# Representing and comparing probabilities with kernels: Part 3 

Arthur Gretton

Gatsby Computational Neuroscience Unit, University College London

MLSS Madrid, 2018

## Training GANs with MMD

## What is a Generative Adversarial Network (GAN)?

- Generator (student)

- Task: critic must teach generator to draw images (here dogs)
- Critic (teacher)




## What is a Generative Adversarial Network (GAN)?



## What is a Generative Adversarial Network (GAN)?



## What is a Generative Adversarial Network (GAN)?



## Why is classification not enough?



## Classification not enough! Need to compare sets

(otherwise student can just produce the same dog over and over)

## MMD for GAN critic

## Can you use MMD as a critic to train GANs?

From ICML 2015:

## Generative Moment Matching Networks

Yujia Li ${ }^{1}$
Kevin Swersky ${ }^{1}$
Richard Zemel ${ }^{1,2}$
${ }^{1}$ Department of Computer Science, University of Toronto, Toronto, ON, CANADA
${ }^{2}$ Canadian Institute for Advanced Research, Toronto, ON, CANADA

YUJIALI@CS.TORONTO.EDU
KSWERSKY@CS.TORONTO.EDU
ZEMEL@CS.TORONTO.EDU

## From UAI 2015:

# Training generative neural networks via Maximum Mean Discrepancy optimization 

University of Cambridge

## Daniel M. Roy

University of Toronto

## Zoubin Ghahramani

University of Cambridge

## MMD for GAN critic

Can you use MMD as a critic to train GANs?


Need better image features.

## How to improve the critic witness

■ Add convolutional features!

- The critic (teacher) also needs to be trained.

■ How to regularise?


MMD GAN Li et al., [NIPS 2017]
Coulomb GAN Unterthiner et al., [ICLR 2018]

## WGAN-GP

Wasserstein GAN Arjovsky et al. [ICML 2017] WGAN-GP Gukrajani et al. [NIPS 2017]


## WGAN-GP

## Wasserstein GAN Arjovsky et al. [ICML 2017] WGAN-GP Gukrajani et al. [NIPS 2017]

Given a generator $G_{\theta}$ with parameters $\theta$ to be trained. Samples $Y \sim G_{\theta}(Z)$ where $Z \sim R$


Given critic features $h_{\psi}$ with parameters $\psi$ to be trained. $f_{\psi}$ a linear function of $h_{\psi}$.

## WGAN-GP

## Wasserstein GAN Arjovsky et al. [ICML 2017]

WGAN-GP Gukrajani et al. [NIPS 2017]

- . 14 Given a generator $G_{\theta}$ with parameters $\theta$ to be trained. Samples $Y \sim G_{\theta}(Z)$ where $Z \sim R$


Given critic features $h_{\psi}$ with parameters $\psi$ to be trained. $f_{\psi}$ a linear function of $h_{\psi}$.
WGAN-GP gradient penalty:

$$
\max _{\psi} \mathbf{E}_{X \sim P} f_{\psi}(X)-\mathbf{E}_{Z \sim R} f_{\psi}\left(G_{\theta}(Z)\right)+\lambda \mathbf{E}_{\tilde{X}}\left(\left\|\nabla_{\tilde{X}} f_{\theta}(\widetilde{X})\right\|-1\right)^{2}
$$

where

$$
\begin{aligned}
\widetilde{X} & =\gamma x_{i}+(1-\gamma) G_{\psi}\left(z_{j}\right) \\
\gamma & \sim \mathcal{U}([0,1]) \quad x_{i} \in\left\{x_{\ell}\right\}_{\ell=1}^{m} \quad z_{j} \in\left\{z_{\ell}\right\}_{\ell=1}^{n}
\end{aligned}
$$

## The (W)MMD

Train MMD critic features with the witness function gradient penalty Binkowski, Sutherland, Arbel, G. [ICLR 2018], Bellemare et al. [2017] for energy distance:

$$
\max _{\psi} M M D^{2}\left(h_{\psi}(X), h_{\psi}\left(G_{\theta}(Z)\right)\right)+\lambda \mathbf{E}_{\widetilde{X}}\left(\left\|\nabla_{\widetilde{X}} f_{\psi}(\widetilde{X})\right\|-1\right)^{2}
$$

where

$$
\begin{aligned}
& \begin{array}{c}
f_{\psi}(\cdot)=\frac{1}{m} \sum_{i=1}^{m} k\left(h_{\psi}\left(x_{i}\right), \cdot\right)-\frac{1}{n} \sum_{j=1}^{n} k\left(h_{\psi}\left(G_{\theta}\left(z_{j}\right)\right), \cdot\right) \\
\text { New }
\end{array} \\
& \widetilde{X}=\gamma x_{i}+(1-\gamma) G_{\psi}\left(z_{j}\right) \\
& \gamma \sim \mathcal{U}([0,1]) \quad x_{i} \in\left\{x_{\ell}\right\}_{\ell=1}^{m} \quad z_{j} \in\left\{z_{\ell}\right\}_{\ell=1}^{n}
\end{aligned}
$$

Remark by Bottou et al. (2017): gradient penalty modifies the function class. So crit $8 /$ /\$ 1 not an MMD in RKHS $\mathcal{F}$.

## MMD for GAN critic: revisited

From ICLR 2018:

# DEMYSTIFYING MMD GANS 

Mikołaj Bińkowski*<br>Department of Mathematics<br>Imperial College London<br>mikbinkowski@gmail.com<br>Dougal J. Sutherland, Michael Arbel \& Arthur Gretton<br>Gatsby Computational Neuroscience Unit<br>University College London<br>\{dougal,michael.n.arbel, arthur.gretton\}@gmail.com

## MMD for GAN critic: revisited



Samples are better!

## MMD for GAN critic: revisited



Samples are better!
Can we do better still?

## Convergence issues for WGAN-GP penalty

WGAN-GP style gradient penalty may not converge near solution
Nagarajan and Kolter [NIPS 2017], Mescheder et al. [ICML 2018], Balduzzi et al. [ICML 2018]
The Dirac-GAN

$$
P=\delta_{0} \quad Q=\delta_{\theta} \quad f_{\psi}(x)=\psi \cdot x
$$



## Convergence issues for WGAN-GP penalty

WGAN-GP style gradient penalty may not converge near solution
Nagarajan and Kolter [NIPS 2017], Mescheder et al. [ICML 2018], Balduzzi et al. [ICML 2018]
The Dirac-GAN

$$
P=\delta_{0} \quad Q=\delta_{\theta} \quad f_{\psi}(x)=\psi \cdot x
$$



## A better gradient penalty

■ New MMD GAN witness regulariser (just accepted, NIPS 2018) Arbel, Sutherland, Binkowski, G. [NIPS 2018]
■ Based on semi-supervised learning regulariser Bousquet et al. [NIPS 2004]

- Related to Sobolev GAN Mroueh et al. [ICLR 2018]

```
arXiv.org > stat > arXiv.1805.11565
    Statistics > Machine Learning
    On gradient regularizers for MMD GANs
    Michael Arbel, Dougal J. Sutherland, Mikołaj Bińkowski, Arthur Gretton
    (Submitted on 29 May 2018)
```


## A better gradient penalty

■ New MMD GAN witness regulariser (just accepted, NIPS 2018)
Arbel, Sutherland, Binkowski, G. [NIPS 2018]
■ Based on semi-supervised learning regulariser Bousquet et al. [NIPS 2004]
■ Related to Sobolev GAN Mroueh et al. [ICLR 2018]


## A better gradient penalty

■ New MMD GAN witness regulariser (just accepted, NIPS 2018)

## Arbel, Sutherland, Binkowski, G. [NIPS 2018]

■ Based on semi-supervised learning regulariser Bousquet et al. [NIPS 2004]
■ Related to Sobolev GAN Mroueh et al. [ICLR 2018]
Modified witness function:

$$
\widetilde{M M D}:=\sup _{\|f\|_{S} \leq 1}\left[\mathbb{E}_{P} f(X)-\mathbb{E}_{Q} f(Y)\right]
$$

where

$$
\begin{aligned}
\|f\|_{S}^{2} & =\|f\|_{L_{2}(P)}^{2}+\|\nabla f\|_{L_{2}(P)}^{2}+\lambda\|f\|_{k}^{2} \\
& \begin{array}{c}
L_{2} \text { norm } \\
\text { control }
\end{array} \\
\begin{array}{c}
\text { Gradient } \\
\text { control }
\end{array} & \begin{array}{c}
\text { RKHS } \\
\text { smoothness }
\end{array}
\end{aligned}
$$

## A better gradient penalty

■ New MMD GAN witness regulariser (just accepted, NIPS 2018)

## Arbel, Sutherland, Binkowski, G. [NIPS 2018]

■ Based on semi-supervised learning regulariser Bousquet et al. [NIPS 2004]
■ Related to Sobolev GAN Mroueh et al. [ICLR 2018]
Modified witness function:

$$
\widetilde{M M D}:=\sup _{\|f\|_{S} \leq 1}\left[\mathbb{E}_{P} f(X)-\mathbb{E}_{Q} f(Y)\right]
$$

where

$$
\begin{aligned}
\|f\|_{S}^{2}= & \|f\|_{L_{2}(P)}^{2}+\|\nabla f\|_{L_{2}(P)}^{2}+\lambda\|f\|_{k}^{2} \\
& \begin{array}{c}
\mathrm{L}_{2} \text { norm } \\
\text { control }
\end{array} \\
\begin{array}{c}
\text { Gradient } \\
\text { control }
\end{array} & \begin{array}{c}
\text { RKHS } \\
\text { smoothness }
\end{array}
\end{aligned}
$$

Problem: not computationally feasible: $O\left(n^{3}\right)$ per iteration.

## A better gradient penalty

■ New MMD GAN witness regulariser (just accepted, NIPS 2018)

## Arbel, Sutherland, Binkowski, G. [NIPS 2018]

■ Based on semi-supervised learning regulariser Bousquet et al. [NIPS 2004]
■ Related to Sobolev GAN Mroueh et al. [ICLR 2018]
The scaled MMD:

$$
S M M D=\sigma_{k, P, \lambda} M M D
$$

where

$$
\sigma_{k, P, \lambda}=\left(\lambda+\int k(x, x) d P(x)+\sum_{i=1}^{d} \int \partial_{i} \partial_{i+d} k(x, x) d P(x)\right)^{-1 / 2}
$$

Replace expensive constraint with cheap upper bound:

$$
\|f\|_{S}^{2} \leq \sigma_{k, P, \lambda}^{-1}\|f\|_{k}^{2}
$$

## A better gradient penalty

■ New MMD GAN witness regulariser (just accepted, NIPS 2018)

## Arbel, Sutherland, Binkowski, G. [NIPS 2018]

■ Based on semi-supervised learning regulariser Bousquet et al. [NIPS 2004]
■ Related to Sobolev GAN Mroueh et al. [ICLR 2018]
The scaled MMD:

$$
S M M D=\sigma_{k, P, \lambda} M M D
$$

where

$$
\sigma_{k, P, \lambda}=\left(\lambda+\int k(x, x) d P(x)+\sum_{i=1}^{d} \int \partial_{i} \partial_{i+d} k(x, x) d P(x)\right)^{-1 / 2}
$$

Replace expensive constraint with cheap upper bound:

$$
\|f\|_{S}^{2} \leq \sigma_{k, P, \lambda}^{-1}\|f\|_{k}^{2}
$$

Idea: rather than regularise the critic or witness function, regularise features directly

## Evaluation and experiments

## Evaluation of GANs

The inception score? Salimans et al. [NIPS 2016]
Based on the classification output $p(y \mid x)$ of the inception model szegedy et al. [ICLR 2014],

$$
E_{X} \exp K L(P(y \mid X) \| P(y))
$$

High when:

- predictive label distribution $P(y \mid x)$ has low entropy (good quality images)
■ label entropy $P(y)$ is high (good variety).


## Evaluation of GANs

The inception score? Salimans et al. [NIPS 2016]
Based on the classification output $p(y \mid x)$ of the inception model szegedy
et al. [ICLR 2014],

$$
E_{X} \exp K L(P(y \mid X) \| P(y))
$$

High when:

- predictive label distribution $P(y \mid x)$ has low entropy (good quality images)
■ label entropy $P(y)$ is high (good variety).

Problem: relies on a trained classifier! Can't be used on new categories (celeb, bedroom...)

## Evaluation of GANs

The Frechet inception distance? Heusel et al. [NiPs 2017]
Fits Gaussians to features in the inception architecture (pool3 layer):

$$
F I D(P, Q)=\left\|\mu_{P}-\mu_{Q}\right\|^{2}+\operatorname{tr}\left(\Sigma_{P}\right)+\operatorname{tr}\left(\Sigma_{Q}\right)-2 \operatorname{tr}\left(\left(\Sigma_{P} \Sigma_{Q}\right)^{\frac{1}{2}}\right)
$$

where $\mu_{P}$ and $\Sigma_{P}$ are the feature mean and covariance of $P$

## Evaluation of GANs

The Frechet inception distance? Heusel et al. [NIPS 2017]
Fits Gaussians to features in the inception architecture (pool3 layer):

$$
F I D(P, Q)=\left\|\mu_{P}-\mu_{Q}\right\|^{2}+\operatorname{tr}\left(\Sigma_{P}\right)+\operatorname{tr}\left(\Sigma_{Q}\right)-2 \operatorname{tr}\left(\left(\Sigma_{P} \Sigma_{Q}\right)^{\frac{1}{2}}\right)
$$

where $\mu_{P}$ and $\Sigma_{P}$ are the feature mean and covariance of $P$

Problem: bias. For finite samples can consistently give incorrect answer.

- Bias demo, CIFAR-10 train vs test



## Evaluation of GANs

The FID can give the wrong answer in theory.
Assume $m$ samples from $P$ and $n \rightarrow \infty$ samples from $Q$.
Given two alternatives:

$$
P_{1} \sim \mathcal{N}\left(0,\left(1-m^{-1}\right)^{2}\right) \quad P_{2} \sim \mathcal{N}(0,1) \quad Q \sim \mathcal{N}(0,1) .
$$

Clearly,

$$
\operatorname{FID}\left(P_{1}, Q\right)=\frac{1}{m^{2}}>\operatorname{FID}\left(P_{2}, Q\right)=0
$$

Given $m$ samples from $P_{1}$ and $P_{2}$,

$$
F I D\left(\widehat{P_{1}}, Q\right)<F I D\left(\widehat{P_{2}}, Q\right)
$$

## Evaluation of GANs

The FID can give the wrong answer in theory.
Assume $m$ samples from $P$ and $n \rightarrow \infty$ samples from $Q$.
Given two alternatives:

$$
P_{1} \sim \mathcal{N}\left(0,\left(1-m^{-1}\right)^{2}\right) \quad P_{2} \sim \mathcal{N}(0,1) \quad Q \sim \mathcal{N}(0,1) .
$$

Clearly,


Given $m$ samples from $P_{1}$ and $P_{2}$,
$F I D\left(\widehat{P_{1}}, Q\right)<\operatorname{FID}\left(\widehat{P_{2}}, Q\right)$.

## Evaluation of GANs

The FID can give the wrong answer in theory.
Assume $m$ samples from $P$ and $n \rightarrow \infty$ samples from $Q$.
Given two alternatives:

$$
P_{1} \sim \mathcal{N}\left(0,\left(1-m^{-1}\right)^{2}\right) \quad P_{2} \sim \mathcal{N}(0,1) \quad Q \sim \mathcal{N}(0,1) .
$$

Clearly,

$$
\operatorname{FID}\left(P_{1}, Q\right)=\frac{1}{m^{2}}>\operatorname{FID}\left(P_{2}, Q\right)=0
$$

Given $m$ samples from $P_{1}$ and $P_{2}$,
$\operatorname{FID}\left(\widehat{P_{1}}, Q\right)<\operatorname{FID}\left(\widehat{P_{2}}, Q\right)$.

## Evaluation of GANs

The FID can give the wrong answer in theory.
Assume $m$ samples from $P$ and $n \rightarrow \infty$ samples from $Q$.
Given two alternatives:

$$
P_{1} \sim \mathcal{N}\left(0,\left(1-m^{-1}\right)^{2}\right) \quad P_{2} \sim \mathcal{N}(0,1) \quad Q \sim \mathcal{N}(0,1) .
$$

Clearly,

$$
\operatorname{FID}\left(P_{1}, Q\right)=\frac{1}{m^{2}}>\operatorname{FID}\left(P_{2}, Q\right)=0
$$

Given $m$ samples from $P_{1}$ and $P_{2}$,

$$
F I D\left(\widehat{P_{1}}, Q\right)<F I D\left(\widehat{P_{2}}, Q\right)
$$

## Evaluation of GANs

The FID can give the wrong answer in practice.
Let $d=2048$, and define
$P_{1}=\operatorname{relu}\left(\mathcal{N}\left(0, I_{d}\right)\right) \quad P_{2}=\operatorname{relu}\left(\mathcal{N}\left(1, .8 \Sigma+.2 I_{d}\right)\right) \quad Q=\operatorname{relu}\left(\mathcal{N}\left(1, I_{d}\right)\right)$
where $\Sigma=\frac{4}{d} C C^{T}$, with $C$ a $d \times d$ matrix with iid standard normal
entries.
For a random draw of $C$ :

$$
F I D\left(P_{1}, Q\right) \approx 1123.0>1114.8 \approx F I D\left(P_{2}, Q\right)
$$

With $m=50000$ samples,

$$
F I D\left(\widehat{P_{1}}, Q\right) \approx 1133.7<1136.2 \approx F I D\left(\widehat{P_{2}}, Q\right)
$$

At $m=100000$ samples, the ordering of the estimates is correct.

## Evaluation of GANs

The FID can give the wrong answer in practice.
Let $d=2048$, and define
$P_{1}=\operatorname{relu}\left(\mathcal{N}\left(0, I_{d}\right)\right) \quad P_{2}=\operatorname{relu}\left(\mathcal{N}\left(1, .8 \Sigma+.2 I_{d}\right)\right) \quad Q=\operatorname{relu}\left(\mathcal{N}\left(1, I_{d}\right)\right)$
where $\Sigma=\frac{4}{d} C C^{T}$, with $C$ a $d \times d$ matrix with iid standard normal entries.

For a random draw of $C$ :


With $m=50000$ samples,


## Evaluation of GANs

The FID can give the wrong answer in practice.
Let $d=2048$, and define
$P_{1}=\operatorname{relu}\left(\mathcal{N}\left(0, I_{d}\right)\right) \quad P_{2}=\operatorname{relu}\left(\mathcal{N}\left(1, .8 \Sigma+.2 I_{d}\right)\right) \quad Q=\operatorname{relu}\left(\mathcal{N}\left(1, I_{d}\right)\right)$
where $\Sigma=\frac{4}{d} C C^{T}$, with $C$ a $d \times d$ matrix with iid standard normal entries.
For a random draw of $C$ :

$$
F I D\left(P_{1}, Q\right) \approx 1123.0>1114.8 \approx F I D\left(P_{2}, Q\right)
$$

With $m=50000$ samples,

$$
F I D\left(\widehat{P_{1}}, Q\right) \approx 1133.7<1136.2 \approx F I D\left(\widehat{P_{2}}, Q\right)
$$

## Evaluation of GANs

The FID can give the wrong answer in practice.
Let $d=2048$, and define
$P_{1}=\operatorname{relu}\left(\mathcal{N}\left(0, I_{d}\right)\right) \quad P_{2}=\operatorname{relu}\left(\mathcal{N}\left(1, .8 \Sigma+.2 I_{d}\right)\right) \quad Q=\operatorname{relu}\left(\mathcal{N}\left(1, I_{d}\right)\right)$ where $\Sigma=\frac{4}{d} C C^{T}$, with $C$ a $d \times d$ matrix with iid standard normal entries.
For a random draw of $C$ :

$$
F I D\left(P_{1}, Q\right) \approx 1123.0>1114.8 \approx F I D\left(P_{2}, Q\right)
$$

With $m=50000$ samples,

$$
F I D\left(\widehat{P_{1}}, Q\right) \approx 1133.7<1136.2 \approx F I D\left(\widehat{P_{2}}, Q\right)
$$

At $m=100000$ samples, the ordering of the estimates is correct. This behavior is similar for other random draws of $C$.

## The kernel inception distance (KID)

The Kernel inception distance Binkowski, Sutherland, Arbel, G. [ICLR 2018] Measures similarity of the samples' representations in the inception architecture (pool3 layer) MMD with kernel

$$
k(x, y)=\left(\frac{1}{d} x^{\top} y+1\right)^{3}
$$

- Checks match for feature means, variances, skewness
■ Unbiased: eg CIFAR-10 train/test



## The kernel inception distance (KID)

The Kernel inception distance Binkowski, Sutherland, Arbel, G. [ICLR 2018]
Measures similarity of the samples' representations in the inception architecture (pool3 layer)
MMD with kernel

$$
k(x, y)=\left(\frac{1}{d} x^{\top} y+1\right)^{3}
$$

- Checks match for feature means, variances, skewness
■ Unbiased : eg CIFAR-10 train/test

..."but isn't KID is computationally costly?"


## The kernel inception distance (KID)

The Kernel inception distance Binkowski, Sutherland, Arbel, G. [ICLR 2018] Measures similarity of the samples' representations in the inception architecture (pool3 layer)
MMD with kernel

$$
k(x, y)=\left(\frac{1}{d} x^{\top} y+1\right)^{3} .
$$

- Checks match for feature means, variances, skewness

■ Unbiased : eg CIFAR-10 train/test

..."but isn't KID is computationally costly?"
"Block" KID implementation is cheaper than FID: see paper (or use Tensorflow implementation)!

## The kernel inception distance (KID)

The Kernel inception distance Binkowski, Sutherland, Arbel, G. [ICLR 2018]
Measures similarity of the samples' representations in the inception architecture (pool3 layer)
MMD with kernel

$$
k(x, y)=\left(\frac{1}{d} x^{\top} y+1\right)^{3} .
$$

- Checks match for feature means, variances, skewness

■ Unbiased : eg CIFAR-10 train/test


Also used for automatic learning rate adjustment: if $K I D\left(\widehat{P}_{t+1}, Q\right)$ not significantly better than $K I D\left(\widehat{P}_{t}, Q\right)$ then reduce learning rate.
[Bounliphone et al. ICLR 2016]

Related: "An empirical study on evaluation metrics of generative adversarial networks", Xu et al. $\mathbf{1 7} / \mathbf{7} \mathbf{7 1}$, June 2018]

## Benchmarks for comparison (all from ICLR 2018)

## Spectral Normalization <br> for Generative Adversarial Networks



BOUNDARY-SEEKING
Generative Adversarial Networks

R Devon Hjelm*
MILA, University of Montréal, IVADO erroneus?gmail.com

## Tong Che

MILA, University of Montróal
tong, che?umontreal.ca

Kyunghyun Cho
New York University,
CIFAR Azrieli Global Scholar
kyunghyun.chognyu.edu

Athul Paul Jacob ${ }^{-}$
MILA, MSR, University of Waterloo
apjacob?edu. uxaterloo.ca

## Adam Trischler

MSR
adam.trischleramicrosoft.com

Yoshua Bengio
MILA, University of Montréal, CIFAR, IVADO
yoshua.bengioßumont real.ca

## Results: what does MMD buy you?

- Critic features from DCGAN: an $f$-filter critic has $f, 2 f, 4 f$ and $8 f$ convolutional filters in layers $1-4$. LSUN $64 \times 64$.


MMD GAN samples, $f=64$, $\mathrm{FID}=32, \mathrm{KID}=3$


WGAN samples, $f=64$, $\mathrm{FID}=41, \mathrm{KID}=4$

## Results: what does MMD buy you?

- Critic features from DCGAN: an $f$-filter critic has $f, 2 f, 4 f$ and $8 f$ convolutional filters in layers $1-4$. LSUN $64 \times 64$.


MMD GAN samples, $f=16$, FID $=86, \mathrm{KID}=9$


WGAN samples, $f=16$, $f=64, \mathrm{FID}=293, \mathrm{KID}=39{ }^{19}$

## The kernel inception distance (KID)

Faster training: performance scores vs generator iterations on MNIST


## Results: celebrity faces $160 \times 160$

## KID (FID)

scores:

■ Sobolev GAN:

$$
14 \text { (20) }
$$

- SN-GAN:

18 (28)
■ Old MMD GAN:
13 (21)

- SMMD GAN:

6 (12)
202599 face images, resized and cropped to 160 $\times 160$


## Results: imagenet $64 \times 64$

## KID (FID)

scores:

- BGAN: 47 (44)
- SN-GAN:

$$
44 \text { (48) }
$$

■ SMMD GAN:

$$
35 \text { (37) }
$$

ILSVRC2012 (ImageNet) dataset, 1281167 images, resized to $64 \times 64$. Around 20000 classes.


## Results: imagenet $64 \times 64$

## KID (FID)

scores:

- BGAN:

$$
47 \text { (44) }
$$

- SN-GAN:

44 (48)

- SMMD GAN: 35 (37)

ILSVRC2012 (ImageNet) dataset, 1281167 images, resized to $64 \times 64$. Around 20000 classes.


## Results: imagenet $64 \times 64$

## KID (FID) <br> scores:

- BGAN:

47 (44)

- SN-GAN:

$$
44 \text { (48) }
$$

- SMMD GAN: 35 (37)

ILSVRC2012 (ImageNet) dataset, 1281167 images, resized to $64 \times 64$. Around 20000 classes.


## Summary

■ MMD critic gives state-of-the-art performance for GAN training (FID and KID)

- use convolutional input features
- train with new gradient regulariser

■ Faster training, simpler critic network
■ Reasons for good performance:

- Unlike WGAN-GP, MMD loss still a valid critic when features not optimal
- Kernel features do some of the "work", so simpler $h_{\psi}$ features possible.
- Better gradient/feature regulariser gives better critic

Code for "Demystifying MMD GANs," ICLR 2018, including KID score: https://github.com/mbinkowski/MMD-GAN
Code for new SMMD:
https://github.com/MichaelArbel/Scaled-MMD-GAN

## Testing against a probabilistic model

## Statistical model criticism


$f^{*}(x)$ is the witness function
Can we compute MMD with samples from $Q$ and a model $P$ ?
Problem: usualy can't compute $E_{p} f$ in closed form.

## Stein idea

To get rid of $E_{p} f$ in

$$
\sup _{\|f\|_{\mathcal{F} \leq 1}}\left[E_{q} f-E_{p} f\right]
$$

we define the Stein operator

$$
\left[T_{p} f\right](x)=\frac{1}{p(x)} \frac{d}{d x}(f(x) p(x))
$$

Then

$$
E_{P} T_{P} f=0
$$

subject to appropriate boundary conditions. (Oates, Girolami, Chopin, 2016)

## Stein idea: proof

$$
\begin{aligned}
E_{p}\left[T_{p} f\right] & =\int\left[\frac{1}{p(x)} \frac{d}{d x}(f(x) p(x))\right] p(x) d x \\
& \int\left[\frac{d}{d x}(f(x) p(x))\right] d x \\
& =[f(x) p(x)]_{-\infty}^{\infty}
\end{aligned}
$$

## Stein idea: proof

$$
E_{p}\left[T_{p} f\right]=\int\left[\frac{1}{p(x)} \frac{d}{d x}(f(x) p(x))\right] p(x) d x
$$

## Stein idea: proof

$$
\begin{gathered}
E_{p}\left[T_{p} f\right]=\int\left[\frac{1}{p(x)} \frac{d}{d x}(f(x) p(x))\right] p(x) d x \\
\int\left[\frac{d}{d x}(f(x) p(x))\right] d x
\end{gathered}
$$

## Stein idea: proof

$$
\begin{aligned}
E_{p}\left[T_{p} f\right] & =\int\left[\frac{1}{p(x)} \frac{d}{d x}(f(x) p(x))\right] p(x) d x \\
& \int\left[\frac{d}{d x}(f(x) p(x))\right] d x \\
& =[f(x) p(x)]_{-\infty}^{\infty}
\end{aligned}
$$

## Stein idea: proof

$$
\begin{aligned}
E_{p}\left[T_{p} f\right] & =\int\left[\frac{1}{p(x)} \frac{d}{d x}(f(x) p(x))\right] p(x) d x \\
& \int\left[\frac{d}{d x}(f(x) p(x))\right] d x \\
& =[f(x) p(x)]_{-\infty}^{\infty} \\
& =0
\end{aligned}
$$

## Kernel Stein Discrepancy Stein operator <br> $$
T_{p} f=\frac{1}{p(x)} \frac{d}{d x}(f(x) p(x))
$$

Kernel Stein Discrepancy (KSD)

$$
K S D(p, q, \mathcal{F})=\sup _{\|g\|_{\mathcal{F} \leq 1}} E_{q} T_{p} g-E_{p} T_{p} g
$$

## Kernel Stein Discrepancy

Stein operator

$$
T_{p} f=\frac{1}{p(x)} \frac{d}{d x}(f(x) p(x))
$$

Kernel Stein Discrepancy (KSD)

$$
K S D(p, q, \mathcal{F})=\sup _{\|g\|_{\mathcal{F} \leq 1}} E_{q} T_{p} g-E_{p} T_{p g}=\sup _{\|g\|_{\mathcal{F}} \leq 1} E_{q} T_{p} g
$$

## Kernel Stein Discrepancy

Stein operator

$$
T_{p} f=\frac{1}{p(x)} \frac{d}{d x}(f(x) p(x))
$$

Kernel Stein Discrepancy (KSD)

$$
K S D(p, q, \mathcal{F})=\sup _{\|g\|_{\mathcal{F} \leq 1}} E_{q} T_{p} g-E_{p} T_{p g}=\sup _{\|g\|_{\mathcal{F}} \leq 1} E_{q} T_{p} g
$$



## Kernel Stein Discrepancy

Stein operator

$$
T_{p} f=\frac{1}{p(x)} \frac{d}{d x}(f(x) p(x))
$$

Kernel Stein Discrepancy (KSD)

$$
K S D(p, q, \mathcal{F})=\sup _{\|g\|_{\mathcal{F}} \leq 1} E_{q} T_{p} g-E_{p} T_{p g}=\sup _{\|g\|_{\mathcal{F}} \leq 1} E_{q} T_{p} g
$$



## Kernel stein discrepancy

Closed-form expression for KSD: given $Z, Z^{\prime} \sim q$, then (Chwialkowski, Strathmann, G., ICML 2016) (Liu, Lee, Jordan ICML 2016)

$$
\operatorname{KSD}(p, q, \mathcal{F})=E_{q} h_{p}\left(Z, Z^{\prime}\right)
$$

where

$$
\begin{aligned}
h_{p}(x, y) & :=\partial_{x} \log p(x) \partial_{x} \log p(y) k(x, y) \\
& +\partial_{y} \log p(y) \partial_{x} k(x, y) \\
& +\partial_{x} \log p(x) \partial_{y} k(x, y) \\
& +\partial_{x} \partial_{y} k(x, y)
\end{aligned}
$$

and $k$ is RKHS kernel for $\mathcal{F}$
Only depends on kernel and $\partial_{x} \log p(x)$. Do not need to normalize $p$, or sample from it.

## Statistical model criticism



Chicago crime data

## Statistical model criticism



Chicago crime data
Model is Gaussian mixture with two components.

## Statistical model criticism



Chicago crime data
Model is Gaussian mixture with two components Stein witness function

## Statistical model criticism



Chicago crime data
Model is Gaussian mixture with ten components.

## Statistical model criticism



## Chicago crime data

Model is Gaussian mixture with ten components Stein witness function
Code: https://github.com/karlnapf/kernel_goodness_of_fit

## Kernel stein discrepancy

Further applications:

- Evaluation of approximate MCMC methods.
(Chwialkowski, Strathmann, G., ICML 2016; Gorham, Mackey, ICML 2017)


## What kernel to use?

■ The inverse multiquadric kernel,

$$
k(x, y)=\left(c+\|x-y\|_{2}^{2}\right)^{\beta}
$$

for $\beta \in(-1,0)$.

```
arXiv.org > stat > arXiv:1703.01717
Statistics > Machine Learning
Measuring Sample Quality with Kernels
Jackson Gorham, Lester Mackey
ICML 2017
(Submitted on 6 Mar 2017 (v1), last revised 3 Aug 2017 (this version, v6))
```


# Testing statistical dependence 

## Dependence testing

■ Given: Samples from a distribution $P_{X Y}$
$\square$ Goal: Are $X$ and $Y$ independent?


## MMD as a dependence measure?

Could we use MMD?

$$
M M D(\underbrace{P_{X Y}}_{P}, \underbrace{P_{X} P_{Y}}_{Q}, \mathcal{H}_{K})
$$

## We don't have samples from $Q:=P_{X} P_{Y}$, only pairs

- Solution: simulate $Q$ with pairs $\left(x_{i}, y_{j}\right)$ for $j \neq i$.


## What kernel $\kappa$ to use for the RKHS $\mathcal{H}_{\kappa}$ ?

## MMD as a dependence measure?

Could we use MMD?

$$
M M D(\underbrace{P_{X Y}}_{P}, \underbrace{P_{X} P_{Y}}_{Q}, \mathcal{H}_{\kappa})
$$

■ We don't have samples from $Q:=P_{X} P_{Y}$, only pairs $\left\{\left(x_{i}, y_{i}\right\}_{i=1}^{n} \stackrel{\text { i.i.d. }}{\sim} P_{X Y}\right.$

- Solution: simulate $Q$ with pairs $\left(x_{i}, y_{j}\right)$ for $j \neq i$.

$$
\text { What kernel } \kappa \text { to use for the RKHS } \mathcal{H}_{\kappa} \text { ? }
$$

## MMD as a dependence measure?

Could we use MMD?

$$
M M D(\underbrace{P_{X Y}}_{P}, \underbrace{P_{X} P_{Y}}_{Q}, \mathcal{H}_{\kappa})
$$

■ We don't have samples from $Q:=P_{X} P_{Y}$, only pairs $\left\{\left(x_{i}, y_{i}\right\}_{i=1}^{n} \stackrel{\text { i.i.d. }}{\sim} P_{X Y}\right.$

- Solution: simulate $Q$ with pairs $\left(x_{i}, y_{j}\right)$ for $j \neq i$.

■ What kernel $\kappa$ to use for the RKHS $\mathcal{H}_{\kappa}$ ?

## MMD as a dependence measure

Kernel $k$ on images with feature space $\mathcal{F}$,

$$
K(F, j)
$$

Kernel $l$ on captions with feature space $\mathcal{G}$,


## MMD as a dependence measure

Kernel $k$ on images with feature space $\mathcal{F}$,


Kernel $l$ on captions with feature space $\mathcal{G}$,


Kernel $\kappa$ on image-text pairs: are images and captions similar?


$$
=k(\pi, \pi) \times l(\approx, \equiv)
$$

## MMD as a dependence measure

- Given: Samples from a distribution $P_{X Y}$
- Goal: Are $X$ and $Y$ independent?

$$
\begin{aligned}
& M M D^{2}\left(\widehat{P}_{X Y}, \widehat{P}_{X} \widehat{P}_{Y}, \mathcal{H}_{\kappa}\right):=\frac{1}{n^{2}} \operatorname{trace}(K L) \\
& (\mathrm{K}, \text { L column centered })
\end{aligned}
$$

## MMD as a dependence measure

■ Given: Samples from a distribution $P_{X Y}$
■ Goal: Are $X$ and $Y$ independent?

$$
M M D^{2}\left(\widehat{P}_{X Y}, \widehat{P}_{X} \widehat{P}_{Y}, \mathcal{H}_{\kappa}\right):=\frac{1}{n^{2}} \operatorname{trace}(K L)
$$



## MMD as a dependence measure

Two questions:
■ Why the product kernel? Many ways to combine kernels - why not eg a sum?

- Is there a more interpretable way of defining this dependence measure?


## Illustration: dependence $\neq$ correlation

■ Given: Samples from a distribution $P_{X Y}$

- Goal: Are $X$ and $Y$ dependent?



## Illustration: dependence $\neq$ correlation

■ Given: Samples from a distribution $P_{X Y}$

- Goal: Are $X$ and $Y$ dependent?

Correlation: 0.07


## Illustration: dependence $\neq$ correlation

- Given: Samples from a distribution $P_{X Y}$
- Goal: Are $X$ and $Y$ dependent?



## Finding covariance with smooth transformations

Illustration: two variables with no correlation but strong dependence.


## Finding covariance with smooth transformations

Illustration: two variables with no correlation but strong dependence.


## Finding covariance with smooth transformations

Illustration: two variables with no correlation but strong dependence.


## Define two spaces, one for each witness

Function in $\mathcal{F}$

$$
f(x)=\sum_{j=1}^{\infty} f_{j} \varphi_{j}(x)
$$

Feature map

| $\varphi(x)=$ | $\left[\varphi_{1}(x) \bigcap \bigcap\right.$ |
| :---: | :---: |
|  | ${ }^{\varphi_{2}(x)}$ ¢ |
|  | $\varphi_{3}(x)$ |

Kernel for RKHS $\mathcal{F}$ on $\mathcal{X}$ :

$$
k\left(x, x^{\prime}\right)=\left\langle\varphi(x), \varphi\left(x^{\prime}\right)\right\rangle_{\mathcal{F}}
$$

Function in $\mathcal{G}$

$$
g(y)=\sum_{j=1}^{\infty} g_{j} \phi_{j}(y)
$$

Feature map


Kernel for RKHS $\mathcal{G}$ on $\mathcal{Y}$ :

$$
l\left(x, x^{\prime}\right)=\left\langle\phi(y), \phi\left(y^{\prime}\right)\right\rangle_{\mathcal{G}}
$$

## The constrained covariance

The constrained covariance is

$$
\operatorname{COCO}\left(P_{X Y}\right)=\sup ^{\|f\|_{\mathcal{F}} \leq 1} \operatorname{cov}[f(x) g(y)]
$$



## The constrained covariance

The constrained covariance is

$$
\operatorname{COCO}\left(P_{X Y}\right)=\sup _{\substack{\|f\|_{\mathcal{F}} \leq 1 \\\|g\|_{\mathcal{G}} \leq 1}} \operatorname{cov}\left[\left(\sum_{j=1}^{\infty} f_{j} \varphi_{j}(x)\right)\left(\sum_{j=1}^{\infty} g_{j} \phi_{j}(y)\right)\right]
$$

## The constrained covariance

The constrained covariance is

$$
\operatorname{COCO}\left(P_{X Y}\right)=\sup _{\substack{\|f\|_{\mathcal{F}} \leq 1 \\\|g\|_{\mathcal{G}} \leq 1}} E_{x y}\left[\left(\sum_{j=1}^{\infty} f_{j} \varphi_{j}(x)\right)\left(\sum_{j=1}^{\infty} g_{j} \phi_{j}(y)\right)\right]
$$

Fine print: feature mappings $\varphi(x)$ and $\phi(y)$ assumed to have zero mean.

## The constrained covariance

The constrained covariance is

$$
\operatorname{COCO}\left(P_{X Y}\right)=\sup _{\substack{\|f\|_{\mathcal{F}} \leq 1 \\\|g\|_{\mathcal{G}} \leq 1}} E_{x y}\left[\left(\sum_{j=1}^{\infty} f_{j} \varphi_{j}(x)\right)\left(\sum_{j=1}^{\infty} g_{j} \phi_{j}(y)\right)\right]
$$

Fine print: feature mappings $\varphi(x)$ and $\phi(y)$ assumed to have zero mean. Rewriting:

$$
\begin{aligned}
& E_{x y}[f(x) g(y)] \\
& =\left[\begin{array}{c}
f_{1} \\
f_{2} \\
\vdots
\end{array}\right]^{\top} \underbrace{\mathbf{E}_{x y}\left(\left[\begin{array}{c}
\varphi_{1}(x) \\
\varphi_{2}(x) \\
\vdots
\end{array}\right]\left[\begin{array}{lll}
\phi_{1}(y) & \phi_{2}(y) & \ldots
\end{array}\right]\right)}_{C_{\varphi(x) \phi(y)}}\left[\begin{array}{c}
g_{1} \\
g_{2} \\
\vdots
\end{array}\right]
\end{aligned}
$$

## The constrained covariance

The constrained covariance is

$$
\operatorname{COCO}\left(P_{X Y}\right)=\sup _{\substack{\|f\|_{\mathcal{F}} \leq 1 \\ \\\|g\|_{\mathcal{G}} \leq 1}} E_{x y}\left[\left(\sum_{j=1}^{\infty} f_{j} \varphi_{j}(x)\right)\left(\sum_{j=1}^{\infty} g_{j} \phi_{j}(y)\right)\right]
$$

Fine print: feature mappings $\varphi(x)$ and $\phi(y)$ assumed to have zero mean.
Rewriting:

$$
\begin{aligned}
& E_{x y}[f(x) g(y)] \\
& =\left[\begin{array}{c}
f_{1} \\
f_{2} \\
\vdots
\end{array}\right]^{\top} \underbrace{\mathbf{E}_{x y}\left(\left[\begin{array}{c}
\varphi_{1}(x) \\
\varphi_{2}(x) \\
\vdots
\end{array}\right]\left[\begin{array}{lll}
\phi_{1}(y) & \phi_{2}(y) & \ldots
\end{array}\right]\right)}_{C_{\varphi(x) \phi(y)}}\left[\begin{array}{c}
g_{1} \\
g_{2} \\
\vdots
\end{array}\right]
\end{aligned}
$$

COCO: max singular value of feature covariance $C_{\varphi(x) \phi\left(y_{41}\right.} /_{71}$

## Computing COCO in practice

Given sample $\left\{\left(x_{i}, y_{i}\right)\right\}_{i=1}^{n} \stackrel{\text { i.i.d. }}{\sim} P_{X Y}$, what is empirical $\widehat{C O C O}$ ?

## Computing COCO in practice

Given sample $\left\{\left(x_{i}, y_{i}\right)\right\}_{i=1}^{n} \stackrel{\text { i.i.d. }}{\sim} P_{X Y}$, what is empirical $\widehat{C O C O}$ ?
$\widehat{C O C O}$ is largest eigenvalue $\gamma_{\max }$ of

$$
\left[\begin{array}{cc}
0 & \frac{1}{n} K L \\
\frac{1}{n} L K & 0
\end{array}\right]\left[\begin{array}{l}
\alpha \\
\beta
\end{array}\right]=\gamma\left[\begin{array}{cc}
K & 0 \\
0 & L
\end{array}\right]\left[\begin{array}{l}
\alpha \\
\beta
\end{array}\right] .
$$

$K_{i j}=k\left(x_{i}, x_{j}\right)$ and $L_{i j}=l\left(y_{i}, y_{j}\right)$.

Fine print: kernels are computed with empirically centered features $\varphi(x)-\frac{1}{n} \sum_{i=1}^{n} \varphi\left(x_{i}\right)$ and $\phi(y)-\frac{1}{n} \sum_{i=1}^{n} \phi\left(y_{i}\right)$.
G., Smola., Bousquet, Herbrich, Belitski, Augath, Murayama, Pauls, Schoelkopf, and Logothetis, AISTATS'05

## Computing COCO in practice

Given sample $\left\{\left(x_{i}, y_{i}\right)\right\}_{i=1}^{n} \stackrel{\text { i.i.d. }}{\sim} P_{X Y}$, what is empirical $\widehat{C O C O}$ ?
$\widehat{C O C O}$ is largest eigenvalue $\gamma_{\text {max }}$ of

$$
\left[\begin{array}{cc}
0 & \frac{1}{n} K L \\
\frac{1}{n} L K & 0
\end{array}\right]\left[\begin{array}{l}
\alpha \\
\beta
\end{array}\right]=\gamma\left[\begin{array}{cc}
K & 0 \\
0 & L
\end{array}\right]\left[\begin{array}{l}
\alpha \\
\beta
\end{array}\right] .
$$

$K_{i j}=k\left(x_{i}, x_{j}\right)$ and $L_{i j}=l\left(y_{i}, y_{j}\right)$.
Witness functions (singular vectors):

$$
f(x) \propto \sum_{i=1}^{n} \alpha_{i} k\left(x_{i}, x\right) \quad g(y) \propto \sum_{i=1}^{n} \beta_{i} l\left(y_{i}, y\right)
$$

Fine print: kernels are computed with empirically centered features $\varphi(x)-\frac{1}{n} \sum_{i=1}^{n} \varphi\left(x_{i}\right)$ and $\phi(y)-\frac{1}{n} \sum_{i=1}^{n} \phi\left(y_{i}\right)$.
G., Smola., Bousquet, Herbrich, Belitski, Augath, Murayama, Pauls, Schoelkopf, and Logothetis, AISTATS'05

## Empirical COCO: proof (1)

The Lagrangian is

$$
\mathcal{L}(f, g, \lambda, \gamma)=\underbrace{\frac{1}{n} \sum_{i=1}^{n}\left[f\left(x_{i}\right) g\left(y_{i}\right)\right]}_{\text {covariance }}-\underbrace{\frac{\lambda}{2}\left(\|f\|_{\mathcal{F}}^{2}-1\right)-\frac{\gamma}{2}\left(\|g\|_{\mathcal{G}}^{2}-1\right)}_{\text {smoothness constraints }}
$$

Fine print: $f\left(x_{i}\right) g\left(y_{i}\right)$ centered to have zero empirical mean.

## Empirical COCO: proof (1)

The Lagrangian is

$$
\mathcal{L}(f, g, \lambda, \gamma)=\underbrace{\frac{1}{n} \sum_{i=1}^{n}\left[f\left(x_{i}\right) g\left(y_{i}\right)\right]}_{\text {covariance }}-\underbrace{\frac{\lambda}{2}\left(\|f\|_{\mathcal{F}}^{2}-1\right)-\frac{\gamma}{2}\left(\|g\|_{\mathcal{G}}^{2}-1\right)}_{\text {smoothness constraints }}
$$

Fine print: $f\left(x_{i}\right) g\left(y_{i}\right)$ centered to have zero empirical mean.
Assume (cf representer theorem):

$$
f=\sum_{i=1}^{n} \alpha_{i} \varphi\left(x_{i}\right) \quad g=\sum_{i=1}^{n} \beta_{i} \psi\left(y_{i}\right)
$$

for centered $\varphi\left(x_{i}\right), \phi\left(y_{i}\right)$.

## Empirical COCO: proof (1)

The Lagrangian is

$$
\mathcal{L}(f, g, \lambda, \gamma)=\underbrace{\frac{1}{n} \sum_{i=1}^{n}\left[f\left(x_{i}\right) g\left(y_{i}\right)\right]}_{\text {covariance }}-\underbrace{\frac{\lambda}{2}\left(\|f\|_{\mathcal{F}}^{2}-1\right)-\frac{\gamma}{2}\left(\|g\|_{\mathcal{G}}^{2}-1\right)}_{\text {smoothness constraints }} .
$$

Fine print: $f\left(x_{i}\right) g\left(y_{i}\right)$ centered to have zero empirical mean.
Assume (cf representer theorem):

$$
f=\sum_{i=1}^{n} \alpha_{i} \varphi\left(x_{i}\right) \quad g=\sum_{i=1}^{n} \beta_{i} \psi\left(y_{i}\right)
$$

for centered $\varphi\left(x_{i}\right), \phi\left(y_{i}\right)$.
First step is smoothness constraint:

$$
\|f\|_{\mathcal{F}}^{2}-1=\langle f, f\rangle_{\mathcal{F}}-1
$$

## Empirical COCO: proof (1)

The Lagrangian is

$$
\mathcal{L}(f, g, \lambda, \gamma)=\underbrace{\frac{1}{n} \sum_{i=1}^{n}\left[f\left(x_{i}\right) g\left(y_{i}\right)\right]}_{\text {covariance }}-\underbrace{\frac{\lambda}{2}\left(\|f\|_{\mathcal{F}}^{2}-1\right)-\frac{\gamma}{2}\left(\|g\|_{\mathcal{G}}^{2}-1\right)}_{\text {smoothness constraints }} .
$$

Fine print: $f\left(x_{i}\right) g\left(y_{i}\right)$ centered to have zero empirical mean.
Assume (cf representer theorem):

$$
f=\sum_{i=1}^{n} \alpha_{i} \varphi\left(x_{i}\right) \quad g=\sum_{i=1}^{n} \beta_{i} \psi\left(y_{i}\right)
$$

for centered $\varphi\left(x_{i}\right), \phi\left(y_{i}\right)$.
First step is smoothness constraint:

$$
\begin{aligned}
\|f\|_{\mathcal{F}}^{2}-1 & =\langle f, f\rangle_{\mathcal{F}}-1 \\
& =\left\langle\sum_{i=1}^{n} \alpha_{i} \varphi\left(x_{i}\right), \sum_{i=1}^{n} \alpha_{i} \varphi\left(x_{i}\right)\right\rangle_{\mathcal{F}}-1
\end{aligned}
$$

## Empirical COCO: proof (1)

The Lagrangian is

$$
\mathcal{L}(f, g, \lambda, \gamma)=\underbrace{\frac{1}{n} \sum_{i=1}^{n}\left[f\left(x_{i}\right) g\left(y_{i}\right)\right]}_{\text {covariance }}-\underbrace{\frac{\lambda}{2}\left(\|f\|_{\mathcal{F}}^{2}-1\right)-\frac{\gamma}{2}\left(\|g\|_{\mathcal{G}}^{2}-1\right)}_{\text {smoothness constraints }} .
$$

Fine print: $f\left(x_{i}\right) g\left(y_{i}\right)$ centered to have zero empirical mean.
Assume (cf representer theorem):

$$
f=\sum_{i=1}^{n} \alpha_{i} \varphi\left(x_{i}\right) \quad g=\sum_{i=1}^{n} \beta_{i} \psi\left(y_{i}\right)
$$

for centered $\varphi\left(x_{i}\right), \phi\left(y_{i}\right)$.
First step is smoothness constraint:

$$
\begin{aligned}
\|f\|_{\mathcal{F}}^{2}-1 & =\langle f, f\rangle_{\mathcal{F}}-1 \\
& =\left\langle\sum_{i=1}^{n} \alpha_{i} \varphi\left(x_{i}\right), \sum_{i=1}^{n} \alpha_{i} \varphi\left(x_{i}\right)\right\rangle_{\mathcal{F}}-1 \\
& =\alpha^{\top} K \alpha-1
\end{aligned}
$$

## Proof sketch (2)

Second step is covariance:

$$
\frac{1}{n} \sum_{i=1}^{n}\left[f\left(x_{i}\right) g\left(y_{i}\right)\right]=\frac{1}{n} \sum_{i=1}^{n}\left\langle f, \varphi\left(x_{i}\right)\right\rangle_{\mathcal{F}}\left\langle g, \varphi\left(y_{i}\right)\right\rangle_{\mathcal{G}}
$$

## Proof sketch (2)

Second step is covariance:

$$
\begin{aligned}
\frac{1}{n} \sum_{i=1}^{n}\left[f\left(x_{i}\right) g\left(y_{i}\right)\right] & =\frac{1}{n} \sum_{i=1}^{n}\left\langle f, \varphi\left(x_{i}\right)\right\rangle_{\mathcal{F}}\left\langle g, \varphi\left(y_{i}\right)\right\rangle_{\mathcal{G}} \\
& =\frac{1}{n} \sum_{i=1}^{n}\left\langle\sum_{\ell=1}^{n} \alpha_{\ell} \varphi\left(x_{\ell}\right), \varphi\left(x_{i}\right)\right\rangle_{\mathcal{F}}\left\langle g, \varphi\left(y_{i}\right)\right\rangle_{\mathcal{G}}
\end{aligned}
$$

## Proof sketch (2)

Second step is covariance:

$$
\begin{aligned}
\frac{1}{n} \sum_{i=1}^{n}\left[f\left(x_{i}\right) g\left(y_{i}\right)\right] & =\frac{1}{n} \sum_{i=1}^{n}\left\langle f, \varphi\left(x_{i}\right)\right\rangle_{\mathcal{F}}\left\langle g, \varphi\left(y_{i}\right)\right\rangle_{\mathcal{G}} \\
& =\frac{1}{n} \sum_{i=1}^{n}\left\langle\sum_{\ell=1}^{n} \alpha_{\ell} \varphi\left(x_{\ell}\right), \varphi\left(x_{i}\right)\right\rangle_{\mathcal{F}}\left\langle g, \varphi\left(y_{i}\right)\right\rangle_{\mathcal{G}} \\
& =\frac{1}{n} \alpha^{\top} K L \beta
\end{aligned}
$$

where $K_{i j}=k\left(x_{i}, x_{j}\right)=\left\langle\varphi\left(x_{i}\right), \varphi\left(x_{j}\right)\right\rangle_{\mathcal{F}} \quad L_{i j}=l\left(y_{i}, y_{j}\right)$.

## Proof sketch (2)

Second step is covariance:

$$
\begin{aligned}
\frac{1}{n} \sum_{i=1}^{n}\left[f\left(x_{i}\right) g\left(y_{i}\right)\right] & =\frac{1}{n} \sum_{i=1}^{n}\left\langle f, \varphi\left(x_{i}\right)\right\rangle_{\mathcal{F}}\left\langle g, \varphi\left(y_{i}\right)\right\rangle_{\mathcal{G}} \\
& =\frac{1}{n} \sum_{i=1}^{n}\left\langle\sum_{\ell=1}^{n} \alpha_{\ell} \varphi\left(x_{\ell}\right), \varphi\left(x_{i}\right)\right\rangle_{\mathcal{F}}\left\langle g, \varphi\left(y_{i}\right)\right\rangle_{\mathcal{G}} \\
& =\frac{1}{n} \alpha^{\top} K L \beta \\
\text { where } K_{i j}=k\left(x_{i}, x_{j}\right)= & \left\langle\varphi\left(x_{i}\right), \varphi\left(x_{j}\right)\right\rangle_{\mathcal{F}} \quad L_{i j}=l\left(y_{i}, y_{j}\right) .
\end{aligned}
$$

The Lagranian is now:

$$
\mathcal{L}(f, g, \lambda, \gamma)=\frac{1}{n} \alpha^{\top} K L \beta-\frac{\lambda}{2}\left(\alpha^{\top} K \alpha-1\right)-\frac{\gamma}{2}\left(\beta^{\top} L \beta-1\right)
$$

## What is a large dependence with COCO?



500 Samples, smooth density


Rough density


500 samples, rough density


Density takes the form:

$$
P_{X Y} \propto 1+\sin (\omega x) \sin (\omega y)
$$

Which of these is the more "dependent"?

## Finding covariance with smooth transformations

Case of $\omega=1$ :




Correlation: $\mathbf{0 . 5 0}$ COCO: 0.09


## Finding covariance with smooth transformations

Case of $\omega=2$ :



Correlation: 0.54


## Finding covariance with smooth transformations

Case of $\omega=3$ :




## Finding covariance with smooth transformations

Case of $\omega=4$ :




Correlation: 0.25 COCO: 0.02


## Finding covariance with smooth transformations

Case of $\omega=$ ??:



Correlation: 0.14 COCO: 0.02


## Finding covariance with smooth transformations

Case of $\omega=0$ : uniform noise! (shows bias)


## Dependence largest when at "low" frequencies

- As dependence is encoded at higher frequencies, the smooth mappings $f, g$ achieve lower linear dependence.
■ Even for independent variables, COCO will not be zero at finite sample sizes, since some mild linear dependence will be found by f,g (bias)
■ This bias will decrease with increasing sample size.


## Can we do better than COCO?

A second example with zero correlation.
First singular value of feature covariance $C_{\varphi(x) \phi(y)}$ :


## Can we do better than COCO?

A second example with zero correlation.
Second singular value of feature covariance $C_{\varphi(x) \phi(y)}$ :


## Can we do better than COCO?

A second example with zero correlation.
Second singular value of feature covariance $C_{\varphi(x) \phi(y)}$ :


## The Hilbert-Schmidt Independence Criterion

Writing the $i$ th singular value of the feature covariance $C_{\varphi(x) \phi(y)}$ as

$$
\gamma_{i}:=\operatorname{COCO}_{i}\left(P_{X Y} ; \mathcal{F}, \mathcal{G}\right)
$$

define Hilbert-Schmidt Independence Criterion (HSIC)

$$
\operatorname{HSIC}^{2}\left(P_{X Y} ; \mathcal{F}, \mathcal{G}\right)=\sum_{i=1}^{\infty} \gamma_{i}^{2}
$$

G, Bousquet , Smola., and Schoelkopf, ALT05; G.., Fukumizu, Teo., Song., Schoelkopf., and Smola, NIPS 2007,.

## The Hilbert-Schmidt Independence Criterion

Writing the $i$ th singular value of the feature covariance $C_{\varphi(x) \phi(y)}$ as

$$
\gamma_{i}:=\operatorname{COCO}_{i}\left(P_{X Y} ; \mathcal{F}, \mathcal{G}\right)
$$

define Hilbert-Schmidt Independence Criterion (HSIC)

$$
\operatorname{HSIC}^{2}\left(P_{X Y} ; \mathcal{F}, \mathcal{G}\right)=\sum_{i=1}^{\infty} \gamma_{i}^{2}
$$

G, Bousquet , Smola., and Schoelkopf, ALT05; G.., Fukumizu, Teo., Song., Schoelkopf., and Smola, NIPS 2007,.
HSIC is MMD with product kernel!

$$
H S I C^{2}\left(P_{X Y} ; \mathcal{F}, \mathcal{G}\right)=M M D^{2}\left(P_{X Y}, P_{X} P_{Y} ; \mathcal{H}_{\kappa}\right)
$$

where $\kappa\left((x, y),\left(x^{\prime}, y^{\prime}\right)\right)=k\left(x, x^{\prime}\right) l\left(y, y^{\prime}\right)$.

## Asymptotics of HSIC under independence

- Given sample $\left\{\left(x_{i}, y_{i}\right\}_{i=1}^{n} \stackrel{\text { i.i.d. }}{\sim} P_{X Y}\right.$, what is empirical $\widehat{H S I C}$ ?
- Empirical HSIC (biased)

$$
\widehat{\text { HSIC }}=\frac{1}{n^{2}} \operatorname{trace}(K L)
$$

$K_{i j}=k\left(x_{i}, x_{j}\right)$ and $L_{i j}=l\left(y_{i} y_{j}\right) \quad$ ( $K$ and $L$ computed with empirically centered features)

- Statistical testing: given $P_{X Y}=P_{X} P_{Y}$, what is the threshold $c_{\alpha}$ such that $P\left(\overline{H S I C}>c_{\alpha}\right)<\alpha$ for small $\alpha$ ?

where $\lambda_{l} \psi_{l}\left(z_{j}\right)=\int h_{i j q r} \psi_{l}\left(z_{i}\right) d F_{i, q, r}, \quad h_{i j q r}=\frac{1}{4!} \sum_{(t, u, v, w)}^{(i, j, q, r)} k_{t u} l_{t u}+k_{t u} l_{v w}-2 k_{t u} l_{t v}$


## Asymptotics of HSIC under independence

- Given sample $\left\{\left(x_{i}, y_{i}\right\}_{i=1}^{n} \stackrel{\text { i.i.d. }}{\sim} P_{X Y}\right.$, what is empirical $\widehat{H S I C}$ ?
- Empirical HSIC (biased)

$$
\begin{gathered}
\widehat{H S I C}=\frac{1}{n^{2}} \operatorname{trace}(K L) \\
\begin{array}{l}
K_{i j}=k\left(x_{i}, x_{j}\right) \text { and } L_{i j}=l\left(y_{i} y_{j}\right) \quad(K \text { and } L \text { computed with } \\
\text { empirically centered features })
\end{array}
\end{gathered}
$$

Statistical testing: given $P_{X Y}=P_{X} P_{Y}$, what is the threshold $c_{\alpha}$ such that $P\left(\widehat{H S I C}>c_{\alpha}\right)<\alpha$ for small $\alpha$ ?

where $\lambda_{l} \psi_{l}\left(z_{j}\right)=\int h_{i j q r} \psi_{l}\left(z_{i}\right) d F_{i, q, r}, \quad h_{i j q r}=\frac{1}{4!} \sum_{(t, u, v, w)}^{(i, j, q, r)} k_{t u} l_{t u}+k_{t u} l_{v w}-2 k_{t u} l_{t v}$

## Asymptotics of HSIC under independence

- Given sample $\left\{\left(x_{i}, y_{i}\right\}_{i=1}^{n} \stackrel{\text { i.i.d. }}{\sim} P_{X Y}\right.$, what is empirical $\widehat{H S I C}$ ?
- Empirical HSIC (biased)

$$
\begin{gathered}
\widehat{H S I C}=\frac{1}{n^{2}} \operatorname{trace}(K L) \\
K_{i j}=k\left(x_{i}, x_{j}\right) \text { and } L_{i j}=l\left(y_{i} y_{j}\right) \quad(K \text { and } L \text { computed with }
\end{gathered}
$$ empirically centered features)

- Statistical testing: given $P_{X Y}=P_{X} P_{Y}$, what is the threshold $c_{\alpha}$ such that $P\left(\widehat{H S I C}>c_{\alpha}\right)<\alpha$ for small $\alpha$ ?

where $\lambda_{l} \psi_{l}\left(z_{j}\right)=\int h_{i j q r} \psi_{l}\left(z_{i}\right) d F_{i, q, r}, \quad h_{i j q r}=\frac{1}{4!} \sum_{(t, u, v, w)}^{(i, j, q, r)} k_{t u} l_{t u}+k_{t u} l_{v w}-2 k_{t u} l_{t v}$


## Asymptotics of HSIC under independence

- Given sample $\left\{\left(x_{i}, y_{i}\right\}_{i=1}^{n} \stackrel{\text { i.i.d. }}{\sim} P_{X Y}\right.$, what is empirical $\widehat{H S I C}$ ?
- Empirical HSIC (biased)

$$
\widehat{\text { HSIC }}=\frac{1}{n^{2}} \operatorname{trace}(K L)
$$

$K_{i j}=k\left(x_{i}, x_{j}\right)$ and $L_{i j}=l\left(y_{i} y_{j}\right) \quad(K$ and $L$ computed with empirically centered features)

- Statistical testing: given $P_{X Y}=P_{X} P_{Y}$, what is the threshold $c_{\alpha}$ such that $P\left(\widehat{H S I C}>c_{\alpha}\right)<\alpha$ for small $\alpha$ ?
- Asymptotics of $\widehat{H S I C}$ when $P_{X Y}=P_{X} P_{Y}$ :

$$
n \widehat{H S I C} \xrightarrow{D} \sum_{l=1}^{\infty} \lambda_{l} z_{l}^{2}, \quad z_{l} \sim \mathcal{N}(0,1) \text { i.i.d. }
$$

where $\lambda_{l} \psi_{l}\left(z_{j}\right)=\int h_{i j g r} \psi_{l}\left(z_{i}\right) d F_{i, q, r}, \quad h_{i j g r}=\frac{1}{4!} \sum_{(t, u, v, w)}^{(i, j, q, r)} k_{t u} l_{t u}+k_{t u} l_{v w}-2 k_{t u} l_{t v}$

## A statistical test

■ Given $P_{X Y}=P_{X} P_{Y}$, what is the threshold $c_{\alpha}$ such that $P\left(\widehat{H S I C}>c_{\alpha}\right)<\alpha$ for small $\alpha$ (prob. of false positive)?

## A statistical test

■ Given $P_{X Y}=P_{X} P_{Y}$, what is the threshold $c_{\alpha}$ such that $P\left(\widehat{H S I C}>c_{\alpha}\right)<\alpha$ for small $\alpha$ (prob. of false positive)?

- Original time series:

$$
\begin{aligned}
& X_{1} X_{2} X_{3} X_{4} X_{5} X_{6} X_{7} X_{8} X_{9} X_{10} \\
& Y_{1} Y_{2} Y_{3} Y_{4} Y_{5} \quad Y_{6} \quad Y_{7} Y_{8} Y_{9} \quad Y_{10}
\end{aligned}
$$

- Permutation:

$$
\begin{aligned}
& X_{1} X_{2} X_{3} X_{4} X_{5} X_{6} X_{7} X_{8} X_{9} X_{10} \\
& Y_{7} Y_{3} \quad Y_{9} \quad Y_{2} \quad Y_{4} \quad Y_{8} \quad Y_{5} \quad Y_{1} \quad Y_{6} \quad Y_{10}
\end{aligned}
$$

## A statistical test

■ Given $P_{X Y}=P_{X} P_{Y}$, what is the threshold $c_{\alpha}$ such that $P\left(\widehat{H S I C}>c_{\alpha}\right)<\alpha$ for small $\alpha$ (prob. of false positive)?

■ Original time series:

$$
\begin{aligned}
& X_{1} X_{2} X_{3} X_{4} X_{5} X_{6} X_{7} X_{8} X_{9} X_{10} \\
& Y_{1} Y_{2} Y_{3} Y_{4} \quad Y_{5} \quad Y_{6} \quad Y_{7} \quad Y_{8} \quad Y_{9} \quad Y_{10}
\end{aligned}
$$

- Permutation:

$$
\begin{aligned}
& X_{1} X_{2} X_{3} X_{4} X_{5} X_{6} X_{7} X_{8} X_{9} X_{10} \\
& Y_{7} Y_{3} \quad Y_{9} \quad Y_{2} \quad Y_{4} \quad Y_{8} \quad Y_{5} \quad Y_{1} \quad Y_{6} \quad Y_{10}
\end{aligned}
$$

■ Null distribution via permutation

- Compute HSIC for $\left\{x_{i}, y_{\pi(i)}\right\}_{i=1}^{n}$ for random permutation $\pi$ of indices $\{1, \ldots, n\}$. This gives HSIC for independent variables.
- Repeat for many different permutations, get empirical CDF
- Threshold $c_{\alpha}$ is $1-\alpha$ quantile of empirical CDF


## Application: dependence detection across languages

Testing task: detect dependence between English and French text

|  |  |
| :--- | :--- |
| Honourable senators, I have a <br> question for the Leader of the <br> Government in the Senate | Honorables sénateurs, ma question <br> s'adresse au leader du <br> gouvernement au Sénat |
| No doubt there is great pressure <br> on provincial and municipal <br> governments | Les ordres de gouvernements <br> provinciaux et municipaux <br> subissent de fortes pressions |
| In fact, we have increased <br> federal investments for early <br> childhood development. | Au contraire, nous avons augmenté <br> le financement fédéral pour le <br> développement des jeunes |
| • | • |

## Application: dependence detection across languages

Testing task: detect dependence between English and French text $k$-spectrum kernel, $k=10$, sample size $n=10$ Honourable senators, I
have a question for the
Leader of the Government
in the Senate
$\begin{aligned} & \text { No doubt there is great } \\ & \text { pressure on provincial and } \\ & \text { municipal governments } \\ & \text { In fact, we have increased } \\ & \text { federal investments for } \\ & \text { early childhood } \\ & \text { development. } \\ & \text { question s'adresse au leader } \\ & \text { du gouvernement au Sénat }\end{aligned}$
$\begin{aligned} & \text { Les ordres de gouvernements }\end{aligned}$
provinciaux et municipaux
subissent de fortes pressions

[^0]
## Application:Dependence detection across languages

Results (for $\alpha=0.05$ )
■ k-spectrum kernel: average Type II error 0
■ Bag of words kernel: average Type II error 0.18

Settings: Five line extracts, averaged over 300 repetitions, for "Agriculture" transcripts. Similar results for Fisheries and Immigration transcripts.

# Testing higher order interactions 

## Detecting higher order interaction

How to detect V-structures with pairwise weak individual dependence?


## Detecting higher order interaction

How to detect V-structures with pairwise weak individual dependence?

reaction

## Detecting higher order interaction

How to detect V-structures with pairwise weak individual dependence?
$X \Perp Y, Y \Perp Z, X \Perp Z$


$\mathrm{X1}^{*} \mathrm{Y} 1$ vs Z 1



■ $X, Y \stackrel{\text { i.i.d. }}{\sim} \mathcal{N}(0,1)$

- $Z \mid X, Y \sim \operatorname{sign}(X Y) \operatorname{Exp}\left(\frac{1}{\sqrt{2}}\right)$

Fine print: Faithfulness violated here!

## V-structure discovery



Assume $X \Perp Y$ has been established.
V-structure can then be detected by:

■ Consistent CI test: $\mathbf{H}_{\mathbf{0}}: X \Perp Y \mid Z$ [Fukumizu et al. 2008, Zhang et al. 2011]
$■$ Factorisation test: $\mathbf{H}_{0}:(X, Y) \Perp Z \vee(X, Z) \Perp Y \vee(Y, Z) \Perp X$ (multiple standard two-variable tests)

How well do these work?

## Detecting higher order interaction

Generalise earlier example to $p$ dimensions
$X \Perp Y, Y \Perp Z, X \Perp Z$



X1 1 Y 1 vs Z 1



- $X, Y \stackrel{\text { i.i.d. }}{\sim} \mathcal{N}(0,1)$
- $Z \mid X, Y \sim \operatorname{sign}(X Y) \operatorname{Exp}\left(\frac{1}{\sqrt{2}}\right)$
- $X_{2: p}, Y_{2: p}, Z_{2: p} \stackrel{i . i . d .}{\sim} \mathcal{N}\left(0, \mathbf{I}_{p-1}\right)$

Fine print: Faithfulness violated here!

## V-structure discovery



CI test for $X \Perp Y \mid Z$ from zhang et al. (2011), and a factorisation test $_{64 / 71}$ $n=500$

## Lancaster interaction measure

Lancaster interaction measure of $\left(X_{1}, \ldots, X_{D}\right) \sim P$ is a signed measure $\Delta P$ that vanishes whenever $P$ can be factorised non-trivially.

$$
D=2: \quad \Delta_{L} P=P_{X Y}-P_{X} P_{Y}
$$

## Lancaster interaction measure

Lancaster interaction measure of $\left(X_{1}, \ldots, X_{D}\right) \sim P$ is a signed measure $\Delta P$ that vanishes whenever $P$ can be factorised non-trivially.

$$
\begin{array}{ll}
D=2: & \Delta_{L} P=P_{X Y}-P_{X} P_{Y} \\
D=3: & \Delta_{L} P=P_{X Y Z}-P_{X} P_{Y Z}-P_{Y} P_{X Z}-P_{Z} P_{X Y}+2 P_{X} P_{Y} P_{Z}
\end{array}
$$

## Lancaster interaction measure

Lancaster interaction measure of $\left(X_{1}, \ldots, X_{D}\right) \sim P$ is a signed measure $\Delta P$ that vanishes whenever $P$ can be factorised non-trivially.
$D=2: \quad \Delta_{L} P=P_{X Y}-P_{X} P_{Y}$
$D=3: \quad \Delta_{L} P=P_{X Y Z}-P_{X} P_{Y Z}-P_{Y} P_{X Z}-P_{Z} P_{X Y}+2 P_{X} P_{Y} P_{Z}$


## Lancaster interaction measure

Lancaster interaction measure of $\left(X_{1}, \ldots, X_{D}\right) \sim P$ is a signed measure $\Delta P$ that vanishes whenever $P$ can be factorised non-trivially.
$D=2: \quad \Delta_{L} P=P_{X Y}-P_{X} P_{Y}$
$D=3: \quad \Delta_{L} P=P_{X Y Z}-P_{X} P_{Y Z}-P_{Y} P_{X Z}-P_{Z} P_{X Y}+2 P_{X} P_{Y} P_{Z}$


Case of $P_{X} \Perp P_{Y Z}$

## Lancaster interaction measure

Lancaster interaction measure of $\left(X_{1}, \ldots, X_{D}\right) \sim P$ is a signed measure $\Delta P$ that vanishes whenever $P$ can be factorised non-trivially.
$D=2: \quad \Delta_{L} P=P_{X Y}-P_{X} P_{Y}$
$D=3: \quad \Delta_{L} P=P_{X Y Z}-P_{X} P_{Y Z}-P_{Y} P_{X Z}-P_{Z} P_{X Y}+2 P_{X} P_{Y} P_{Z}$
$(X, Y) \Perp Z \vee(X, Z) \Perp Y \vee(Y, Z) \Perp X \Rightarrow \Delta_{L} P=0$.
...so what might be missed?

## Lancaster interaction measure

Lancaster interaction measure of $\left(X_{1}, \ldots, X_{D}\right) \sim P$ is a signed measure $\Delta P$ that vanishes whenever $P$ can be factorised non-trivially.
$D=2: \quad \Delta_{L} P=P_{X Y}-P_{X} P_{Y}$
$D=3: \quad \Delta_{L} P=P_{X Y Z}-P_{X} P_{Y Z}-P_{Y} P_{X Z}-P_{Z} P_{X Y}+2 P_{X} P_{Y} P_{Z}$

$$
\Delta_{L} P=0 \nRightarrow(X, Y) \Perp Z \vee(X, Z) \Perp Y \vee(Y, Z) \Perp X
$$

Example:

$$
\begin{array}{|l|l|l|l|}
\hline P(0,0,0)=0.2 & P(0,0,1)=0.1 & P(1,0,0)=0.1 & P(1,0,1)=0.1 \\
\hline P(0,1,0)=0.1 & P(0,1,1)=0.1 & P(1,1,0)=0.1 & P(1,1,1)=0.2 \\
\hline
\end{array}
$$

## A kernel test statistic using Lancaster Measure

Construct a test by estimating $\left\|\mu_{\kappa}\left(\Delta_{L} P\right)\right\|_{\mathcal{H}_{\kappa}}^{2}$, where $\kappa=k \otimes l \otimes m$ :

$$
\begin{aligned}
& \left\|\mu_{\kappa}\left(P_{X Y Z}-P_{X Y} P_{Z}-\cdots\right)\right\|_{\mathcal{H}_{\kappa}}^{2}= \\
& \left\langle\mu_{\kappa} P_{X Y Z}, \mu_{\kappa} P_{X Y Z}\right\rangle_{\mathcal{H}_{\kappa}}-2\left\langle\mu_{\kappa} P_{X Y Z}, \mu_{\kappa} P_{X Y} P_{Z}\right\rangle_{\mathcal{H}_{\kappa}} \ldots
\end{aligned}
$$

## A kernel test statistic using Lancaster Measure

| $\nu \backslash \nu^{\prime}$ | $P_{X Y Z}$ | $P_{X Y Y} P_{Z}$ | $P_{X Z} P_{Y}$ | $P_{Y Z} P_{X}$ | $P_{X} P_{Y} P_{Z}$ |
| :---: | :---: | :---: | :---: | :---: | :---: |
| $P_{X Y Y}$ | $(\mathrm{K} \circ \mathbf{L} \circ \mathbf{M})_{++}$ | $((\mathbf{K} \circ \mathrm{L}) \mathrm{M})_{++}$ | $((\mathrm{K} \circ \mathrm{M}) \mathrm{L})_{++}$ | $((\mathrm{M} \circ \mathrm{L}) \mathrm{K})_{++}$ | $\operatorname{tr}\left(\mathrm{K}_{+} \circ \mathrm{L}_{+} \circ \mathrm{M}_{+}\right)$ |
| $P_{X Y} P_{Z}$ |  | $(\mathrm{K} \circ \mathrm{L})_{++} \mathrm{M}_{++}$ | $(\mathrm{MKL})_{++}$ | (KLM) ${ }_{++}$ | $(\mathrm{KL})_{++} \mathrm{M}_{++}$ |
| $P_{X Z} P_{Y}$ |  |  | $(\mathbf{K} \circ \mathbf{M})_{++} \mathbf{L}_{++}$ | (KML) ${ }_{++}$ | (KM) ${ }_{++} \mathbf{L}_{++}$ |
| $P_{\boldsymbol{Y Z}} P_{X}$ |  |  |  | $(\mathbf{L} \circ \mathbf{M})_{++} \mathbf{K}_{++}$ | $(\mathrm{LM})_{++} \mathrm{K}_{++}$ |
| $P_{X} P_{Y} P_{Z}$ |  |  |  |  | $\mathbf{K}_{++} \mathbf{L}_{++} \mathbf{M}_{++}$ |

Table: $V$-statistic estimators of $\left\langle\mu_{\kappa} \nu, \mu_{\kappa} \nu^{\prime}\right\rangle_{\mathcal{H}_{\kappa}}$ (without terms $P_{X} P_{Y} P_{Z}$ ). $H$ is centering matrix $I-n^{-1}$

## A kernel test statistic using Lancaster Measure

| $\nu \backslash \nu^{\prime}$ | P ${ }_{\text {XYZ }}$ | $P_{X Y Y} P_{Z}$ | $P_{X Z} P_{Y}$ | $P_{Y Z} P_{X}$ | $P_{X} P_{Y} P_{Z}$ |
| :---: | :---: | :---: | :---: | :---: | :---: |
| $P_{X Y Z}$ | $(\mathrm{K} \circ \mathbf{L} \circ \mathbf{M})_{++}$ | $((\mathrm{K} \circ \mathrm{L}) \mathrm{M})_{++}$ | $((\mathrm{K} \circ \mathrm{M}) \mathrm{L})_{++}$ | $((\mathrm{M} \circ \mathrm{L}) \mathrm{K})_{++}$ | $\operatorname{tr}\left(\mathrm{K}_{+} \circ \mathrm{L}_{+} \circ \mathrm{M}_{+}\right)$ |
| $P_{X Y Y} P_{Z}$ |  | $(\mathrm{K} \circ \mathrm{L})_{++} \mathrm{M}_{++}$ | $(\mathrm{MKL})_{++}$ | $(\mathrm{KLM})_{++}$ | $(\mathrm{KL})_{++} \mathrm{M}_{++}$ |
| $P_{X X Z} P_{Y}$ |  |  | $(\mathbf{K} \circ \mathbf{M})_{++} \mathbf{L}_{++}$ | (KML) ${ }_{++}$ | (KM) ++ $^{\mathbf{L}_{++}}$ |
| $P_{Y Z} P_{X}$ |  |  |  | $(\mathbf{L} \circ \mathbf{M})_{++} \mathbf{K}_{++}$ | $(\mathrm{LM})_{++} \mathrm{K}_{++}$ |
| $P_{X} P_{Y} P_{Z}$ |  |  |  |  | $\mathbf{K}_{++} \mathbf{L}_{++} \mathbf{M}_{++}$ |

Table: $V$-statistic estimators of $\left\langle\mu_{\kappa} \nu, \mu_{\kappa} \nu^{\prime}\right\rangle_{\mathcal{H}_{\kappa}}$ (without terms $P_{X} P_{Y} P_{Z}$ ). $H$ is centering matrix $I-n^{-1}$

Lancaster interaction statistic: Sejdinovic, G, Bergsma, NIPS13

$$
\left\|\mu_{\kappa}\left(\Delta_{L} P\right)\right\|_{\mathcal{H}_{\kappa}}^{2}=\frac{1}{n^{2}}(H \mathbf{K} H \circ H \mathrm{~L} H \circ H \mathrm{M} H)_{++}
$$

Empirical joint central moment in the feature space

## V-structure discovery



Lancaster test, CI test for $X \Perp Y \mid Z$ from zhang et al. (2011), and a factorisation test, $n=500$

## Interaction for $D \geq 4$

- Interaction measure valid for all $D$ :
(Streitberg, 1990)

$$
\Delta_{S} P=\sum_{\pi}(-1)^{|\pi|-1}(|\pi|-1)!J_{\pi} P
$$

- For a partition $\pi, J_{\pi}$ associates to the joint the corresponding factorisation,

$$
\text { e.g., } J_{13|2| 4} P=P_{X_{1} X_{3}} P_{X_{2}} P_{X_{4}} \text {. }
$$

## Interaction for $D \geq 4$

- Interaction measure valid for all $D$ :
(Streitberg, 1990)

$$
\Delta_{S} P=\sum_{\pi}(-1)^{|\pi|-1}(|\pi|-1)!J_{\pi} P
$$

- For a partition $\pi, J_{\pi}$ associates to the joint the corresponding factorisation, e.g., $J_{13|2| 4} P=P_{X_{1} X_{3}} P_{X_{2}} P_{X_{4}}$.



## Interaction for $D \geq 4$

- Interaction measure valid for all $D$ :
(Streitberg, 1990)

$$
\Delta_{S} P=\sum_{\pi}(-1)^{|\pi|-1}(|\pi|-1)!J_{\pi} P
$$

- For a partition $\pi, J_{\pi}$ associates to the joint the corresponding factorisation,

$$
\text { e.g., } J_{13|2| 4} P=P_{X_{1} X_{3}} P_{X_{2}} P_{X_{4}} \text {. }
$$



## Co-authors

## From Gatsby:

■ Mikolaj Binkowski
■ Kacper Chwialkowski

- Wittawat Jitkrittum
- Heiko Strathmann

■ Dougal Sutherland
■ Wenkai Xu

## External collaborators:

■ Kenji Fukumizu

- Bernhard Schoelkopf
- Dino Sejdinovic

■ Bharath Sriperumbudur

- Alex Smola
- Zoltan Szabo


## Questions?




[^0]:    ( $K$ and $L$ column centered)

