

Advances in Machine Learning for Molecules

José Miguel Hernández-Lobato

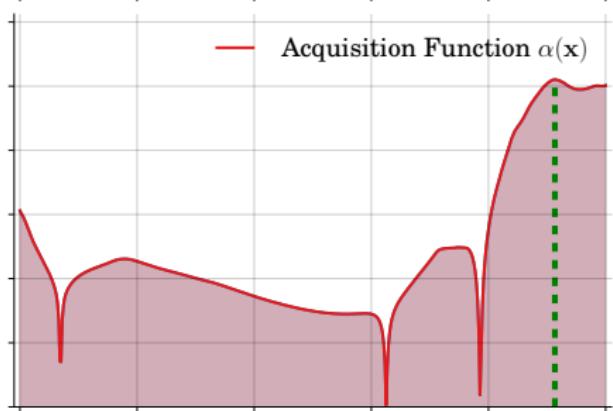
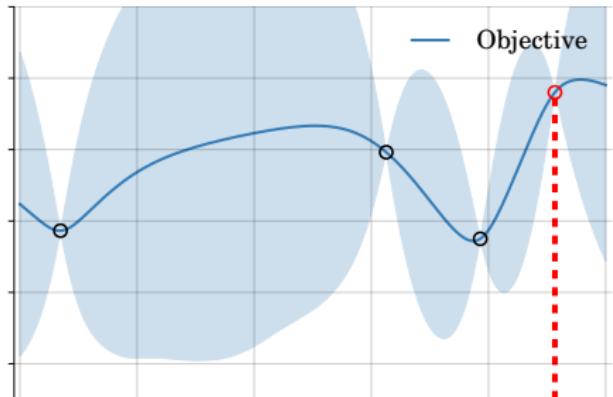
Department of Engineering
University of Cambridge,
Microsoft Research Cambridge,
Alan Turing Institute

<http://jmhl.org>, jmh233@cam.ac.uk

Machine Learning Summer School, Madrid, 2018.

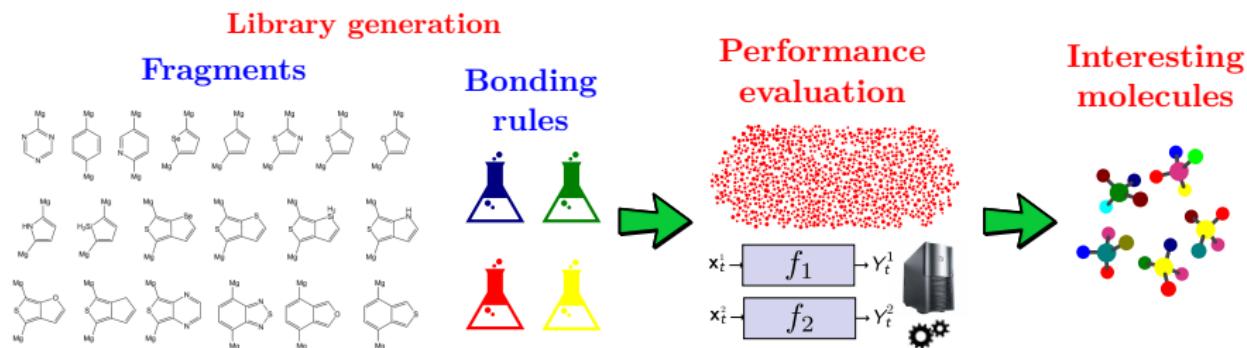
Part III: Bayesian Optimization of Molecules

Bayesian optimization

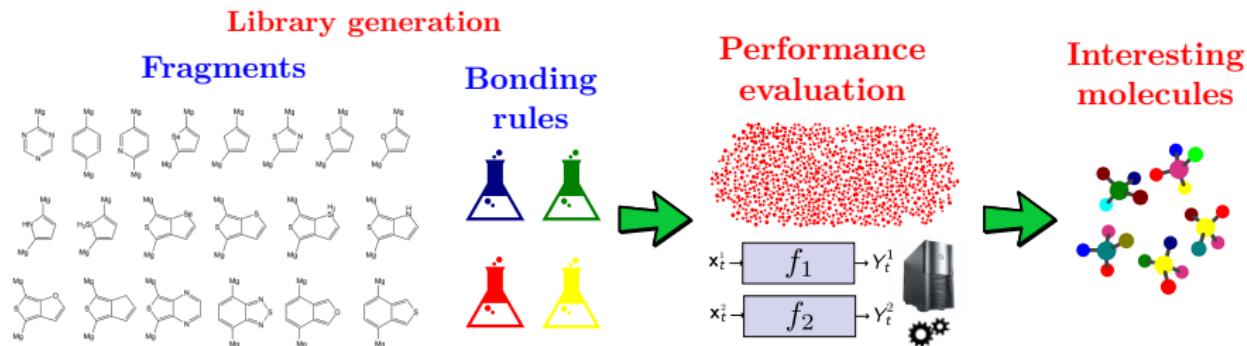


- ① Get initial sample.
- ② Fit a model to the data:
 $p(y|\mathbf{x}, \mathcal{D}_n)$.
- ③ Select data collection strategy:
 $\alpha(\mathbf{x}) = \mathbf{E}_{p(y|\mathbf{x}, \mathcal{D}_n)}[U(y|\mathbf{x}, \mathcal{D}_n)]$.
- ④ Optimize acquisition function $\alpha(\mathbf{x})$.
- ⑤ Collect data and update model.
- ⑥ Repeat!

Discovering new optimal molecules

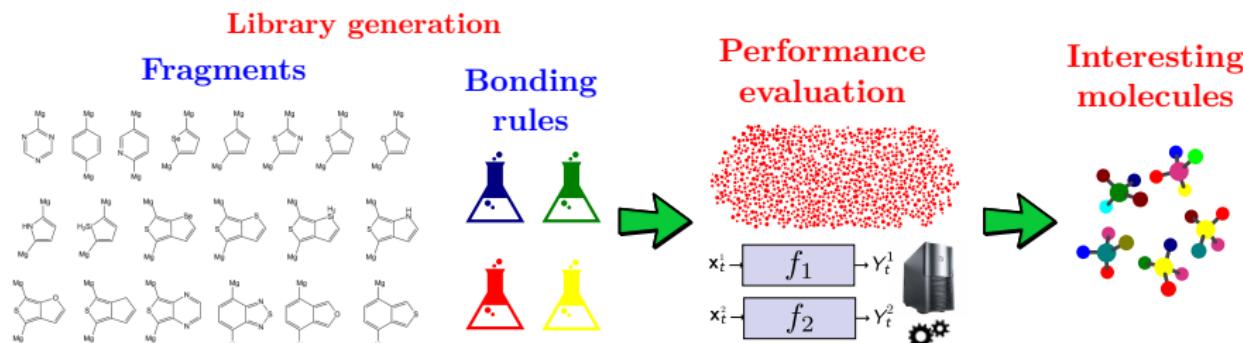


Discovering new optimal molecules



Bayesian optimization (BO) can accelerate the search!

Discovering new optimal molecules

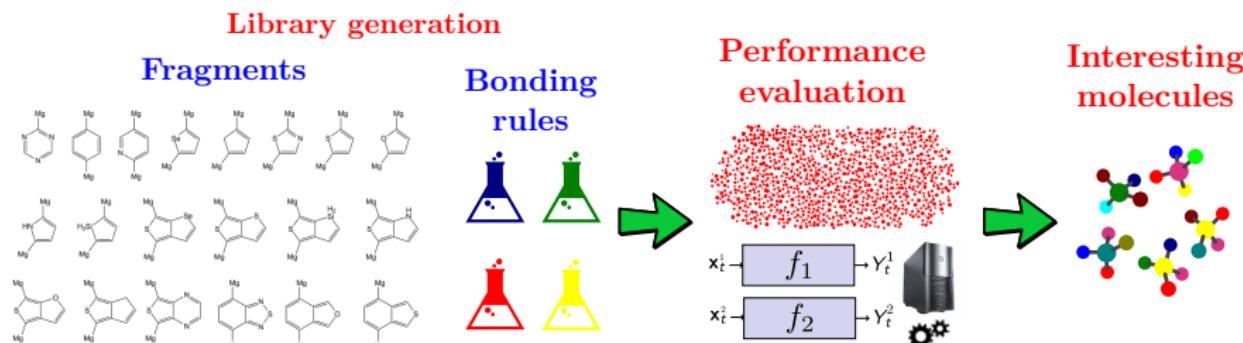


Bayesian optimization (BO) can accelerate the search!

Challenges:

- The search space is discrete and structured.
- BO output should exhibit regularities found in real-world molecules.

Discovering new optimal molecules



Bayesian optimization (BO) can accelerate the search!

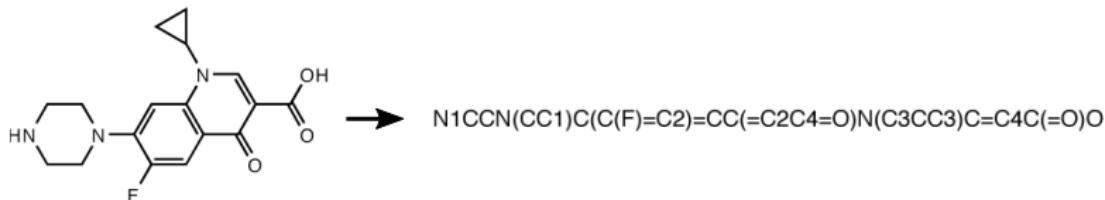
Challenges:

- The search space is discrete and structured.
- BO output should exhibit regularities found in real-world molecules.

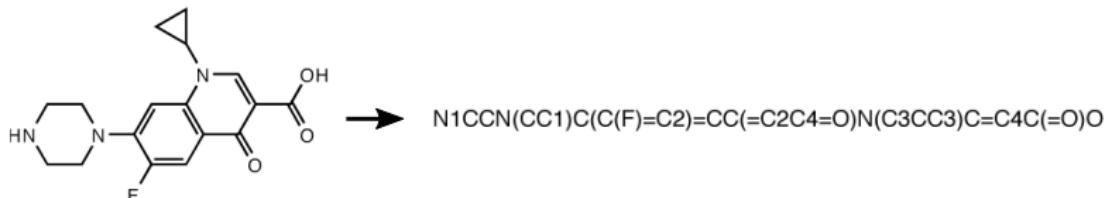
Solution: combine BO methods with generative models of molecules.

Gómez-Bombarelli*, Wei*, Duvenaud*, Hernández-Lobato*, Sánchez-Lengeling, Sheberla, Aguilera-Iparraguirre, Hirzel, Adams and Aspuru-Guzik, 2018. (* equal contributors).

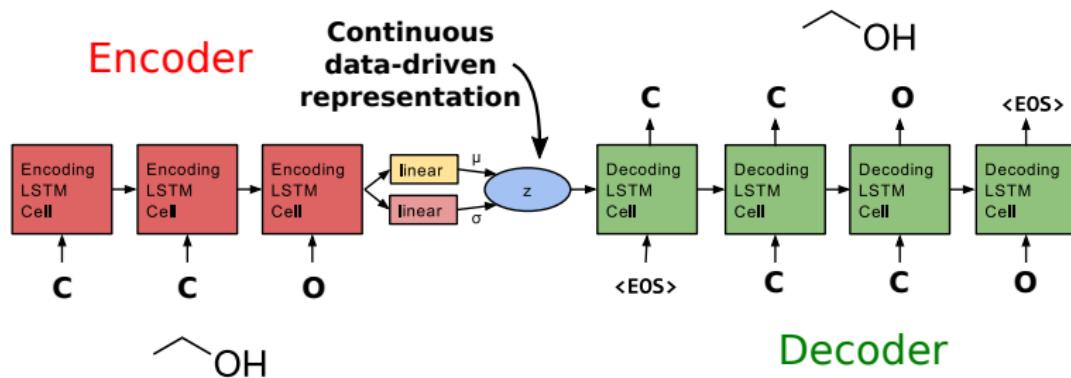
Molecules encoded as **character strings** using SMILES language.



Molecules encoded as **character strings** using SMILES language.



A **generative model** on SMILES is obtained using a **seq2seq VAE**.



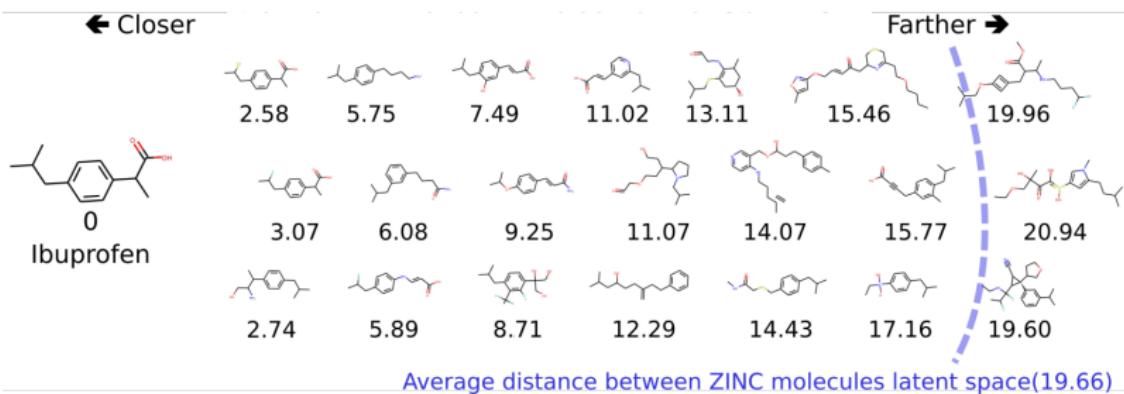
Sampled molecules have **statistics similar** to those of real molecules:

source	data set	logP	SAS	QED
Data	ZINC	2.46 (1.43)	3.05 (0.83)	0.73 (0.14)
GA	ZINC	2.84 (1.86)	3.80 (1.01)	0.57 (0.20)
VAE	ZINC	2.67 (1.46)	3.18 (0.86)	0.70 (0.14)
Data	QM9	0.30 (1.00)	4.25 (0.94)	0.48 (0.07)
GA	QM9	0.96 (1.53)	4.47 (1.01)	0.53 (0.13)
VAE	QM9	0.30 (0.97)	4.34 (0.98)	0.47 (0.08)

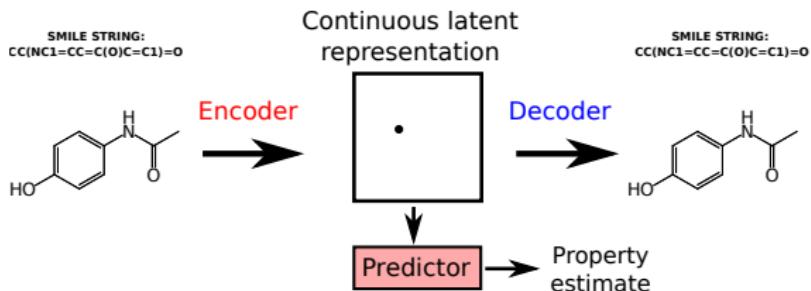
Sampled molecules have **statistics similar** to those of real molecules:

source	data set	logP	SAS	QED
Data	ZINC	2.46 (1.43)	3.05 (0.83)	0.73 (0.14)
GA	ZINC	2.84 (1.86)	3.80 (1.01)	0.57 (0.20)
VAE	ZINC	2.67 (1.46)	3.18 (0.86)	0.70 (0.14)
Data	QM9	0.30 (1.00)	4.25 (0.94)	0.48 (0.07)
GA	QM9	0.96 (1.53)	4.47 (1.01)	0.53 (0.13)
VAE	QM9	0.30 (0.97)	4.34 (0.98)	0.47 (0.08)

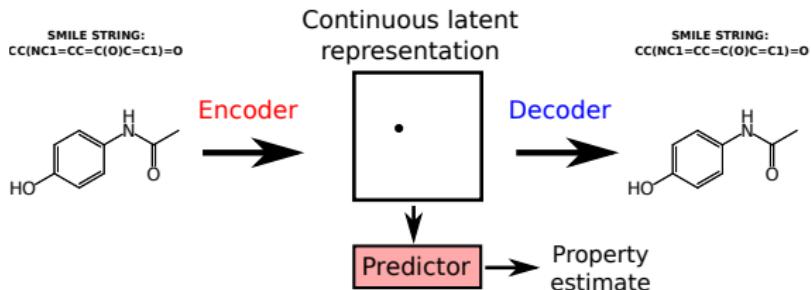
Nearby **latent representations** decode into similar molecules:



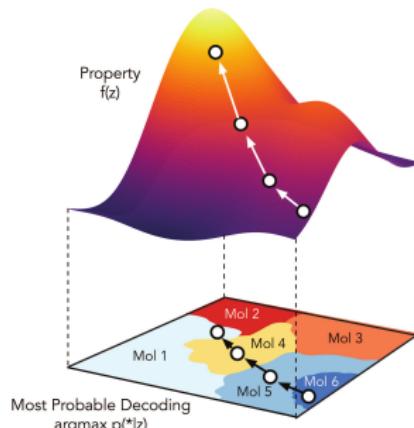
Easy to add a **surrogate model** from latent space to property.



Easy to add a **surrogate model** from latent space to property.



Gradient-based optimization can be used in latent space.



Local optimization of OLED molecules



We optimize the delayed fluorescence decay rate κ_{TADF} , as estimated from TDDFT computations on 150,000 molecules.



Local optimization of OLED molecules

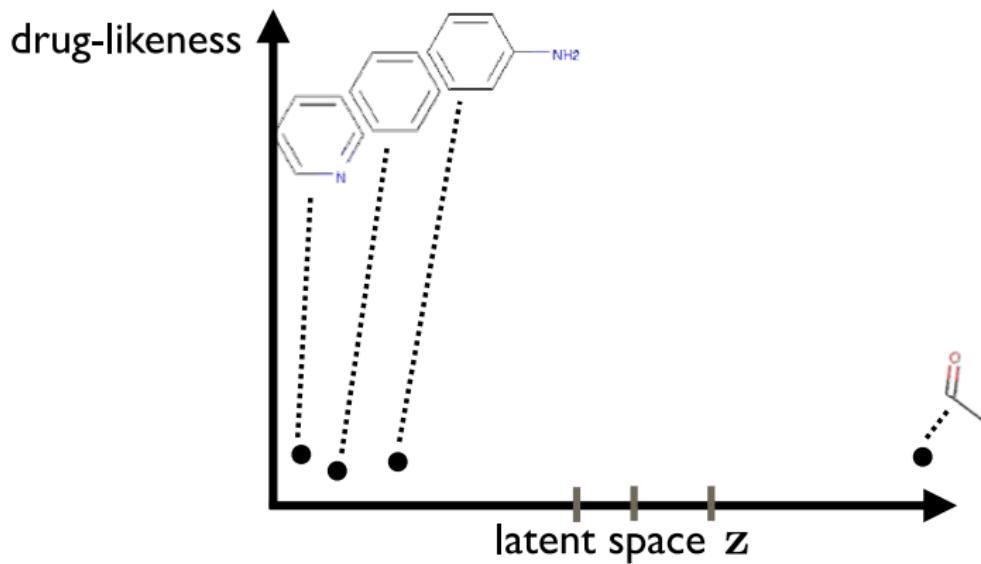


We optimize the delayed fluorescence decay rate κ_{TADF} , as estimated from TDDFT computations on 150,000 molecules.

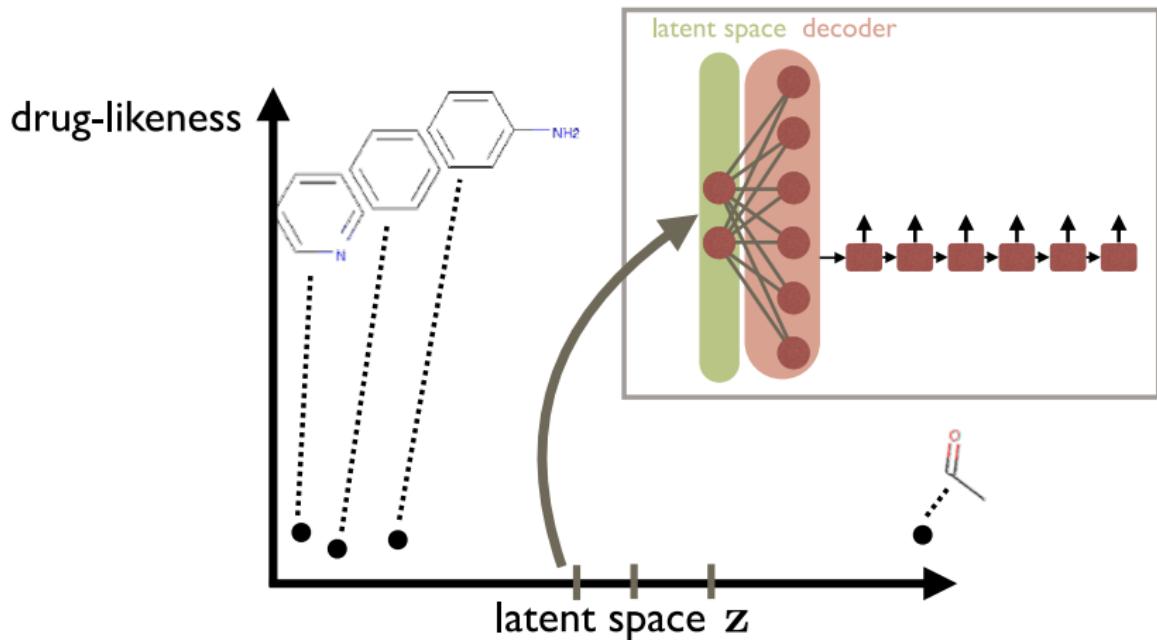


However, many of the sampled SMILES strings are not valid molecules.

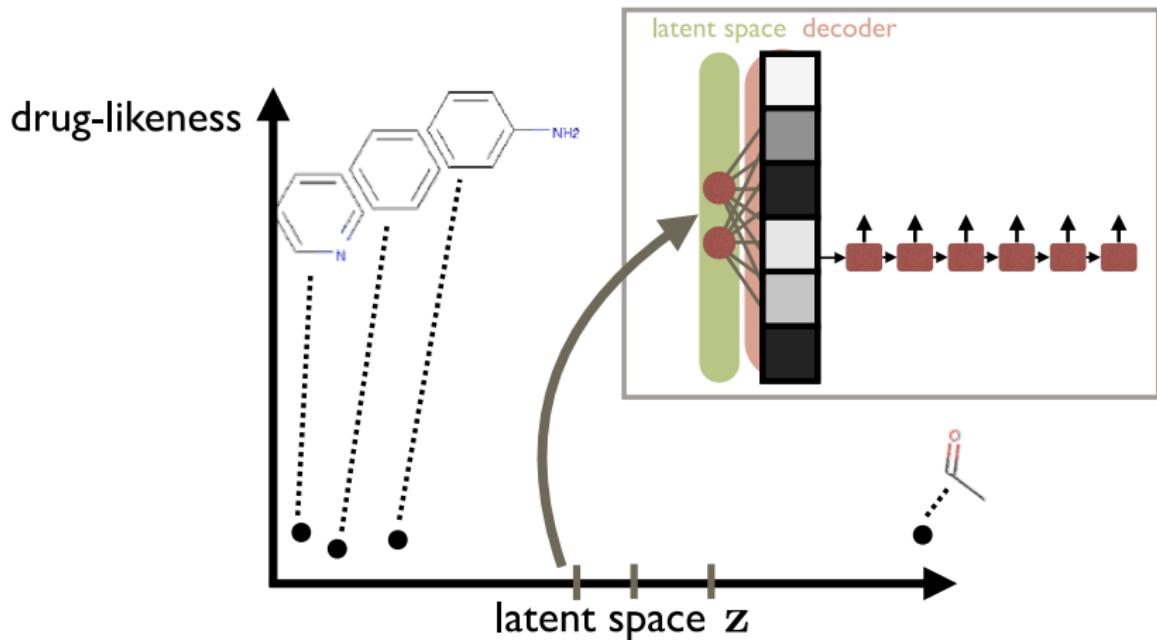
THE PROBLEM



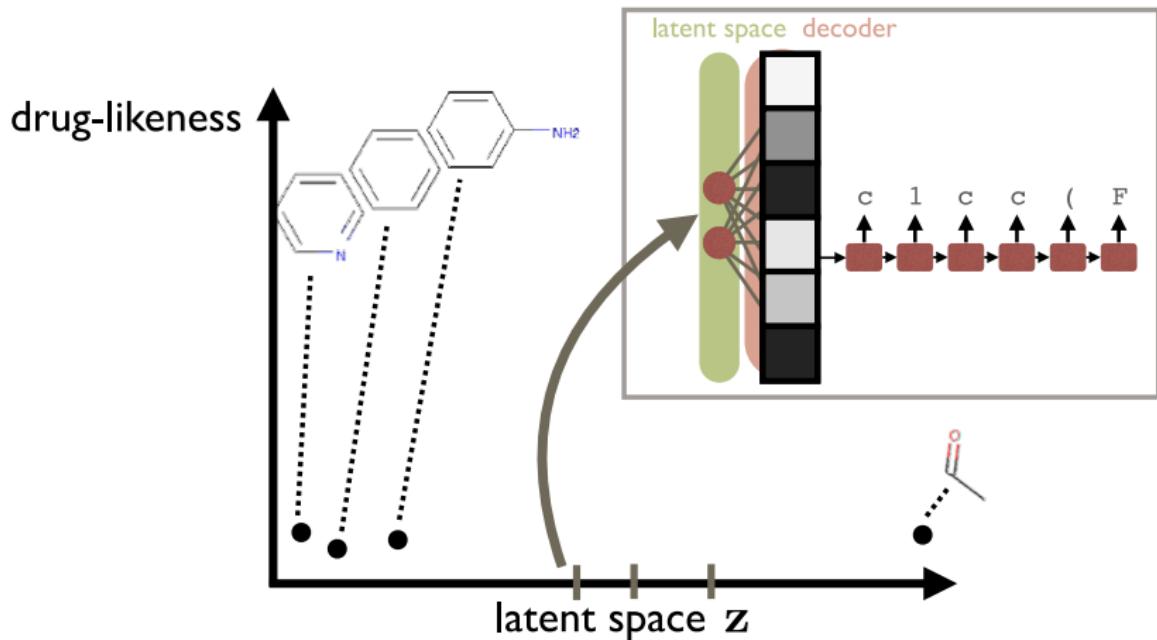
THE PROBLEM



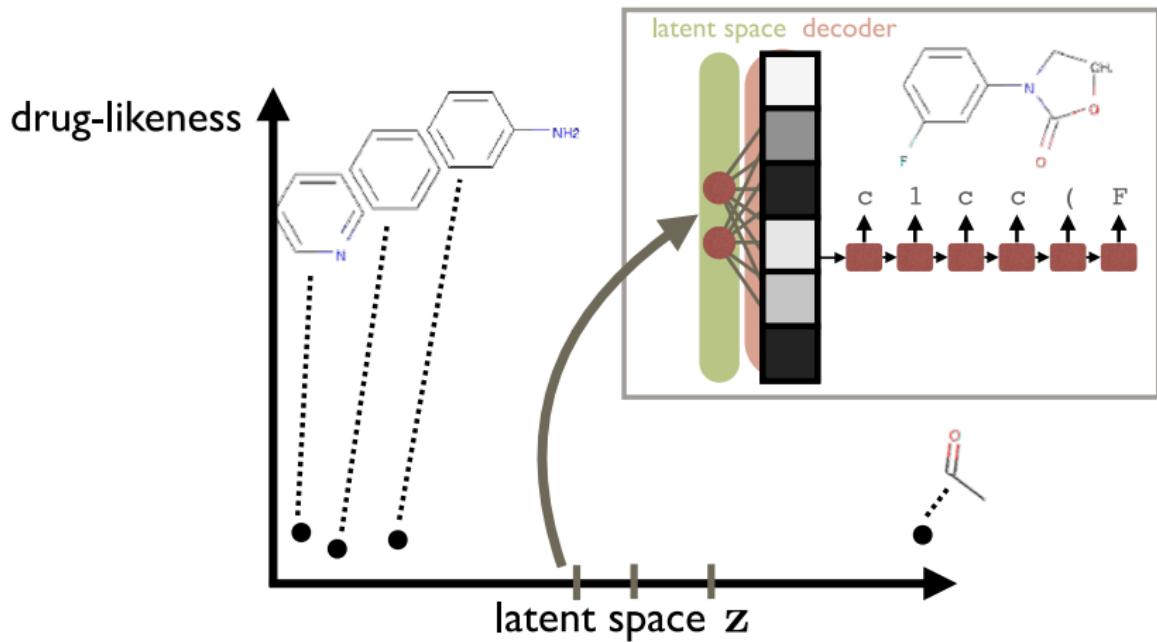
THE PROBLEM



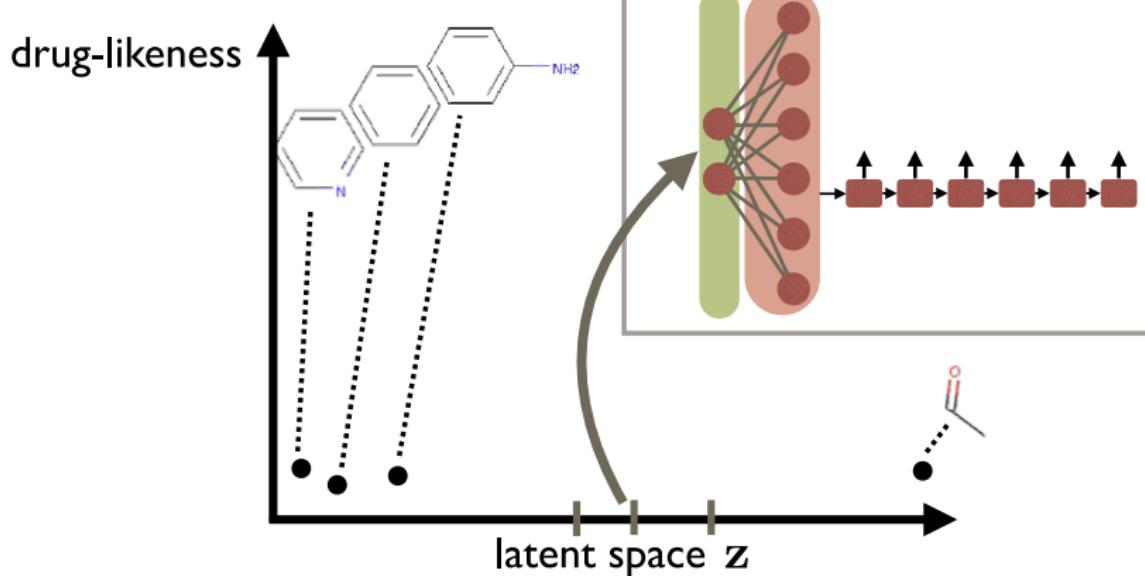
THE PROBLEM



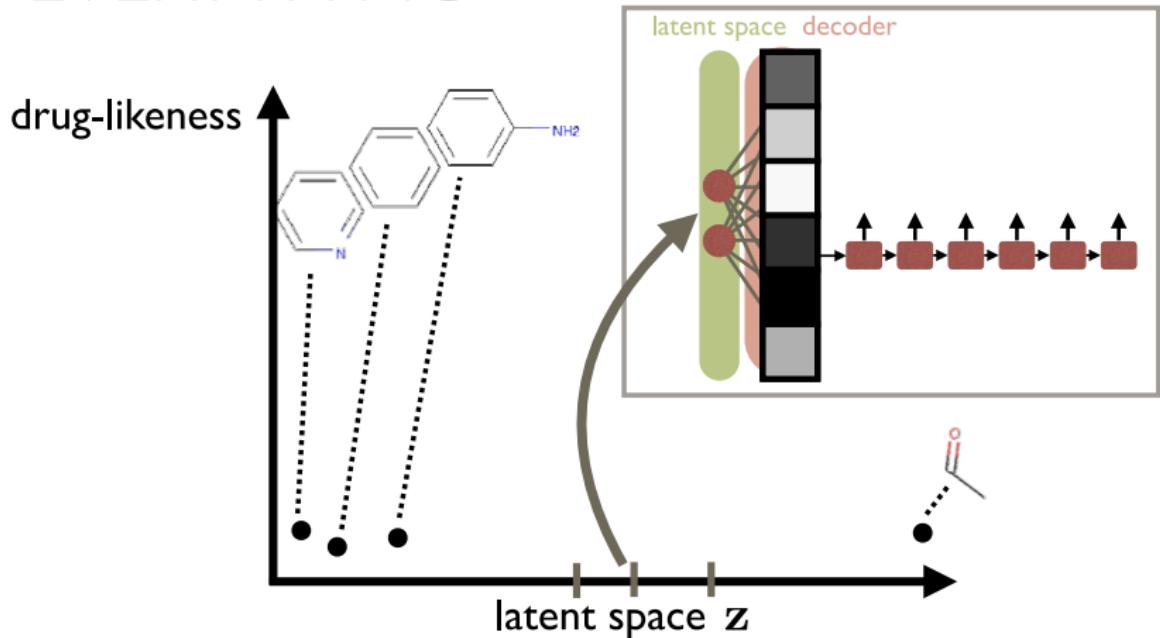
THE PROBLEM



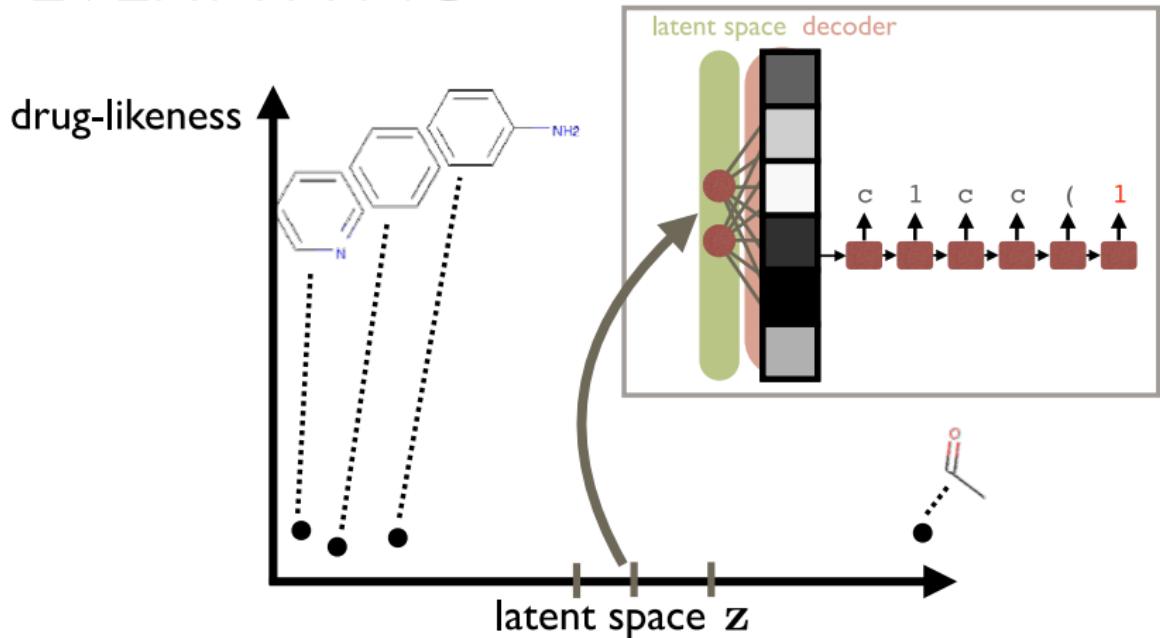
SMALL CHANGES RUIN EVERYTHING



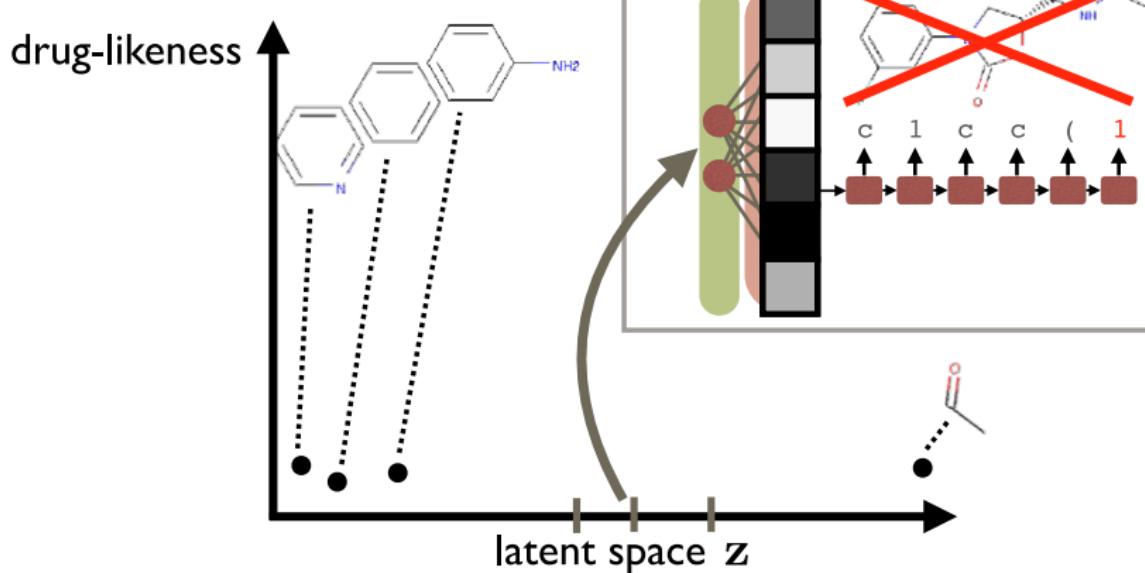
SMALL CHANGES RUIN EVERYTHING



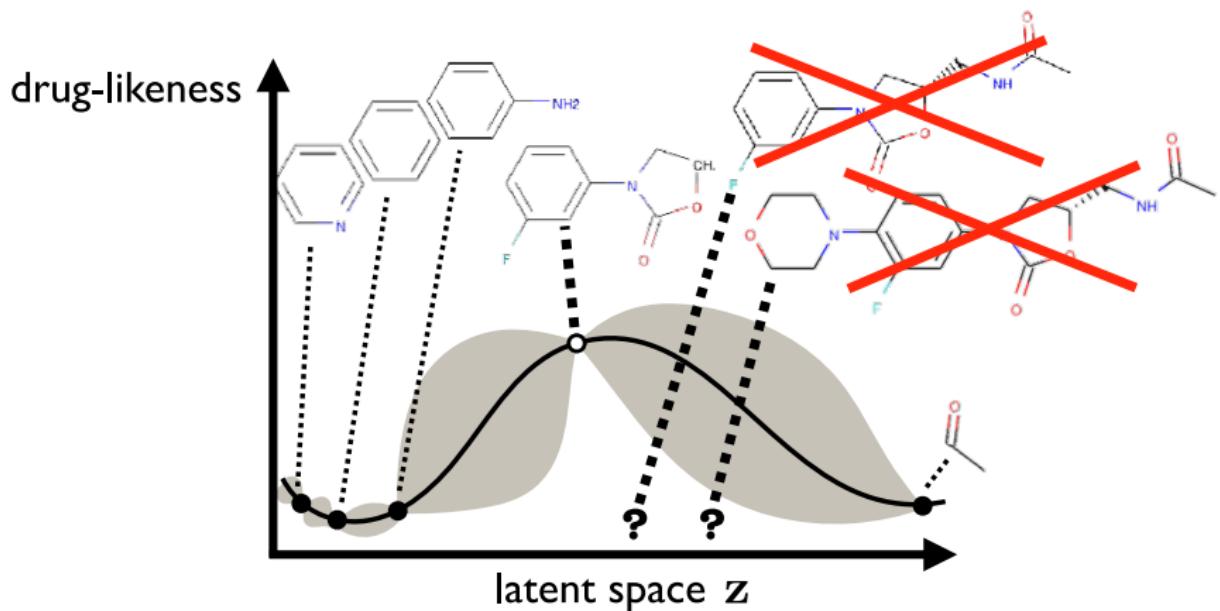
SMALL CHANGES RUIN EVERYTHING



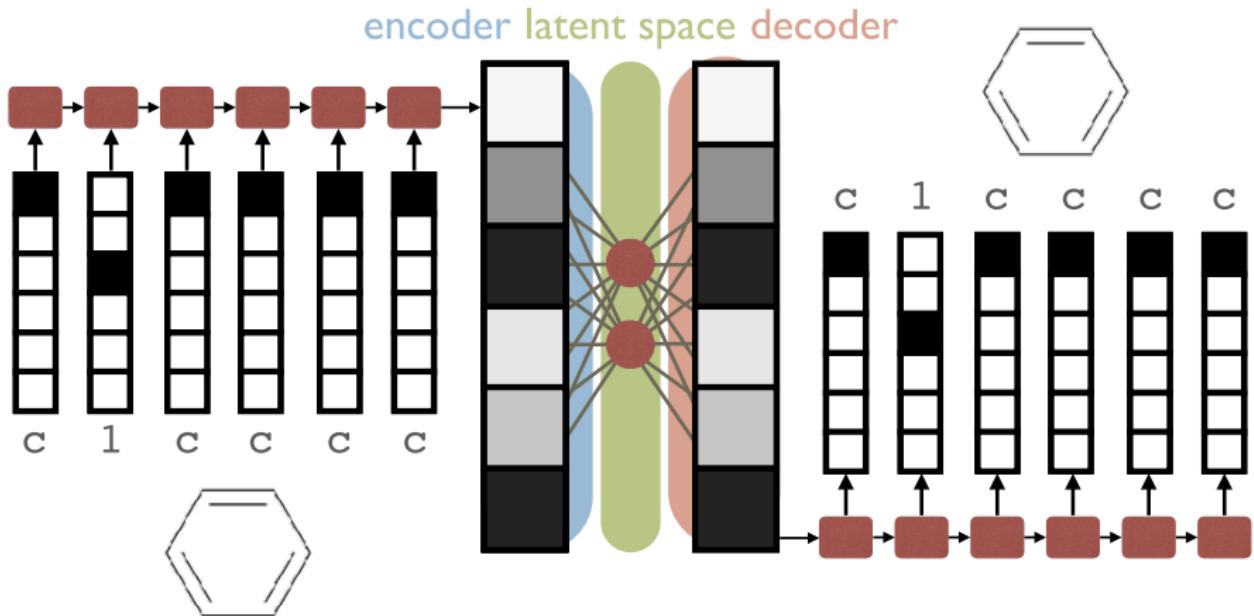
SMALL CHANGES RUIN EVERYTHING



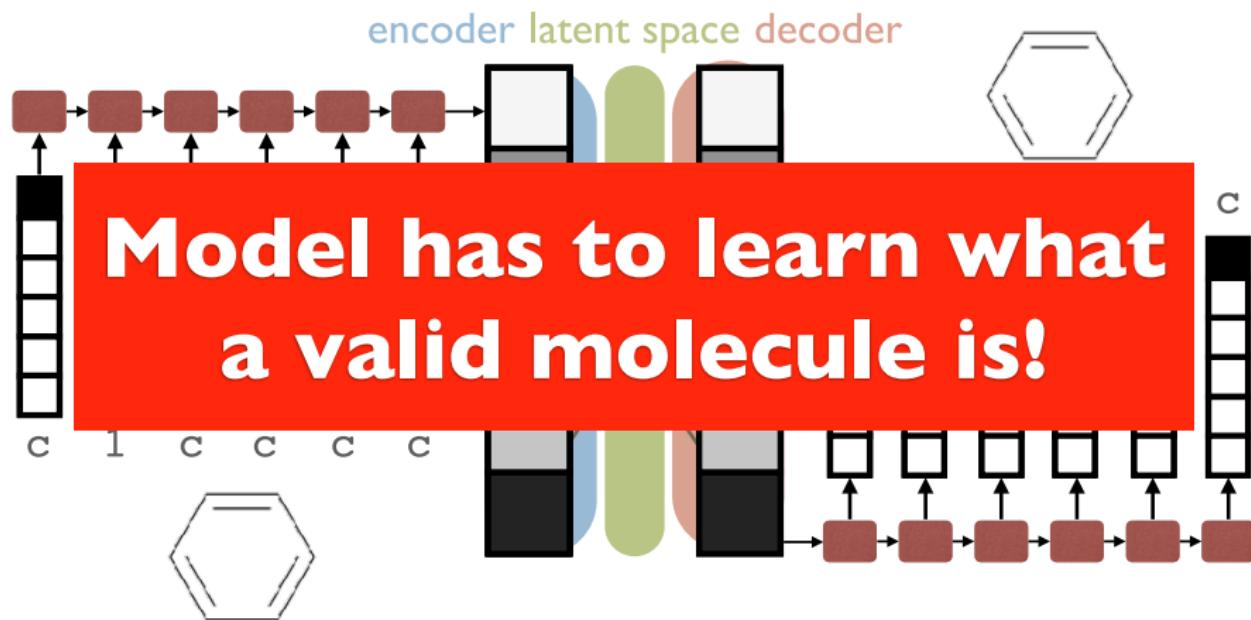
SMALL CHANGES RUIN EVERYTHING



DISCRETE GENERATIVE MODELS: CAN WE DO BETTER?



DISCRETE GENERATIVE MODELS: CAN WE DO BETTER?



A NEW REPRESENTATION

[Kusner, Paige, Hernández-Lobato, 2017]

Context-Free Grammar

[James et al., 2015]

smiles \longrightarrow chain

chain \longrightarrow chain, branched atom

chain \longrightarrow branched atom

branched atom \longrightarrow atom, ringbond

branched atom \longrightarrow atom

atom \longrightarrow aromatic organic

atom \longrightarrow aliphatic organic

ringbond \longrightarrow digit

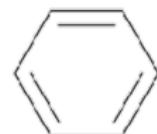
aromatic organic \longrightarrow 'c'

aliphatic organic \longrightarrow 'C'

aliphatic organic \longrightarrow 'N'

digit \longrightarrow '1'

digit \longrightarrow '2'



c 1 c c c c c 1

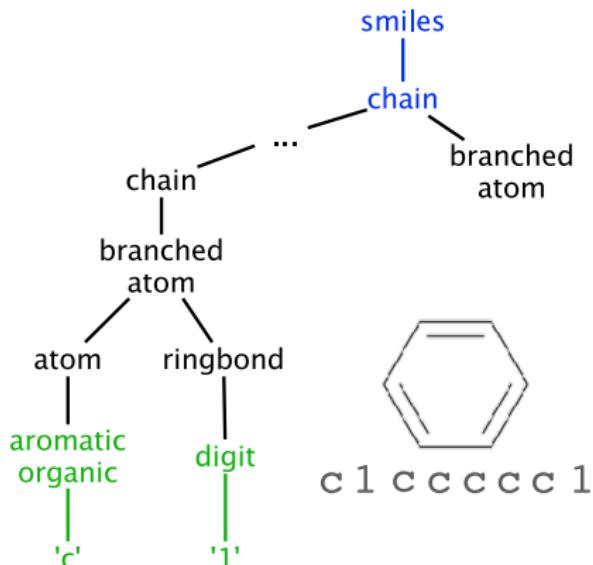
A NEW REPRESENTATION

[Kusner, Paige, Hernández-Lobato, 2017]

Context-Free Grammar

[James et al., 2015]

smiles → chain
chain → chain, branched atom
chain → branched atom
branched atom → atom, ringbond
branched atom → atom
atom → aromatic organic
atom → aliphatic organic
ringbond → digit
aromatic organic → 'c'
aliphatic organic → 'C'
aliphatic organic → 'N'
digit → '1'
digit → '2'



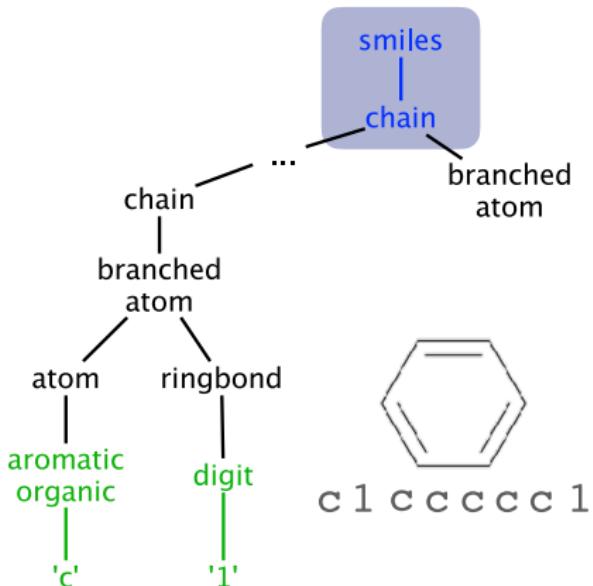
A NEW SEQUENCE!

[Kusner, Paige, Hernández-Lobato, 2017]

Context-Free Grammar

[James et al., 2015]

smiles \rightarrow chain
chain \rightarrow chain, branched atom
chain \rightarrow branched atom
branched atom \rightarrow atom, ringbond
branched atom \rightarrow atom
atom \rightarrow aromatic organic
atom \rightarrow aliphatic organic
ringbond \rightarrow digit
aromatic organic \rightarrow 'c'
aliphatic organic \rightarrow 'C'
aliphatic organic \rightarrow 'N'
digit \rightarrow '1'
digit \rightarrow '2'



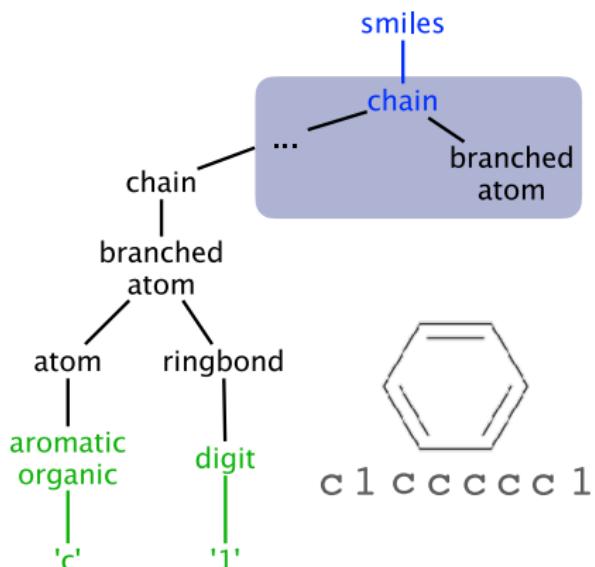
A NEW SEQUENCE!

[Kusner, Paige, Hernández-Lobato, 2017]

Context-Free Grammar

[James et al., 2015]

smiles → chain
chain → chain, branched atom
chain → branched atom
branched atom → atom, ringbond
branched atom → atom
atom → aromatic organic
atom → aliphatic organic
ringbond → digit
aromatic organic → 'c'
aliphatic organic → 'C'
aliphatic organic → 'N'
digit → '1'
digit → '2'



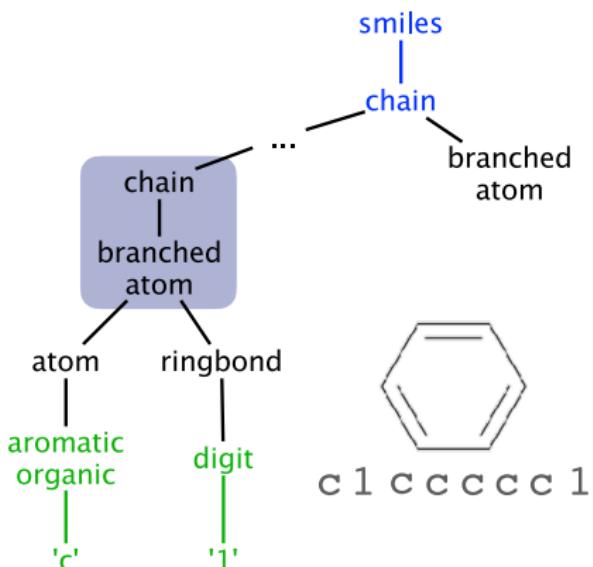
A NEW SEQUENCE!

[Kusner, Paige, Hernández-Lobato, 2017]

Context-Free Grammar

[James et al., 2015]

smiles \rightarrow chain
chain \rightarrow chain, branched atom
chain \rightarrow branched atom
branched atom \rightarrow atom, ringbond
branched atom \rightarrow atom
atom \rightarrow aromatic organic
atom \rightarrow aliphatic organic
ringbond \rightarrow digit
aromatic organic \rightarrow 'c'
aliphatic organic \rightarrow 'C'
aliphatic organic \rightarrow 'N'
digit \rightarrow '1'
digit \rightarrow '2'



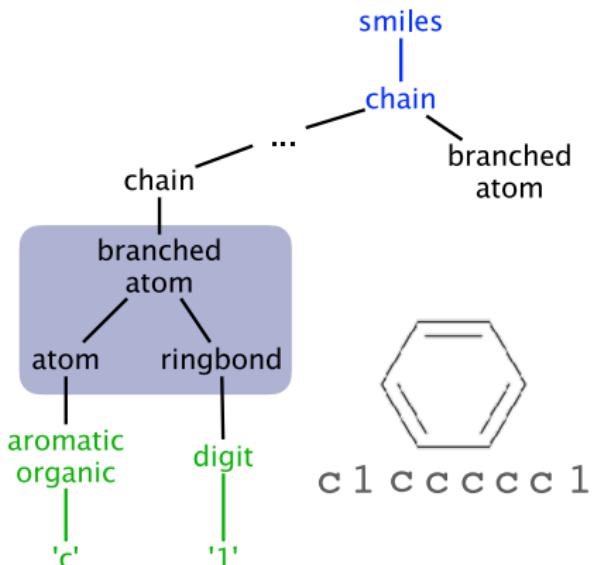
A NEW SEQUENCE!

[Kusner, Paige, Hernández-Lobato, 2017]

Context-Free Grammar

[James et al., 2015]

smiles \rightarrow chain
chain \rightarrow chain, branched atom
chain \rightarrow branched atom
branched atom \rightarrow atom, ringbond
branched atom \rightarrow atom
atom \rightarrow aromatic organic
atom \rightarrow aliphatic organic
ringbond \rightarrow digit
aromatic organic \rightarrow 'c'
aliphatic organic \rightarrow 'C'
aliphatic organic \rightarrow 'N'
digit \rightarrow '1'
digit \rightarrow '2'



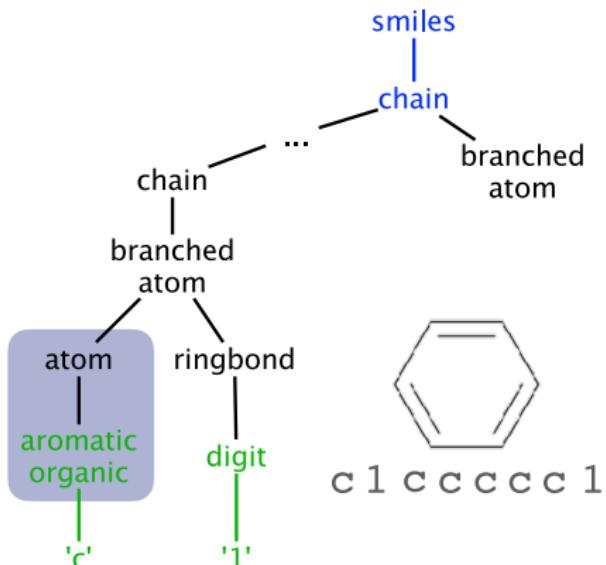
A NEW SEQUENCE!

[Kusner, Paige, Hernández-Lobato, 2017]

Context-Free Grammar

[James et al., 2015]

smiles \rightarrow chain
chain \rightarrow chain, branched atom
chain \rightarrow branched atom
branched atom \rightarrow atom, ringbond
branched atom \rightarrow atom
atom \rightarrow aromatic organic
atom \rightarrow aliphatic organic
ringbond \rightarrow digit
aromatic organic \rightarrow 'c'
aliphatic organic \rightarrow 'C'
aliphatic organic \rightarrow 'N'
digit \rightarrow '1'
digit \rightarrow '2'



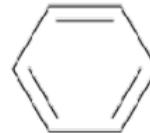
A NEW SEQUENCE!

[Kusner, Paige, Hernández-Lobato, 2017]

Context-Free Grammar

[James et al., 2015]

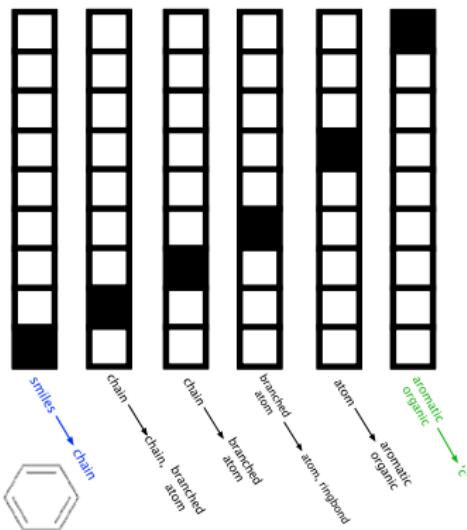
smiles \longrightarrow chain
chain \longrightarrow chain, branched atom
chain \longrightarrow branched atom
branched atom \longrightarrow atom, ringbond
branched atom \longrightarrow atom
atom \longrightarrow aromatic organic
atom \longrightarrow aliphatic organic
ringbond \longrightarrow digit
aromatic organic \longrightarrow 'c'
aliphatic organic \longrightarrow 'C'
aliphatic organic \longrightarrow 'N'
digit \longrightarrow '1'
digit \longrightarrow '2'



smiles \longrightarrow chain
chain \longrightarrow chain, branched atom
...
chain \longrightarrow branched atom
branched atom \longrightarrow atom, ringbond
atom \longrightarrow aromatic organic
aromatic organic \longrightarrow 'c'
ringbond \longrightarrow digit
digit \longrightarrow '1'
...

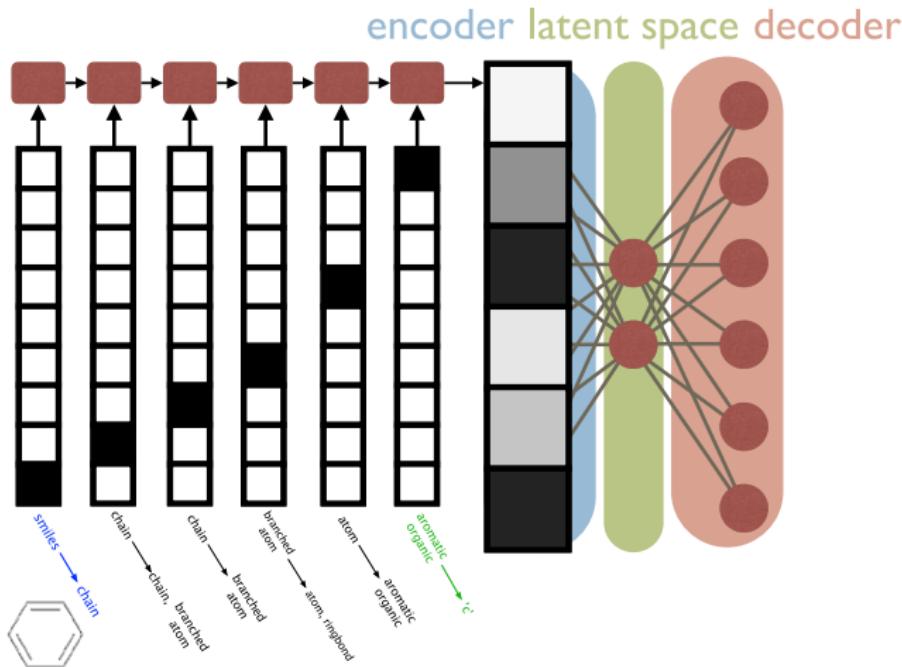
ENCODER

[Kusner, Paige, Hernández-Lobato, 2017]



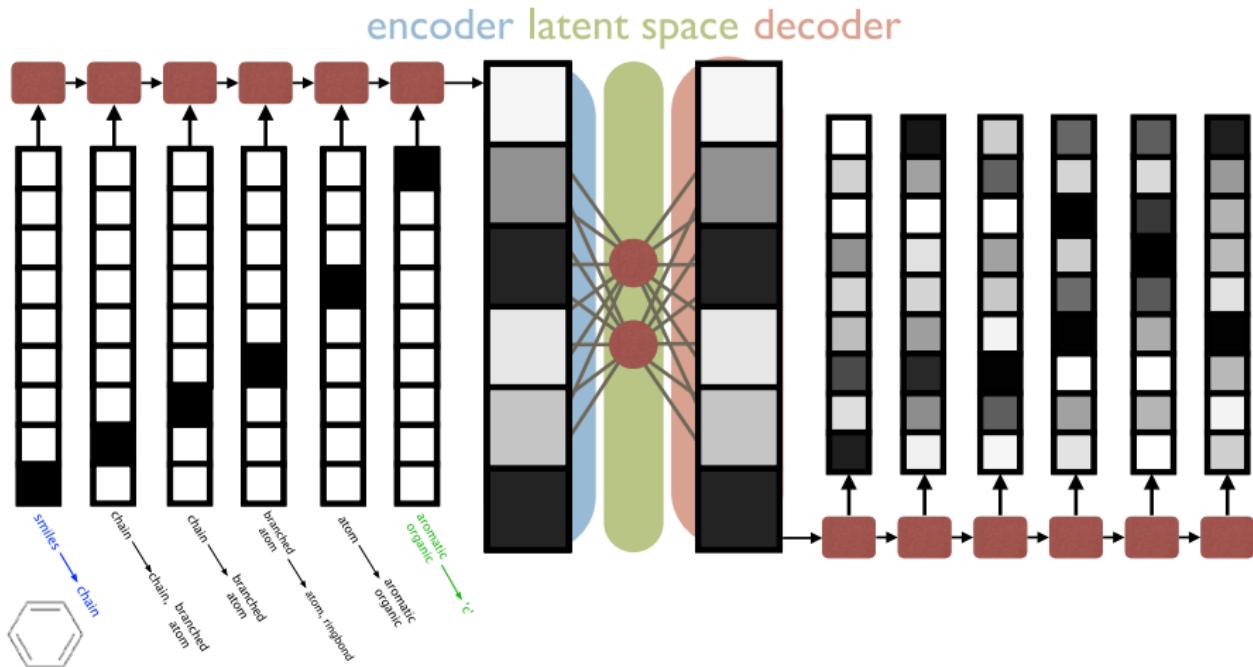
ENCODER

[Kusner, Paige, Hernández-Lobato, 2017]



DECODER

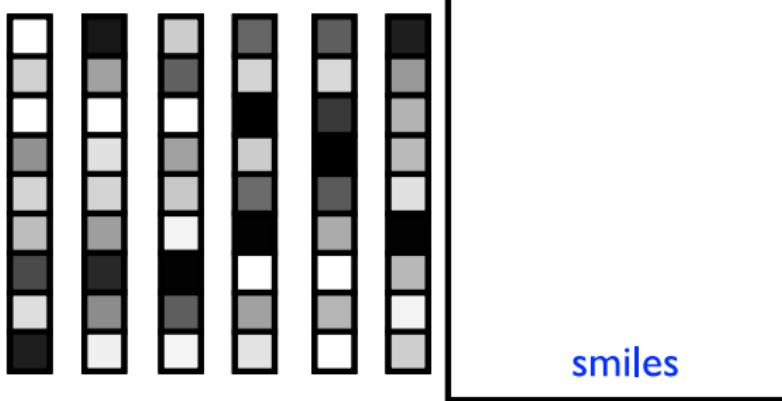
[Kusner, Paige, Hernández-Lobato, 2017]



DECODER

[Kusner, Paige, Hernández-Lobato, 2017]

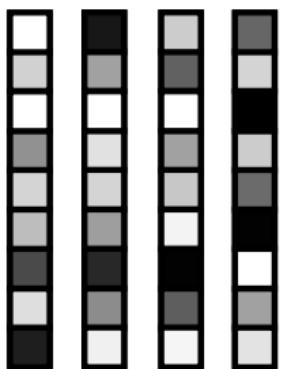
LIFO stack



DECODER

[Kusner, Paige, Hernández-Lobato, 2017]

LIFO stack

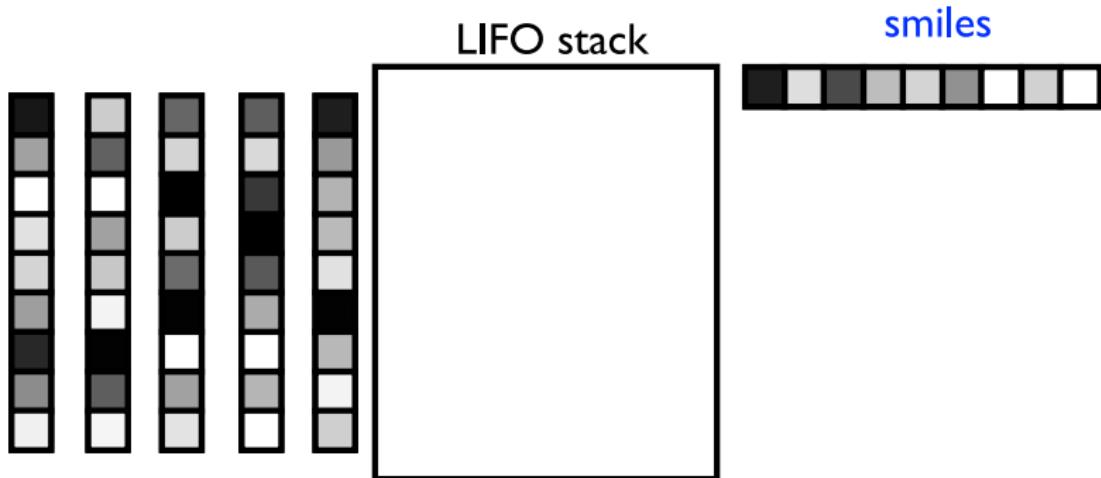


smiles



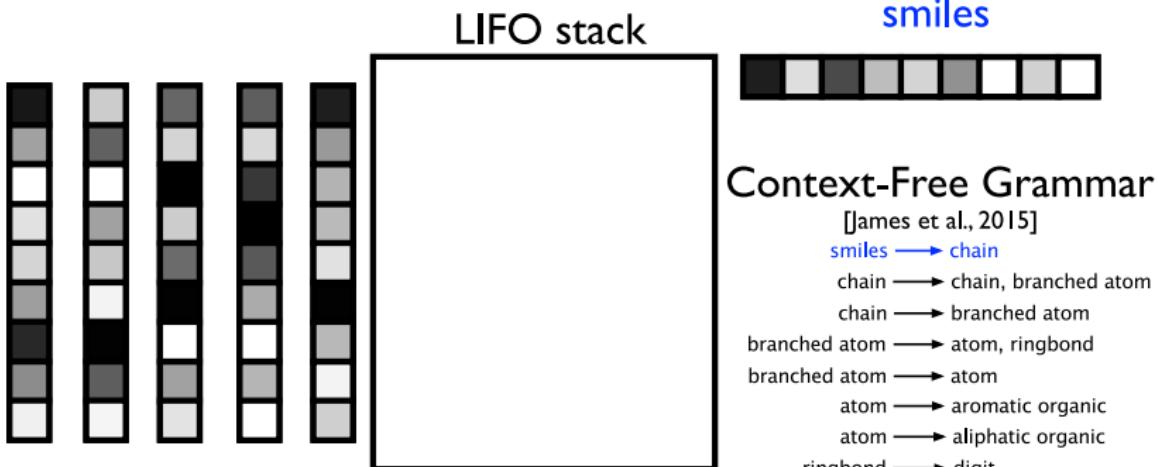
DECODER

[Kusner, Paige, Hernández-Lobato, 2017]



DECODER

[Kusner, Paige, Hernández-Lobato, 2017]



Context-Free Grammar

[James et al., 2015]

smiles \longrightarrow chain

chain \longrightarrow chain, branched atom

chain \longrightarrow branched atom

branched atom \longrightarrow atom, ringbond

branched atom \longrightarrow atom

atom \longrightarrow aromatic organic

atom \longrightarrow aliphatic organic

ringbond \longrightarrow digit

aromatic organic \longrightarrow 'c'

aliphatic organic \longrightarrow 'C'

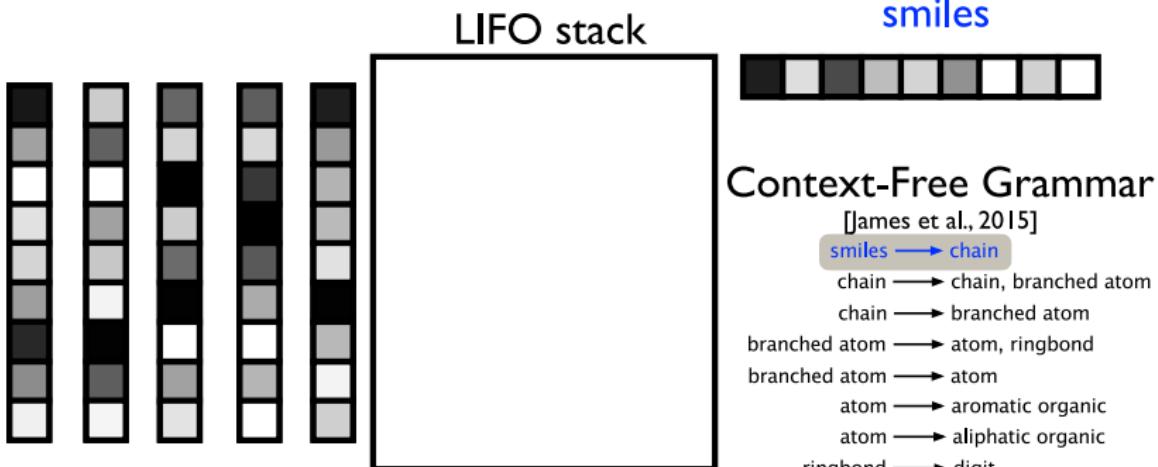
aliphatic organic \longrightarrow 'N'

digit \longrightarrow '1'

digit \longrightarrow '2'

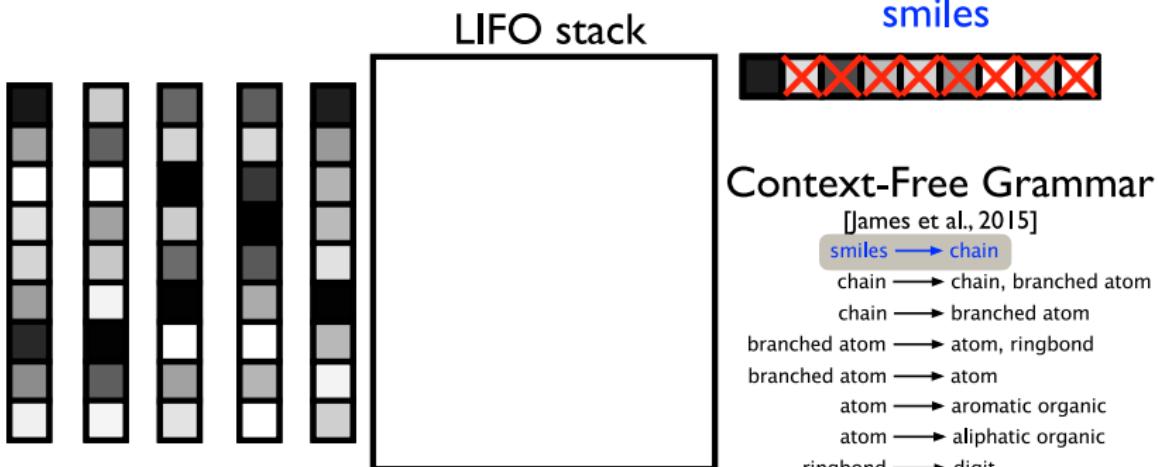
DECODER

[Kusner, Paige, Hernández-Lobato, 2017]



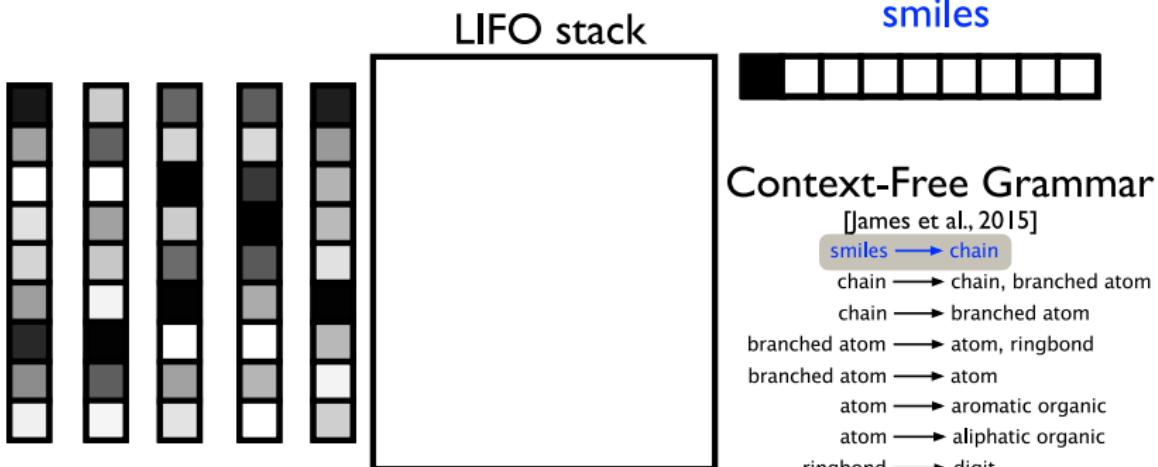
DECODER

[Kusner, Paige, Hernández-Lobato, 2017]



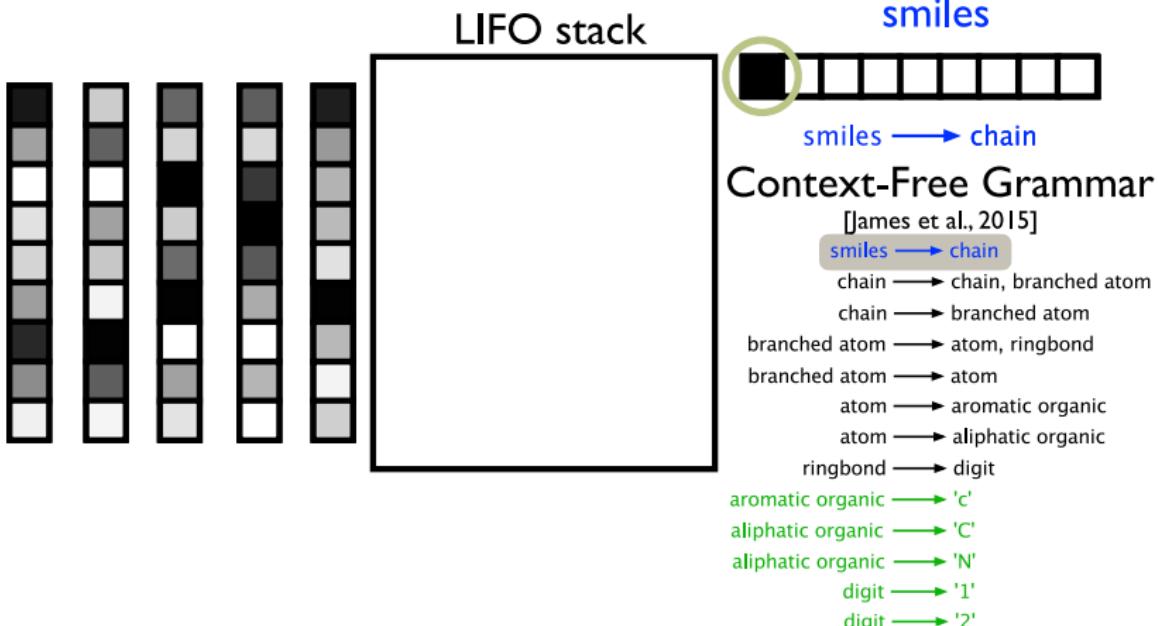
DECODER

[Kusner, Paige, Hernández-Lobato, 2017]



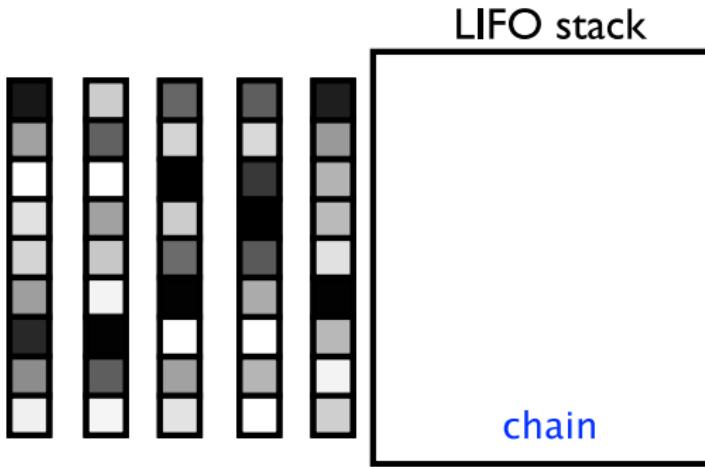
DECODER

[Kusner, Paige, Hernández-Lobato, 2017]



DECODER

[Kusner, Paige, Hernández-Lobato, 2017]



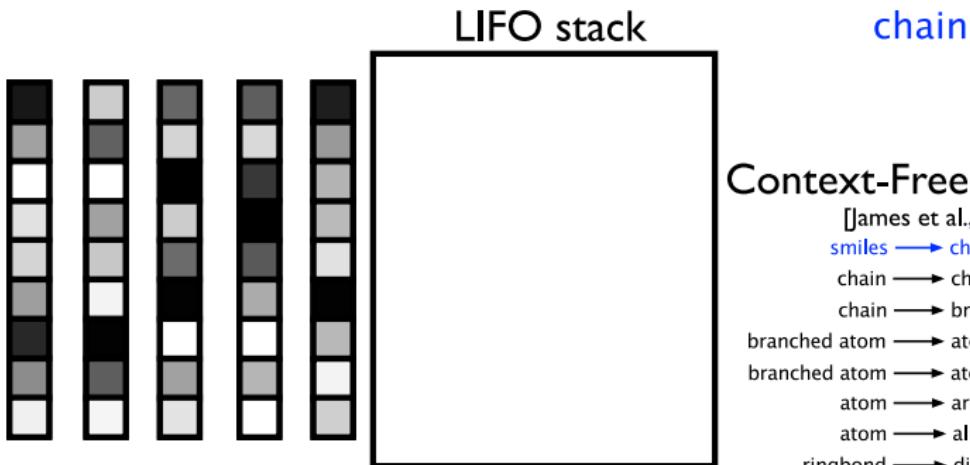
Context-Free Grammar

[James et al., 2015]

- smiles \longrightarrow chain
- chain \longrightarrow chain, branched atom
- chain \longrightarrow branched atom
- branched atom \longrightarrow atom, ringbond
- branched atom \longrightarrow atom
 - atom \longrightarrow aromatic organic
 - atom \longrightarrow aliphatic organic
- ringbond \longrightarrow digit
- aromatic organic \longrightarrow 'c'
- aliphatic organic \longrightarrow 'C'
- aliphatic organic \longrightarrow 'N'
- digit \longrightarrow '1'
- digit \longrightarrow '2'

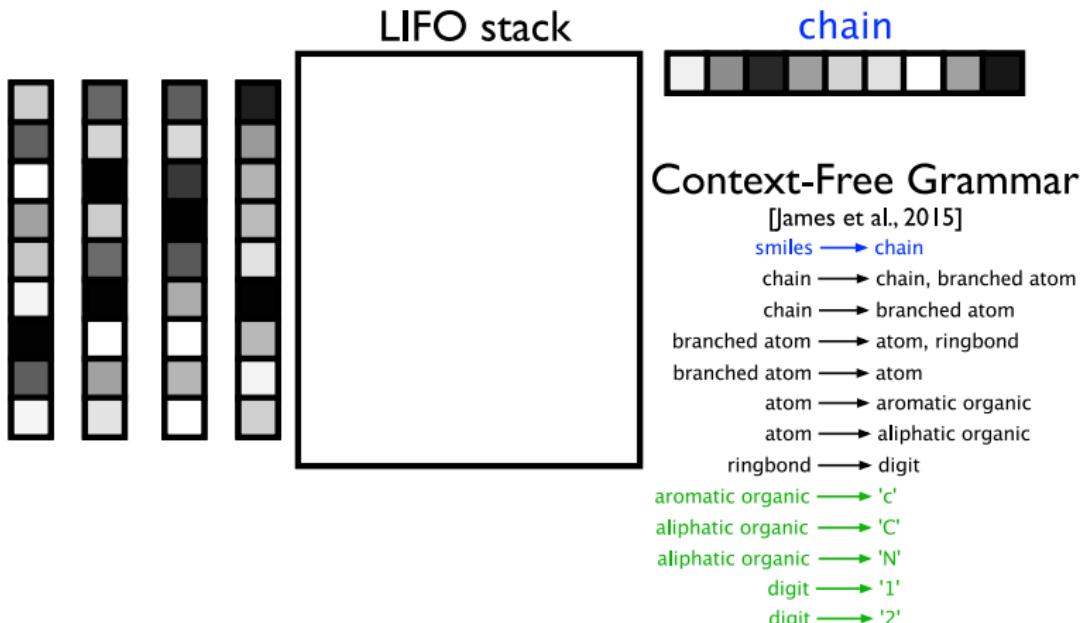
DECODER

[Kusner, Paige, Hernández-Lobato, 2017]



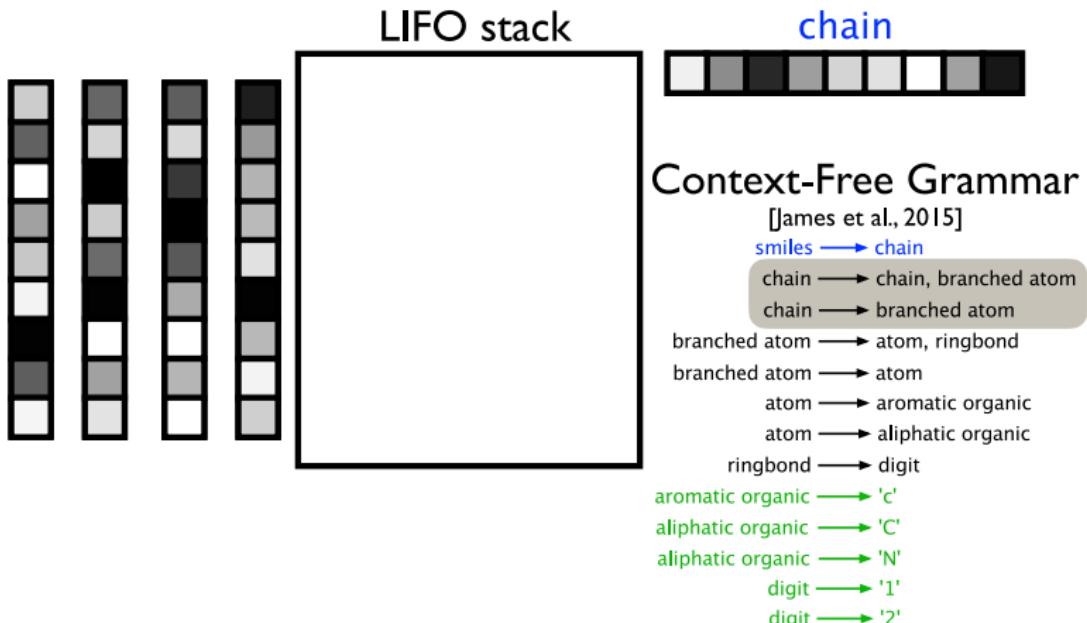
DECODER

[Kusner, Paige, Hernández-Lobato, 2017]



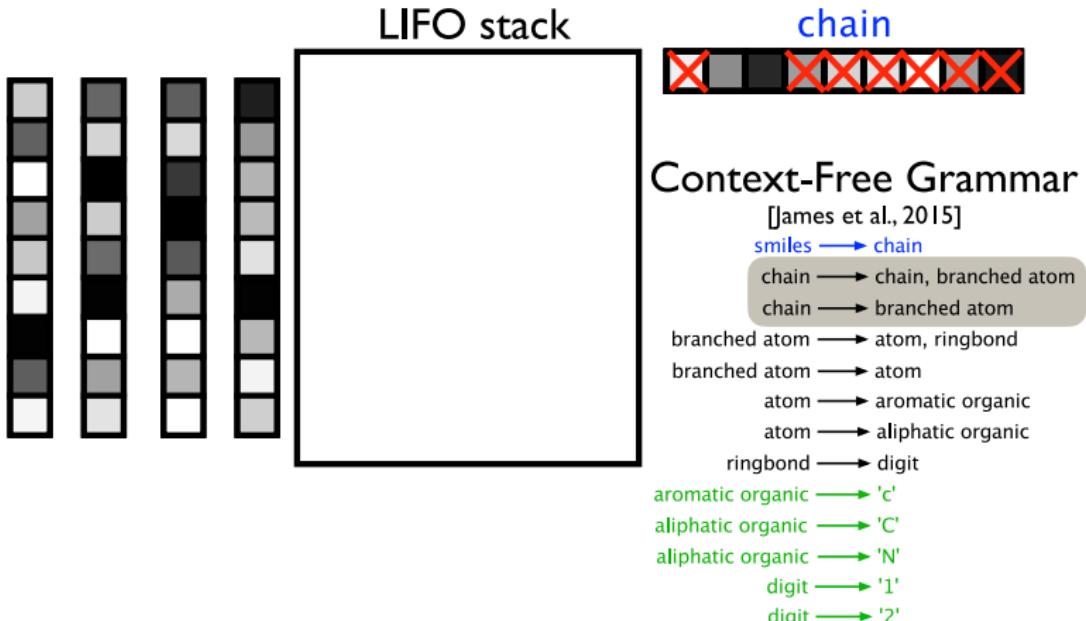
DECODER

[Kusner, Paige, Hernández-Lobato, 2017]



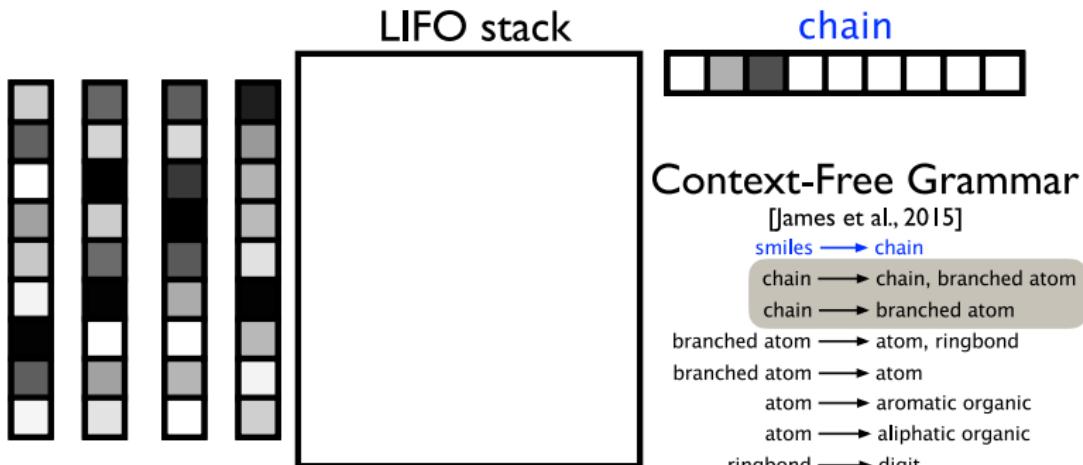
DECODER

[Kusner, Paige, Hernández-Lobato, 2017]



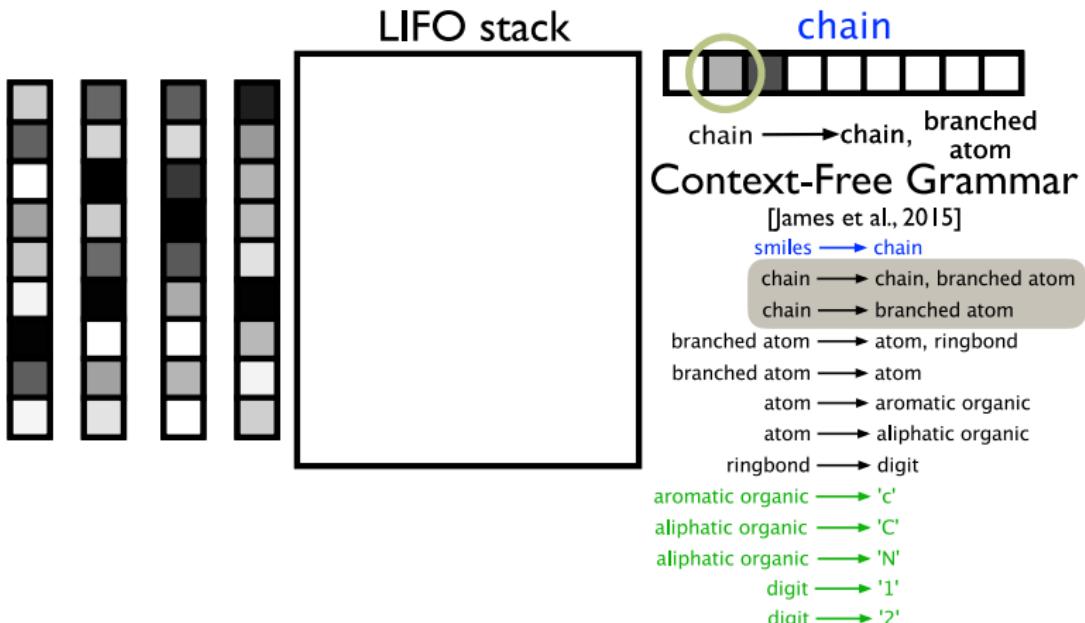
DECODER

[Kusner, Paige, Hernández-Lobato, 2017]



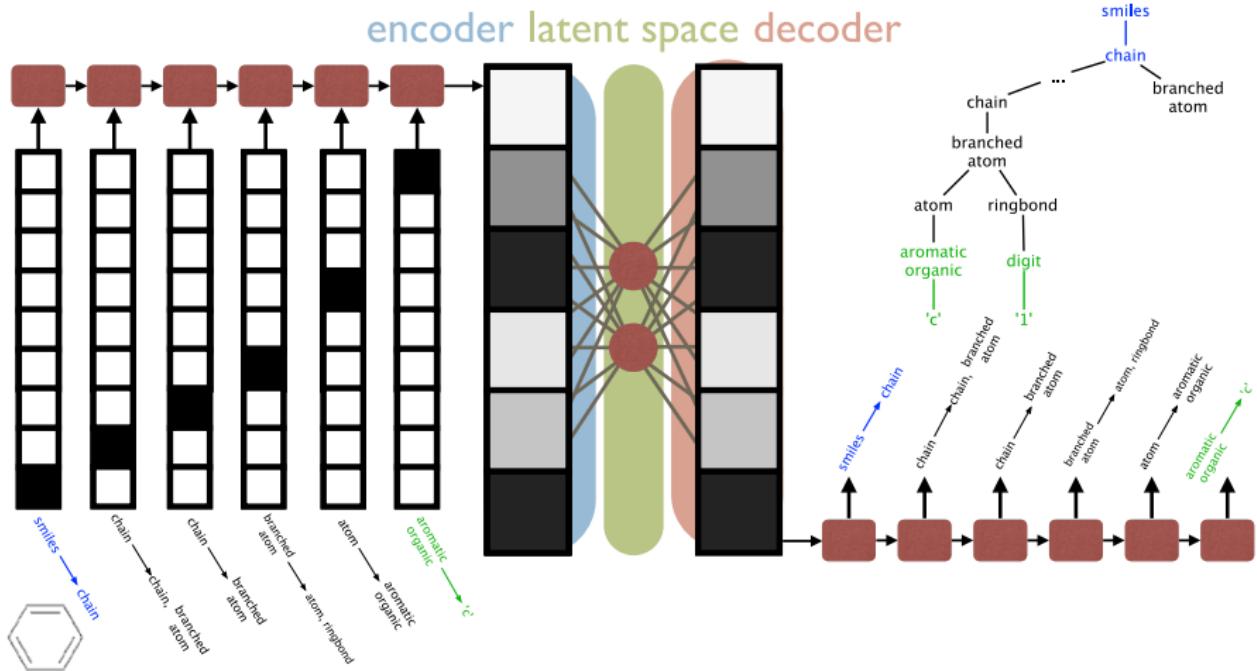
DECODER

[Kusner, Paige, Hernández-Lobato, 2017]



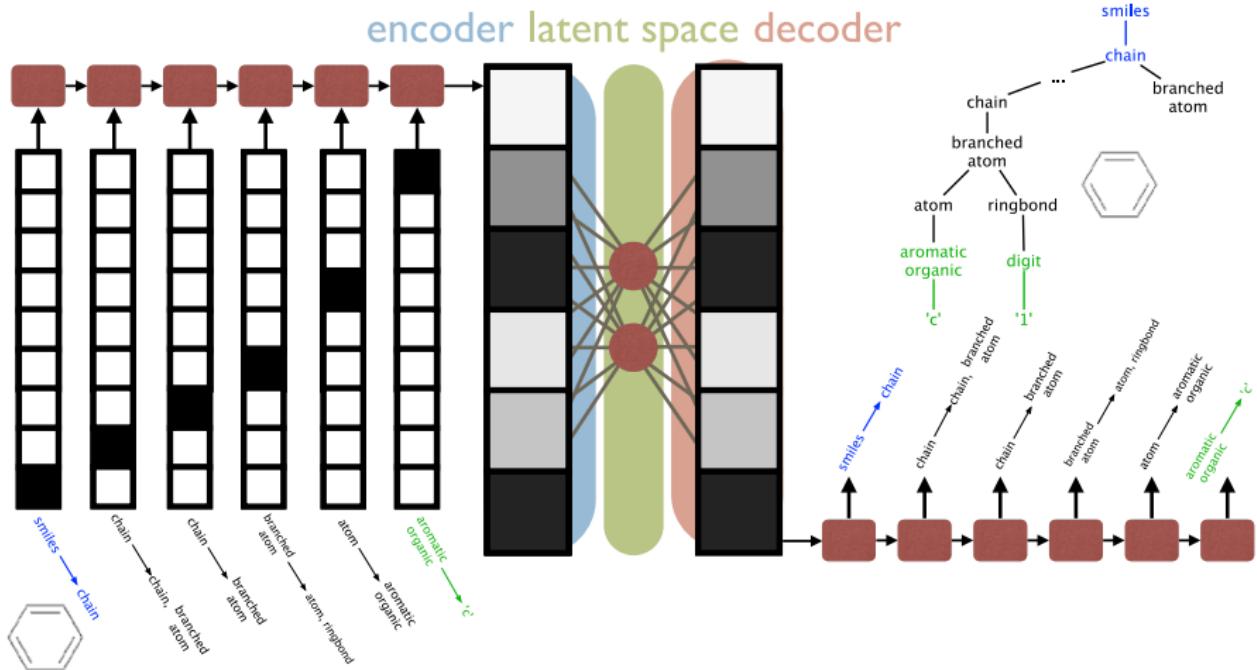
DECODER

[Kusner, Paige, Hernández-Lobato, 2017]



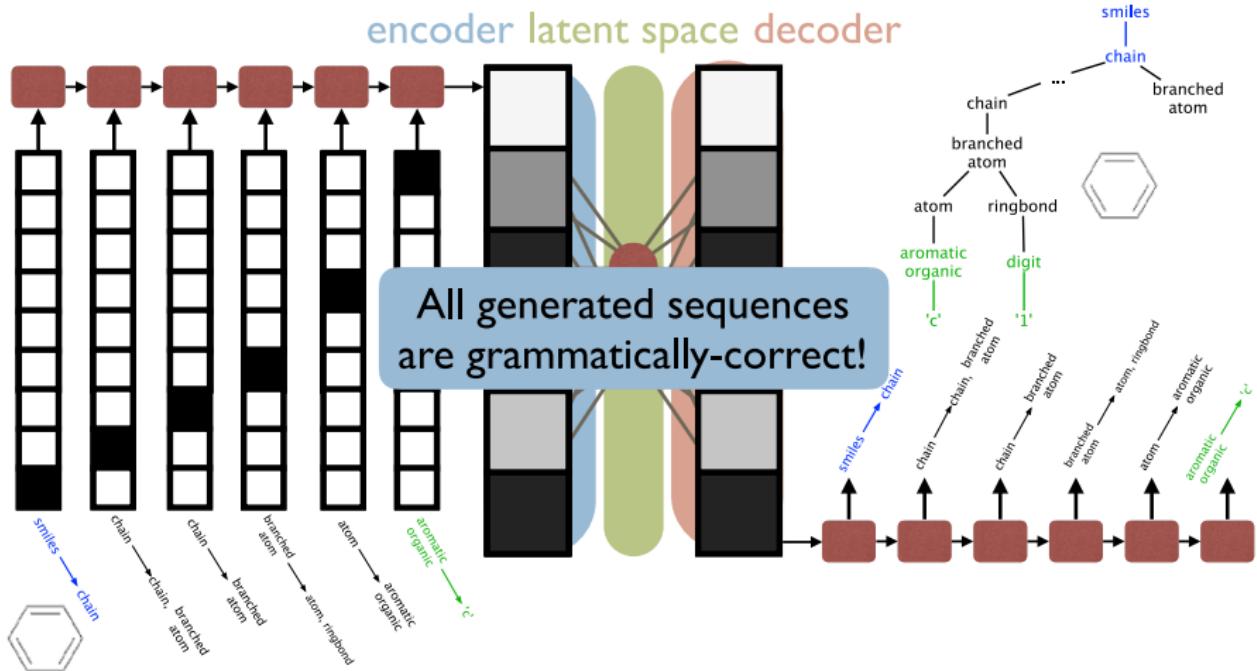
DECODER

[Kusner, Paige, Hernández-Lobato, 2017]



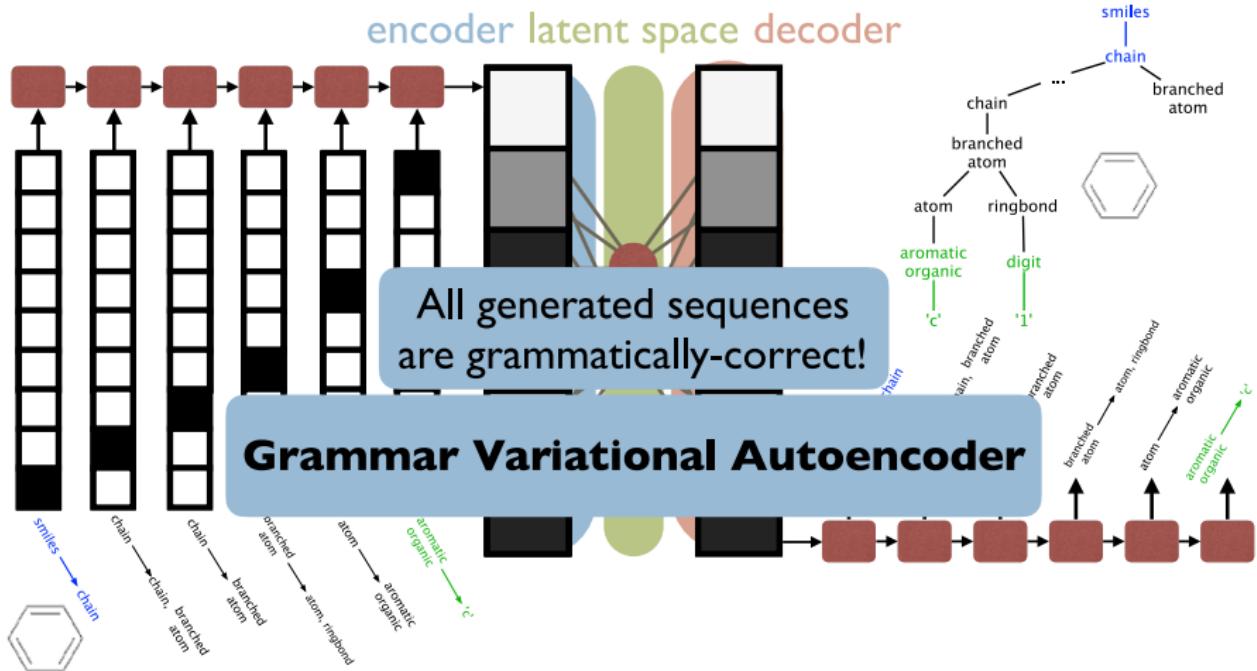
DECODER

[Kusner, Paige, Hernández-Lobato, 2017]



DECODER

[Kusner, Paige, Hernández-Lobato, 2017]



HOW WELL DOES THIS WORK?

Reconstruction accuracy and sample validity results.

Method	% Reconstruct	% Prior Valid
GVAE	53.7	7.2
CVAE	44.6	0.70

Results finding best molecule

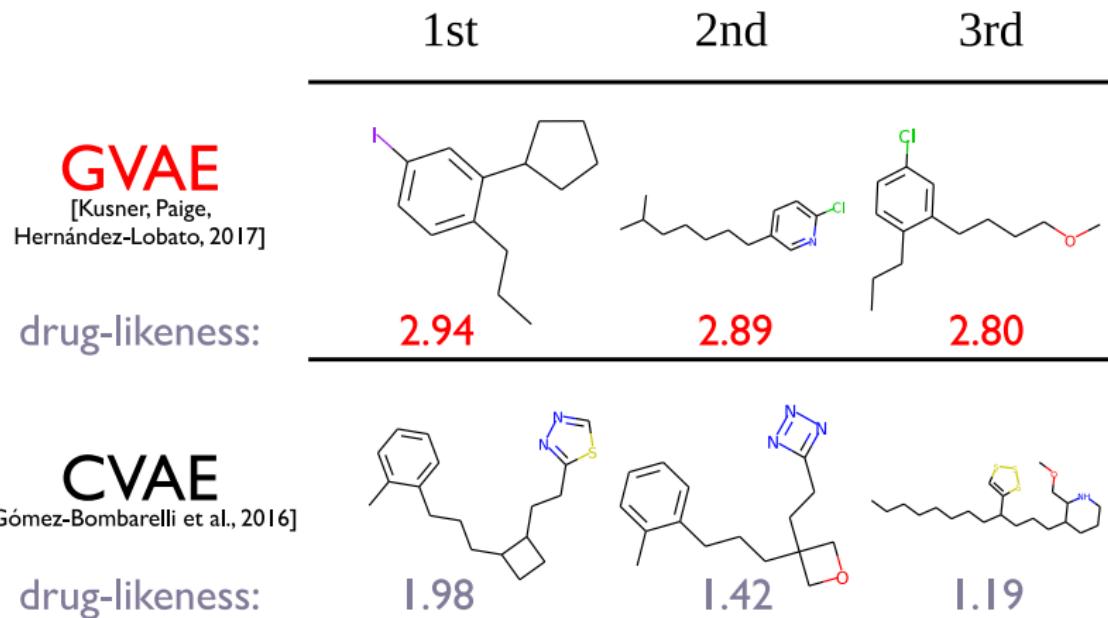
Method	Frac. valid	Avg. score
GVAE	0.31±0.07	-9.57 ±1.77
CVAE	0.17±0.05	-54.66±2.66

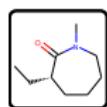
Test Log-likelihood (LL) and RMSE for predictions from latent space

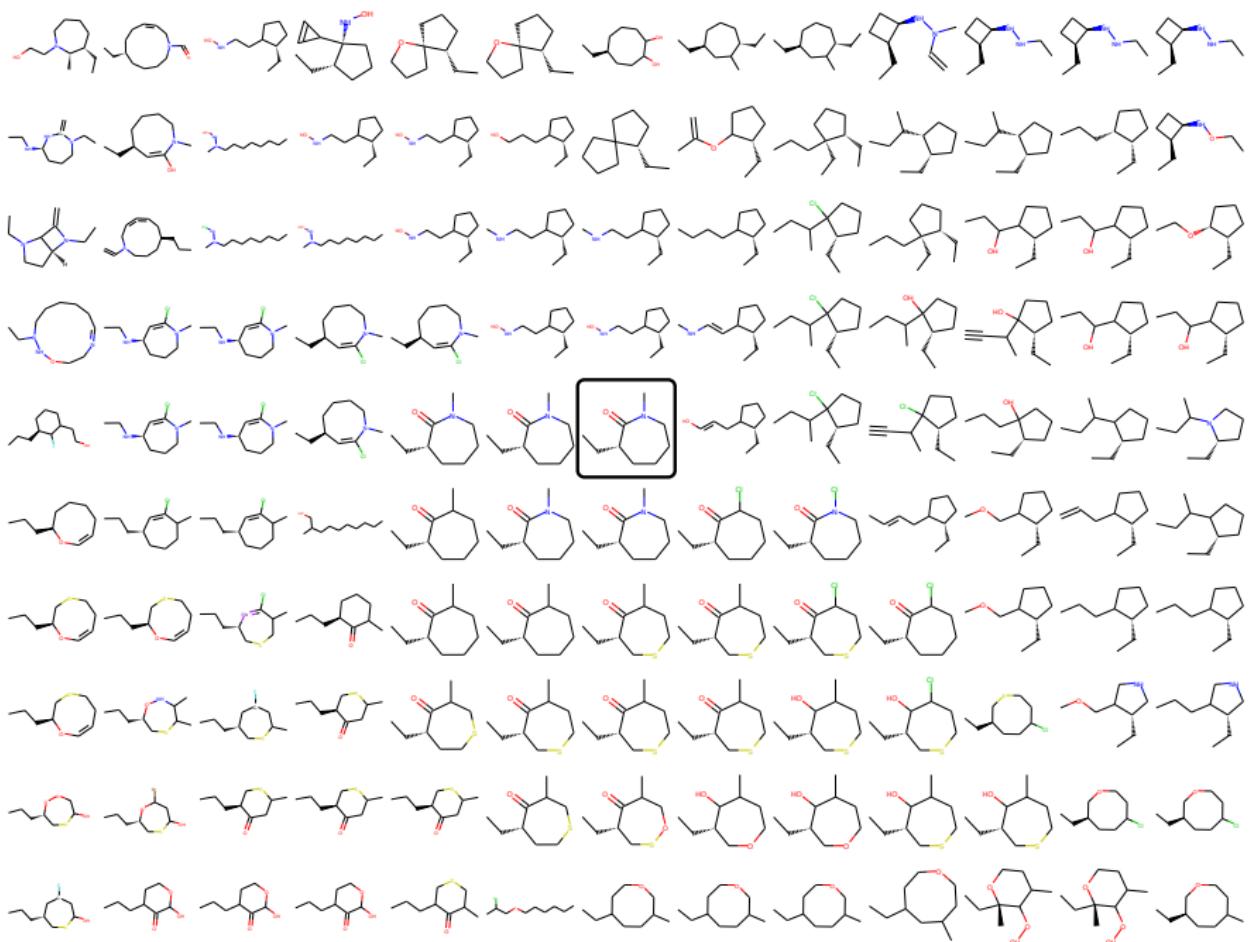
Objective	Method	Molecules
LL	GVAE	-1.739 ±0.004
	CVAE	-1.812±0.004
RMSE	GVAE	1.404 ±0.006
	CVAE	1.504±0.006

DRUG SEARCH

Trained on **ZINC**
250K molecules
maximum 120 chars







Take home messages

Generative models of molecules...

- ① **Data-driven approach** to molecule generation, no expertise needed.
- ② Sampled molecules are **realistic**, unlike those generated with rules.
- ③ Create a **continuous latent space** which is useful for optimization.

Grammar variational autoencoder...

- ① Produces a **larger fraction of valid molecules** when decoding.
- ② Produces **better predictions of molecule properties** from latent space.
- ③ **Molecule optimization** results are **improved**.

A lot of work going on this area...

Adding semantic constraints into GVAE:

- Dai, Tian, Dai, Skiena, and Song, 2018.

Generative models of graphs:

- Li, Vinyals, Dyer, Pascanu and Battaglia, 2018.
- Liu, Allamanis, Brockschmidt and Gaunt, 2018.
- Jin, Barzilay and Jaakkola, 2018.

Semi-supervised generative models:

- Kang and Cho, 2018.

Using GANs instead of VAEs:

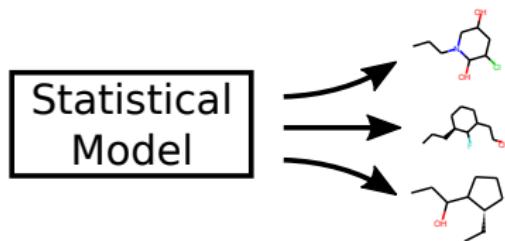
- Guimaraes, Sanchez-Lengeling, Outeiral, Farias and Aspuru-Guzik, 2017.
- De Cao and Kipf, 2018.

and many more!

Part IV: Predicting Electron Paths

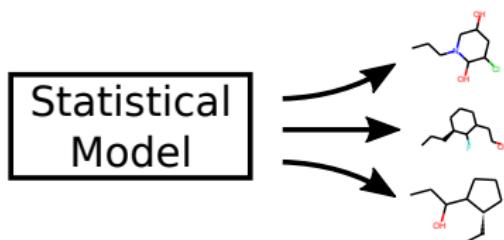
Motivation

Current generative models for molecules have **limitations**: will not tell you how to **synthesize** the generated molecules.

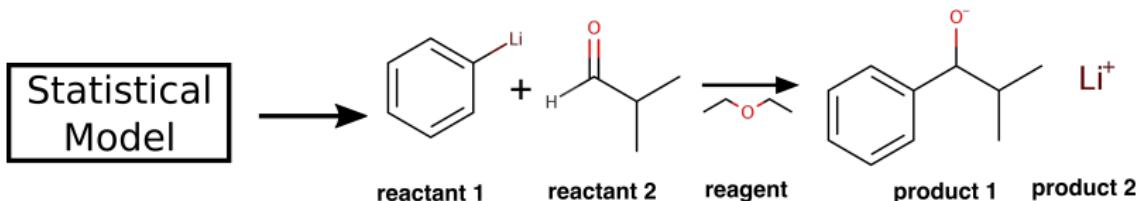


Motivation

Current generative models for molecules have **limitations**: will not tell you how to **synthesize** the generated molecules.

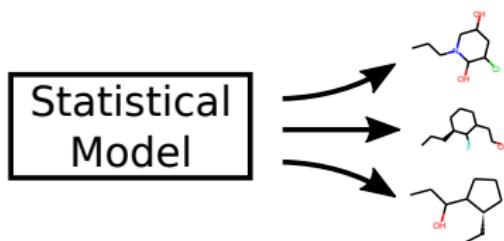


We want models that generate **synthetic routes** instead of just molecules!

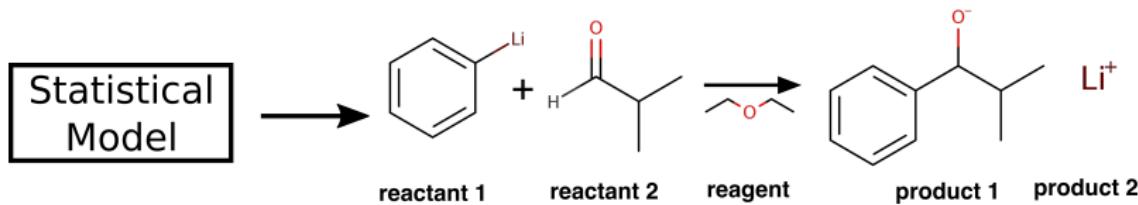


Motivation

Current generative models for molecules have **limitations**: will not tell you how to **synthesize** the generated molecules.



We want models that generate **synthetic routes** instead of just molecules!

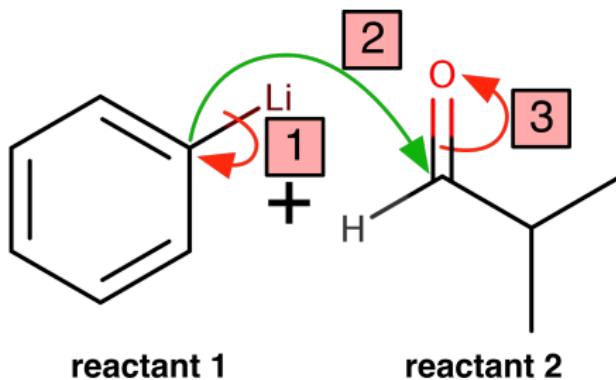


Perhaps too challenging. We just focus for now on reaction prediction.

Electron path

A chemical reaction is characterized by a **sequential movement of electrons** alternating between breaking and forming bonds.

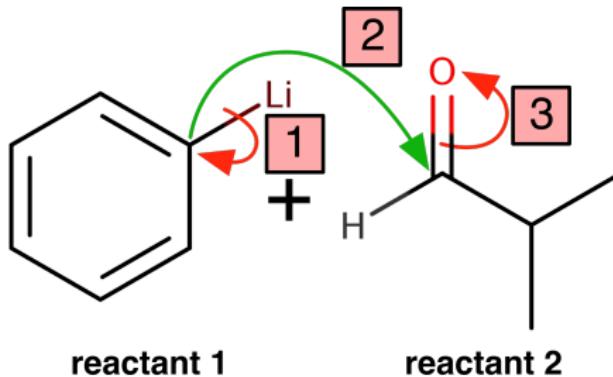
Represented by an **arrow pushing diagram** showing the **path followed by electrons** (detailed steps of the reaction):



Electron path

A chemical reaction is characterized by a **sequential movement of electrons** alternating between breaking and forming bonds.

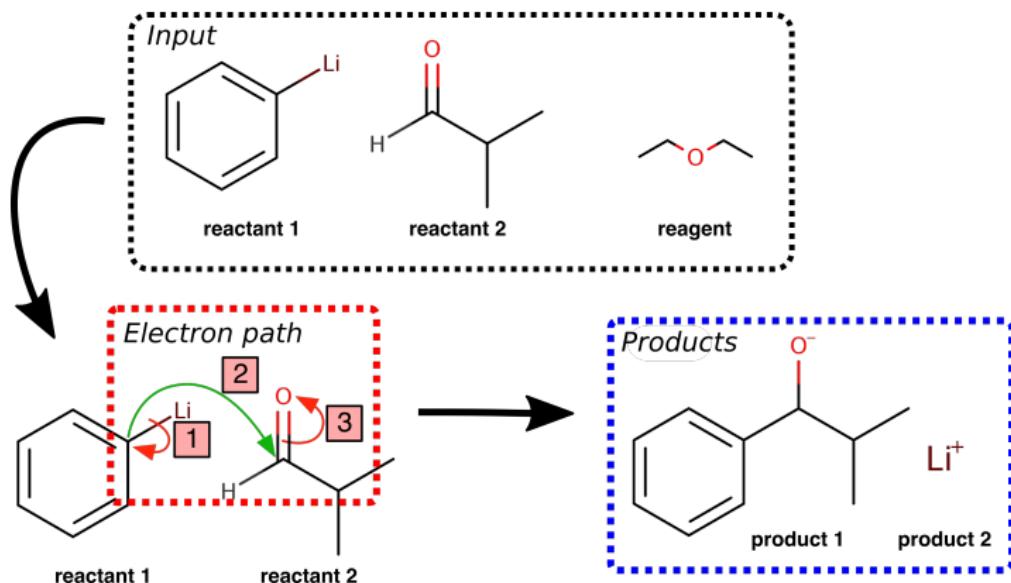
Represented by an **arrow pushing diagram** showing the **path followed by electrons** (detailed steps of the reaction):



Goal: predict the **electron path** from input molecules.

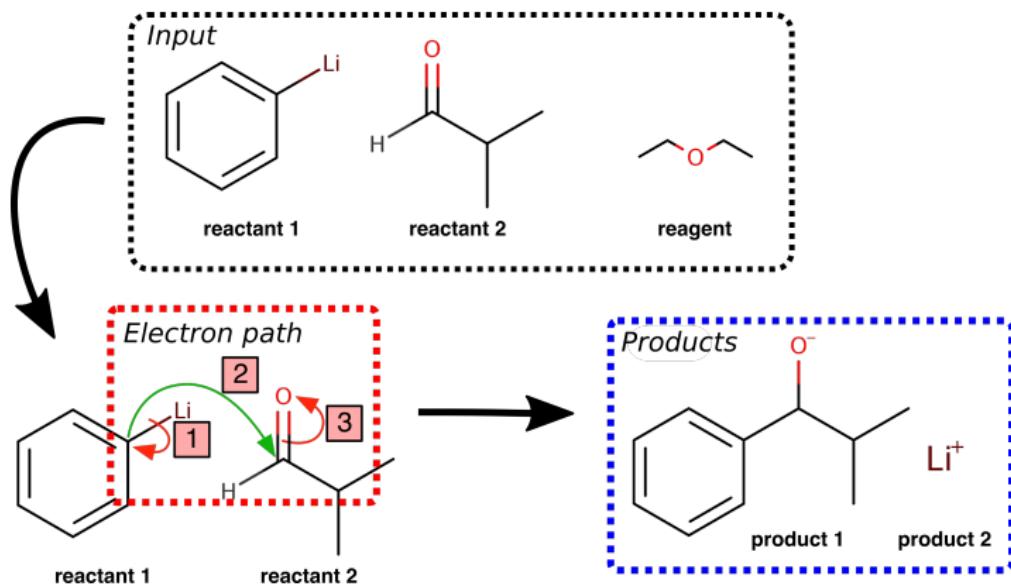
Prediction problem

Given the **electron path**, we directly obtain the final **products** of the reaction.



Prediction problem

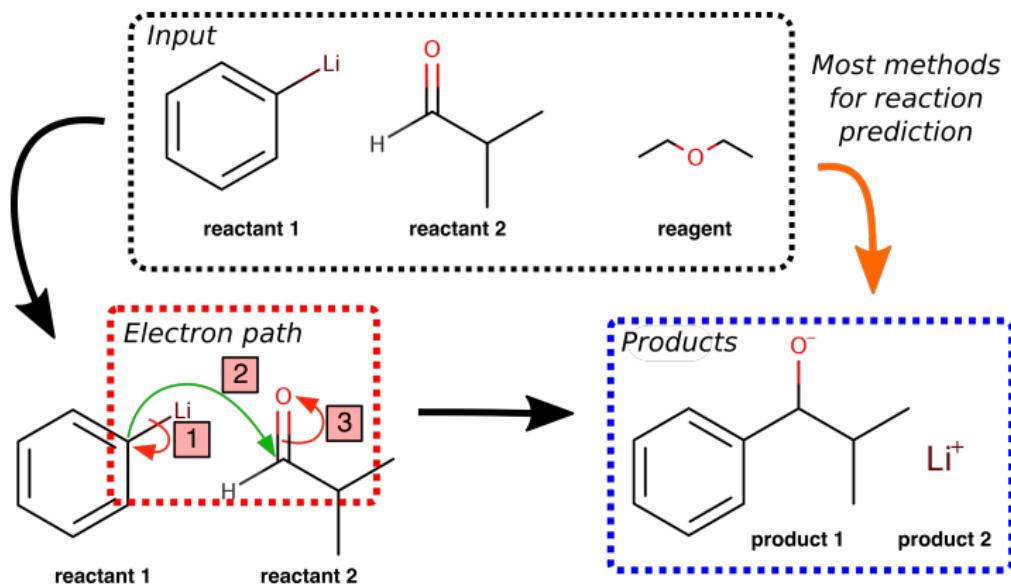
Given the **electron path**, we directly obtain the final **products** of the reaction.



Most reaction prediction approaches directly **predict products** from **input**.

Prediction problem

Given the **electron path**, we directly obtain the final **products** of the reaction.



Most reaction prediction approaches directly **predict products** from **input**.

Why predict electron paths?

- **Easy to interpret**: if the model makes a mistake, its easy to see where, and maybe even why.
- **Sparse**: typically reactions only affect 3-7 atoms out of 10-50. Modeling an electron path enforces this sparsity.
- **Chemically consistent**: easy to incorporate chemical constraints directly into a model for predicting electron paths.
- **Generalizable**: as electron paths exhibit regularities across different reactions, we naturally generalize to new inputs.

ELECTRO: an algorithm for electron paths

→ Select **starting** atom A.

Repeat:

Select atom B bonded to A and **remove** 2 electrons from bond A-B.

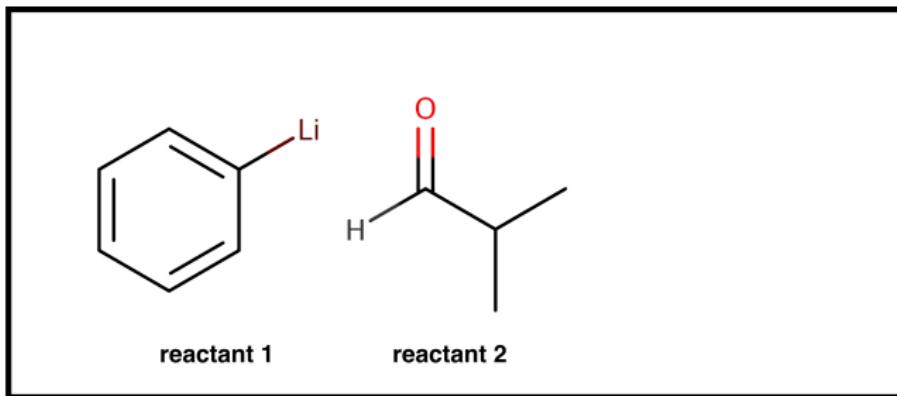
Should we **stop**?

Select new atom A and **add** bond B-A.

Should we **stop**?

Compute resulting electron path.

Compute resulting products.



ELECTRO: an algorithm for electron paths

→ Select **starting** atom A.

Repeat:

Select atom B bonded to A and **remove** 2 electrons from bond A-B.

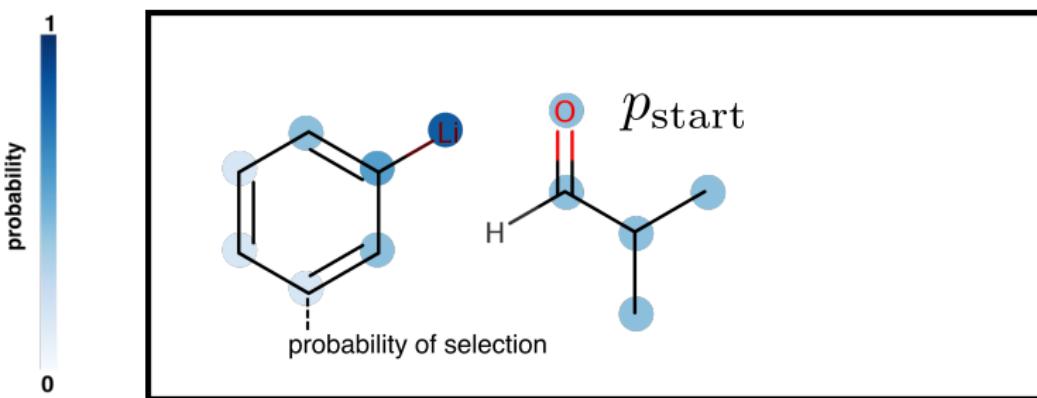
Should we **stop**?

Select new atom A and **add** bond B-A.

Should we **stop**?

Compute resulting electron path.

Compute resulting products.



ELECTRO: an algorithm for electron paths

→ Select **starting** atom A.

Repeat:

Select atom B bonded to A and **remove** 2 electrons from bond A-B.

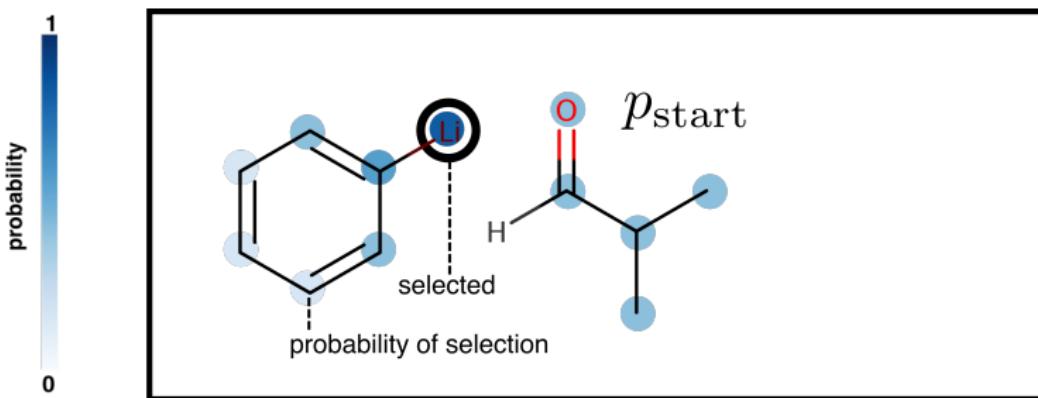
Should we **stop**?

Select new atom A and **add** bond B-A.

Should we **stop**?

Compute resulting electron path.

Compute resulting products.



ELECTRO: an algorithm for electron paths

Select **starting** atom A.

Repeat:

→ Select atom B bonded to A and **remove** 2 electrons from bond A-B.

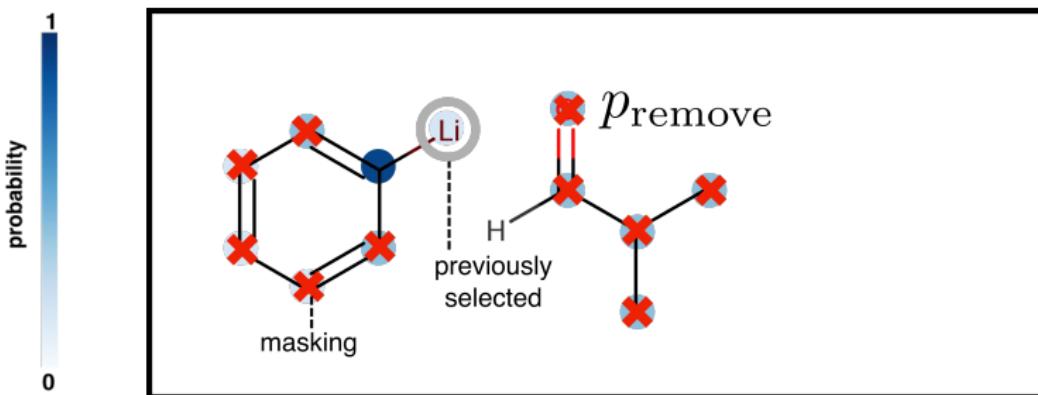
Should we **stop**?

Select new atom A and **add** bond B-A.

Should we **stop**?

Compute resulting electron path.

Compute resulting products.



ELECTRO: an algorithm for electron paths

Select **starting** atom A.

Repeat:

→ Select atom B bonded to A and **remove** 2 electrons from bond A-B.

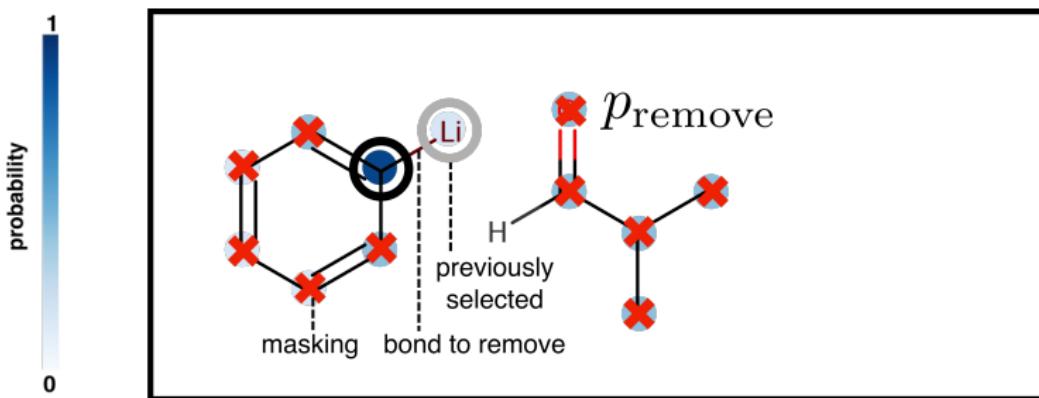
Should we **stop**?

Select new atom A and **add** bond B-A.

Should we **stop**?

Compute resulting electron path.

Compute resulting products.



ELECTRO: an algorithm for electron paths

Select **starting** atom A.

Repeat:

Select atom B bonded to A and **remove** 2 electrons from bond A-B.

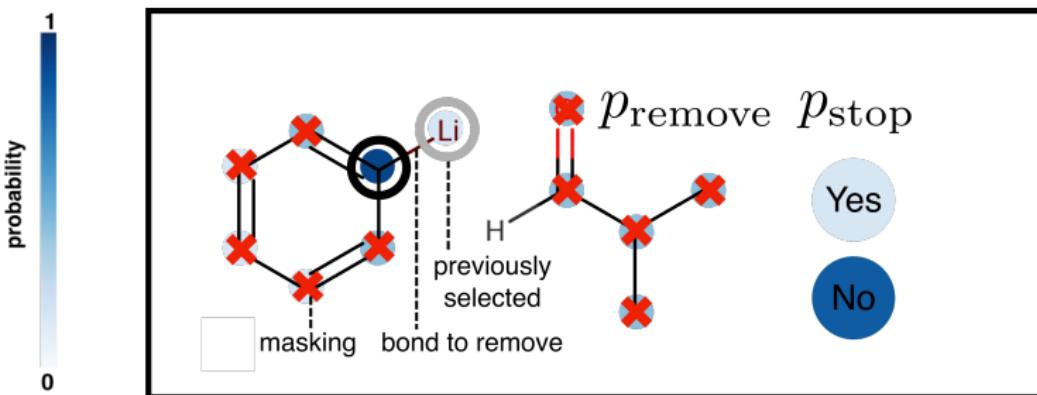
→ Should we **stop**?

Select new atom A and **add** bond B-A.

Should we **stop**?

Compute resulting electron path.

Compute resulting products.



ELECTRO: an algorithm for electron paths

Select **starting** atom A.

Repeat:

Select atom B bonded to A and **remove** 2 electrons from bond A-B.

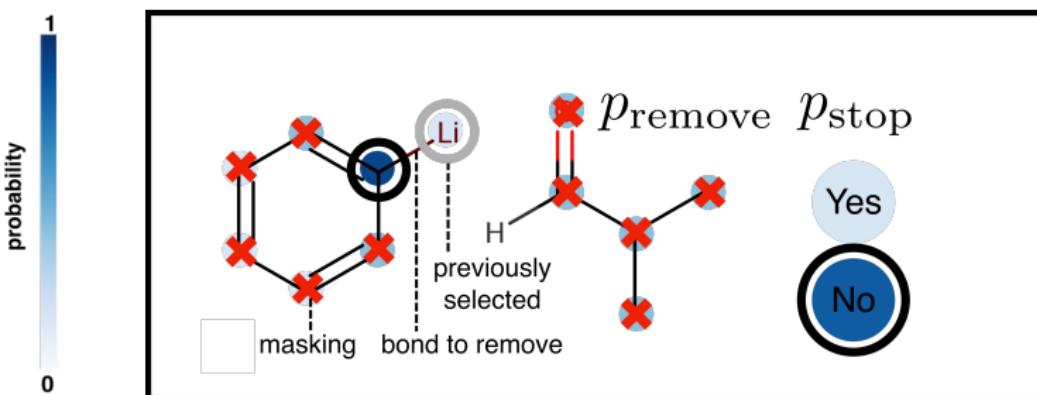
→ Should we **stop**?

Select new atom A and **add** bond B-A.

Should we **stop**?

Compute resulting electron path.

Compute resulting products.



ELECTRO: an algorithm for electron paths

Select **starting** atom A.

Repeat:

Select atom B bonded to A and **remove** 2 electrons from bond A-B.

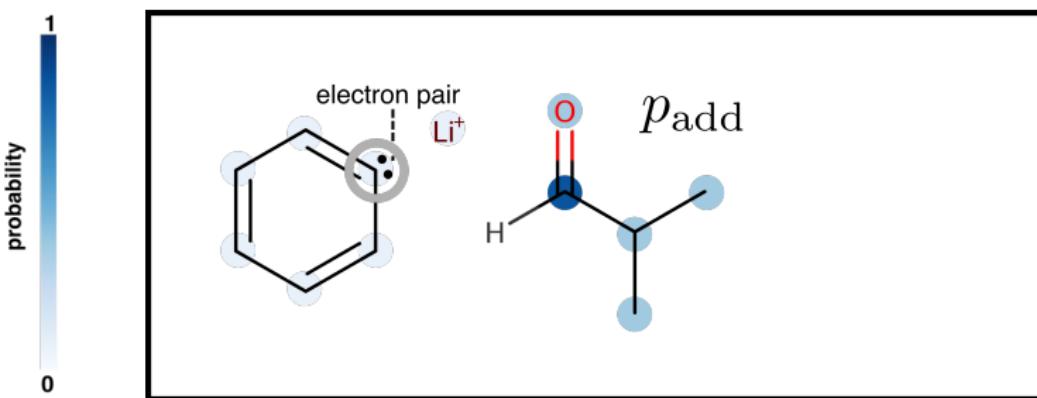
Should we **stop**?

→ Select new atom A and **add** bond B-A.

Should we **stop**?

Compute resulting electron path.

Compute resulting products.



ELECTRO: an algorithm for electron paths

Select **starting** atom A.

Repeat:

Select atom B bonded to A and **remove** 2 electrons from bond A-B.

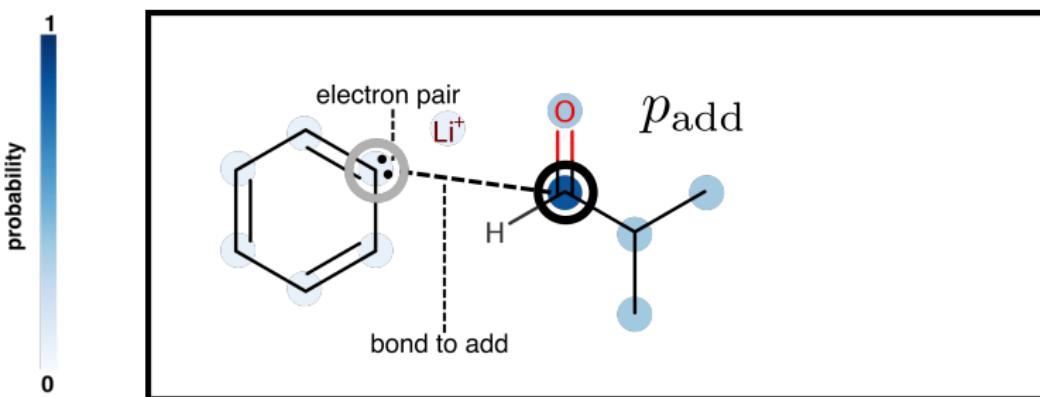
Should we **stop**?

→ Select new atom A and **add** bond B-A.

Should we **stop**?

Compute resulting electron path.

Compute resulting products.



ELECTRO: an algorithm for electron paths

Select **starting** atom A.

Repeat:

Select atom B bonded to A and **remove** 2 electrons from bond A-B.

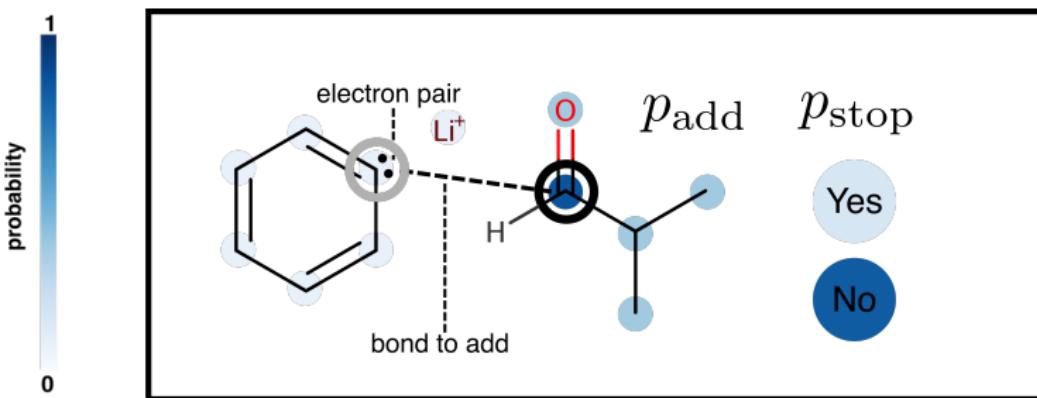
Should we **stop**?

Select new atom A and **add** bond B-A.

→ Should we **stop**?

Compute resulting electron path.

Compute resulting products.



ELECTRO: an algorithm for electron paths

Select **starting** atom A.

Repeat:

Select atom B bonded to A and **remove** 2 electrons from bond A-B.

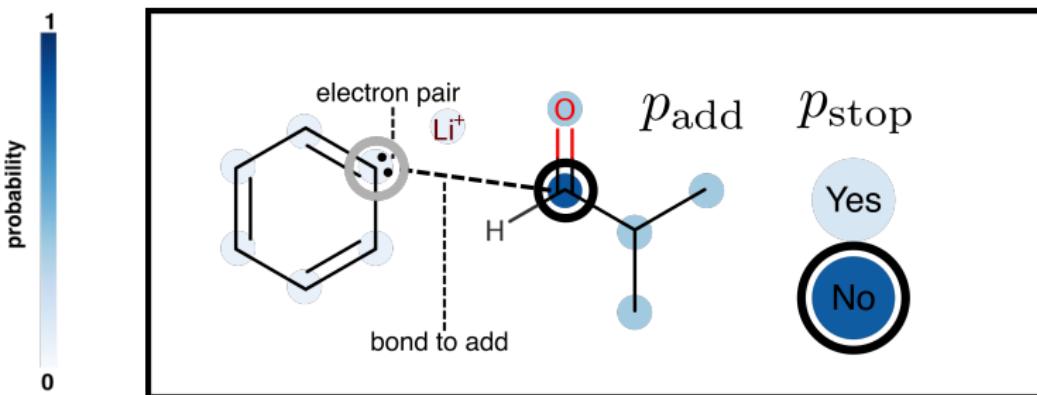
Should we **stop**?

Select new atom A and **add** bond B-A.

→ Should we **stop**?

Compute resulting electron path.

Compute resulting products.



ELECTRO: an algorithm for electron paths

Select **starting** atom A.

Repeat:

→ Select atom B bonded to A and **remove** 2 electrons from bond A-B.

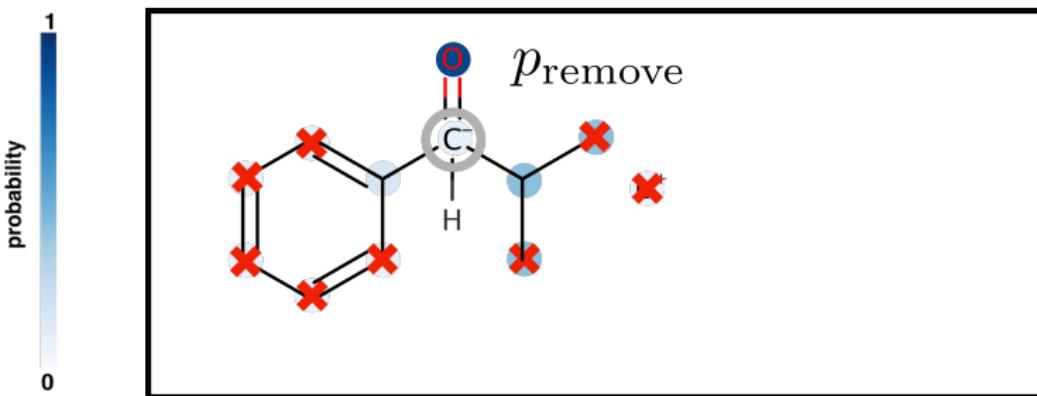
Should we **stop**?

Select new atom A and **add** bond B-A.

Should we **stop**?

Compute resulting electron path.

Compute resulting products.



ELECTRO: an algorithm for electron paths

Select **starting** atom A.

Repeat:

→ Select atom B bonded to A and **remove** 2 electrons from bond A-B.

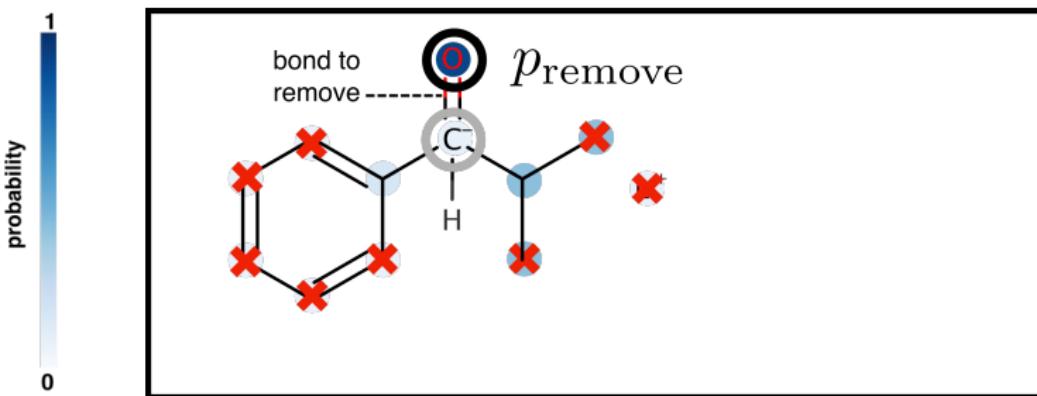
Should we **stop**?

Select new atom A and **add** bond B-A.

Should we **stop**?

Compute resulting electron path.

Compute resulting products.



ELECTRO: an algorithm for electron paths

Select **starting** atom A.

Repeat:

Select atom B bonded to A and **remove** 2 electrons from bond A-B.

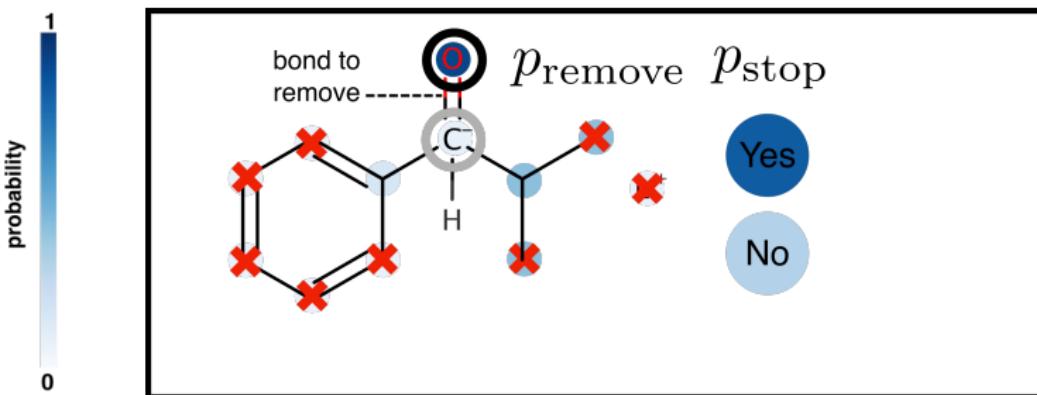
→ Should we **stop**?

Select new atom A and **add** bond B-A.

Should we **stop**?

Compute resulting electron path.

Compute resulting products.



ELECTRO: an algorithm for electron paths

Select **starting** atom A.

Repeat:

Select atom B bonded to A and **remove** 2 electrons from bond A-B.

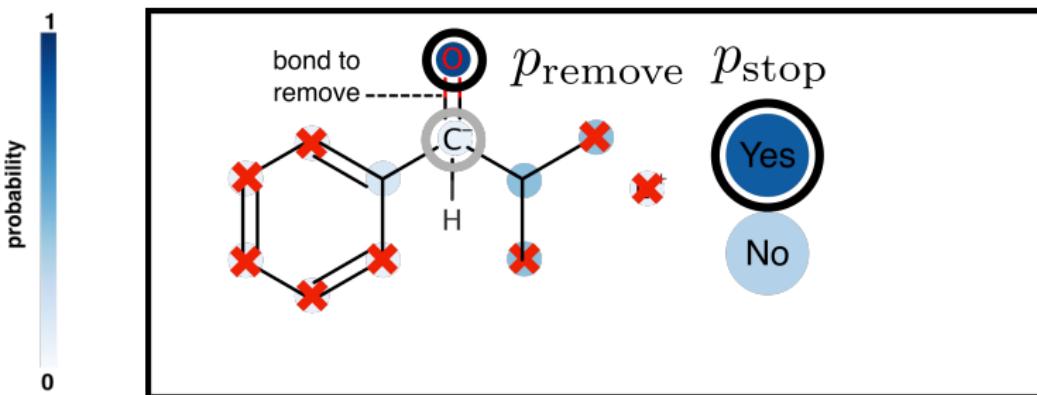
→ Should we **stop**?

Select new atom A and **add** bond B-A.

Should we **stop**?

Compute resulting electron path.

Compute resulting products.



ELECTRO: an algorithm for electron paths

Select **starting** atom A.

Repeat:

Select atom B bonded to A and **remove** 2 electrons from bond A-B.

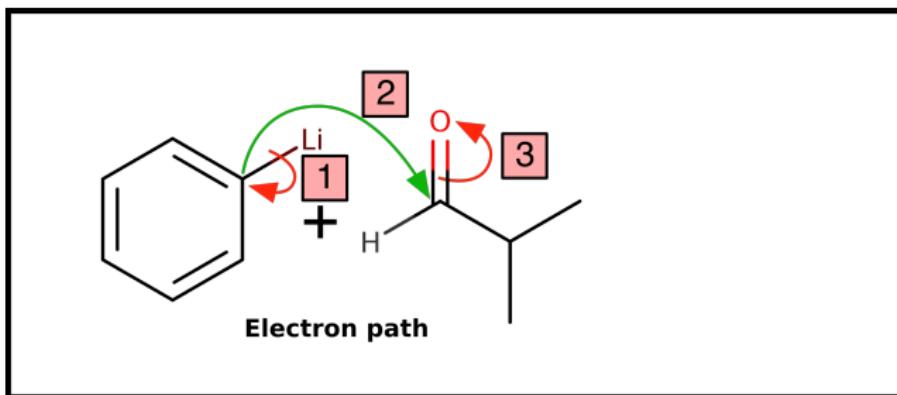
Should we **stop**?

Select new atom A and **add** bond B-A.

Should we **stop**?

→ Compute resulting electron path.

Compute resulting products.



ELECTRO: an algorithm for electron paths

Select **starting** atom A.

Repeat:

Select atom B bonded to A and **remove** 2 electrons from bond A-B.

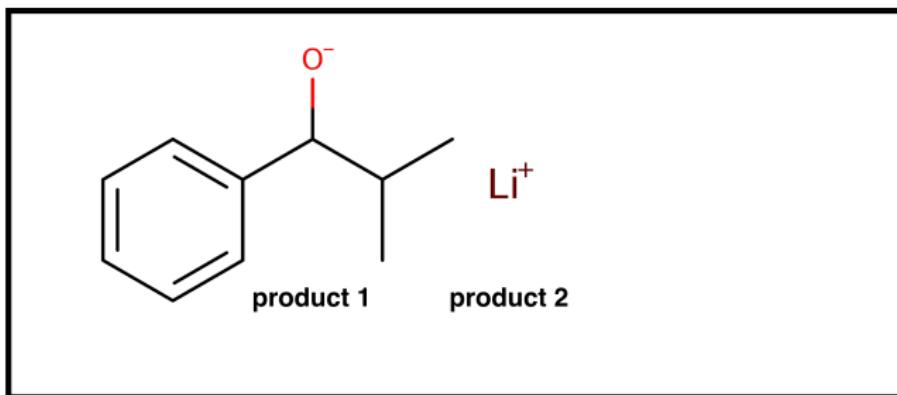
Should we **stop**?

Select new atom A and **add** bond B-A.

Should we **stop**?

Compute resulting electron path.

→ Compute resulting products.



A probabilistic model

Electron path $\mathbf{a} = (a_1, a_2, \dots, a_T)$ defined by **sequence of selected atoms**.

Electrons **removed and added alternatively** from pairs of consecutive atoms.

A probabilistic model

Electron path $\mathbf{a} = (a_1, a_2, \dots, a_T)$ defined by **sequence of selected atoms**.

Electrons **removed and added alternatively** from pairs of consecutive atoms.

We assume \mathbf{a} is sampled from

$$p_{\theta}(\mathbf{a}|\mathcal{M}_1) = \left[\prod_{t=1}^T p_{\theta}(a_t|a_{t-1}, \mathcal{M}_t) p_{\theta}(\text{CONT}|\mathcal{M}_t) \right] p_{\theta}(\text{STOP}|\mathcal{M}_{T+1}),$$

where

- θ are the model parameters.
- \mathcal{M}_t represents the current state of molecules.
- $p_{\theta}(\text{STOP}|\mathcal{M}_t) = 1 - p_{\theta}(\text{CONT}|\mathcal{M}_t)$ is the probability of stopping.

A probabilistic model

Electron path $\mathbf{a} = (a_1, a_2, \dots, a_T)$ defined by **sequence of selected atoms**.

Electrons **removed and added alternatively** from pairs of consecutive atoms.

We assume \mathbf{a} is sampled from

$$p_{\theta}(\mathbf{a}|\mathcal{M}_1) = \left[\prod_{t=1}^T p_{\theta}(a_t|a_{t-1}, \mathcal{M}_t) p_{\theta}(\text{CONT}|\mathcal{M}_t) \right] p_{\theta}(\text{STOP}|\mathcal{M}_{T+1}),$$

where

- θ are the model parameters.
- \mathcal{M}_t represents the current state of molecules.
- $p_{\theta}(\text{STOP}|\mathcal{M}_t) = 1 - p_{\theta}(\text{CONT}|\mathcal{M}_t)$ is the probability of stopping.

How to specify $p_{\theta}(a_t|a_{t-1}, \mathcal{M}_t)$ and $p_{\theta}(\text{CONT}|\mathcal{M}_t)$?

Forming node and graph embeddings

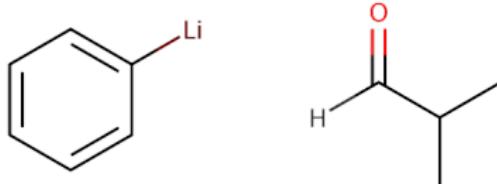
→ Initialize atom features $\mathbf{h}_{0,i}$.

For $t = 1$ to T :

Compute atom messages $\mathbf{m}_{t,i} = \sum_{j \sim i} \mathbf{A}_{\text{type}(j \sim i)} \mathbf{h}_{t-1,j}$.

Update atom features $\mathbf{h}_{t,i} = \text{RNN}(\mathbf{m}_{t,i}, \mathbf{h}_{t-1,i})$.

Compute graph embedding $\mathbf{h}_G = f_{\text{down}}\{\sum_i \sigma[f_{\text{gate}}(\mathbf{h}_{T,i})]f_{\text{up}}(\mathbf{h}_{T,i})\}$.



Forming node and graph embeddings

→ Initialize atom features $\mathbf{h}_{0,i}$.

For $t = 1$ to T :

Compute atom messages $\mathbf{m}_{t,i} = \sum_{j \sim i} \mathbf{A}_{\text{type}(j \sim i)} \mathbf{h}_{t-1,j}$.

Update atom features $\mathbf{h}_{t,i} = \text{RNN}(\mathbf{m}_{t,i}, \mathbf{h}_{t-1,i})$.

Compute graph embedding $\mathbf{h}_G = f_{\text{down}}\{\sum_i \sigma[f_{\text{gate}}(\mathbf{h}_{T,i})]f_{\text{up}}(\mathbf{h}_{T,i})\}$.

Feature

Atom type

Degree

Explicit Valence

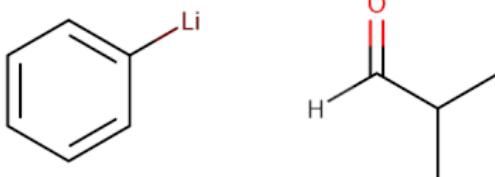
Hybridization

H count

Electronegativity

Atomic number

Part of an aromatic ring



Forming node and graph embeddings

→ Initialize atom features $\mathbf{h}_{0,i}$.

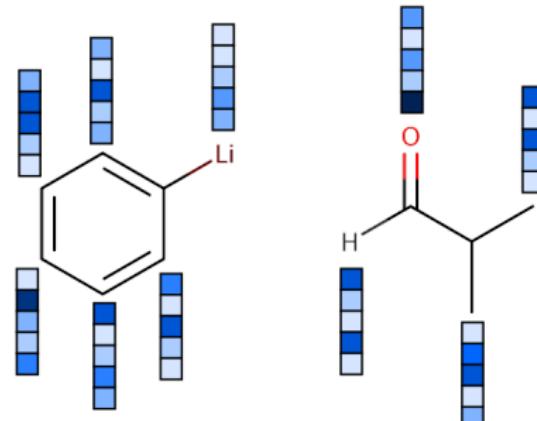
For $t = 1$ to T :

Compute atom messages $\mathbf{m}_{t,i} = \sum_{j \sim i} \mathbf{A}_{\text{type}(j \sim i)} \mathbf{h}_{t-1,j}$.

Update atom features $\mathbf{h}_{t,i} = \text{RNN}(\mathbf{m}_{t,i}, \mathbf{h}_{t-1,i})$.

Compute graph embedding $\mathbf{h}_G = f_{\text{down}}\{\sum_i \sigma[f_{\text{gate}}(\mathbf{h}_{T,i})]f_{\text{up}}(\mathbf{h}_{T,i})\}$.

Feature
Atom type
Degree
Explicit Valence
Hybridization
H count
Electronegativity
Atomic number
Part of an aromatic ring



Forming node and graph embeddings

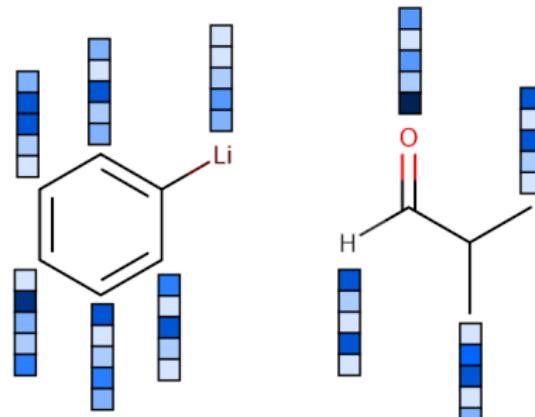
Initialize atom features $\mathbf{h}_{0,i}$.

For $t = 1$ to T :

→ Compute atom messages $\mathbf{m}_{t,i} = \sum_{j \sim i} \mathbf{A}_{\text{type}(j \sim i)} \mathbf{h}_{t-1,j}$.

Update atom features $\mathbf{h}_{t,i} = \text{RNN}(\mathbf{m}_{t,i}, \mathbf{h}_{t-1,i})$.

Compute graph embedding $\mathbf{h}_G = f_{\text{down}}\{\sum_i \sigma[f_{\text{gate}}(\mathbf{h}_{T,i})]f_{\text{up}}(\mathbf{h}_{T,i})\}$.



Forming node and graph embeddings

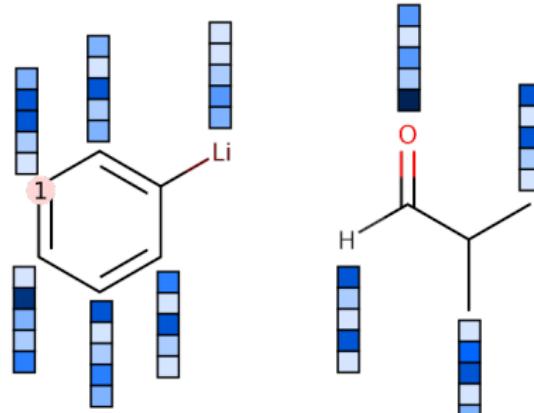
Initialize atom features $\mathbf{h}_{0,i}$.

For $t = 1$ to T :

→ Compute atom messages $\mathbf{m}_{t,i} = \sum_{j \sim i} \mathbf{A}_{\text{type}(j \sim i)} \mathbf{h}_{t-1,j}$.

Update atom features $\mathbf{h}_{t,i} = \text{RNN}(\mathbf{m}_{t,i}, \mathbf{h}_{t-1,i})$.

Compute graph embedding $\mathbf{h}_G = f_{\text{down}}\{\sum_i \sigma[f_{\text{gate}}(\mathbf{h}_{T,i})]f_{\text{up}}(\mathbf{h}_{T,i})\}$.



Forming node and graph embeddings

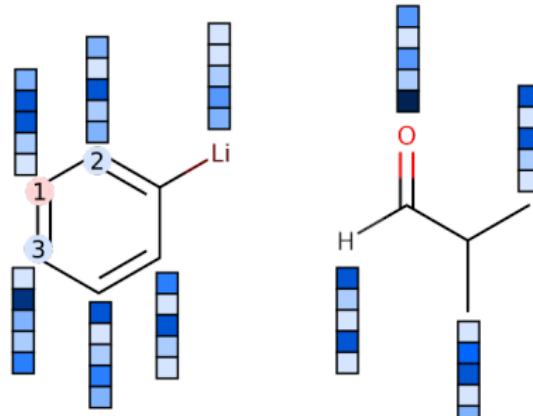
Initialize atom features $\mathbf{h}_{0,i}$.

For $t = 1$ to T :

→ Compute atom messages $\mathbf{m}_{t,i} = \sum_{j \sim i} \mathbf{A}_{\text{type}(j \sim i)} \mathbf{h}_{t-1,j}$.

Update atom features $\mathbf{h}_{t,i} = \text{RNN}(\mathbf{m}_{t,i}, \mathbf{h}_{t-1,i})$.

Compute graph embedding $\mathbf{h}_G = f_{\text{down}}\{\sum_i \sigma[f_{\text{gate}}(\mathbf{h}_{T,i})]f_{\text{up}}(\mathbf{h}_{T,i})\}$.



Forming node and graph embeddings

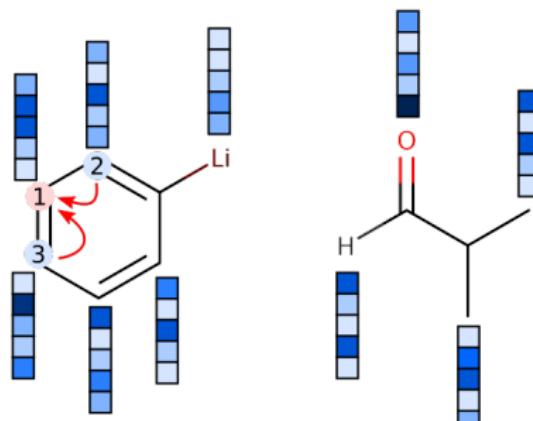
Initialize atom features $\mathbf{h}_{0,i}$.

For $t = 1$ to T :

→ Compute atom messages $\mathbf{m}_{t,i} = \sum_{j \sim i} \mathbf{A}_{\text{type}(j \sim i)} \mathbf{h}_{t-1,j}$.

Update atom features $\mathbf{h}_{t,i} = \text{RNN}(\mathbf{m}_{t,i}, \mathbf{h}_{t-1,i})$.

Compute graph embedding $\mathbf{h}_G = f_{\text{down}}\{\sum_i \sigma[f_{\text{gate}}(\mathbf{h}_{T,i})]f_{\text{up}}(\mathbf{h}_{T,i})\}$.



Forming node and graph embeddings

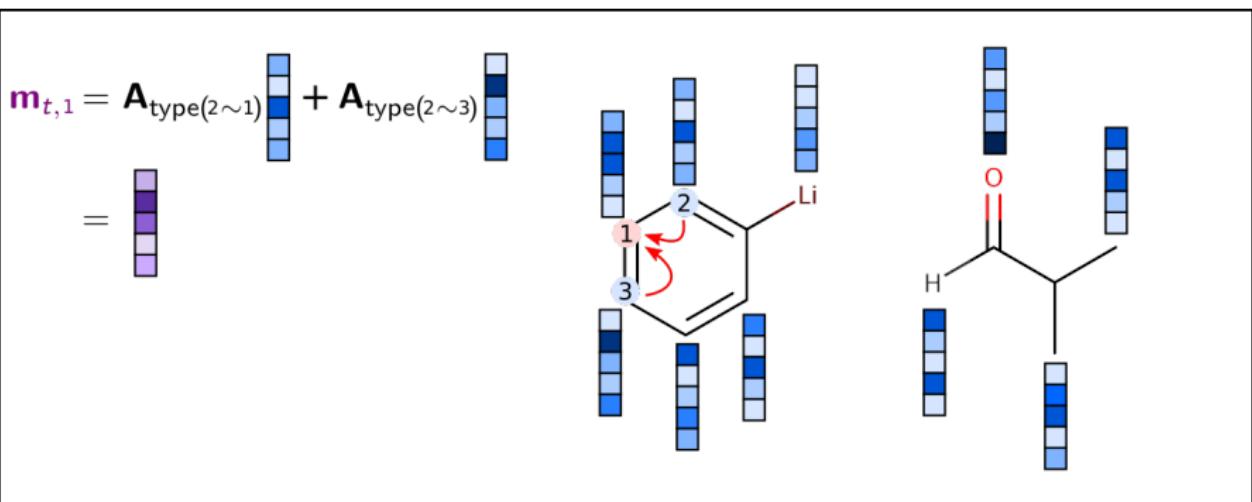
Initialize atom features $\mathbf{h}_{0,i}$.

For $t = 1$ to T :

→ Compute atom messages $\mathbf{m}_{t,i} = \sum_{j \sim i} \mathbf{A}_{\text{type}(j \sim i)} \mathbf{h}_{t-1,j}$.

Update atom features $\mathbf{h}_{t,i} = \text{RNN}(\mathbf{m}_{t,i}, \mathbf{h}_{t-1,i})$.

Compute graph embedding $\mathbf{h}_G = f_{\text{down}}\{\sum_i \sigma[f_{\text{gate}}(\mathbf{h}_{T,i})]f_{\text{up}}(\mathbf{h}_{T,i})\}$.



Forming node and graph embeddings

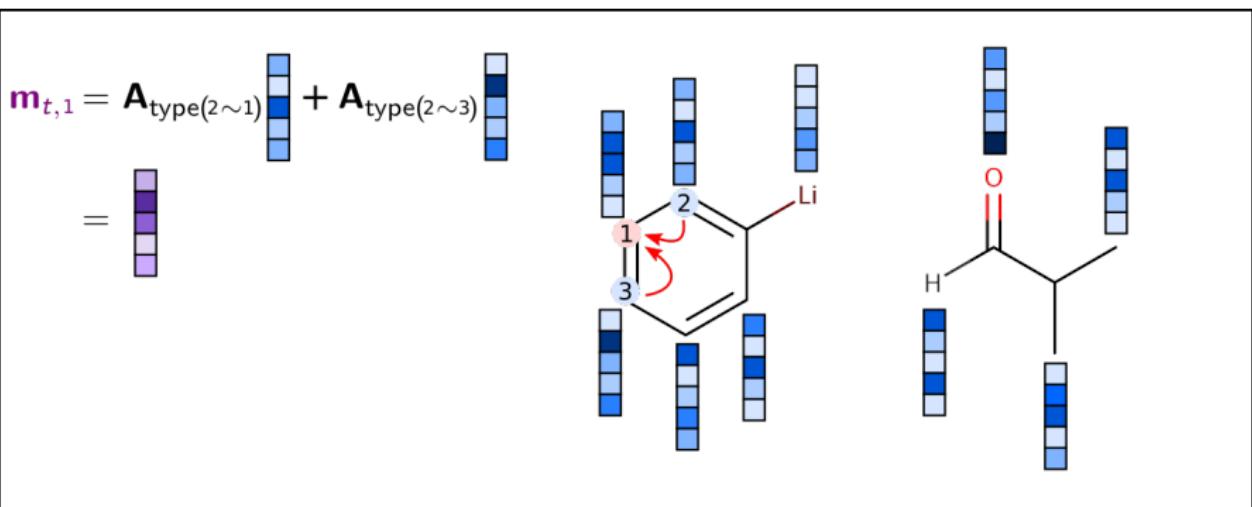
Initialize atom features $\mathbf{h}_{0,i}$.

For $t = 1$ to T :

Compute atom messages $\mathbf{m}_{t,i} = \sum_{j \sim i} \mathbf{A}_{\text{type}(j \sim i)} \mathbf{h}_{t-1,j}$.

→ Update atom features $\mathbf{h}_{t,i} = \text{RNN}(\mathbf{m}_{t,i}, \mathbf{h}_{t-1,i})$.

Compute graph embedding $\mathbf{h}_G = f_{\text{down}}\{\sum_i \sigma[f_{\text{gate}}(\mathbf{h}_{T,i})]f_{\text{up}}(\mathbf{h}_{T,i})\}$.



Forming node and graph embeddings

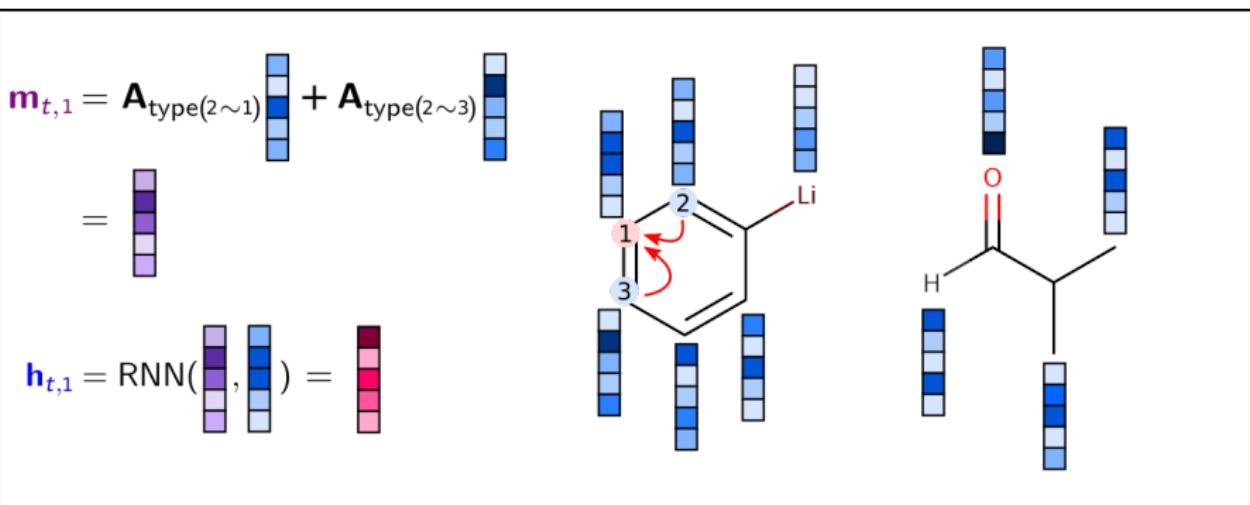
Initialize atom features $\mathbf{h}_{0,i}$.

For $t = 1$ to T :

Compute atom messages $\mathbf{m}_{t,i} = \sum_{j \sim i} \mathbf{A}_{\text{type}(j \sim i)} \mathbf{h}_{t-1,j}$.

→ Update atom features $\mathbf{h}_{t,i} = \text{RNN}(\mathbf{m}_{t,i}, \mathbf{h}_{t-1,i})$.

Compute graph embedding $\mathbf{h}_G = f_{\text{down}}\{\sum_i \sigma[f_{\text{gate}}(\mathbf{h}_{T,i})]f_{\text{up}}(\mathbf{h}_{T,i})\}$.



Forming node and graph embeddings

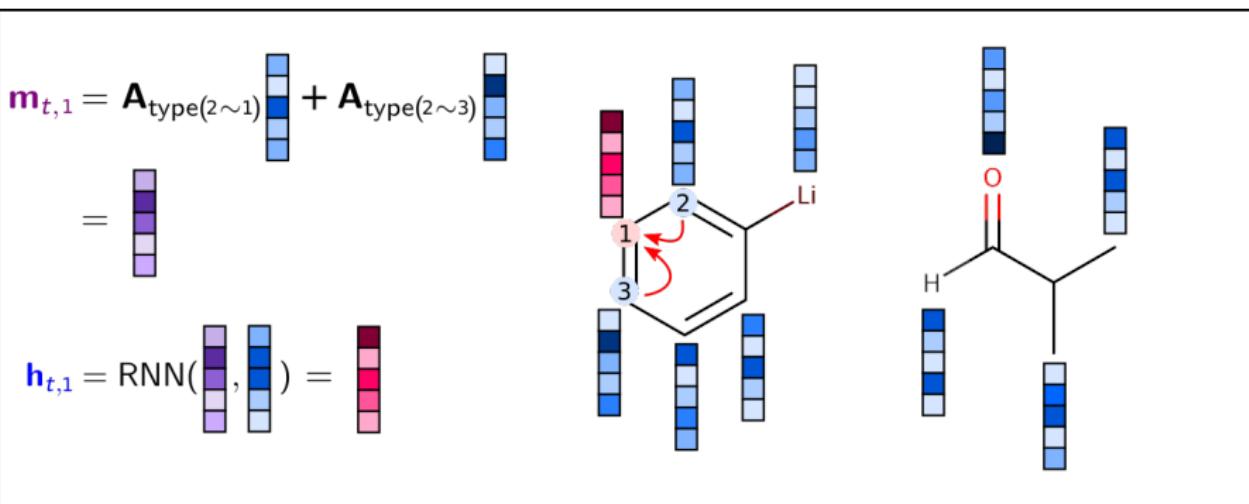
Initialize atom features $\mathbf{h}_{0,i}$.

For $t = 1$ to T :

Compute atom messages $\mathbf{m}_{t,i} = \sum_{j \sim i} \mathbf{A}_{\text{type}(j \sim i)} \mathbf{h}_{t-1,j}$.

→ Update atom features $\mathbf{h}_{t,i} = \text{RNN}(\mathbf{m}_{t,i}, \mathbf{h}_{t-1,i})$.

Compute graph embedding $\mathbf{h}_G = f_{\text{down}}\{\sum_i \sigma[f_{\text{gate}}(\mathbf{h}_{T,i})]f_{\text{up}}(\mathbf{h}_{T,i})\}$.



Forming node and graph embeddings

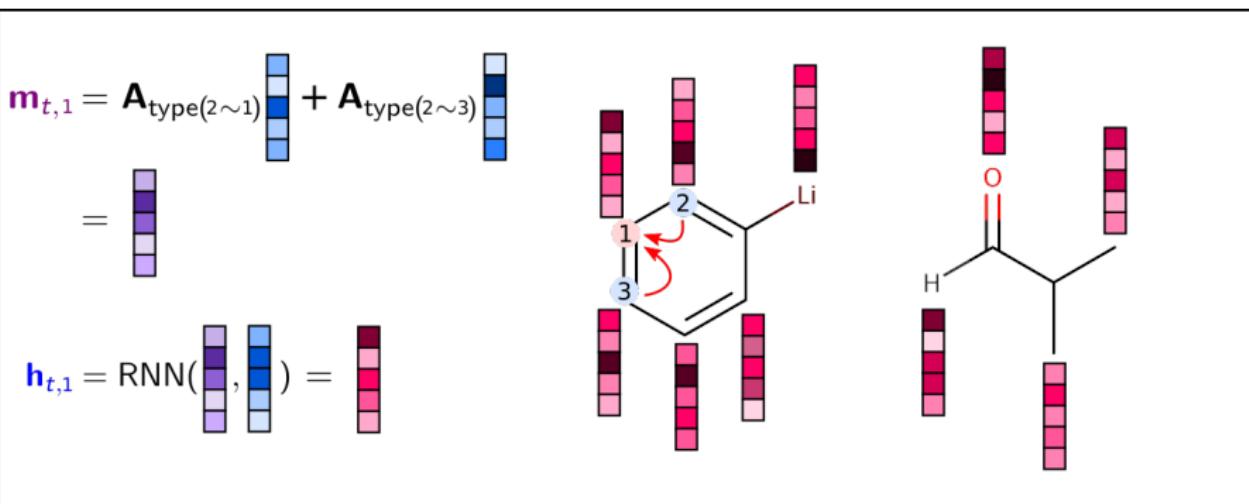
Initialize atom features $\mathbf{h}_{0,i}$.

For $t = 1$ to T :

Compute atom messages $\mathbf{m}_{t,i} = \sum_{j \sim i} \mathbf{A}_{\text{type}(j \sim i)} \mathbf{h}_{t-1,j}$.

→ Update atom features $\mathbf{h}_{t,i} = \text{RNN}(\mathbf{m}_{t,i}, \mathbf{h}_{t-1,i})$.

Compute graph embedding $\mathbf{h}_G = f_{\text{down}}\{\sum_i \sigma[f_{\text{gate}}(\mathbf{h}_{T,i})]f_{\text{up}}(\mathbf{h}_{T,i})\}$.



Forming node and graph embeddings

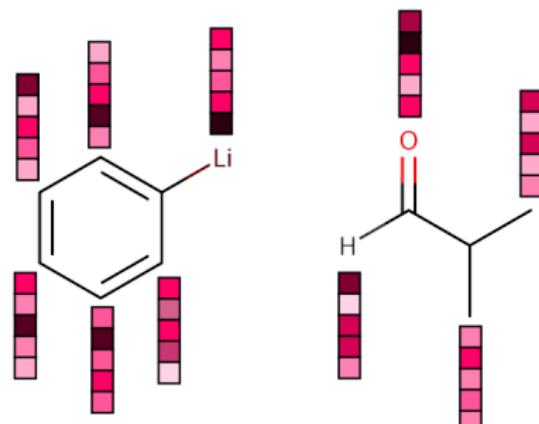
Initialize atom features $\mathbf{h}_{0,i}$.

For $t = 1$ to T :

Compute atom messages $\mathbf{m}_{t,i} = \sum_{j \sim i} \mathbf{A}_{\text{type}(j \sim i)} \mathbf{h}_{t-1,j}$.

Update atom features $\mathbf{h}_{t,i} = \text{RNN}(\mathbf{m}_{t,i}, \mathbf{h}_{t-1,i})$.

→ Compute graph embedding $\mathbf{h}_G = f_{\text{down}}\{\sum_i \sigma[f_{\text{gate}}(\mathbf{h}_{T,i})]f_{\text{up}}(\mathbf{h}_{T,i})\}$.



Forming node and graph embeddings

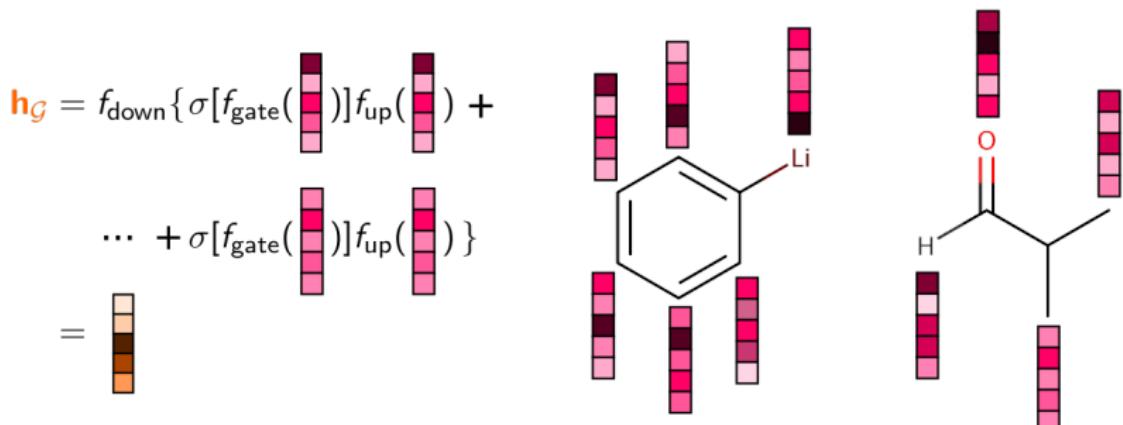
Initialize atom features $\mathbf{h}_{0,i}$.

For $t = 1$ to T :

Compute atom messages $\mathbf{m}_{t,i} = \sum_{j \sim i} \mathbf{A}_{\text{type}(j \sim i)} \mathbf{h}_{t-1,j}$.

Update atom features $\mathbf{h}_{t,i} = \text{RNN}(\mathbf{m}_{t,i}, \mathbf{h}_{t-1,i})$.

→ Compute graph embedding $\mathbf{h}_G = f_{\text{down}}\{\sum_i \sigma[f_{\text{gate}}(\mathbf{h}_{T,i})]f_{\text{up}}(\mathbf{h}_{T,i})\}$.



Stopping probability and next atom in path

The stopping probability is defined as $p_\theta(\text{STOP}|\mathcal{M}_t) = \sigma(f_{\text{stop}}(\mathbf{h}_{\mathcal{G}}))$.

Stopping probability and next atom in path

The stopping probability is defined as $p_\theta(\text{STOP}|\mathcal{M}_t) = \sigma(f_{\text{stop}}(\mathbf{h}_G))$.

For the next atom, we use a **softmax layer** with inputs $s_i = \text{MLP}(\mathbf{h}_{T,i}, \text{CONTEXT})$ and corresponding binary masks $m_i \in \{0, 1\}$ to remove invalid choices:

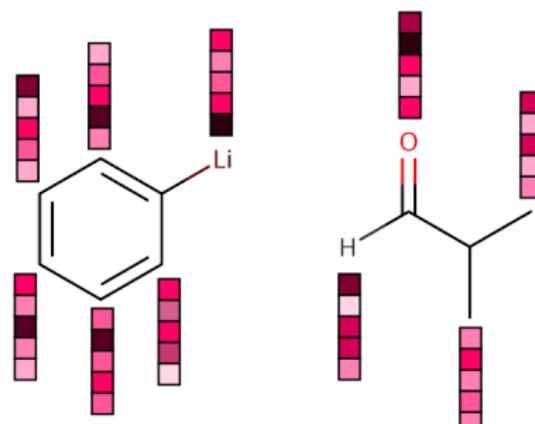
$$p_\theta(a_t = i | a_{t-1}, \mathcal{M}_t) = \frac{m_i \exp(s_i)}{\sum_j m_j \exp(s_j)}.$$

Stopping probability and next atom in path

The stopping probability is defined as $p_\theta(\text{STOP}|\mathcal{M}_t) = \sigma(f_{\text{stop}}(\mathbf{h}_G))$.

For the next atom, we use a **softmax layer** with inputs $s_i = \text{MLP}(\mathbf{h}_{T,i}, \text{CONTEXT})$ and corresponding binary masks $m_i \in \{0, 1\}$ to remove invalid choices:

$$p_\theta(a_t = i | a_{t-1}, \mathcal{M}_t) = \frac{m_i \exp(s_i)}{\sum_j m_j \exp(s_j)}.$$

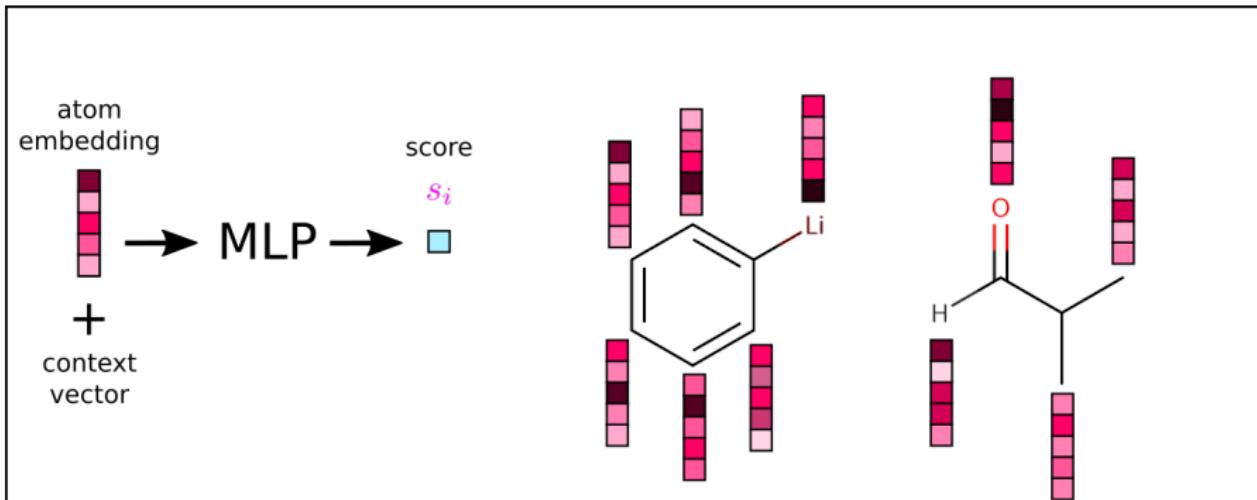


Stopping probability and next atom in path

The stopping probability is defined as $p_\theta(\text{STOP}|\mathcal{M}_t) = \sigma(f_{\text{stop}}(\mathbf{h}_G))$.

For the next atom, we use a **softmax layer** with inputs $s_i = \text{MLP}(\mathbf{h}_{T,i}, \text{CONTEXT})$ and corresponding binary masks $m_i \in \{0, 1\}$ to remove invalid choices:

$$p_\theta(a_t = i | a_{t-1}, \mathcal{M}_t) = \frac{m_i \exp(s_i)}{\sum_j m_j \exp(s_j)}.$$

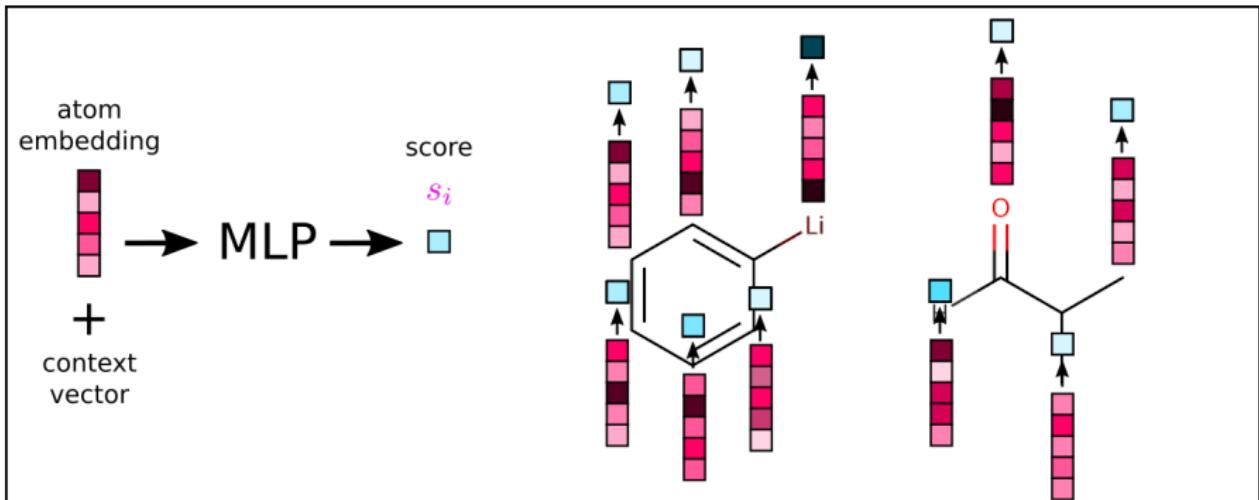


Stopping probability and next atom in path

The stopping probability is defined as $p_\theta(\text{STOP}|\mathcal{M}_t) = \sigma(f_{\text{stop}}(\mathbf{h}_G))$.

For the next atom, we use a **softmax layer** with inputs $s_i = \text{MLP}(\mathbf{h}_{T,i}, \text{CONTEXT})$ and corresponding binary masks $m_i \in \{0, 1\}$ to remove invalid choices:

$$p_\theta(a_t = i | a_{t-1}, \mathcal{M}_t) = \frac{m_i \exp(s_i)}{\sum_j m_j \exp(s_j)}.$$



Training data

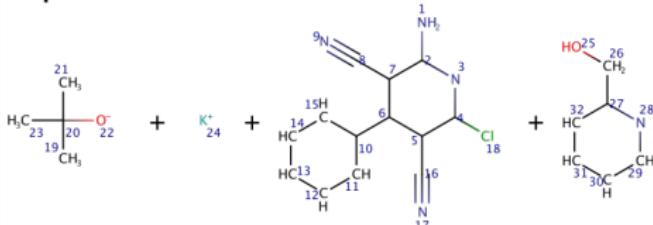
USPTO: 480K chemical reactions extracted from the US patent database.

Atoms have unique ID to easily match them before and after the reaction.

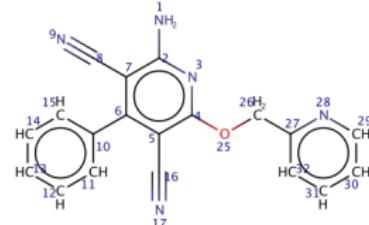
Approximate electron paths obtained for **73%** of USPTO by using simple rules.

```
[CH3:19][C:20]([CH3:21])([O:-22])[CH3:23].[K+:24].[NH2:1][c:2]1[n:3]
[c:4]([Cl:18])[c:5]([C:16]#[N:17])[c:6](-[c:10]2[cH:11][cH:12][cH:13]
[cH:14][cH:15]2)[c:7]1[C:8]#[N:9].[OH:25][CH2:26][c:27]1[n:28][cH:29]
[cH:30][cH:31][cH:32]1>>[NH2:1][c:2]1[n:3][c:4]([O:25][CH2:26][c:
27]2[n:28][cH:29][cH:30][cH:31][cH:32]2)[c:5]([C:16]#[N:17])[c:6](-[c:
10]2[cH:11][cH:12][cH:13][cH:14]2)[c:7]1[C:8]#[N:9]
```

Input:



Target:



Results

Beam search to find top-K predictions for the electron path.

Multiple electron paths can result in same product (harder prediction task).

LITE variant ignores reagent information.

Electron path prediction

Model Name	Accuracies (%)			
	Top-1	Top-2	Top-3	Top-5
ELECTRO-LITE	70.3	82.8	87.7	92.2
ELECTRO	77.8	89.2	92.4	94.7

Product prediction

Model Name	Accuracies (%)			
	Top-1	Top-2	Top-3	Top-5
ELECTRO-LITE	78.2	87.7	91.5	94.4
ELECTRO	87.0	92.6	94.5	95.9

WLDN	84.0	89.2	91.1	92.3
Seq2Seq	80.3*	84.7*	86.2*	87.5*

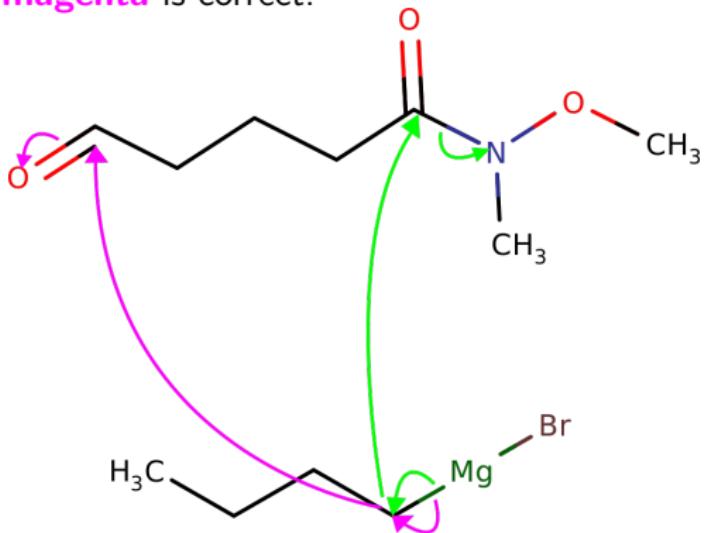
Bonus: **interpretable explanation** for product prediction task.

Qualitative analysis

Easy to identify what went wrong!

First choice in **green** is incorrect, but “chemically reasonable”.

Second choice in **magenta** is correct.



Take home messages

ELECTRO...

- ① predicts electron paths with **high accuracy**.
- ② produces an output that is very easy to **interpret** by chemists.
- ③ exploits that only a **small number of atoms** interact during the reaction.
- ④ also predicts **final products** with high accuracy.

Collaborators



R. P. Adams



J. Aguilera-Ipaguirre



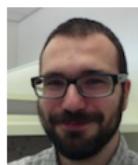
A. Aspuru-Guzik



John Bradshaw



D. Duvenaud



R. Gomez-Bonbardelli



T. Hirzel



M. Kusner



B. Paige



B. Sánchez-Lengeling



Marwin H. S.
Segler



D. Sheberla



J. Wey

Thanks!

Take home messages

Generative models of molecules...

- ① Data-driven approach to molecule generation, no expertise needed.
- ② Sampled molecules are realistic, unlike those generated with rules.
- ③ Create a continuous latent space which is useful for optimization.

Grammar variational autoencoder...

- ① Produces a larger fraction of valid molecules when decoding.
- ② Produces better predictions of molecule properties from latent space.
- ③ Molecule optimization results are improved.

ELECTRO...

- ① predicts electron paths with high accuracy.
- ② produces an output that is very easy to interpret by chemists.
- ③ exploits that only a small number of atoms interact during the reaction.
- ④ also predicts final products with high accuracy.