

Planting the Seeds of Probabilistic Thinking

Foundations | Tricks | Algorithms

Shakir Mohamed

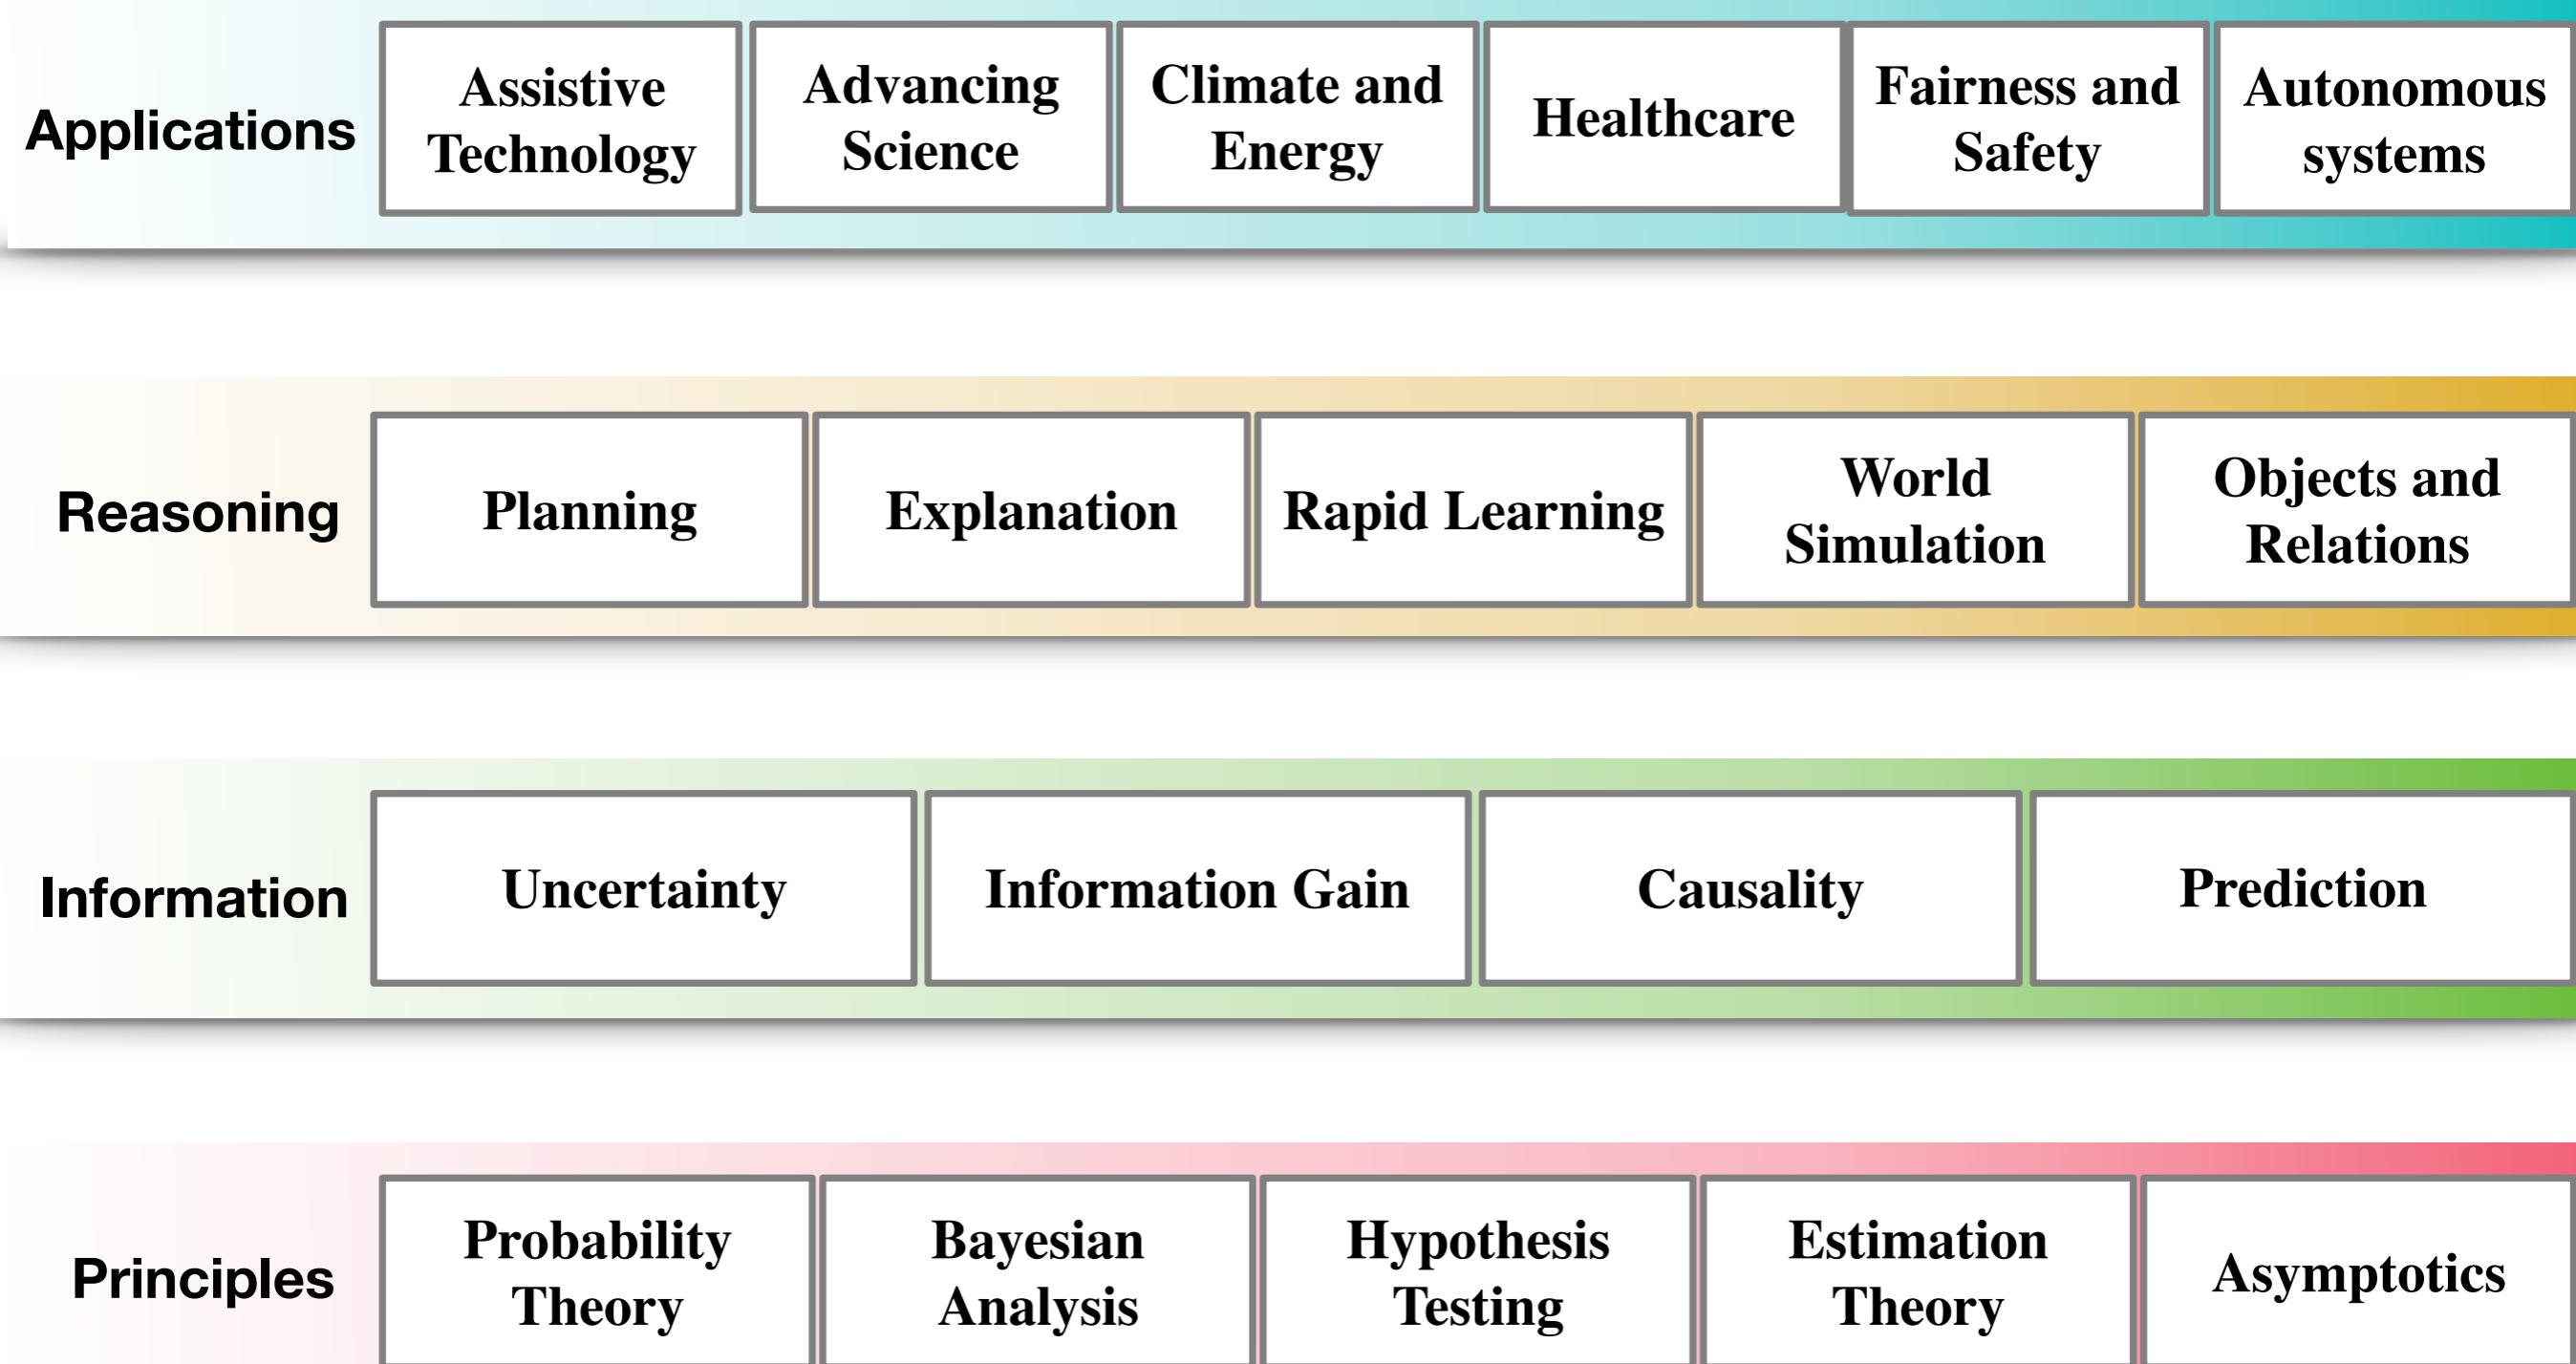
Research Scientist, DeepMind

Planting the Seeds of Probabilistic Thinking: Foundations, Tricks and Algorithms

Probabilistic machine learning approaches task of describing of data, to complex systems or our world using the language and tools of probability. Almost all of machine learning can be viewed in probabilistic terms, making probabilistic thinking fundamental. It is, of course, not the only view. But it is through this view that we can connect what we do in machine learning to every other computational science, whether that be in stochastic optimisation, control theory, operations research, econometrics, information theory, statistical physics or bio-statistics. For this reason alone, mastery of probabilistic thinking is essential.

The aim of this tutorial is to develop flexible and broad tools that will support your probabilistic thinking. Part 1, Foundations looks at the philosophy of machine learning, builds an understanding of the model-inference-algorithm paradigm, and explores fundamental areas of machine learning - we'll look at deep learning, kernels and reinforcement learning. Part 2 Tricks, will look at 6 individual probabilistic problems and a tricks to solve them, using these tricks to develop flexibility in our thinking. Part 3: Algorithms will look at how the foundations and tricks combine to develop machine learning algorithms, with a specific focus on the area of deep generative models.

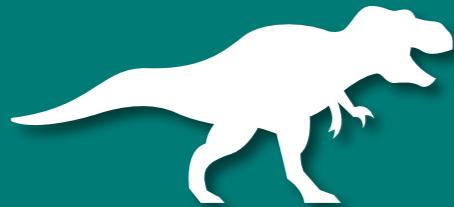
Principles to Products



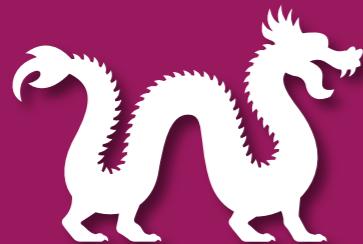
Planting the Seeds of Probabilistic Thinking

Part I: Foundations

Learning Objectives



1. Language to think about the
Philosophy of Machine Learning



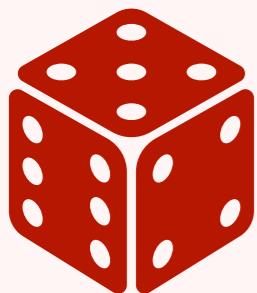
2. Understand the
Model-Inference-Algorithm paradigm



3. Use probabilistic thinking applied to
problems in supervised, unsupervised,
and reinforcement learning.

Probability

Some Definitions for probability



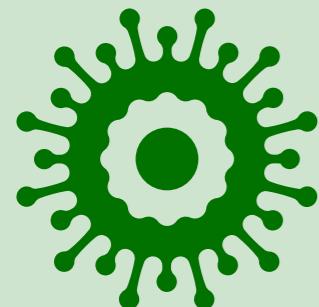
Statistical Probability

Frequency ratio of items



Logical Probability

Degree of confirmation of a hypothesis based on logical analysis



Probability as Propensity

Probability used for predictions



Subjective Probability

Probability as a degree of belief

Probability is sufficient for the task of reasoning under uncertainty

Probability

Probability as a Degree of Belief



Probability is a measure of the belief in a proposition **given** evidence.
A description of a state of knowledge.

No such thing as **the probability** of an event, since the value depends on the evidence used.

Inherently subjective in that it depends on the believer's information

Different observers with different information will have different beliefs.

Probabilistic Quantities

Probability

$$p(\mathbf{x}) \quad p^*(\mathbf{x}) \quad q(\mathbf{x})$$

Conditions

$$p(\mathbf{x}) \geq 0 \quad \int p(\mathbf{x}) d\mathbf{x} = 1$$

Bayes Rule

$$p(\mathbf{z}|\mathbf{x}) = \frac{p(\mathbf{x}|\mathbf{z})p(\mathbf{z})}{p(\mathbf{x})}$$

Parameterisation

$$p_{\theta}(\mathbf{x}|\mathbf{z}) \equiv p(\mathbf{x}|\mathbf{z}; \boldsymbol{\theta})$$

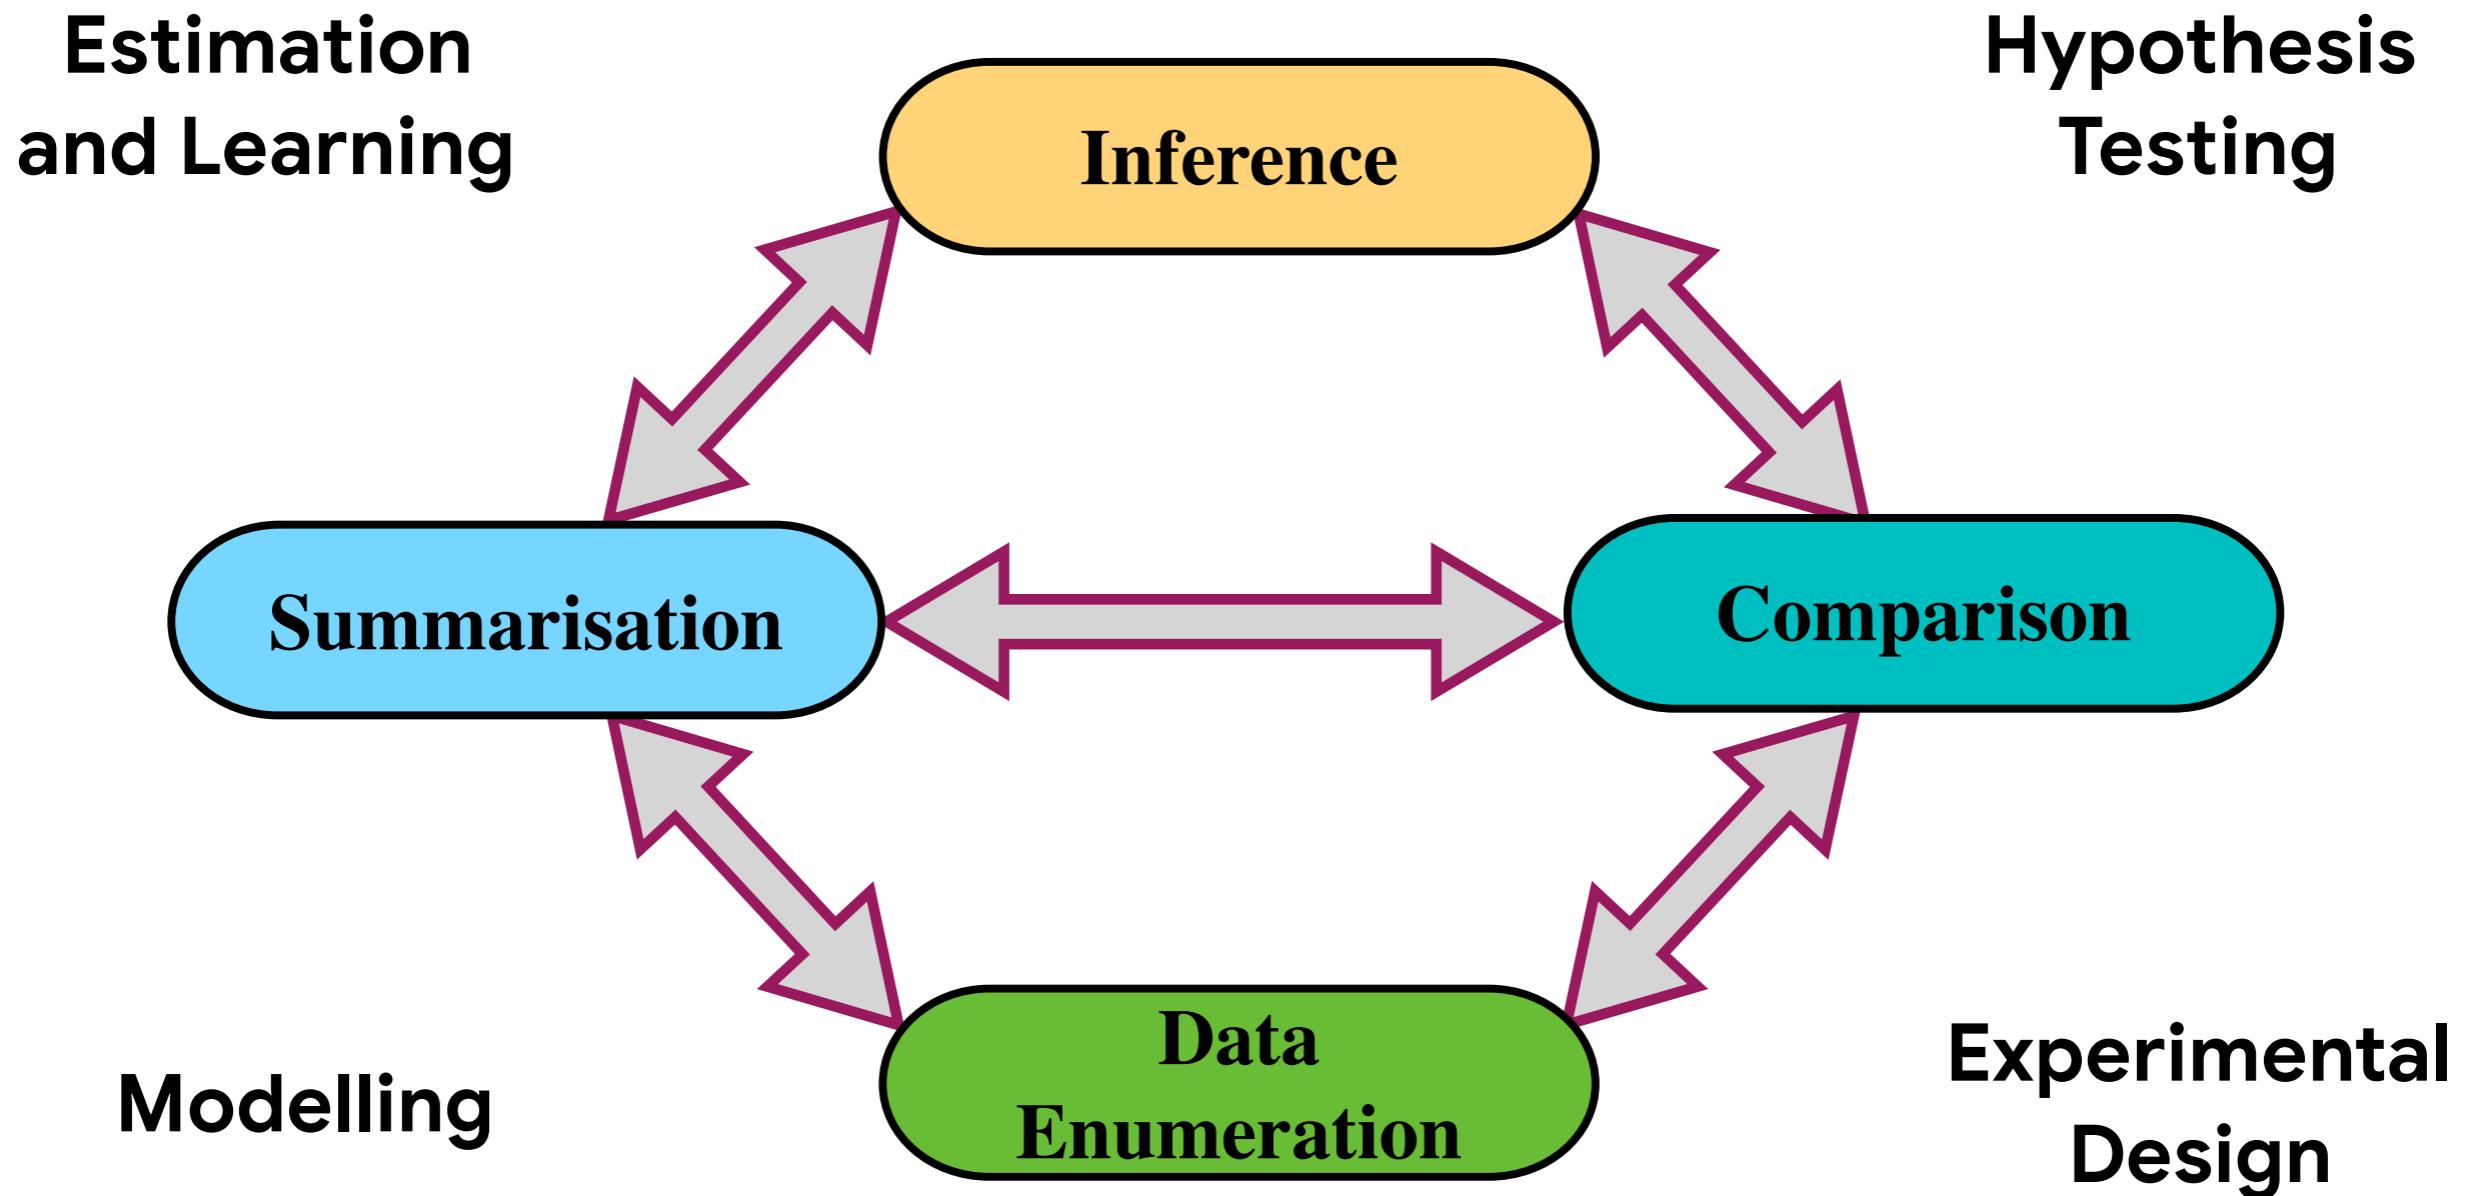
Expectation

$$\mathbb{E}_{p_{\theta}(\mathbf{x}|\mathbf{z})}[f(\mathbf{x}; \boldsymbol{\phi})] = \int p_{\theta}(\mathbf{x}|\mathbf{z})f(\mathbf{x}; \boldsymbol{\phi})d\mathbf{x}$$

Gradient

$$\nabla_{\boldsymbol{\phi}} f(\mathbf{x}; \boldsymbol{\phi}) = \frac{\partial f(\mathbf{x}; \boldsymbol{\phi})}{\partial \boldsymbol{\phi}}$$

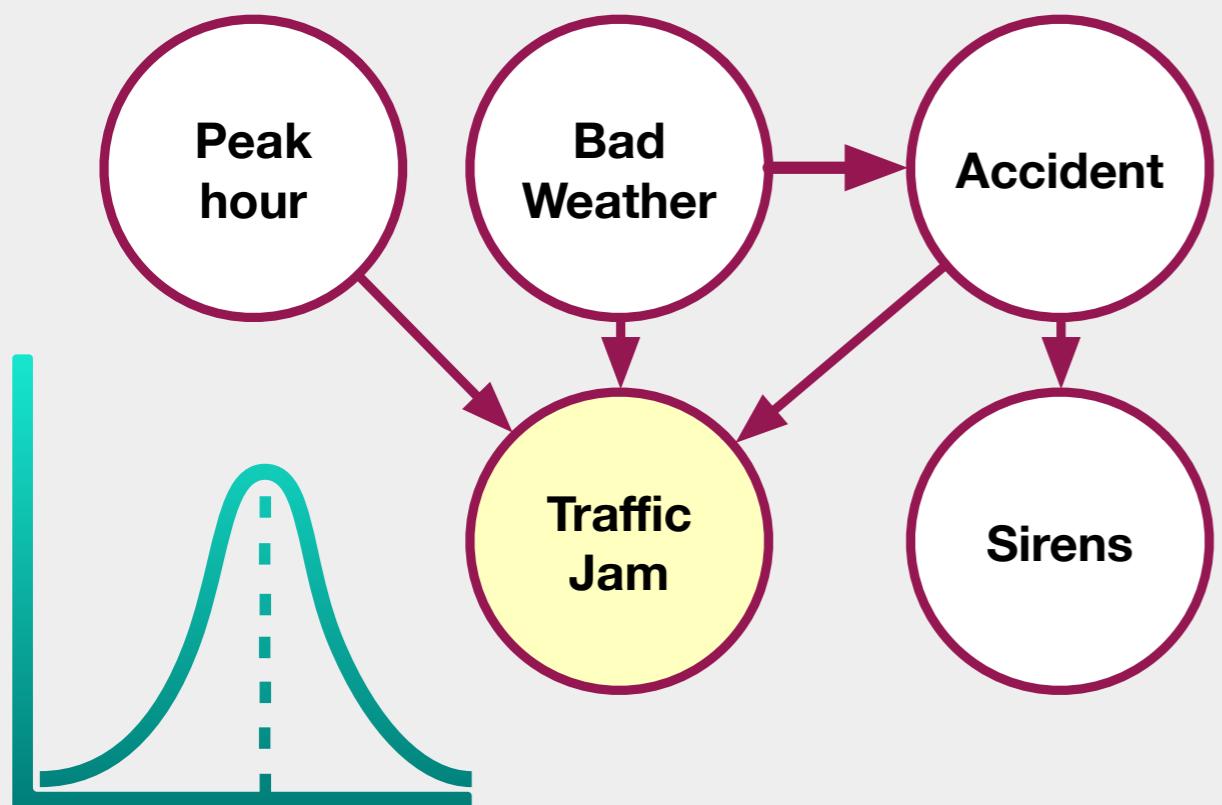
Statistical Operations



Probabilistic Models

Model: Description of the world, of data, of potential scenarios, of processes.

A probabilistic model writes out these models using the language of probability



prob(traffic Jam)

prob(sirens | Accident)

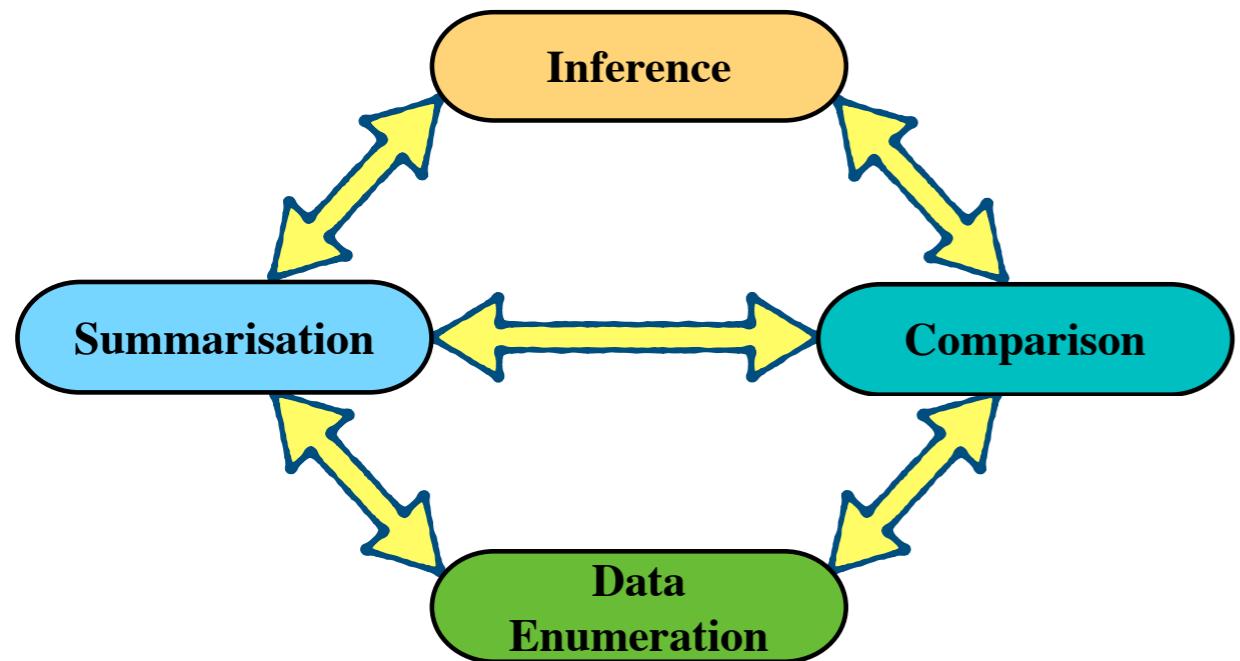
prob(peak hour | Traffic Jam)

Most models in machine learning are probabilistic.

Probabilistic models let you learn probability distributions of data.

You can choose what to learn: Just the mean. Or the entire distribution.

Centrality of Inference



**Artificial General Intelligence
will be the refined instantiation
of these statistical operations.**

**The core questions of
AGI will be those of
probabilistic inference**



Linear Regression

Generalised Linear Regression

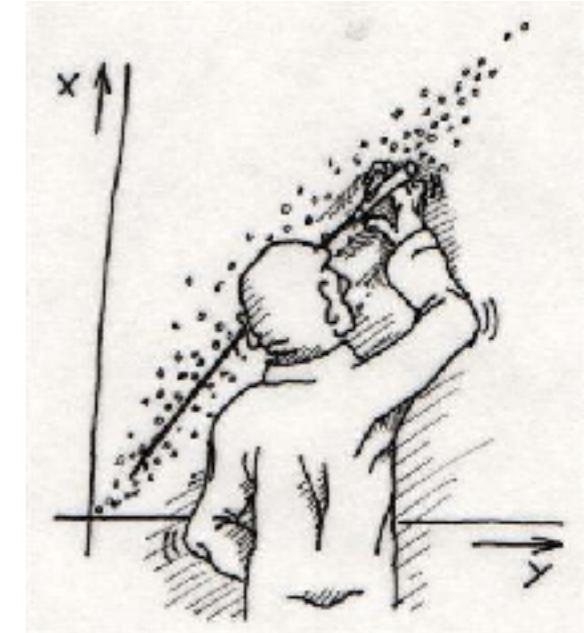
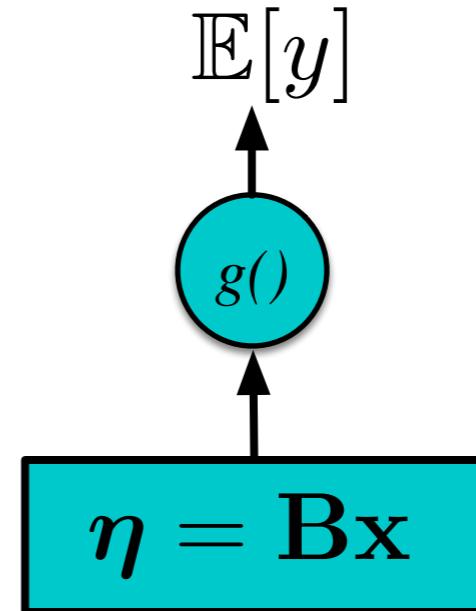
$$\eta = \mathbf{w}^\top \mathbf{x} + b$$

$$p(y|\mathbf{x}) = p(y|g(\eta); \theta)$$

- The basic function can be any linear function, e.g., affine, convolution.
- $g(\cdot)$ is an **inverse link function** that we'll refer to as an activation function.

Optimise the negative log-likelihood

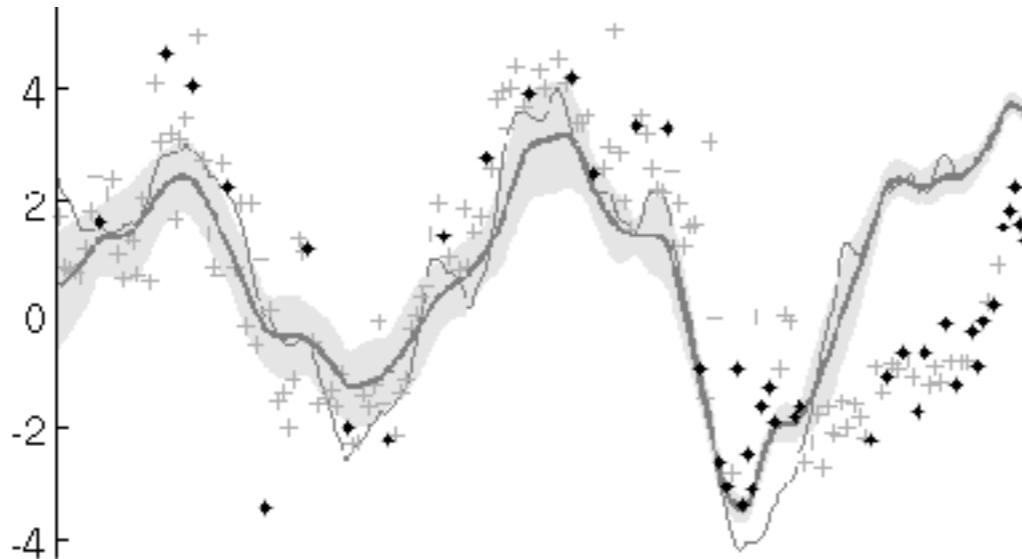
$$\mathcal{L} = -\log p(y|g(\eta); \theta)$$



Target	Regression	Link	Inv link	Activation
Real	Linear	Identity	Identity	
Binary	Logistic	Logit $\log \frac{\mu}{1-\mu}$	Sigmoid $\frac{1}{1+\exp(-\eta)}$	Sigmoid
Binary	Probit	Inv Gauss CDF $\Phi^{-1}(\mu)$	Gauss CDF $\Phi(\eta)$	Probit
Binary	Gumbel	Compl. log-log $\log(-\log(\mu))$	Gumbel CDF $e^{-e^{-x}}$	
Binary	Logistic			Hyperbolic Tangent $\tanh(\eta)$
Categorical	Multinomial		Multin. Logit	Softmax
Counts	Poisson	$\log(\mu)$	$\sum_j \frac{n_j}{n_i}$ $\exp(v)$	
Counts	Poisson	$\sqrt{(\mu)}$	v^2	
Non-neg.	Gamma	Reciprocal	$\frac{1}{v}$	
Sparse	Tobit		max $\max(0; v)$	ReLU
Ordered	Ordinal		Cum. Logit $\sigma(\phi_k - \eta)$	

Deep Networks

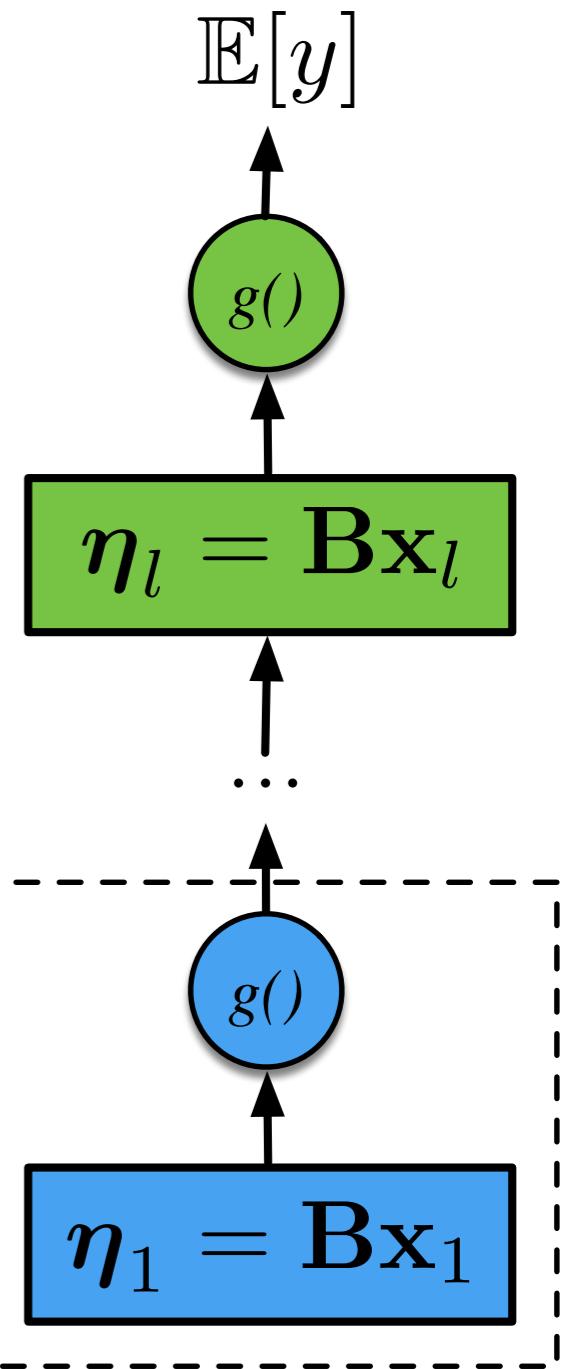
Recursive Generalised Linear Regression



- Recursively compose the basic linear functions.
- Gives a deep neural network.

$$\mathbb{E}[y] = h_L \circ \dots \circ h_l \circ h_0(\mathbf{x})$$

A general, flexible framework for building
non-linear, parametric models



Likelihood

Probabilistic Model

$$p(y|\mathbf{x}) = p(y|h(\mathbf{x}); \boldsymbol{\theta})$$

Likelihood function

$$\mathcal{L}(\boldsymbol{\theta}) = \sum_n \log p(y_n|\mathbf{x}_n; \boldsymbol{\theta})$$

Likelihood of parameters

Efficient Estimators

- Statistically efficient (Cramer-Rao lower bound)
- Asymptotically unbiased, consistent
- Maximum entropy (principle of indifference)

Tests with Good Power

- Likelihood ratio tests
- Can construct small confidence regions

Widely-applicable

- Handle data that is incompletely observed, distorted, samples with bias
- Can offset or correct these issues.

Pool Information

- Combine different data sources
- Knowledge outside the data can be used, like constraints on domain or prior probabilities.

Misspecification: Inefficient estimates; or confidence intervals/tests can fail completely.

Estimation Theory

Probabilistic Model

$$p(y|\mathbf{x}) = p(y|h(\mathbf{x}); \boldsymbol{\theta})$$

Likelihood function

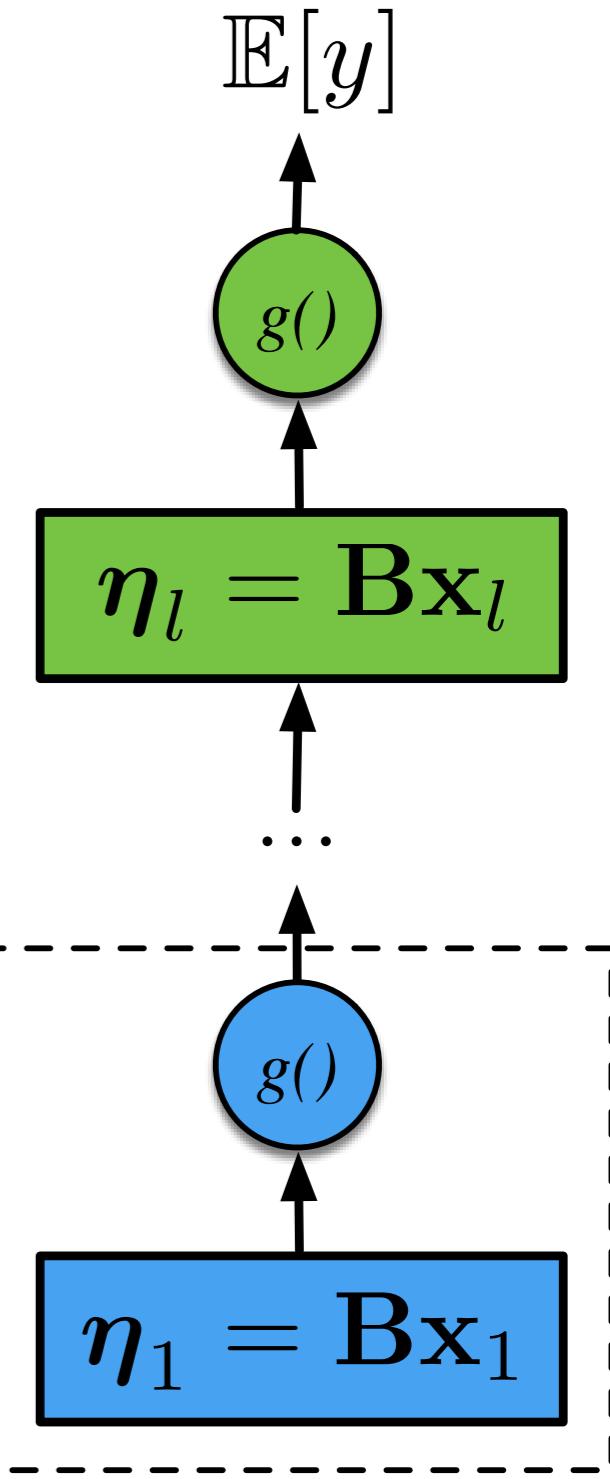
$$\mathcal{L}(\boldsymbol{\theta}) = \sum_n \log p(y_n | \mathbf{x}_n; \boldsymbol{\theta})$$

Optimisation Objective

$$\arg \max_{\boldsymbol{\theta}} \mathcal{L}(\boldsymbol{\theta})$$

- Straightforward and natural way to learn parameters
- Can be biased in finite sample size, e.g., Gaussian variances with N and N-1.
- Easy to observe **overfitting** of parameters.

Maximum Likelihood



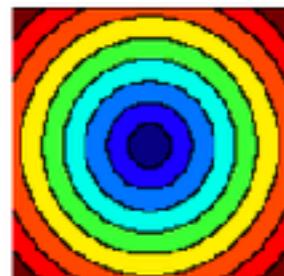
Estimation Theory

Probabilistic Model

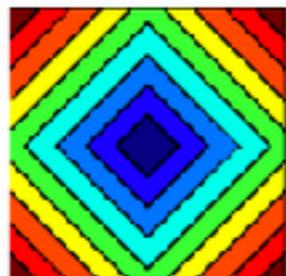
$$p(\boldsymbol{\theta}|y, \mathbf{x}) \propto p(y|h(\mathbf{x}); \boldsymbol{\theta})p(\boldsymbol{\theta})$$

Likelihood function

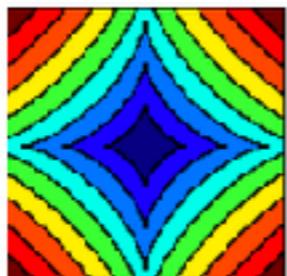
$$\mathcal{L}(\boldsymbol{\theta}) = \sum_n \log p(y_n|\mathbf{x}_n; \boldsymbol{\theta}) + \frac{1}{\lambda} \mathcal{R}(\boldsymbol{\theta})$$



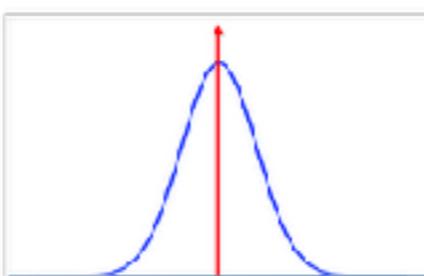
Gaussian (L2)



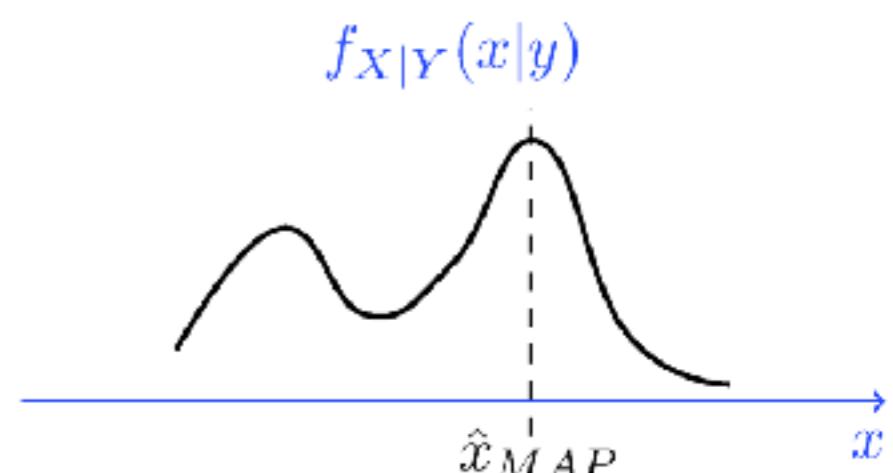
Laplace (L1)



L_p -norm



Spike and Slab



Maximum a Posteriori (MAP)

Optimisation Objective

$$\arg \max_{\boldsymbol{\theta}} \mathcal{L}(\boldsymbol{\theta})$$

- Generalises the MLE (uniform prior)
- **Shrinkage**: shrink parameters back to initial beliefs.
- Not every regulariser corresponds a valid probability distribution.

Regularisation

$$\mathcal{L}(\boldsymbol{\theta}) = \sum_n \log p(y_n | \mathbf{x}_n; \boldsymbol{\theta}) + \frac{1}{\lambda} \mathcal{R}(\boldsymbol{\theta})$$

- ◆ **Regularisation** is essential to overcome the limitations of maximum likelihood estimation.
- ◆ **Other names:** Regularisation, penalised regression, shrinkage.

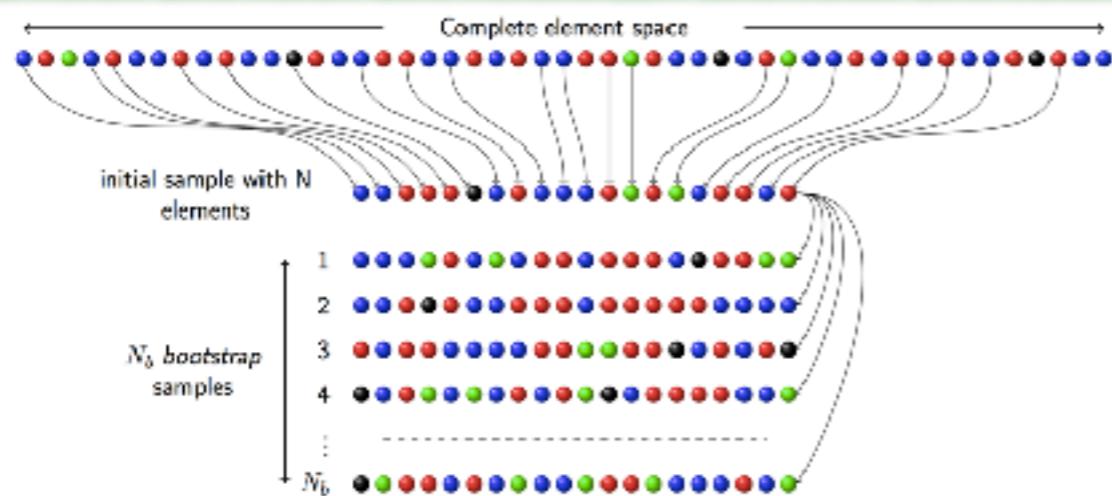
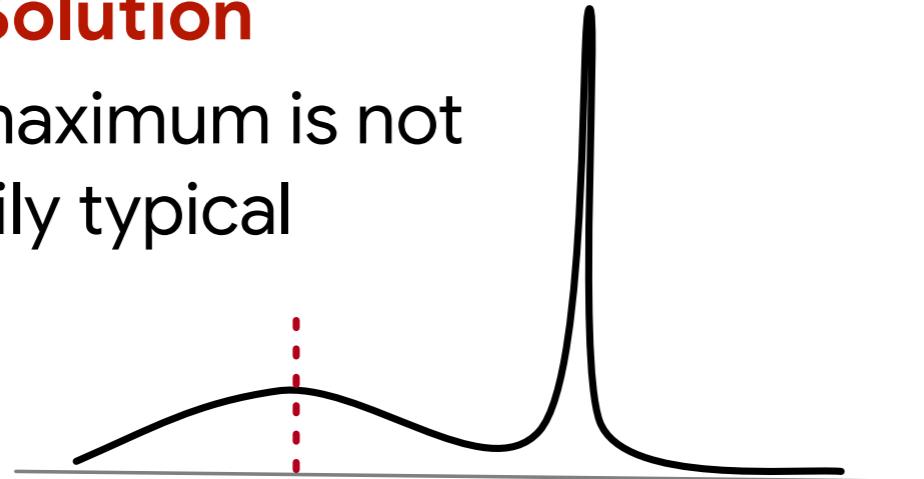
A wide range of available regularisation techniques:

- Large data sets
- Input noise/jittering and data augmentation/expansion.
- L2 /L1 regularisation (Weight decay, Gaussian prior)
- Binary or Gaussian Dropout
- Batch normalisation

MAP Estimation

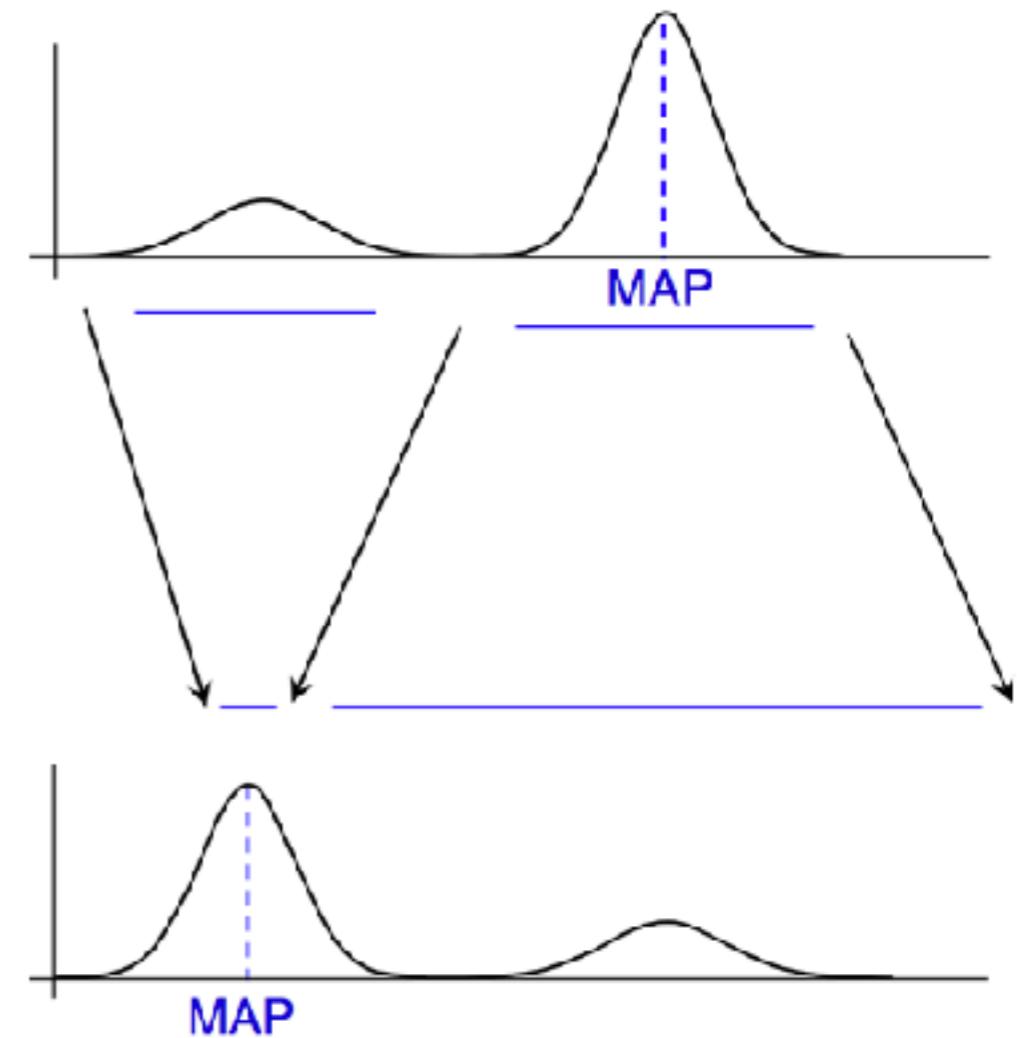
Type of Solution

What is maximum is not necessarily typical



Uncertainty

Can be reported using confidence intervals or bootstrap estimates.



Parameterisation sensitive

Location of max will change depending on parameterisation

Invariant MAP

Change of variables

$$p(\phi) = p(\mu) \left| \frac{d\mu}{d\phi} \right|$$

Popular Example

$$y \in \{0, 1\}; \quad 0 \leq \mu \leq 1$$

Bernoulli

$$p(y=1|\mu) = \mu$$

Uniform

$$p(\mu) = 1$$

Mode of the prior

$$\hat{\phi}_{MAP} = \arg \max_{\phi \in [0,1]} p(\phi)$$

Transform

$$\mu = \phi^2$$

New prior

$$p(\phi) = 2\phi$$

MAP Est.

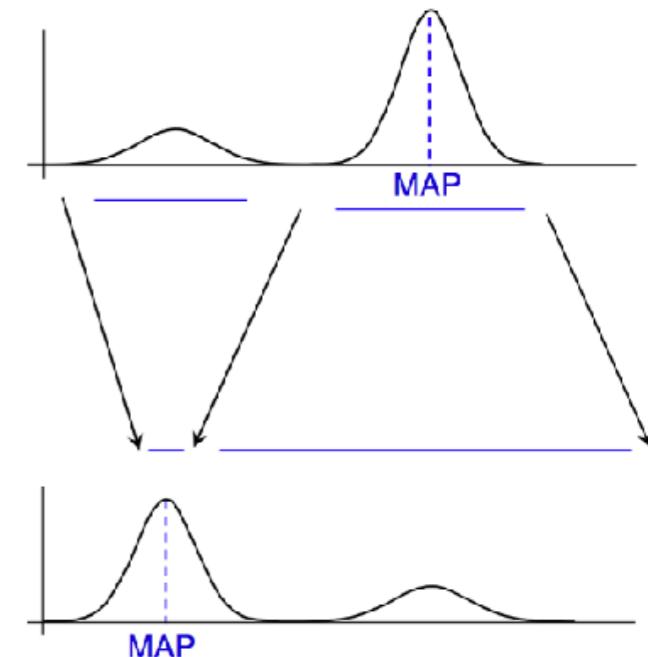
$$\hat{\phi}_{MAP} = 1$$

Parameterisation 2

$$\mu = 1 - (1 - \phi)^2$$

$$p(\phi) = 2(1 - \phi)$$

$$\hat{\phi}_{MAP} = 0$$



Clear sensitivity: Sensitive to units, affects interpretability, affects gradients, learning stability, design of models.

Invariant MAP

Use a modified probabilistic model that removes sensitivity

Invariant MAP

$$p(y|h(\mathbf{x}); \boldsymbol{\theta})p(\boldsymbol{\theta})|\mathcal{I}(\boldsymbol{\theta})|^{\frac{1}{2}}$$

- Use the Fisher information
- Connection to the natural gradients and trust-region optimisation.
- Uninformative priors.

Proposed solutions have not fully dealt with the underlying issues.



Bayesian Analysis

Issues arise as a consequence of:

- Reasoning only about the most likely solution, and
- Not maintaining knowledge of the underlying variability (and averaging over this).

**Motivates learning more than the mean.
This is the core of a Bayesian philosophy.**

$$p(\boldsymbol{\theta}|y, \mathbf{x}) \propto p(y|h(\mathbf{x}); \boldsymbol{\theta})p(\boldsymbol{\theta})$$

Pragmatic Bayesian Approach for
Probabilistic Reasoning in Deep Networks.
(and all of machine learning)

Bayesian reasoning over some, but not all parts of our models (yet).

Bayesian Analysis

Interested in reasoning about two important quantities

Evidence

$$p(y|\mathbf{x}) = \int p(y|h(\mathbf{x}); \boldsymbol{\theta})p(\boldsymbol{\theta})d\boldsymbol{\theta}$$

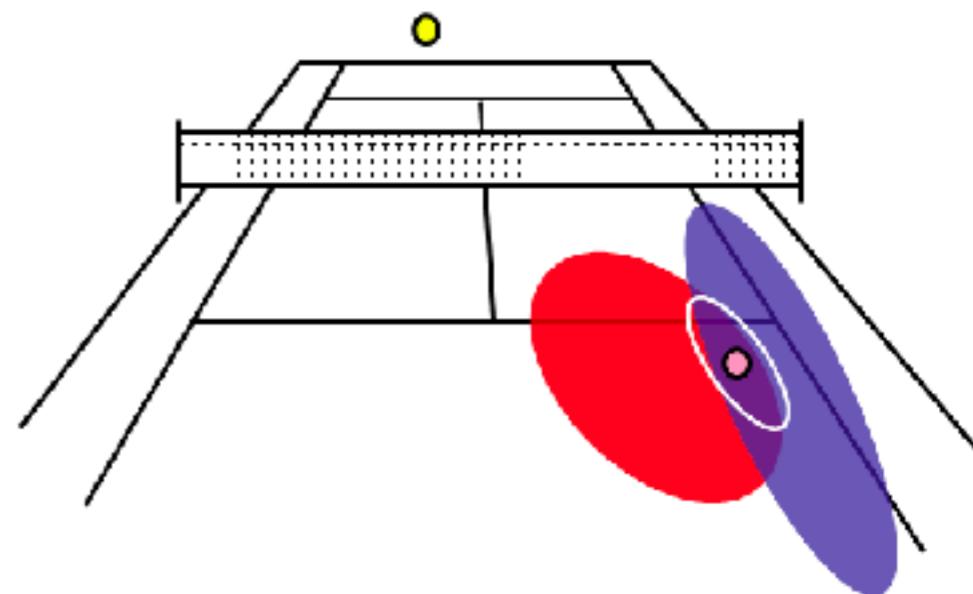
Posterior

$$p(\boldsymbol{\theta}|y, \mathbf{x}) \propto p(y|h(\mathbf{x}); \boldsymbol{\theta})p(\boldsymbol{\theta})$$

- In Bayesian analysis, things that are *not* observed must be integrated over - averaged out.
- This makes computation difficult.
- Integration is the central operation.

Intractable Integrals: Will often see this phrasing.

- Don't know the integral in closed form
- Very high-dimensional quantities and can't compute (e.g., using quadrature)



Learning and Inference

Statistics, no distinction between learning and inference - only inference (or **estimation**).

Machine learning makes a distinction between **inference and learning**:

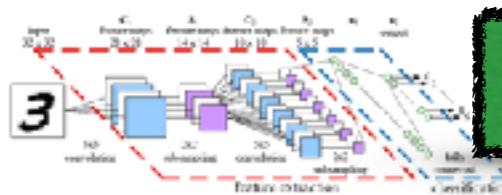
- **Inference**: reason about (and compute) unknown probability distributions.
- **(Parameter) Learning** is finding point estimates of quantities in the model.

Bayesian statistics, all quantities are probability distributions, so there is only the problem of **inference**.

Software engineering, **inference** is the forward evaluation of a trained model (to get predictions).

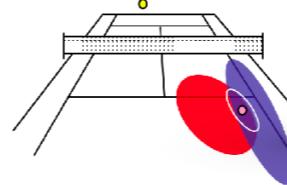
Decision making and AI, refer to **learning** in general as the means of understanding and acting based on past experience (data).

Two Streams of ML



Deep Learning

- + Rich non-linear models for classification and sequence prediction.
- + Scalable learning using stochastic approximation and conceptually simple.
- + Easily composable with other gradient-based methods
- Only point estimates
- Hard to score models, do selection and complexity penalisation.



Bayesian Reasoning

- Mainly conjugate and linear models
- Potentially intractable inference, computationally expensive or long simulation time.
- + Unified framework for model building, inference, prediction and decision making
- + Explicit accounting for uncertainty and variability of outcomes
- + Robust to overfitting; tools for model selection and composition.

Natural to consider the marriage of these approaches: Bayesian Deep Learning

Bayesian Regression

Probabilistic models over functions

Prior

$$p(\boldsymbol{\theta}) = \mathcal{N}(\boldsymbol{\theta} | 0, \mathbf{I})$$

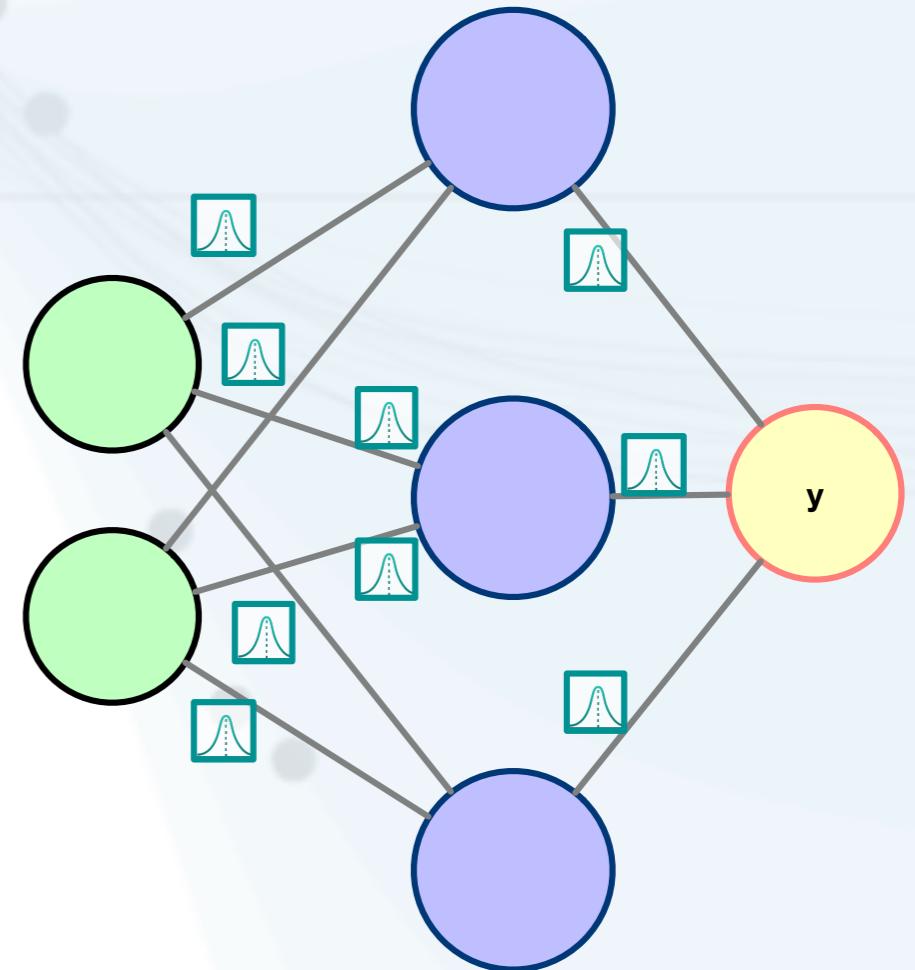
Observation model

$$p(y|\mathbf{x}, \boldsymbol{\theta}) = \text{Categorical}(\pi(\mathbf{x}; \boldsymbol{\theta}))$$

Posterior

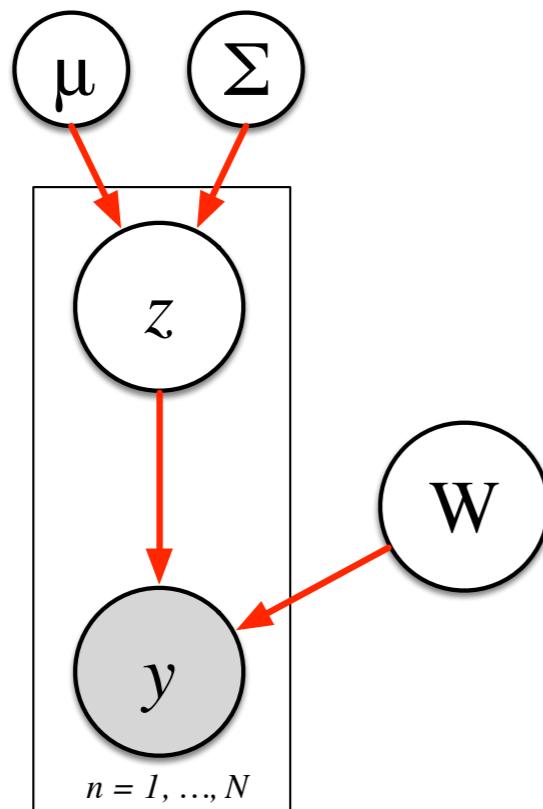
$$p(\boldsymbol{\theta}|y, \mathbf{x})$$

- Ways of learning distributions over functions and maintaining uncertainty over functions.
- Difficult in parametric models (like deep networks) because of high-dimensional parameter space.
- Many ways to learn the posterior distribution. Focus of Part III



Density Estimation

Learn probability distributions over the data itself

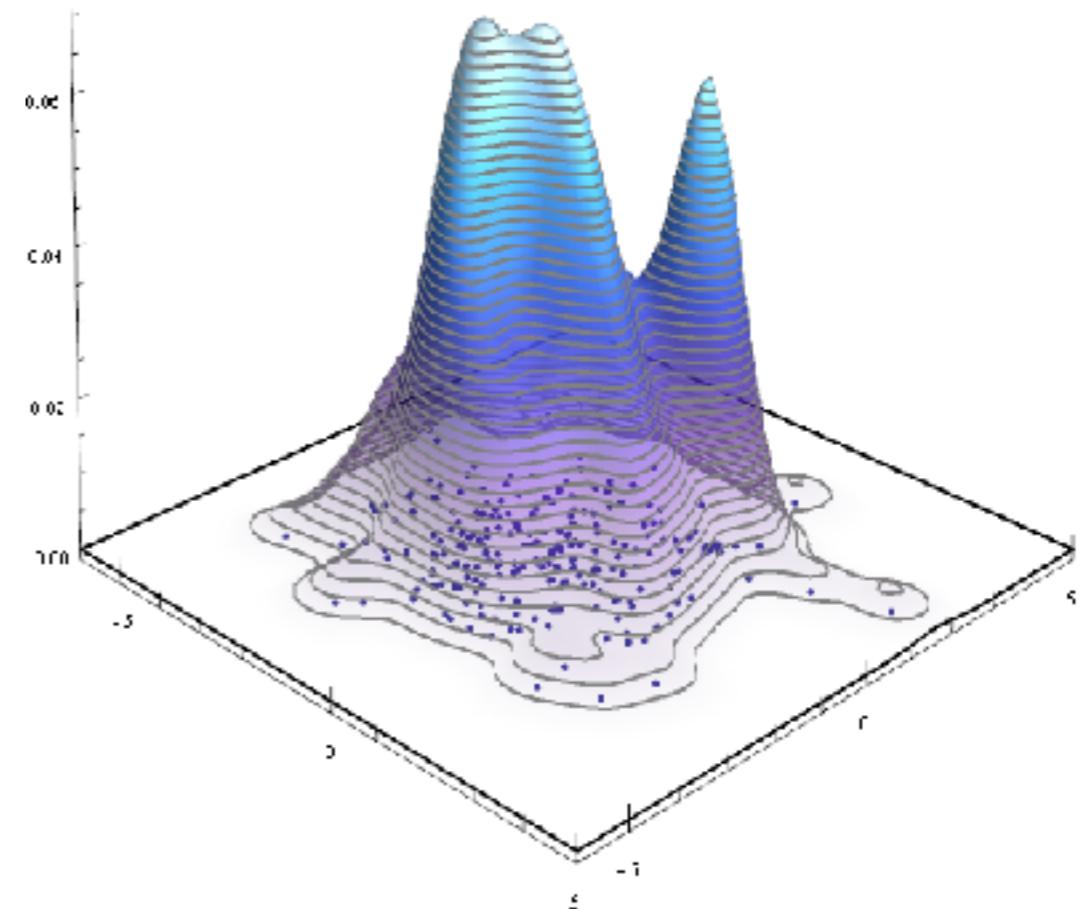


- Can learn distributions of some things and point estimates of others.
- Deep Generative Models and Unsupervised learning - more in Part III

Factor Analysis / PCA

$$z \sim \mathcal{N}(z|\mu, \Sigma)$$

$$y \sim \mathcal{N}(y|Wz, \sigma_y^2 I)$$



Decision-making

Probabilistic models of environments and actions

Prior over actions

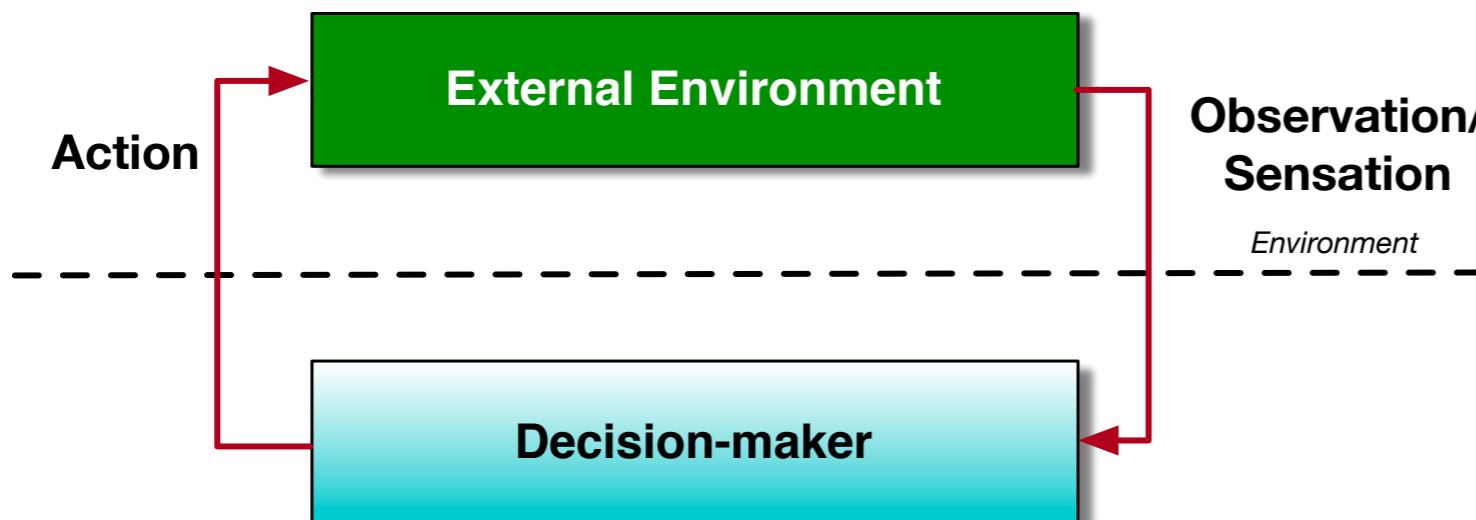
$$a \sim p(a)$$

Interaction only

$$u(s, a) \sim \text{Environment}(a)$$

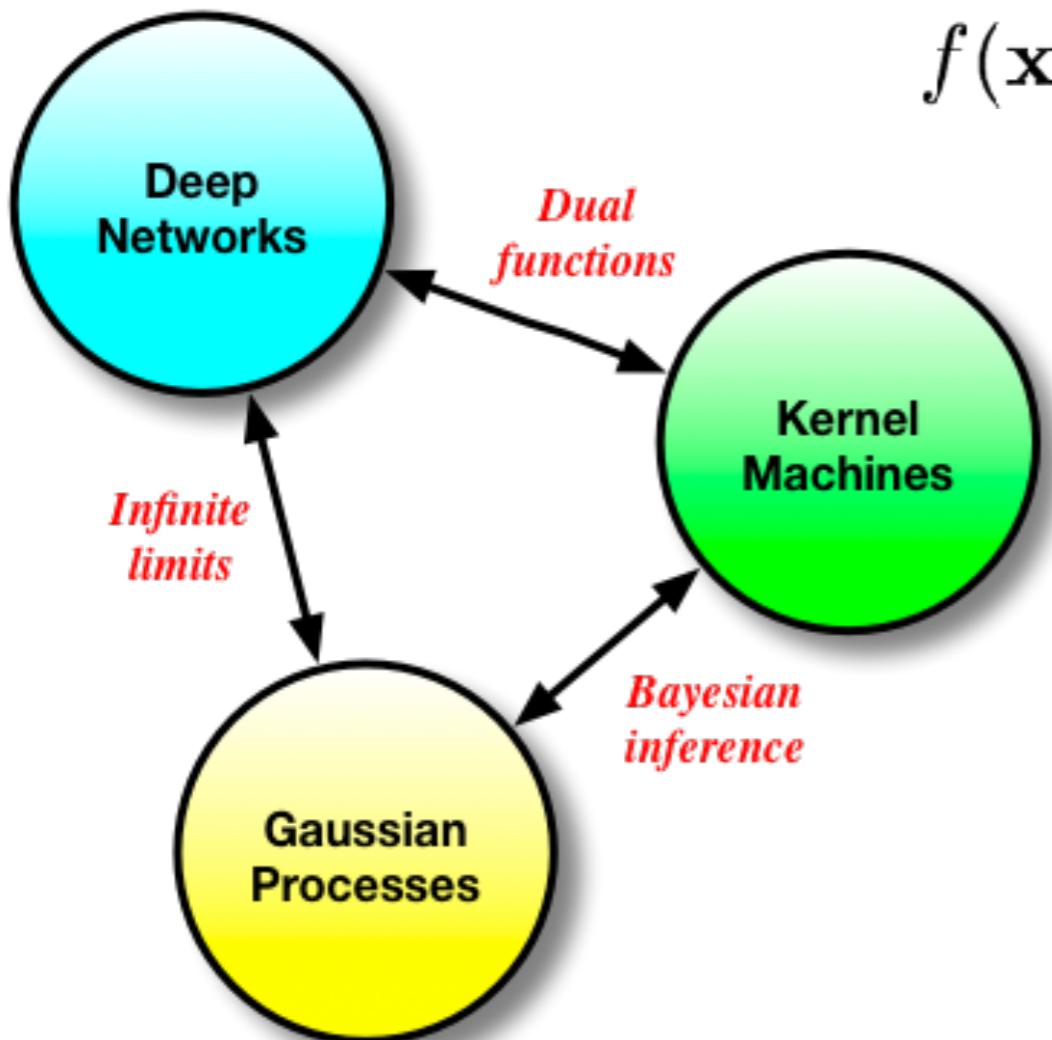
Reward/Utility

$$p(R(s)|a) \propto \exp(u(s, a))$$



Setup is common in experimental design, causal learning, reinforcement learning.

Probabilistic Dualities



Basis Function Regression

$$f(\mathbf{x}) = \mathbf{w}^\top \phi(\mathbf{x}; \boldsymbol{\theta}); \quad \{\mathbf{w}, \boldsymbol{\theta}\} \sim \mathcal{N}(0, \sigma^2 \mathbf{I})$$

$$y = f(\mathbf{x}) + \epsilon; \quad \epsilon \sim \mathcal{N}(0, \sigma_y^2)$$

Move from primal variables to dual variables

$$\mathcal{L}(f) = \frac{1}{2} \sum_n (y_n - f(\mathbf{x}))^2 + \frac{\lambda}{2} \|f\|_{\mathcal{H}}^2$$

Kernel trick and methods

Probability distributions over functions

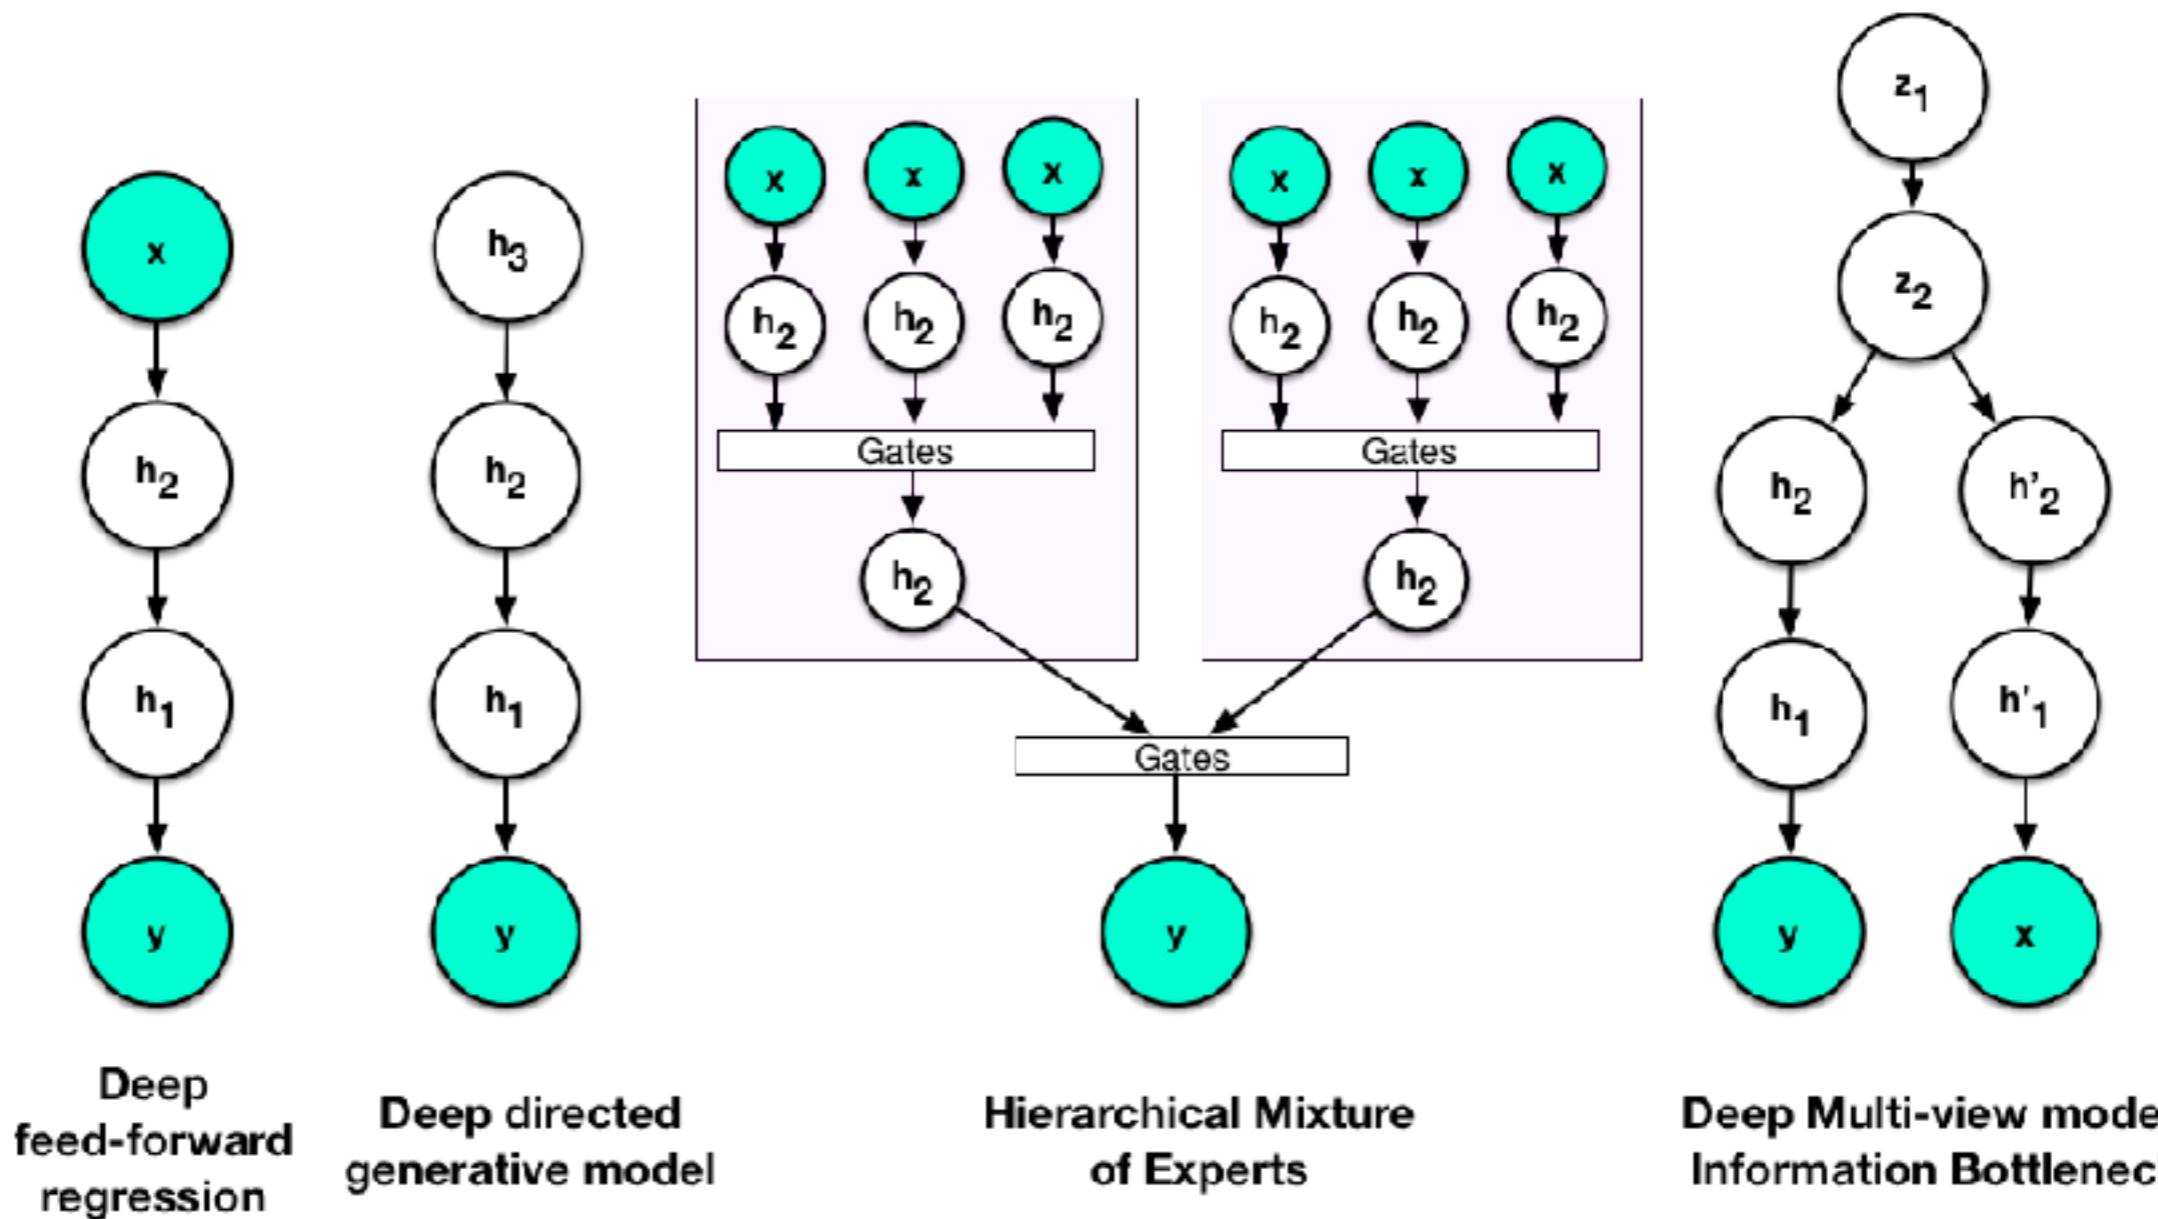
$$p(f) = \mathcal{N}(0, \mathbf{K}) \quad p(y|f) = \mathcal{N}(f, \sigma^2)$$

Gaussian processes

Deep and Hierarchical

Hierarchical Model: models where the (prior) probability distributions can be decomposed into a sequence of conditional distributions

$$p(z) = p(z_1|z_2)p(z_2|z_3)\dots p(z_{L-1}|z_L)p(z_L)$$



Deep
feed-forward
regression

Deep directed
generative model

Hierarchical Mixture
of Experts

Deep Multi-view model/
Information Bottleneck

Foundations

How will you approach your ML research and practice?

In general:

Human-centred,
interdisciplinary approach



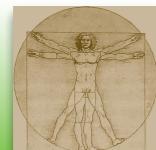
Sociological



Psychological



Componential

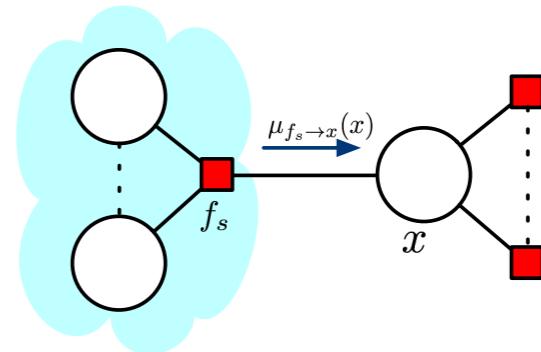


Physiological

**Sun's Phenomenological
Levels**

For the ML Core:

Probabilistic and pragmatic in approach

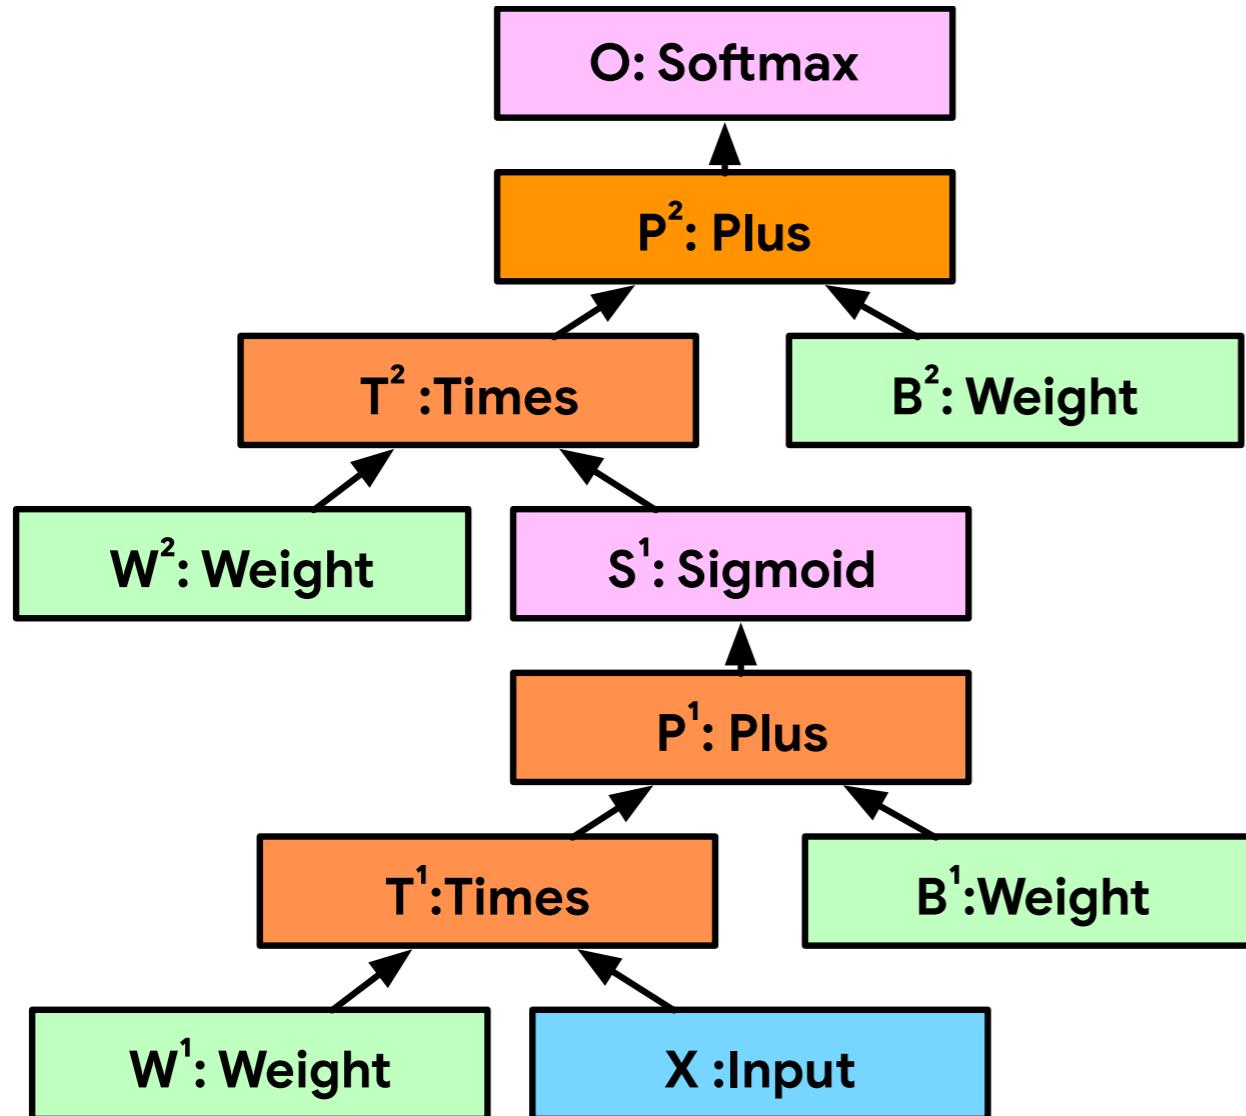


Architecture-Loss

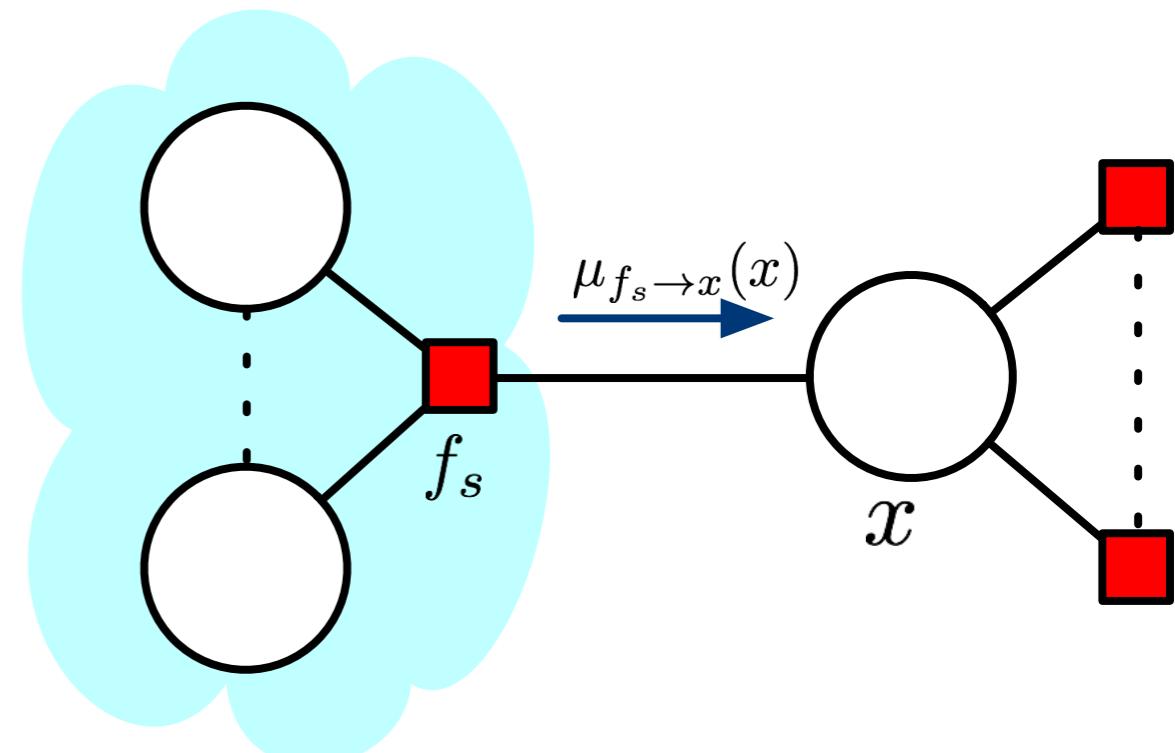


Model-Inference-Algorithm

Architecture-Loss



1. Computational Graphs



2. Error propagation

Model-Inference-Algorithm



3. Algorithms

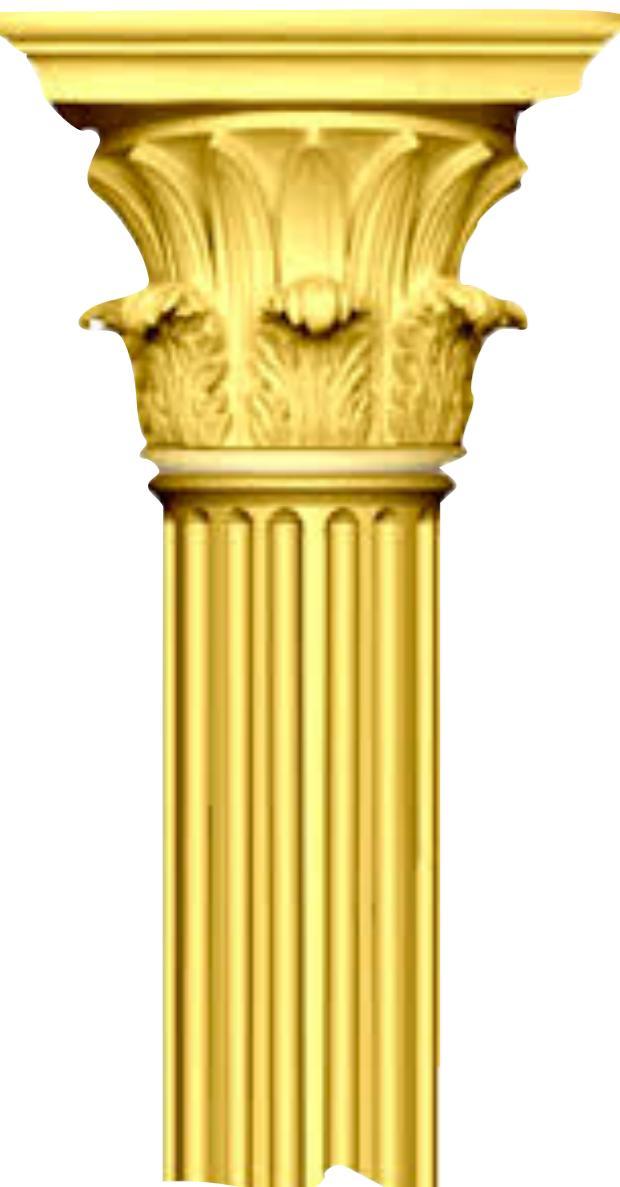


1. Models

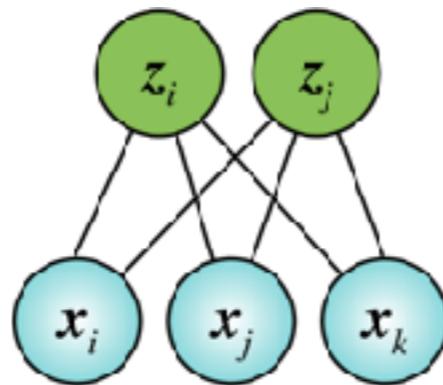


**2. Learning
Principles**

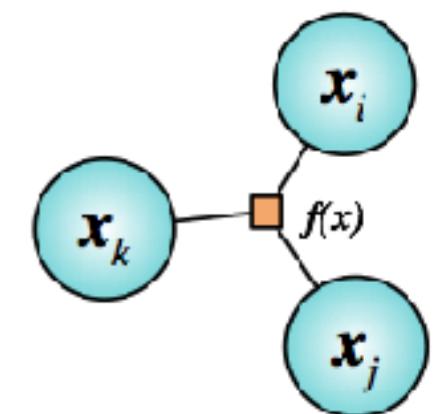
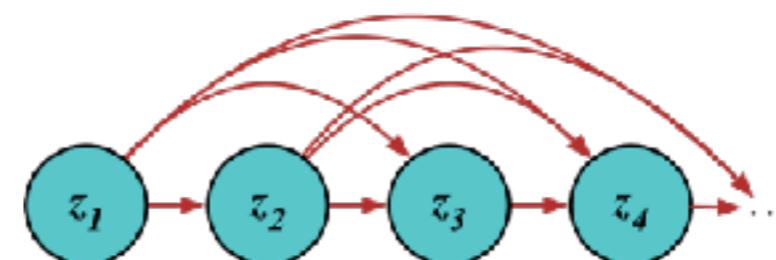
Models



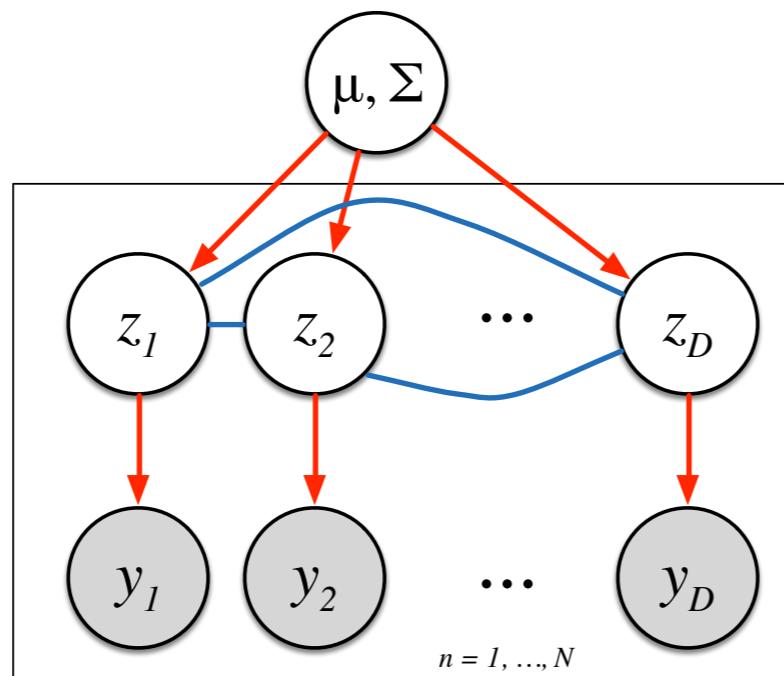
Directed and Undirected



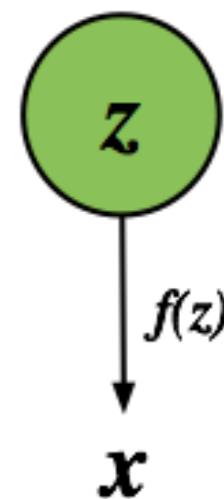
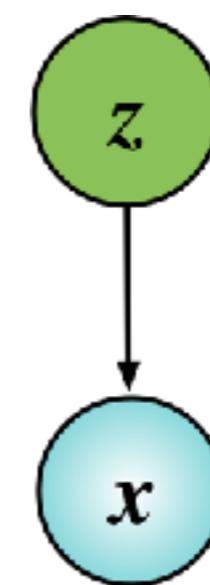
Fully-observed



Parametric, Non-parametric
And semi-parametric

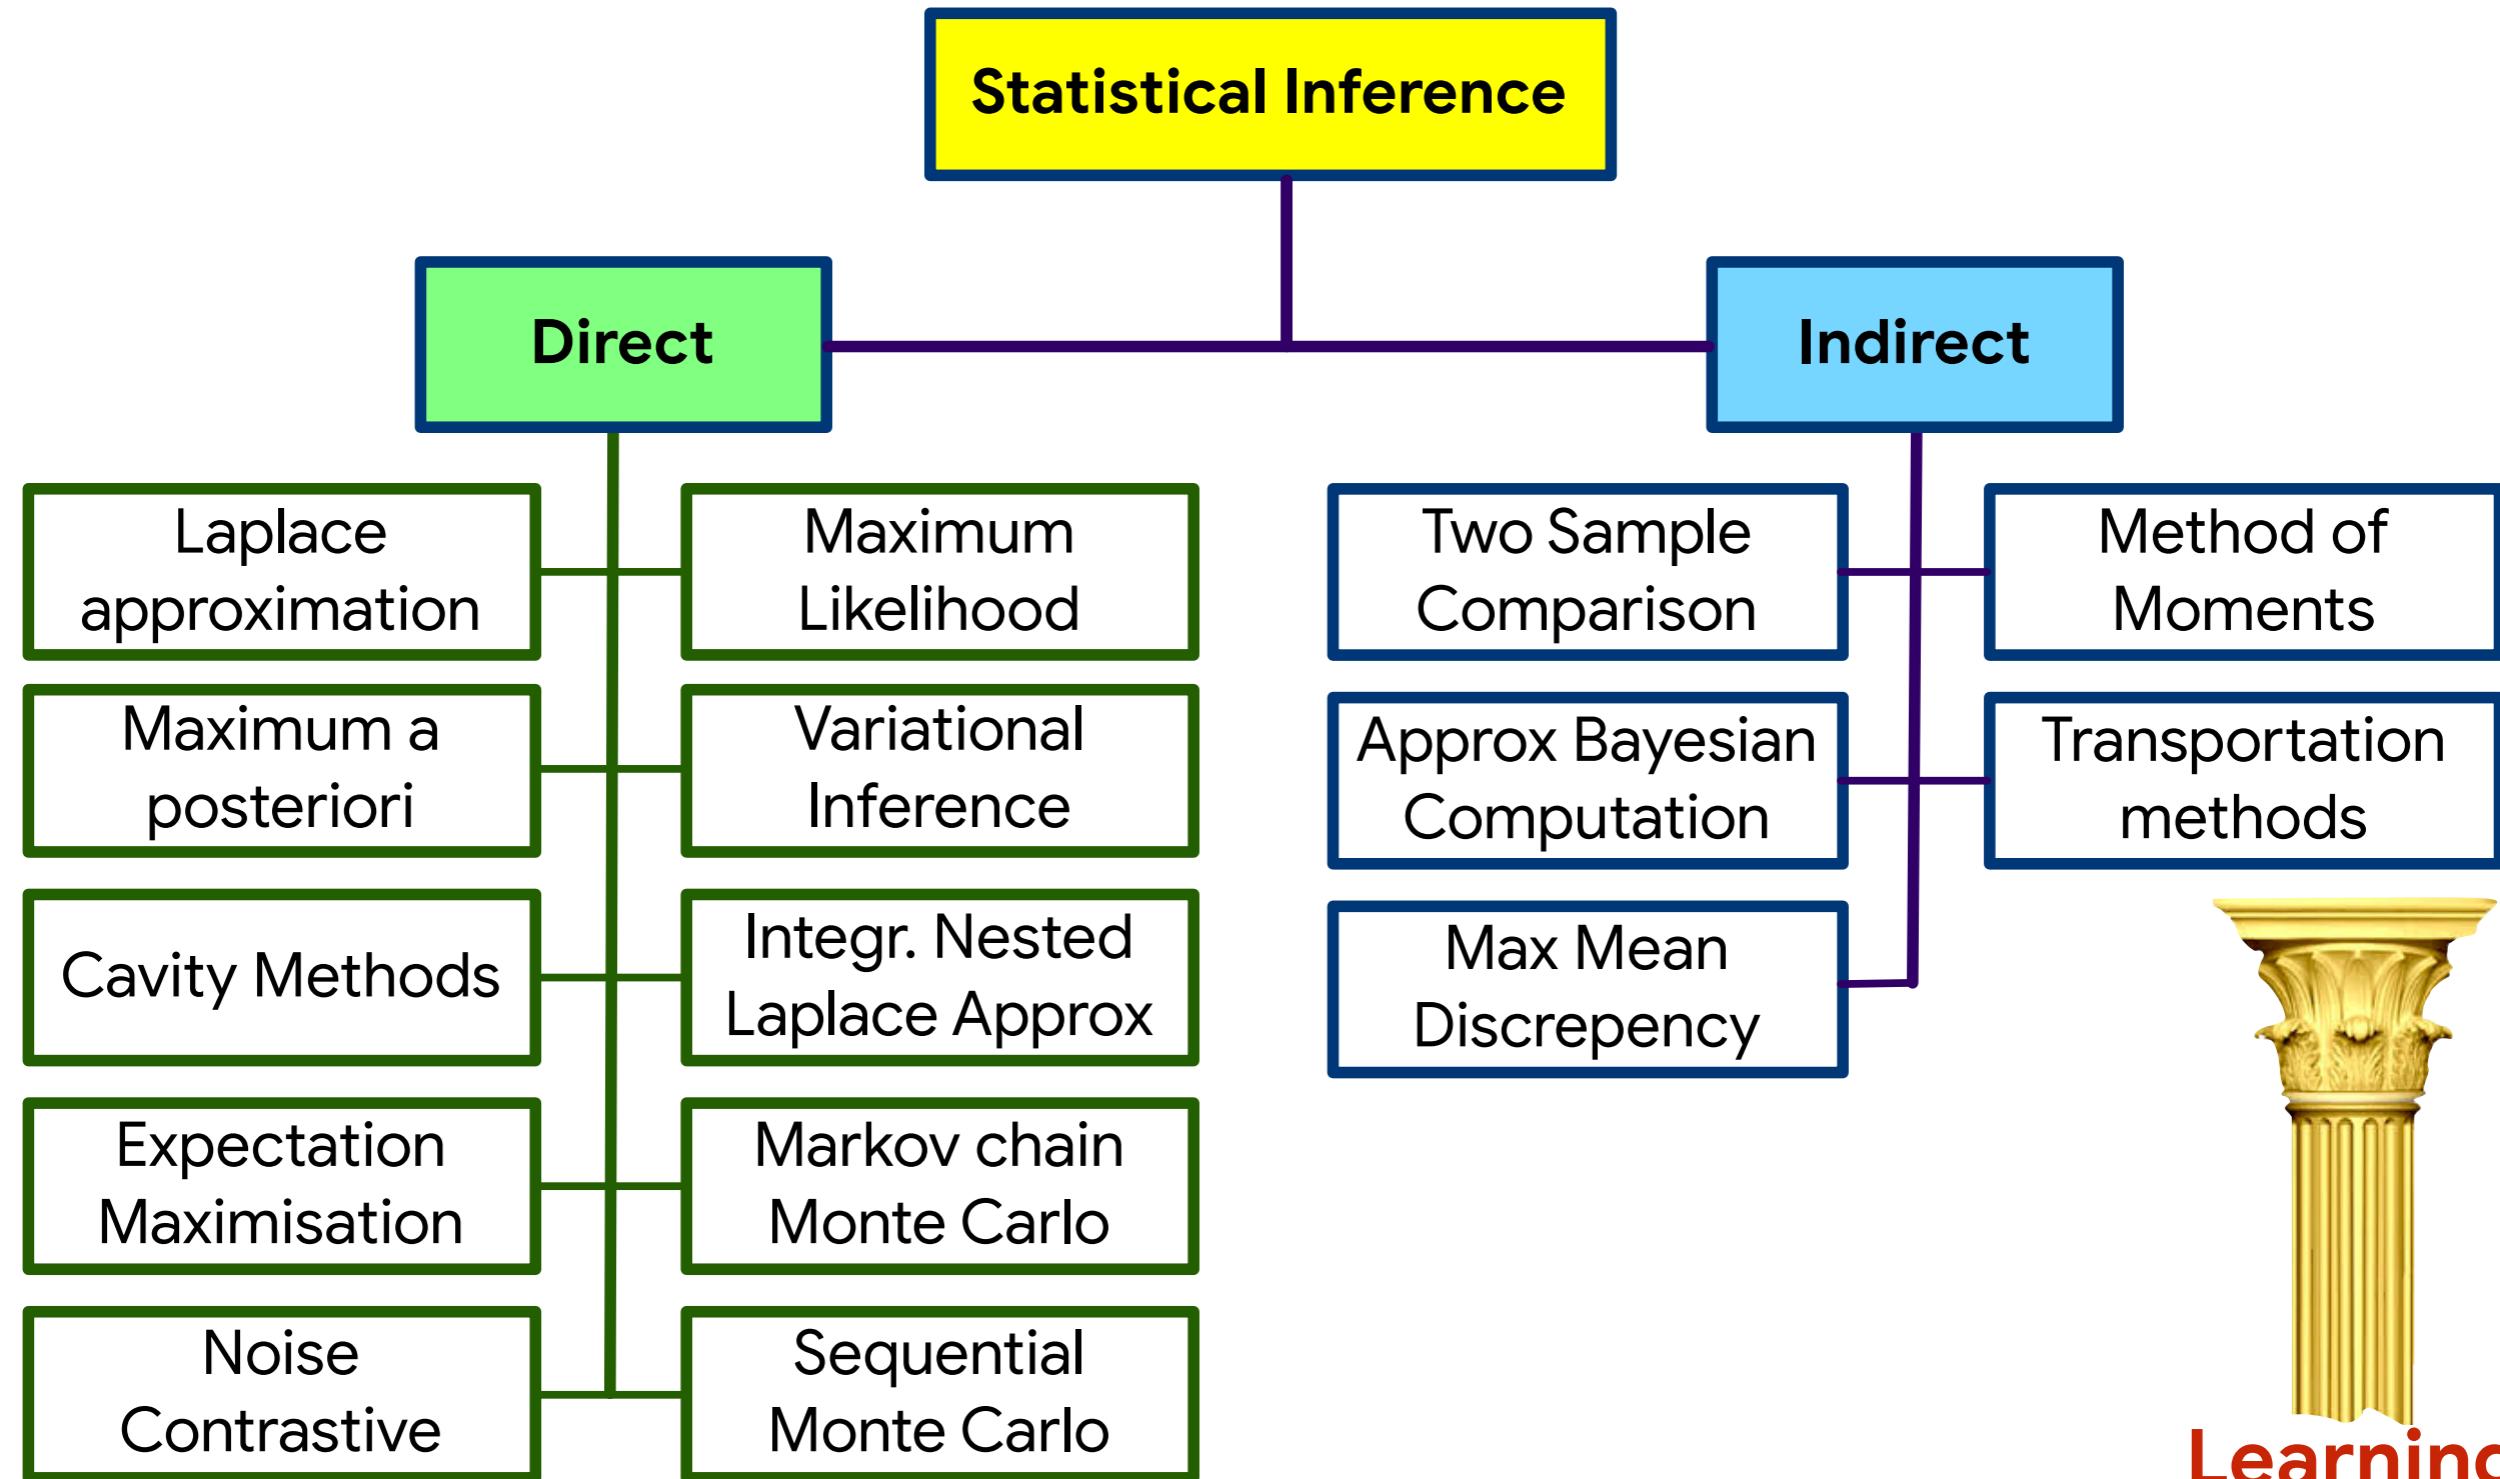


Latent Variable



Models

Learning Principles



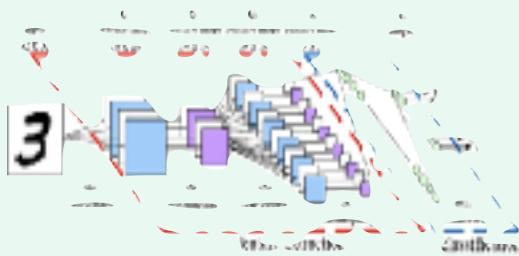
Learning
Principles

Algorithms



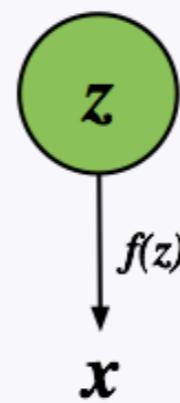
A given model and learning principle can be implemented in many ways.

Convolutional neural network + penalised maximum likelihood



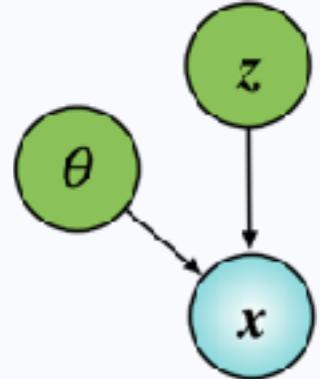
- Optimisation methods (SGD, Adagrad)
- Regularisation (L1, L2, batchnorm, dropout)

Implicit Generative Model + Two-sample testing



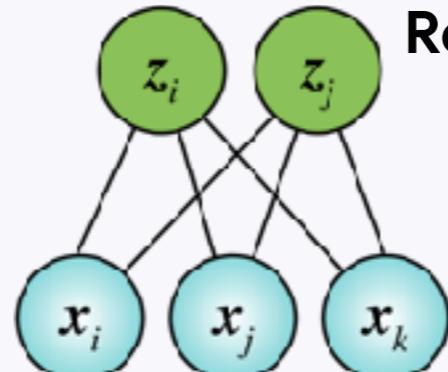
- Unsupervised-as-supervised learning
- Approximate Bayesian Computation (ABC)
- Generative adversarial network (GAN)

Latent variable model + variational inference



- VEM algorithm
- Expectation propagation
- Approximate message passing
- Variational auto-encoders (VAE)

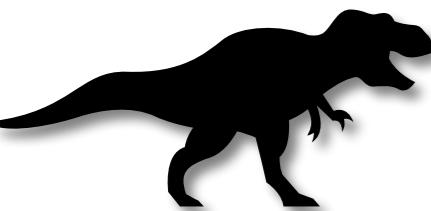
Restricted Boltzmann Machine + maximum likelihood



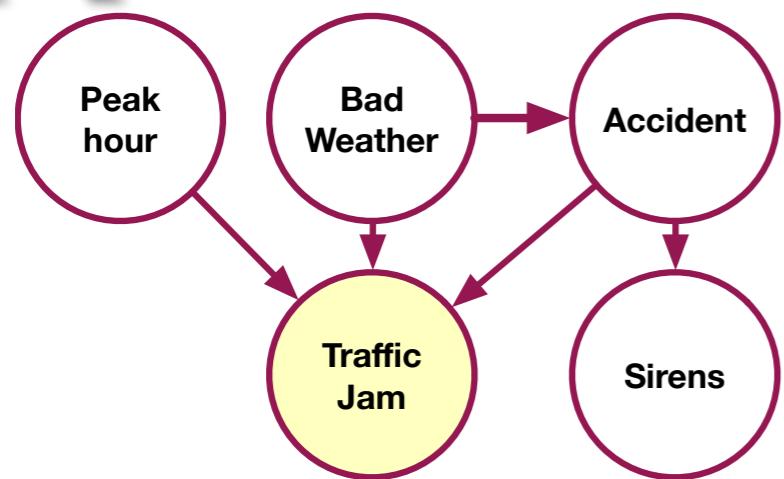
- Contrastive Divergence
- Persistent CD
- Parallel Tempering
- Natural gradients

Final Words

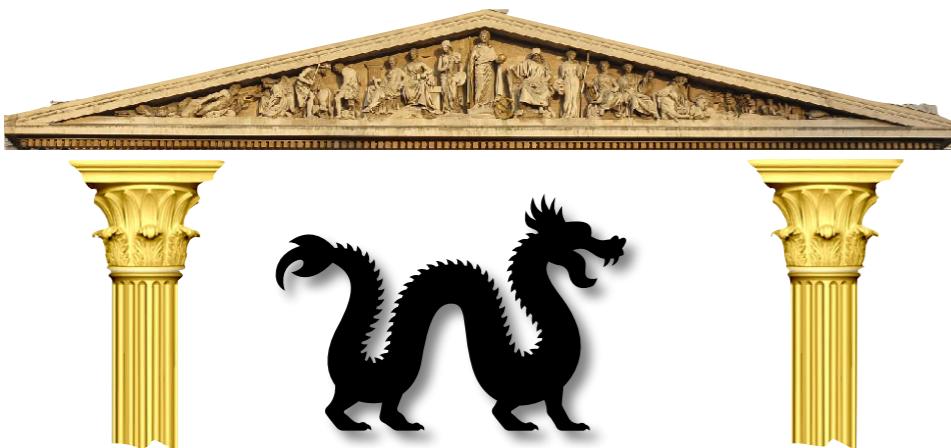
Subjective Probability
Probability as a degree of belief



**Probabilistic descriptions
of systems and data**

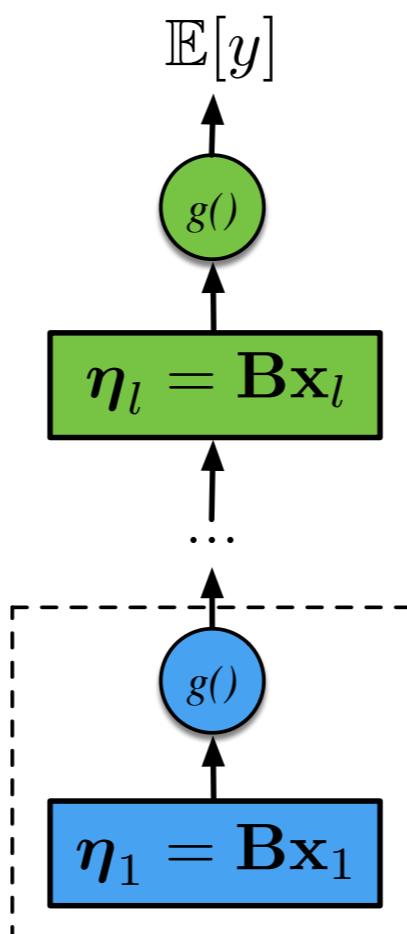


Model-Inference-Algorithm

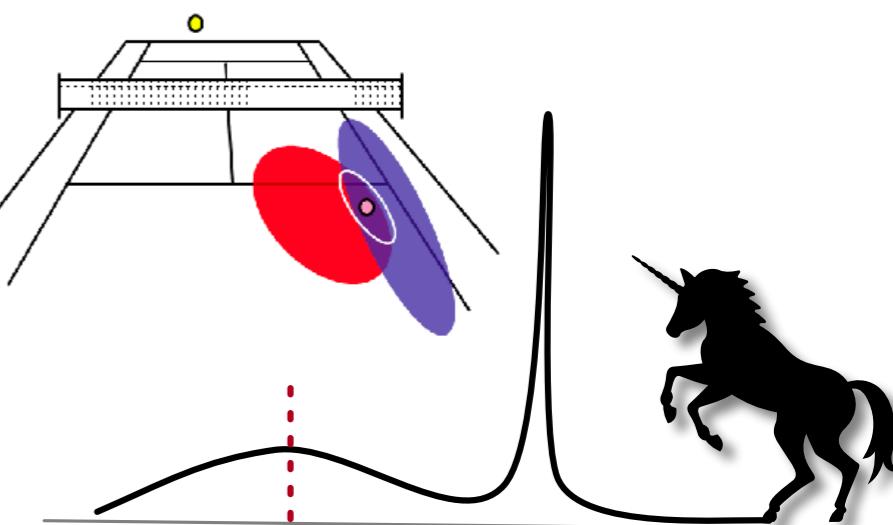
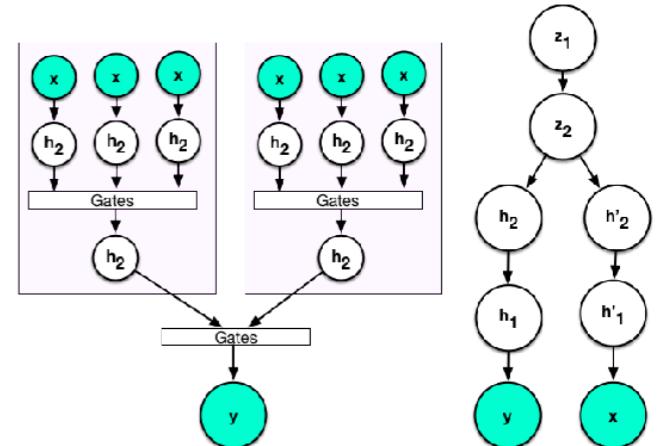


**Deep Learning, Estimation theory,
hierarchical models, dualities**

Probabilistic Model



Likelihood function



Planting the Seeds of Probabilistic Thinking

Foundations | Tricks | Algorithms

Shakir Mohamed

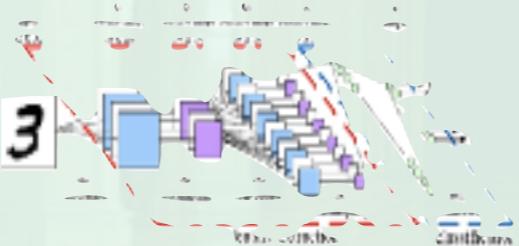
Research Scientist, DeepMind

Last Time ...



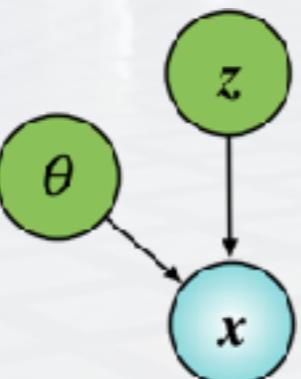
1. Models

**Convolutional neural network
+ penalised maximum likelihood**



- Optimisation methods (SGD, Adagrad)
- Regularisation (L1, L2, batchnorm, dropout)

**Latent variable model
+ variational inference**



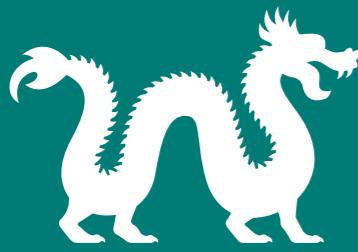
- VEM algorithm
- Expectation propagation
- Approximate message passing
- Variational auto-encoders (VAE)

2. Learning Principles

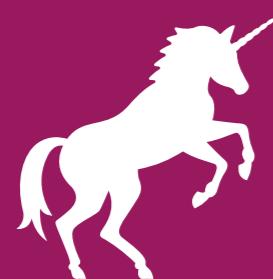
Planting the Seeds of Probabilistic Thinking

Part II: Tricks

Learning Objectives



1. Develop tools to manipulate distributions by studying 6 probability questions.



2. Build connections between concepts in machine learning and those in other computational sciences.

Inferential Questions

Probabilistic dexterity is needed to solve the fundamental problems of machine learning and artificial intelligence.

Evidence
Estimation

$$p(\mathbf{x}) = \int p(\mathbf{x}, \mathbf{z}) d\mathbf{z}$$

Moment
Computation

$$\mathbb{E}[f(\mathbf{z})|\mathbf{x}] = \int f(\mathbf{z})p(\mathbf{z}|\mathbf{x})d\mathbf{z}$$

Parameter
Estimation

$$p(\boldsymbol{\theta}|\mathbf{x}_{0:N})$$

Prediction

$$p(\mathbf{x}_{t+1}|\mathbf{x}_{0:t})$$

Planning

$$\mathcal{J} = \mathbb{E}_p \left[\int_0^\infty C(\mathbf{x}_t) dt | \mathbf{x}_0, \mathbf{u} \right]$$

Hypothesis Testing

$$\mathcal{B} = \log p(\mathbf{x}|H_1) - \log p(\mathbf{x}|H_2)$$

Experimental Design

$$\mathcal{IG} = D[p(\mathbf{x}_{t:T}|u) \| p(\mathbf{x}_{0:t})]$$

Identity Trick

Transform an expectation w.r.t. distribution p ,
into an expectation w.r.t. distribution q .

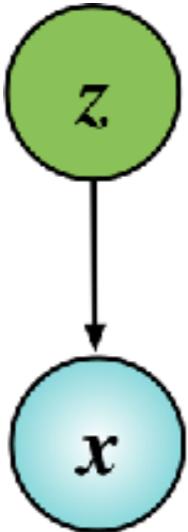
$$\int p(\mathbf{x}) f(\mathbf{x}) d\mathbf{x} = \boxed{\mathbb{E}_{p(\mathbf{x})}[f(\mathbf{x})]}$$



$$\boxed{\mathbb{E}_{q(\mathbf{x})}[g(\mathbf{x}; f)]} = \int q(\mathbf{x}) g(\mathbf{x}, f) d\mathbf{x}$$

Do this by introducing a **probabilistic one** $\frac{p(\mathbf{x})}{p(\mathbf{x})}$

Identity Trick

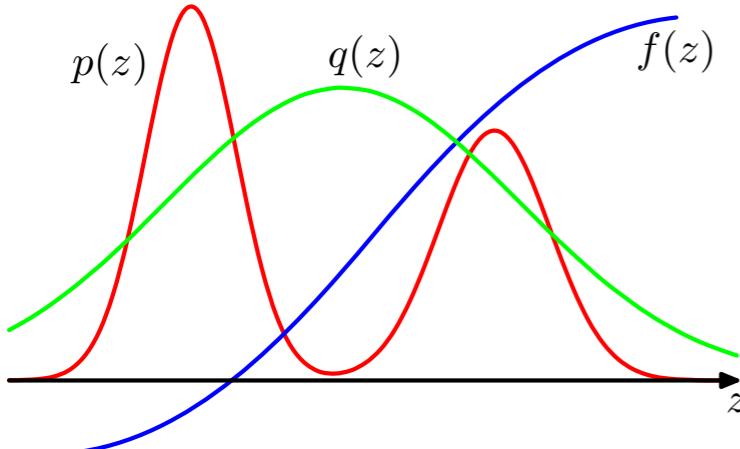


Integral problem

$$p(\mathbf{x}) = \int p(\mathbf{x}|\mathbf{z})p(\mathbf{z})d\mathbf{z}$$

Probabilistic one

$$p(\mathbf{x}) = \int p(\mathbf{x}|\mathbf{z})p(\mathbf{z})\frac{q(\mathbf{z})}{q(\mathbf{z})}d\mathbf{z}$$



Re-group/re-weight

$$p(\mathbf{x}) = \int p(\mathbf{x}|\mathbf{z})\frac{p(\mathbf{z})}{q(\mathbf{z})}q(\mathbf{z})d\mathbf{z}$$

Conditions

- $q(z) > 0$, when $p(\mathbf{x}|z)p(z) \neq 0$.
- $q(z)$ is known/easy to handle.

$$p(\mathbf{x}) = \mathbb{E}_{q(\mathbf{z})} \left[p(\mathbf{x}|\mathbf{z})\frac{p(\mathbf{z})}{q(\mathbf{z})} \right]$$

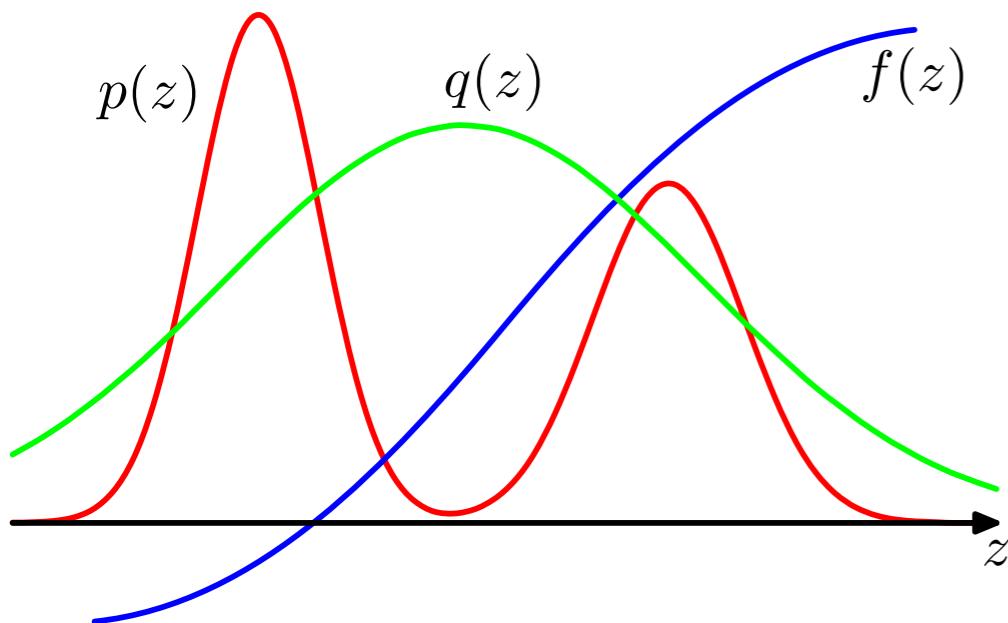
Importance Sampling

$$p(\mathbf{x}) = \mathbb{E}_{q(\mathbf{z})} \left[p(\mathbf{x}|\mathbf{z}) \frac{p(\mathbf{z})}{q(\mathbf{z})} \right]$$

Monte Carlo
Estimator

$$p(\mathbf{x}) = \frac{1}{S} \sum_s w^{(s)} p(\mathbf{x}|\mathbf{z}^{(s)})$$

$$w^{(s)} = \frac{p(z)}{q(z)} \quad z^{(s)} \sim q(z)$$



Identity Trick Elsewhere

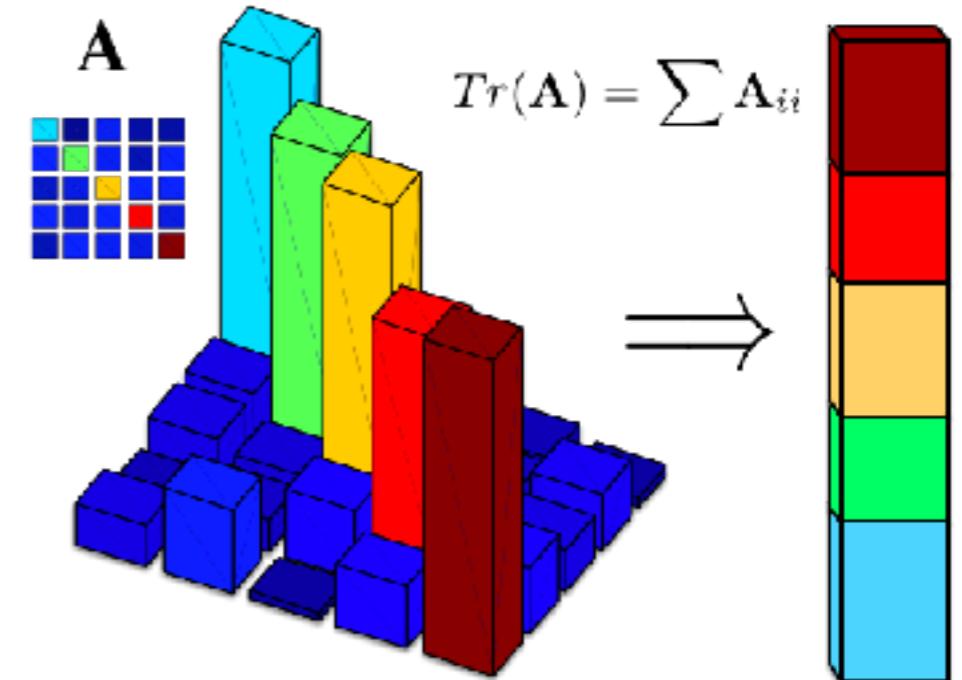
- Manipulate stochastic gradients
- Derive probability bounds
- RL for policy corrections

Hutchinson's Trick

Compute the trace of a matrix:

- KL between two Gaussians.
- Gradient of a log-determinant.

$$\partial(\log \det(\mathbf{X})) = \text{Tr}(\mathbf{X}^{-1} \partial X)$$



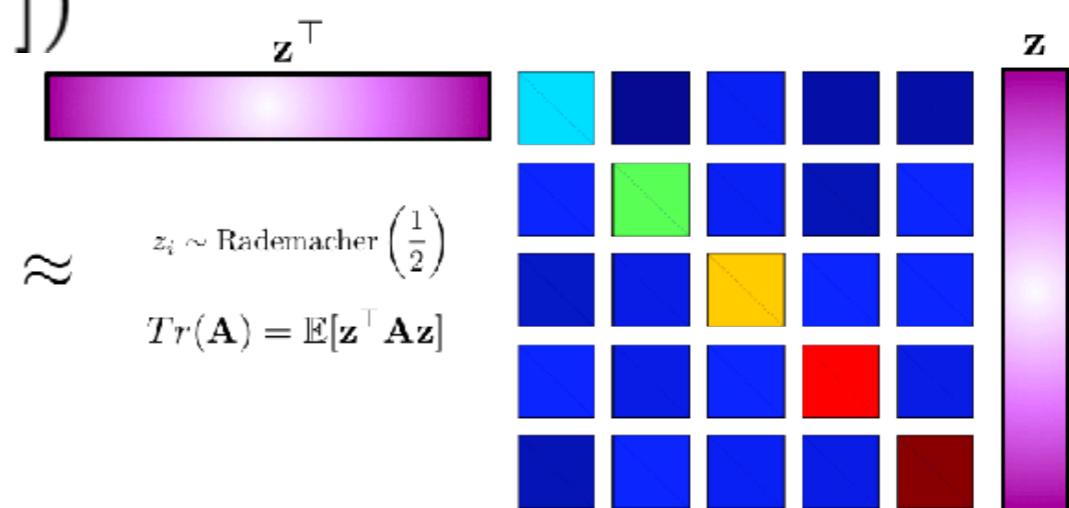
Trace problem $\text{Tr}(\mathbf{A})$

Zero mean unit var
 $(\Sigma + \mathbf{m}\mathbf{m}^\top)$

Identity Trick $\text{Tr}(\mathbf{A}\mathbf{I}) = \text{Tr}(\mathbf{A}\mathbb{E}[\mathbf{z}\mathbf{z}^\top])$

Linear operations $\mathbb{E}[\text{Tr}(\mathbf{A}\mathbf{z}\mathbf{z}^\top)]$

Trace property $\mathbb{E}[\mathbf{z}^\top \mathbf{A}\mathbf{z}]$



Sampling z randomly, compute Trace using linear systems of equations

Probability Flow Tricks

Distribution and sample

$$\hat{x} \sim p(x)$$

Transformation

$$\hat{y} = g(\hat{x}; \theta)$$

Unconscious Statistician

$$\mathbb{E}_{p(x)}[g(x; \theta)]$$

Change of Variables

$$p(\mathbf{y}) = p(\mathbf{x}) \left| \frac{dg}{d\mathbf{x}} \right|^{-1}$$

Makes entropy computation and backpropagation easy.

Begin with a diagonal Gaussian and improve by change of variables.

Triangular Jacobians allow for computational efficiency.

Compute

$$\log \det \left(\frac{\partial f(\mathbf{z})}{\partial \mathbf{z}} \right)$$

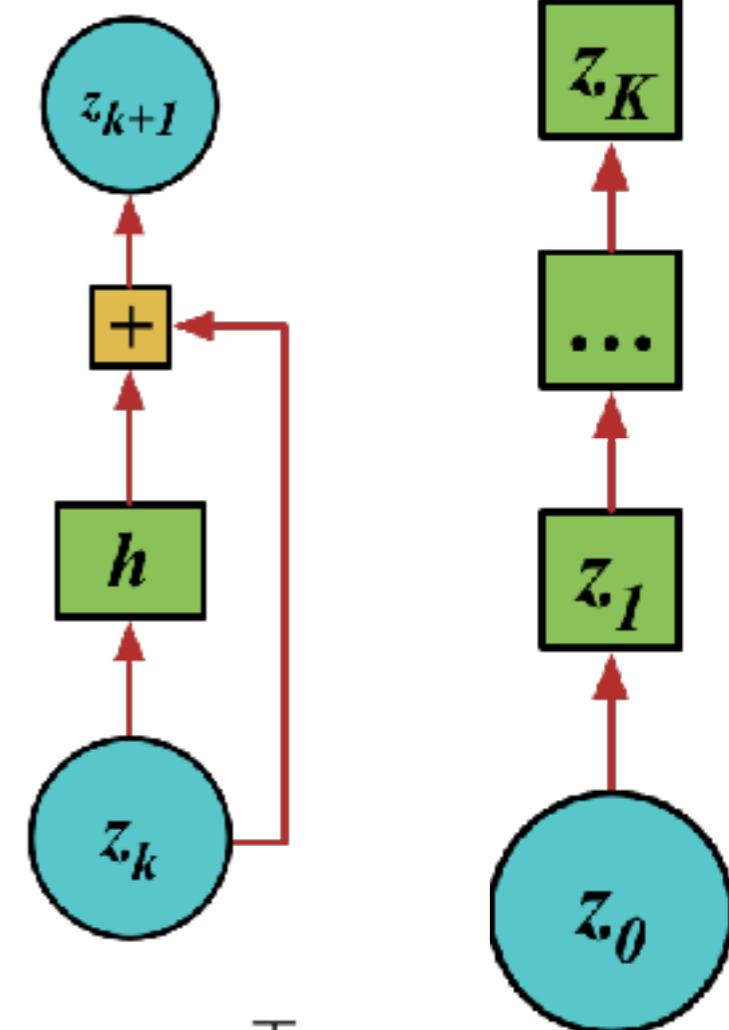
$$\det(\mathbf{I} + \mathbf{u}\mathbf{s}^\top) = (1 + \mathbf{u}^\top \mathbf{s})$$

$$f(\mathbf{z}) = \mathbf{z} + \mathbf{u}h(\mathbf{w}^\top \mathbf{z} + b)$$
$$\det(I + ab^\top) = 1 + a^\top b$$

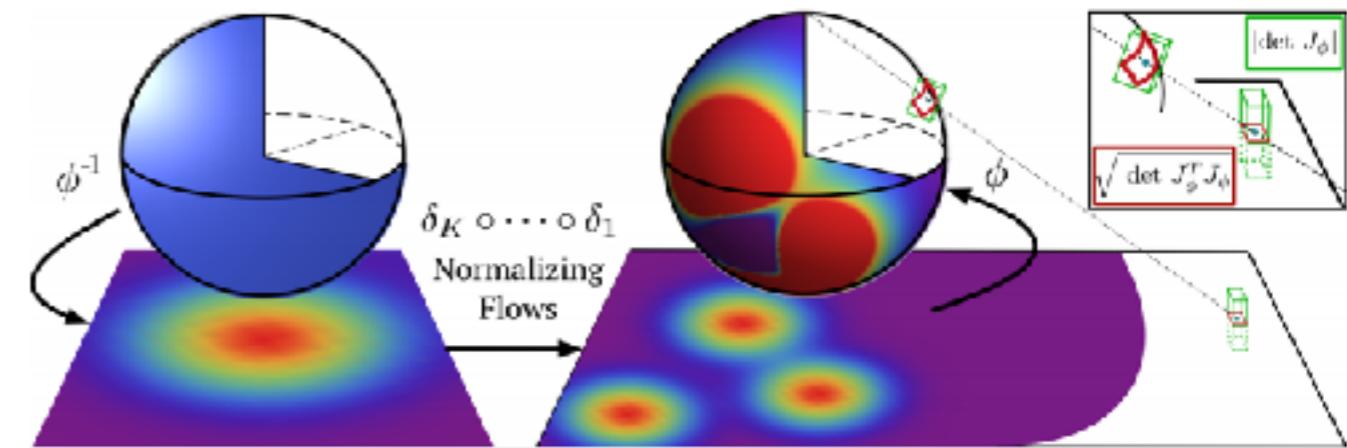
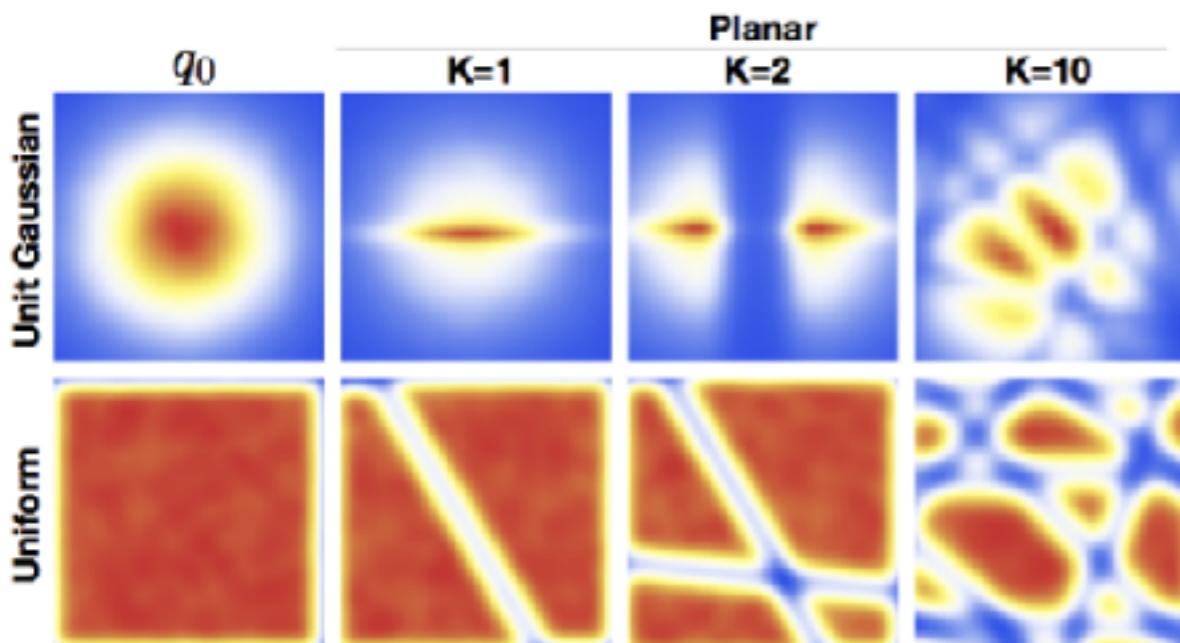
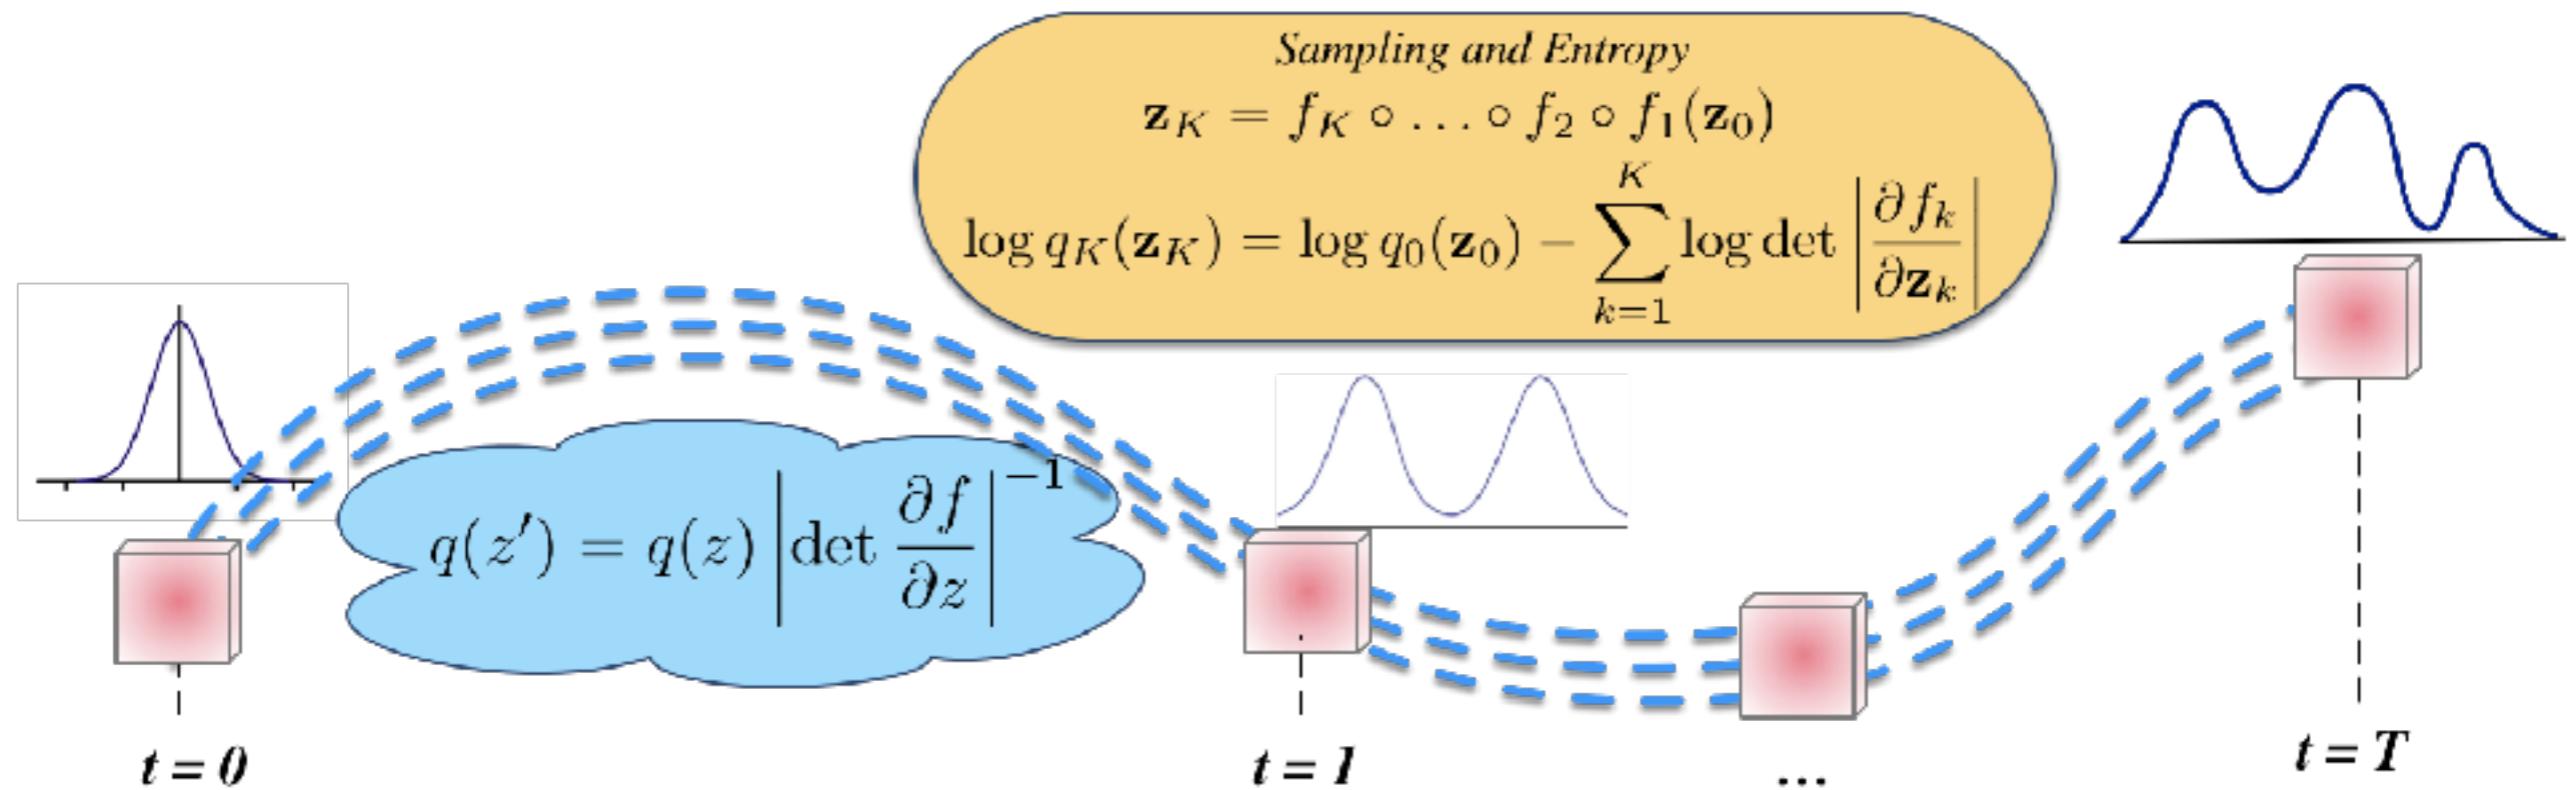
$$\mathbf{s} = h' \mathbf{w}$$

Linear time computation of the determinant and its gradient.

Planar Flow



Normalising Flows



Stochastic Optimisation

Common gradient problem

$$\nabla_{\phi} \mathbb{E}_{q_{\phi}(\mathbf{z})} [f_{\theta}(\mathbf{z})] = \nabla \int q_{\phi}(\mathbf{z}) f_{\theta}(\mathbf{z}) d\mathbf{z}$$

- Don't know this expectation in general.
- Gradient is of the parameters of the distribution w.r.t. which the expectation is taken.

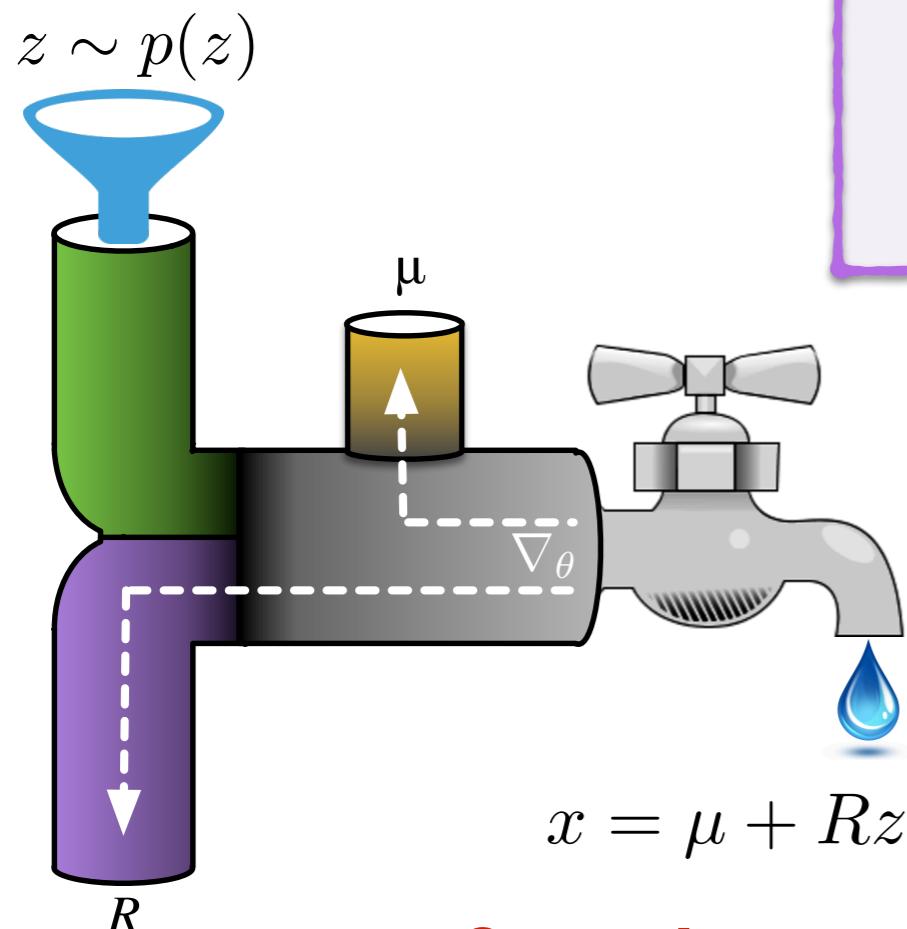
1. **Pathwise estimator**: Differentiate the function $\mathbf{f}(\mathbf{z})$
2. **Score-function estimator**: Differentiate the density $\mathbf{q}(\mathbf{z}|\mathbf{x})$

Typical problem areas

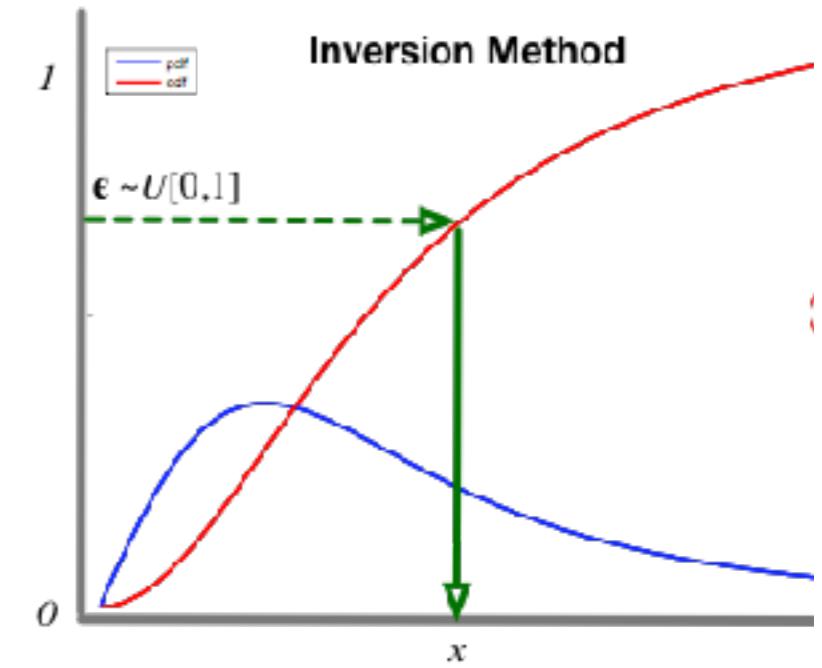
- Sensitivity analysis
- Generative models and inference
- Reinforcement learning and control
- Operations research and inventory control
- Monte Carlo simulation
- Finance and asset pricing

Reparameterisation Tricks

Distributions can be expressed as a transformations of other distributions.



$$\boxed{z \sim q_\phi(\mathbf{z})}$$
$$\mathbf{z} = g(\epsilon, \phi) \quad \epsilon \sim p(\epsilon)$$

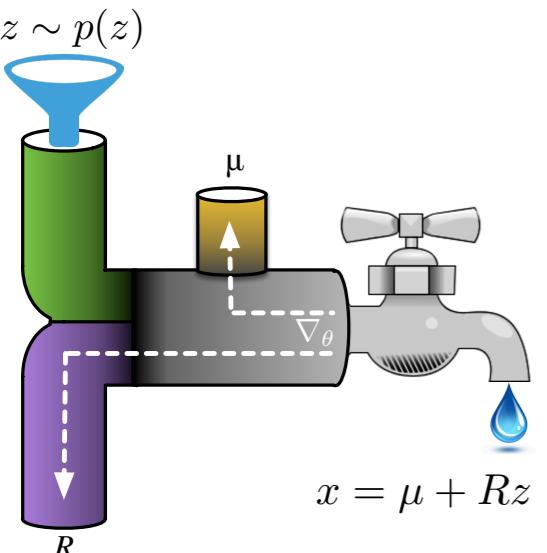


Samplers, one-liners and change-of-variables

$$p(z) = \left| \frac{d\epsilon}{dz} \right| p(\epsilon) \implies |p(z)dz| = |p(\epsilon)d\epsilon|$$

Pathwise Estimator

(Non-rigorous) Derivation

$$\begin{aligned}\nabla_{\phi} \mathbb{E}_{q(z)}[f(\mathbf{z})] &= \nabla_{\phi} \int q_{\phi}(\mathbf{z}) f(\mathbf{z}) d\mathbf{z} && \text{Known transformation} \\ &= \nabla_{\phi} \int p(\epsilon) \frac{d\epsilon}{d\mathbf{z}} f(g(\epsilon, \phi)) g'(\epsilon, \phi) d\epsilon && \text{Change of variables} \\ &= \nabla_{\phi} \mathbb{E}_{p(\epsilon)}[f(g(\phi, \epsilon))] = \mathbb{E}_{p(\epsilon)}[\nabla_{\phi} f(g(\phi, \epsilon))] && \text{Inv fn Thm}\end{aligned}$$

$$\nabla_{\phi} \mathbb{E}_{q_{\phi}(\mathbf{z})}[f_{\theta}(\mathbf{z})] = \mathbb{E}_{p(\epsilon)}[\nabla_{\phi} f_{\theta}(g(\epsilon, \phi))]$$

Other names

- Unconscious statistician
- Stochastic backpropagation
- Perturbation analysis
- Reparameterisation trick
- Affine-independent inference

When to use

- Function f is differentiable
- Density q is known with a suitable transform of a simpler base distribution: inverse CDF, location-scale transform, or other co-ordinate transform.
- Easy to sample from base distribution.

Log-derivative Trick

Score function is the derivative of a log-likelihood function.

$$\nabla_{\phi} \log q_{\phi}(\mathbf{z}) = \frac{\nabla_{\phi} q_{\phi}(\mathbf{z})}{q_{\phi}(\mathbf{z})}$$

Several useful properties

Expected score

$$\mathbb{E}_{q(z)} [\nabla_{\phi} \log q_{\phi}(\mathbf{z})] = 0$$

←Show this

$$\mathbb{E}_{q(z)} [\nabla_{\phi} \log q_{\phi}(\mathbf{z})] = \int q(z) \frac{\nabla_{\phi} q_{\phi}(\mathbf{z})}{q_{\phi}(\mathbf{z})} = \nabla \int q_{\phi}(\mathbf{z}) = \nabla 1 = 0$$

Fisher Information

$$\mathbb{V}[\nabla_{\theta} \log p(\mathbf{x}; \theta)] = \mathcal{I}(\theta) = \mathbb{E}_{p(x; \theta)} [\nabla_{\theta} \log p(\mathbf{x}; \theta) \nabla_{\theta} \log p(\mathbf{x}; \theta)^{\top}]$$

Score-function Estimator

$$\nabla_{\phi} \mathbb{E}_{q_{\phi}(\mathbf{z})}[f_{\theta}(\mathbf{z})] = \nabla \int q_{\phi}(\mathbf{z}) f_{\theta}(\mathbf{z}) d\mathbf{z}$$

Leibnitz integral rule

$$= \int \frac{q_{\phi}(\mathbf{z})}{q_{\phi}(\mathbf{z})} \nabla_{\phi} q_{\phi}(\mathbf{z}) f(\mathbf{z}) d\mathbf{z}$$

Identity

$$= \int q_{\phi}(\mathbf{z}) \nabla_{\phi} \log q_{\phi}(\mathbf{z}) f(\mathbf{z}) d\mathbf{z}$$

Log-deriv

$$= \mathbb{E}_{q_{\phi}(\mathbf{z})} [f(\mathbf{z}) \nabla_{\phi} \log q_{\phi}(\mathbf{z})]$$

Gradient

$$\nabla_{\phi} \mathbb{E}_{q_{\phi}(\mathbf{z})}[f_{\theta}(\mathbf{z})] = \mathbb{E}_{q_{\phi}(\mathbf{z})} [(f(\mathbf{z}) - c) \nabla_{\phi} \log q_{\phi}(\mathbf{z})]$$

Control Variate

Other names

- Likelihood ratio method
- REINFORCE and policy gradients
- Automated & Black-box inference

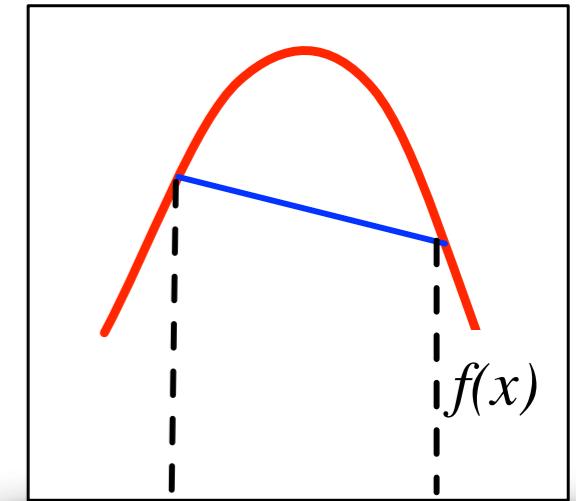
When to use

- Function is not differentiable, not analytical.
- Distribution q is easy to sample from.
- Density q is known and differentiable.

Bounding Tricks

An important result from convex analysis lets us move expectations through a function:

For concave functions $f(\cdot)$
$$f(\mathbb{E}[x]) \geq \mathbb{E}[f(x)]$$



Logarithms are strictly concave allowing us to use Jensen's inequality.

$$\log \int p(x)g(x)dx \geq \int p(x)\log g(x)dx$$

Bounding Trick Elsewhere

Optimisation; Variational Inference; Rao-Blackwell Theorem;

Other Bounding Tricks

- Fenchel duality
- Holder's inequality
- Monge-Kantorovich Inequality

Evidence Bounds

Integral problem

$$p(\mathbf{x}) = \int p(\mathbf{x}|\mathbf{z})p(\mathbf{z})d\mathbf{z}$$

Proposal

$$p(\mathbf{x}) = \int p(\mathbf{x}|\mathbf{z})p(\mathbf{z})\frac{q(\mathbf{z})}{q(\mathbf{z})}d\mathbf{z}$$

Importance Weight

$$p(\mathbf{x}) = \int p(\mathbf{x}|\mathbf{z})\frac{p(\mathbf{z})}{q(\mathbf{z})}q(\mathbf{z})d\mathbf{z}$$

Jensen's inequality

$$\log \int p(x)g(x)dx \geq \int p(x)\log g(x)dx$$

$$\begin{aligned}\log p(\mathbf{x}) &\geq \int q(\mathbf{z}) \log \left(p(\mathbf{x}|\mathbf{z})\frac{p(\mathbf{z})}{q(\mathbf{z})} \right) d\mathbf{z} \\ &= \int q(\mathbf{z}) \log p(\mathbf{x}|\mathbf{z}) - \int q(\mathbf{z}) \log \frac{q(\mathbf{z})}{p(\mathbf{z})}\end{aligned}$$

Lower bound

$$\mathbb{E}_{q(\mathbf{z})}[\log p(\mathbf{x}|\mathbf{z})] - KL[q(\mathbf{z})||p(\mathbf{z})]$$

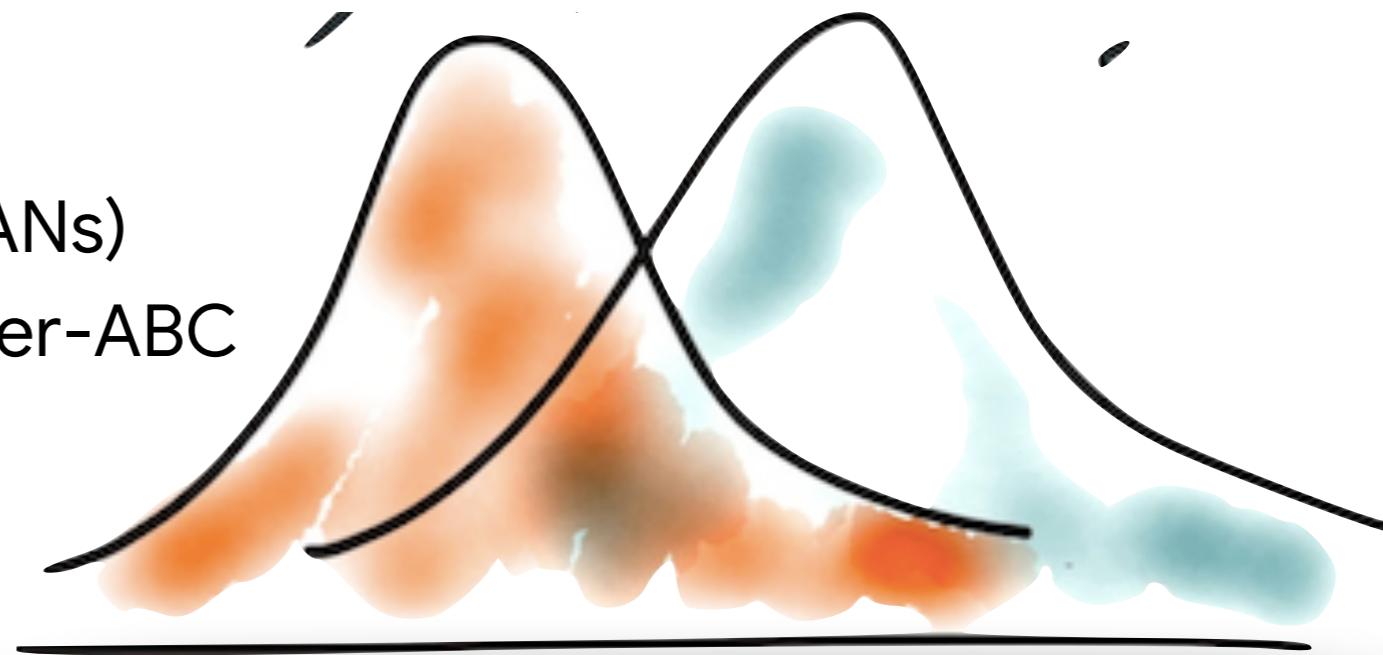
Density Ratio Trick

The ratio of two densities can be computed using a classifier of using samples drawn from the two distributions.

$$\frac{p^*(\mathbf{x})}{q(\mathbf{x})} = \frac{p(y = 1|\mathbf{x})}{p(y = -1|\mathbf{x})}$$

Density Ratio Trick Elsewhere

- Generative Adversarial Networks (GANs)
- Noise contrastive estimation, Classifier-ABC
- Two-sample testing
- Covariate-shift, calibration



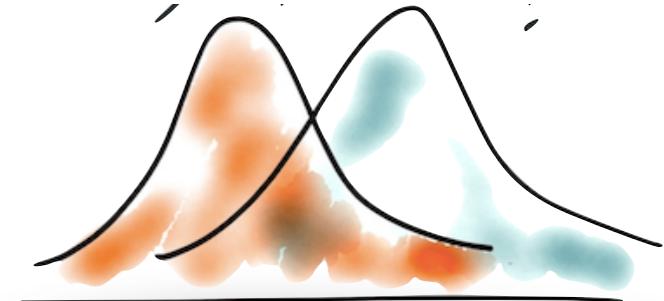
Density Ratio Estimation

Assign labels

$$\{y_1, \dots, y_N\} = \{+1, \dots, +1, -1, \dots, -1\}$$

Equivalence

$$p^*(\mathbf{x}) = p(\mathbf{x}|y=1) \quad q(\mathbf{x}) = p(\mathbf{x}|y=-1)$$



Density Ratio

$$\frac{p^*(\mathbf{x})}{q(\mathbf{x})}$$

Bayes' Rule

$$p(\mathbf{x}|y) = \frac{p(y|\mathbf{x})p(\mathbf{x})}{p(y)}$$

Conditional

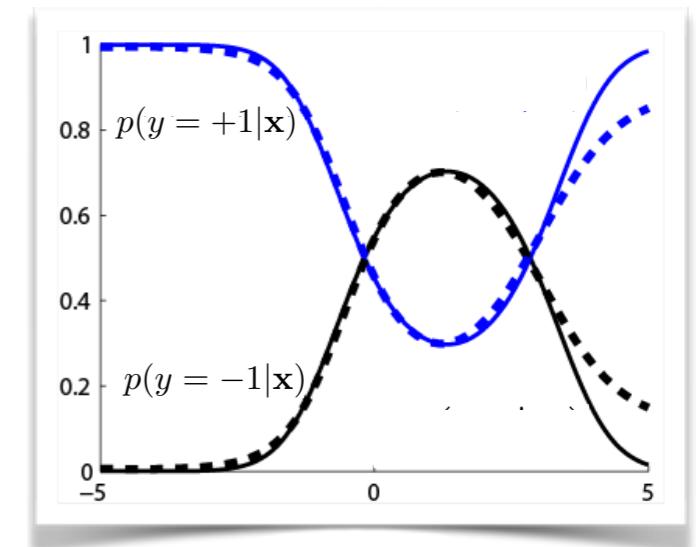
$$\frac{p^*(\mathbf{x})}{q(\mathbf{x})} = \frac{p(\mathbf{x}|y=1)}{p(\mathbf{x}|y=-1)}$$

Bayes' Subst.

$$= \frac{p(y=+1|\mathbf{x})p(\mathbf{x})}{p(y=+1)} \Big/ \frac{p(y=-1|\mathbf{x})p(\mathbf{x})}{p(y=-1)}$$

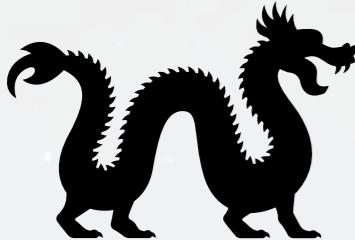
Class probability

$$\frac{p^*(\mathbf{x})}{q(\mathbf{x})} = \frac{p(y=1|\mathbf{x})}{p(y=-1|\mathbf{x})}$$



Computing a density ratio is equivalent to class probability estimation.

Final Words



Strengthen your probabilistic dexterity.

Identity

$$\frac{p(\mathbf{x})}{p(\mathbf{x})}$$

Hutchinson's

$$\mathbb{E}[\mathbf{z}\mathbf{z}^\top] = \mathbf{I}$$

Flows

$$p(\mathbf{y}) = p(\mathbf{x}) \left| \frac{dg}{d\mathbf{x}} \right|^{-1}$$

Log-derivative

$$\nabla_\phi \log q_\phi(\mathbf{z}) = \frac{\nabla_\phi q_\phi(\mathbf{z})}{q_\phi(\mathbf{z})}$$

Reparameterisation

$$\mathbf{z} = g(\epsilon, \phi) \quad \epsilon \sim p(\epsilon)$$

Density Ratio

$$\frac{p^*(\mathbf{x})}{q(\mathbf{x})} = \frac{p(y=1|\mathbf{x})}{p(y=-1|\mathbf{x})}$$



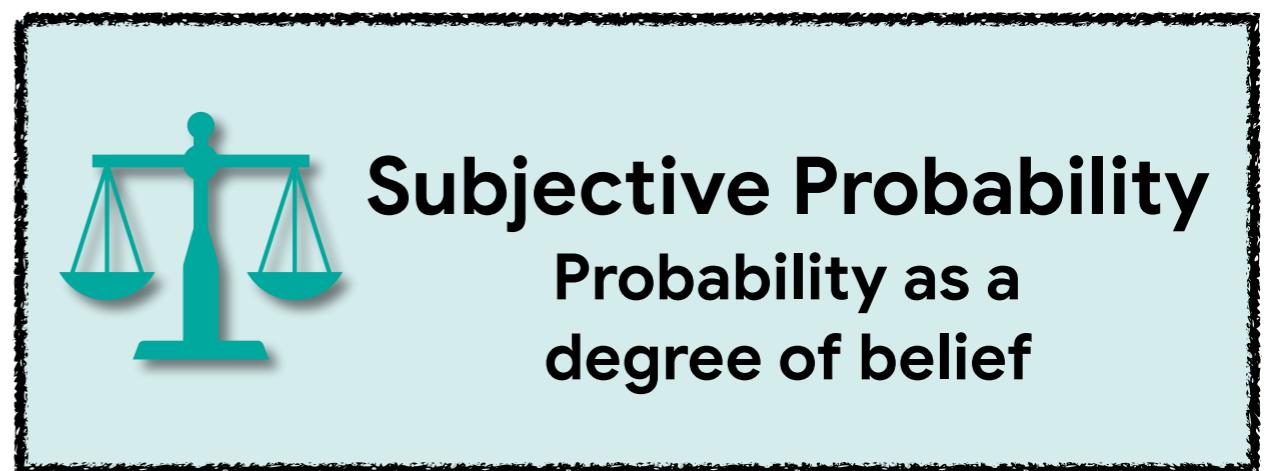
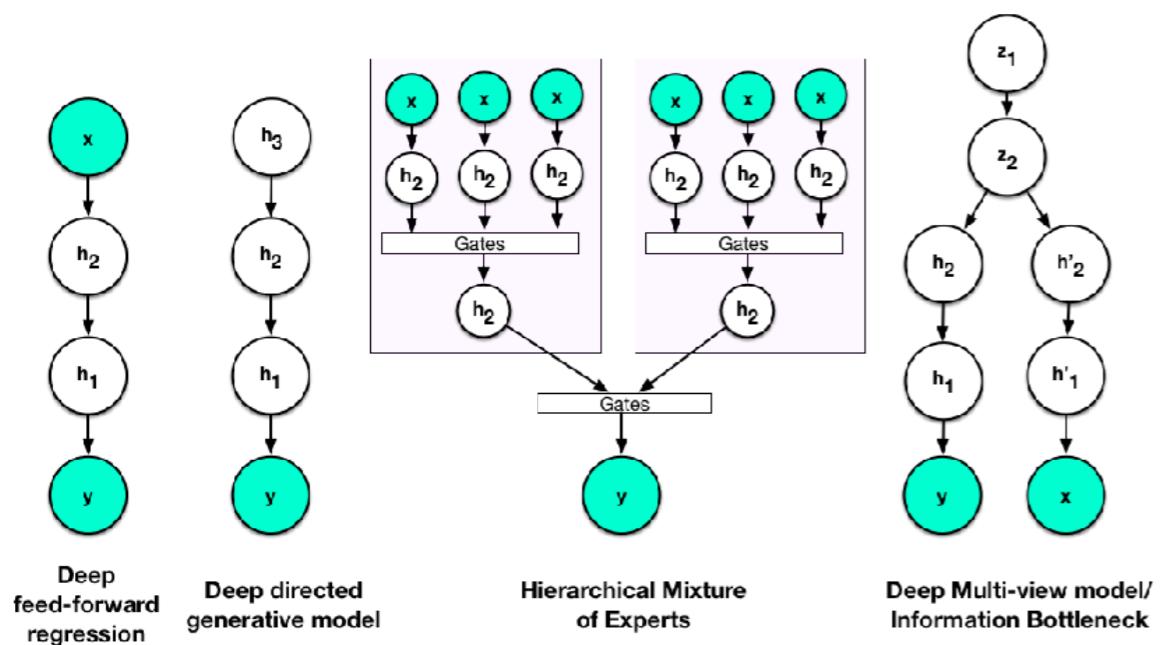
Planting the Seeds of Probabilistic Thinking

Foundations | Tricks | Algorithms

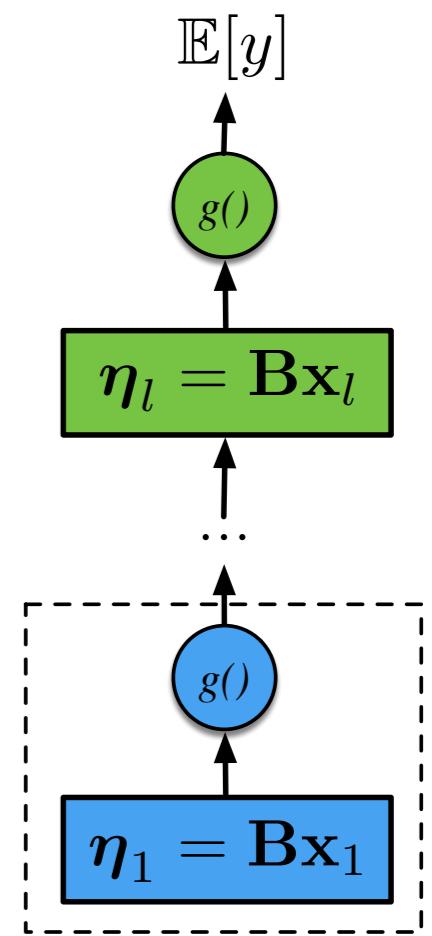
Shakir Mohamed

Research Scientist, DeepMind

Last Time ...



Model-Inference-Algorithm



Manipulating Integrals

Evidence
Estimation

$$p(\mathbf{x}) = \int p(\mathbf{x}, \mathbf{z}) d\mathbf{z}$$

Prediction

$$p(\mathbf{x}_{t+1} | \mathbf{x}_{0:t})$$

Hypothesis Testing

$$\mathcal{B} = \log p(\mathbf{x}|H_1) - \log p(\mathbf{x}|H_2)$$

Moment
Computation

$$\mathbb{E}[f(\mathbf{z})|\mathbf{x}] = \int f(\mathbf{z})p(\mathbf{z}|\mathbf{x})d\mathbf{z}$$



Parameter
Estimation

$$p(\boldsymbol{\theta} | \mathbf{x}_{0:N})$$

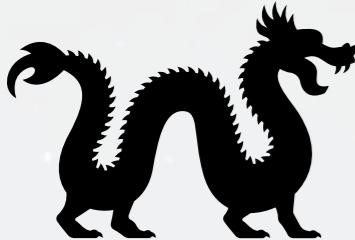
Planning

$$\mathcal{J} = \mathbb{E}_p \left[\int_0^\infty C(\mathbf{x}_t) dt | \mathbf{x}_0, \mathbf{u} \right]$$

Experimental Design

$$\mathcal{IG} = D[p(\mathbf{x}_{t:T} | u) \| p(\mathbf{x}_{0:t})]$$

Your Tricks



Strengthen your probabilistic dexterity.

Identity

$$\frac{p(\mathbf{x})}{p(\mathbf{x})}$$

Hutchinson's

$$\mathbb{E}[\mathbf{z}\mathbf{z}^\top] = \mathbf{I}$$

Flows

$$p(\mathbf{y}) = p(\mathbf{x}) \left| \frac{dg}{d\mathbf{x}} \right|^{-1}$$

Log-derivative

$$\nabla_\phi \log q_\phi(\mathbf{z}) = \frac{\nabla_\phi q_\phi(\mathbf{z})}{q_\phi(\mathbf{z})}$$

Reparameterisation

$$\mathbf{z} = g(\epsilon, \phi) \quad \epsilon \sim p(\epsilon)$$

Density Ratio

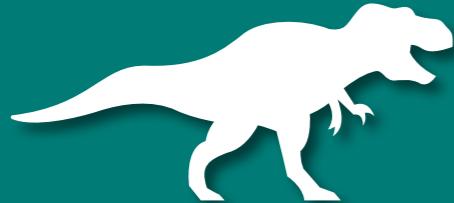
$$\frac{p^*(\mathbf{x})}{q(\mathbf{x})} = \frac{p(y=1|\mathbf{x})}{p(y=-1|\mathbf{x})}$$



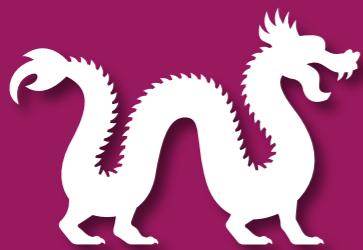
Planting the Seeds of Probabilistic Thinking

Part III: Algorithms

Learning Objectives



1. Have knowledge of different types of probabilistic models for unsupervised learning.



2. Understand generative algorithms (pixelCNN, VAEs, GANs) within the model-inference-algorithm framework.



3. Build awareness of the breadth of applications of generative models.

Beyond Classification

**Move beyond
associating inputs
to outputs**

**Understand and simulate
how the world evolves**

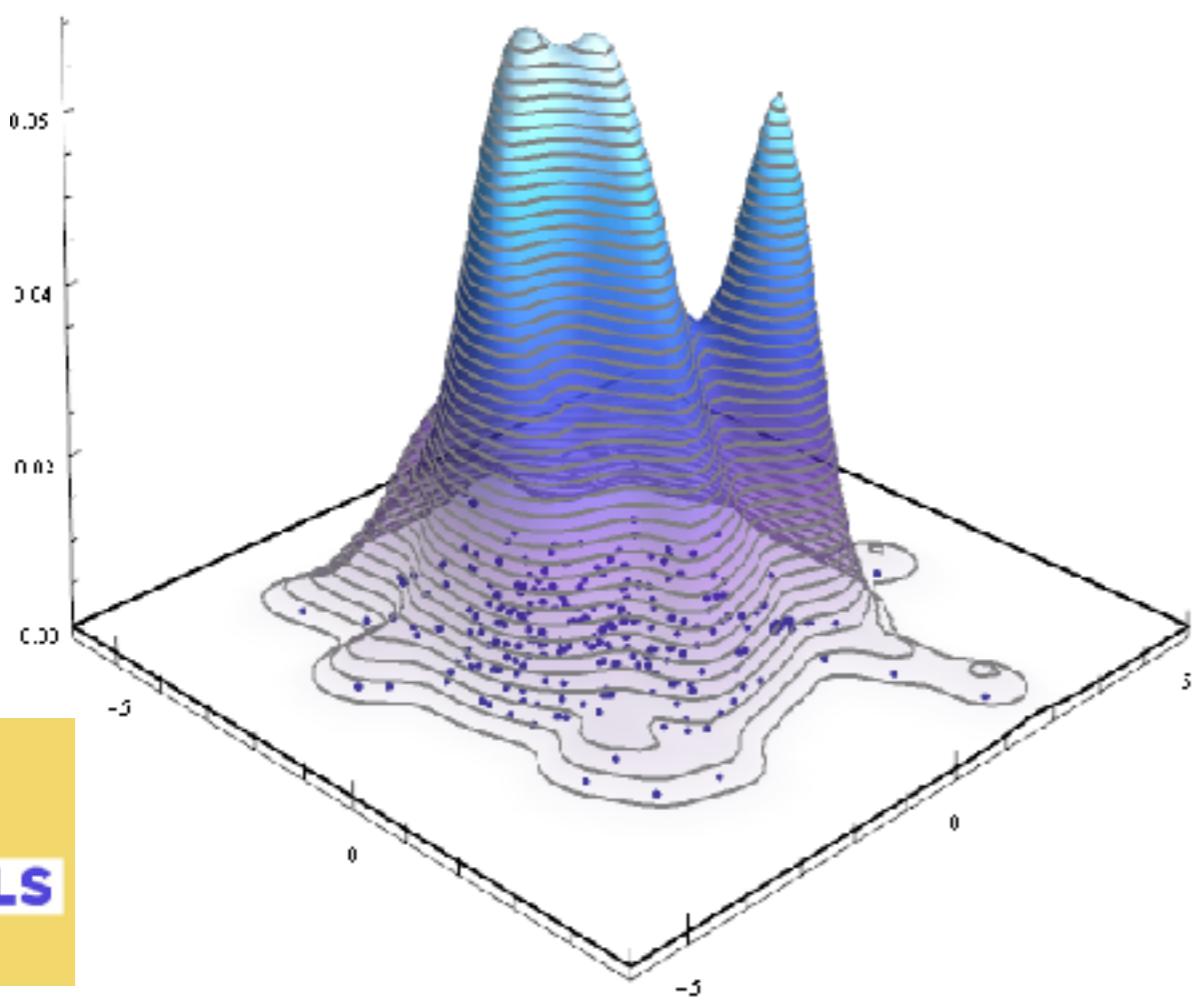
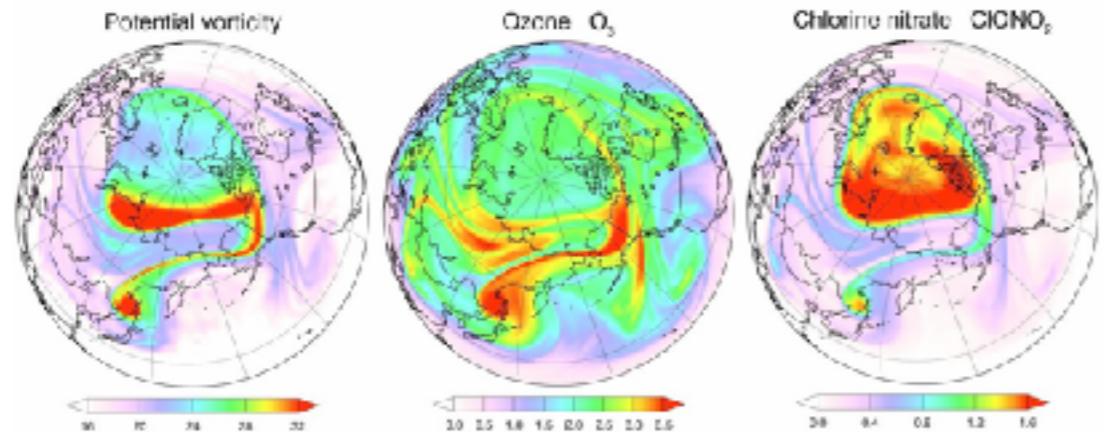
**Recognise objects in the
world and their factors
of variation**

**Detect surprising
events in the world**

**Establish concepts as
useful for reasoning and
decision making**

**Anticipate and
generate rich plans
for the future**

Generative Models



A model that allows us to learn a simulator of data

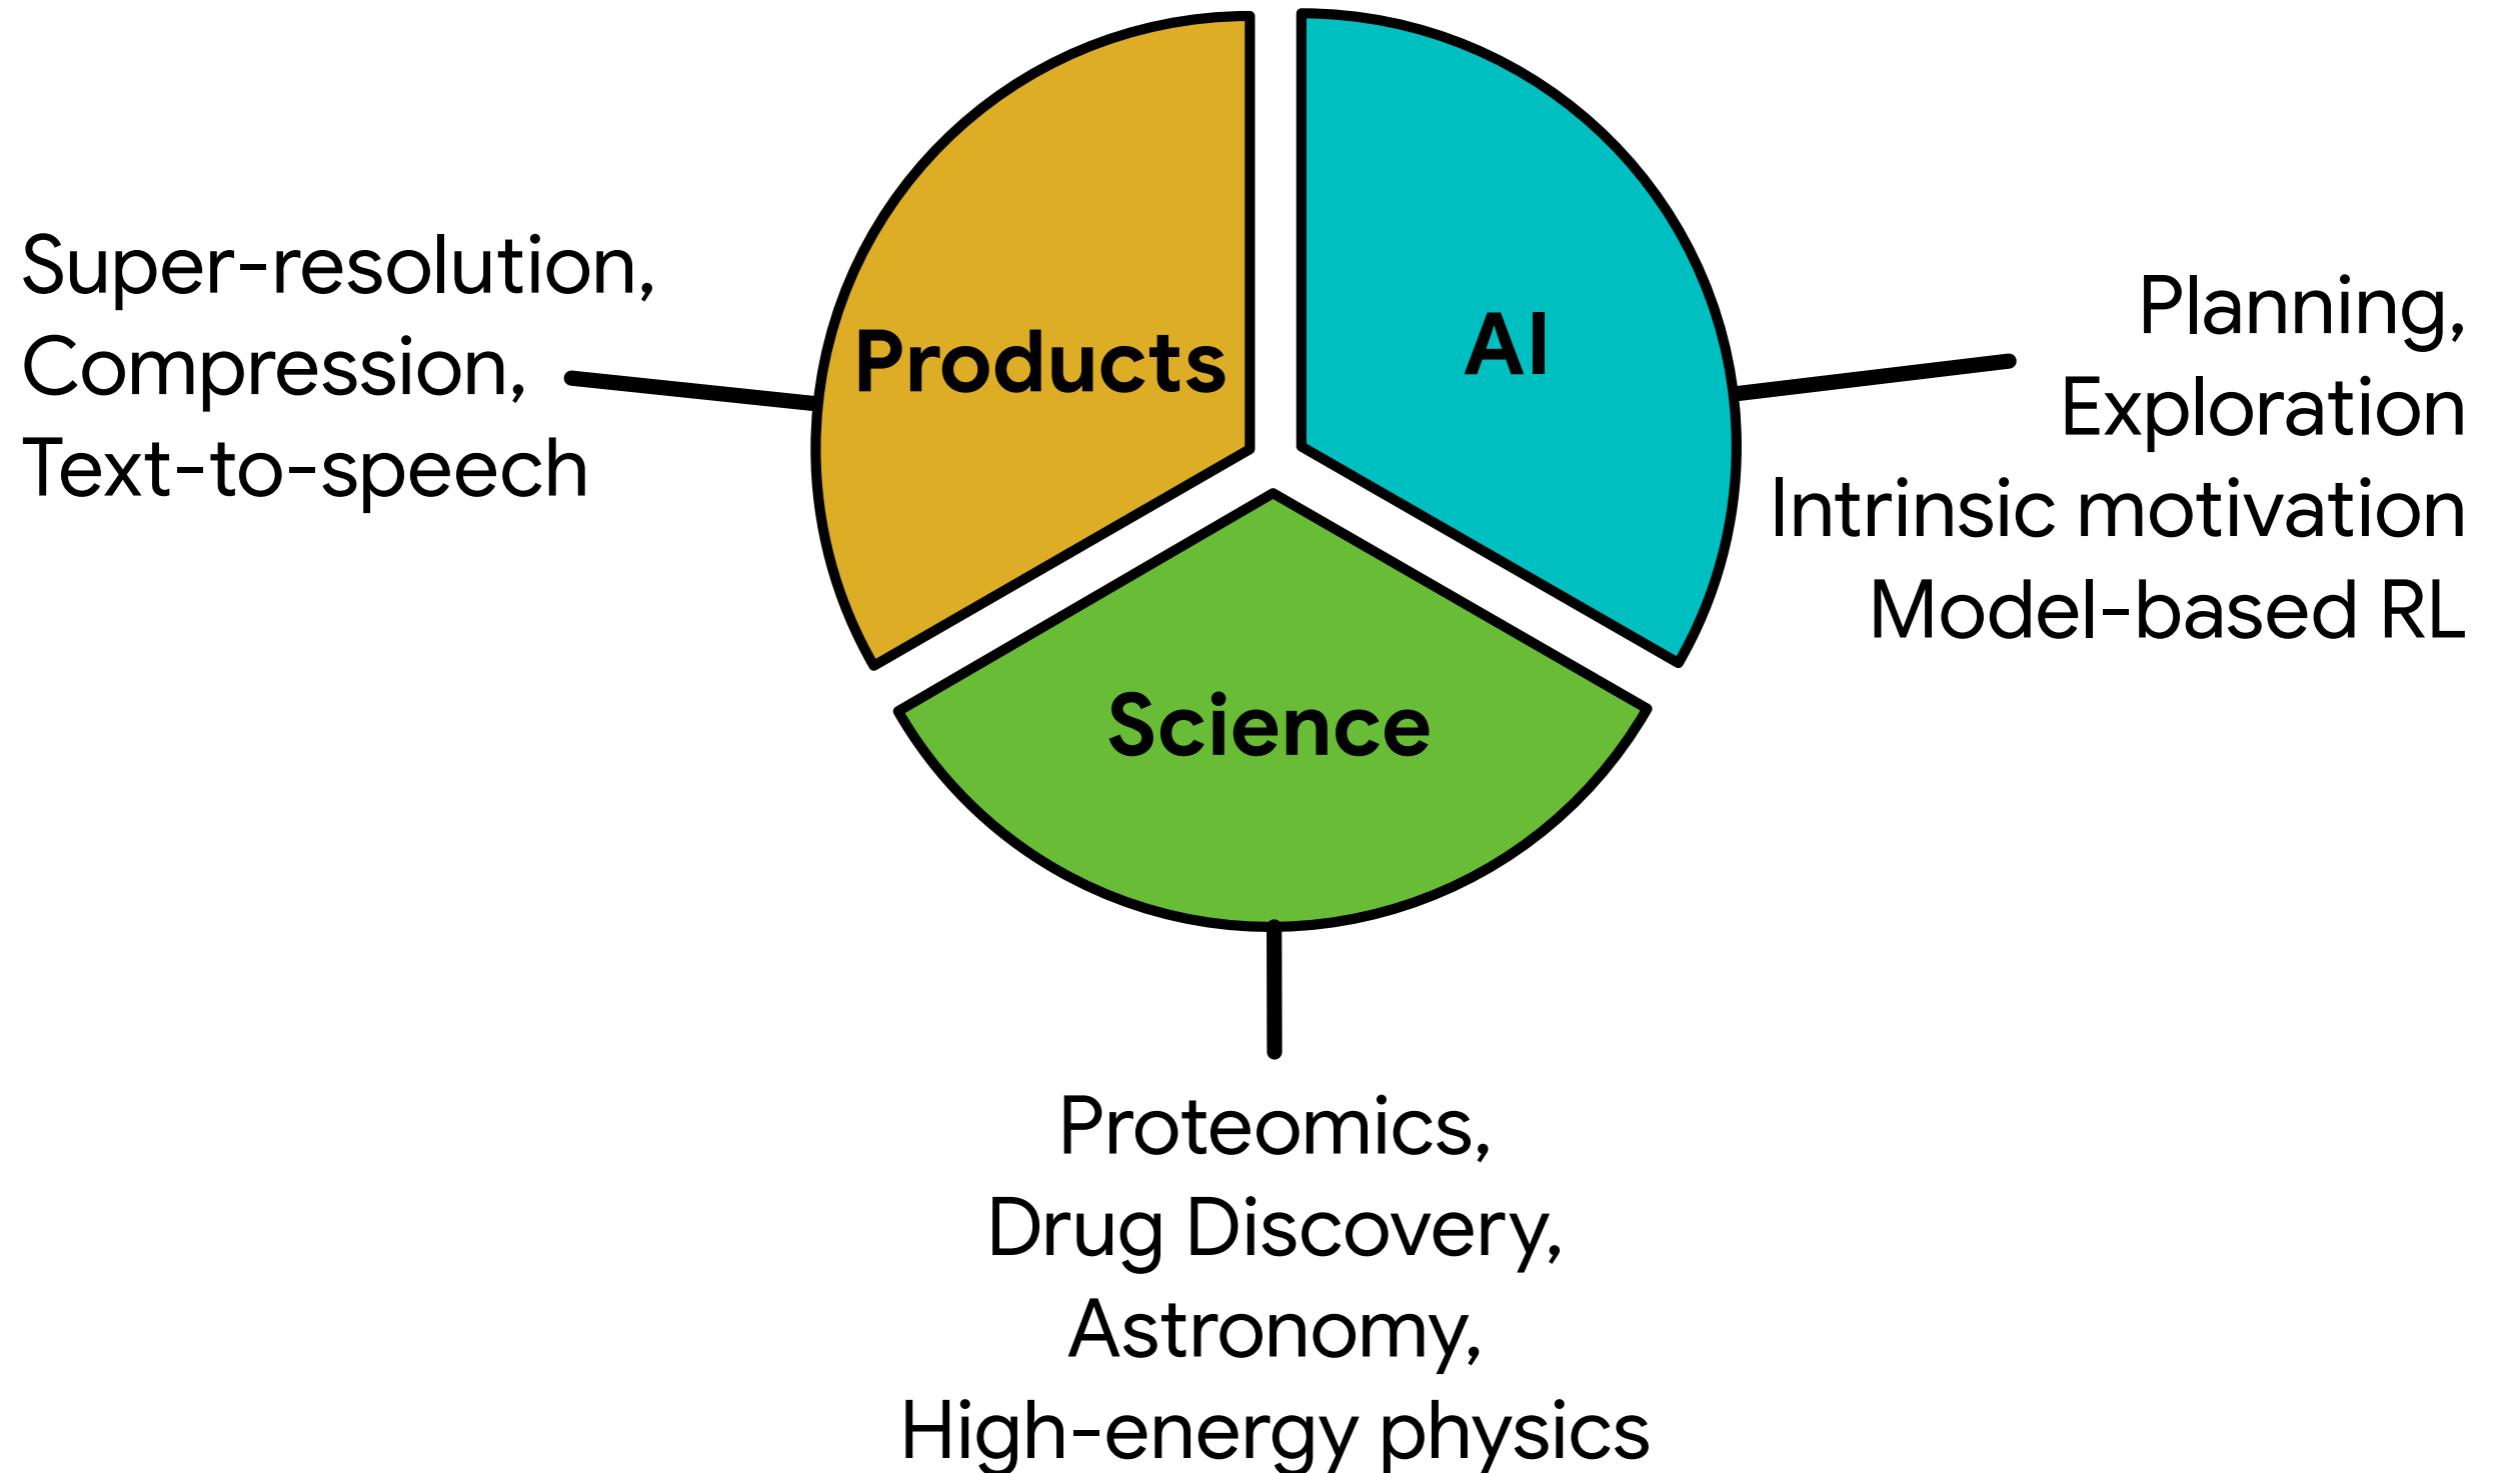
Models that allow for (conditional) density estimation

Approaches for unsupervised learning of data

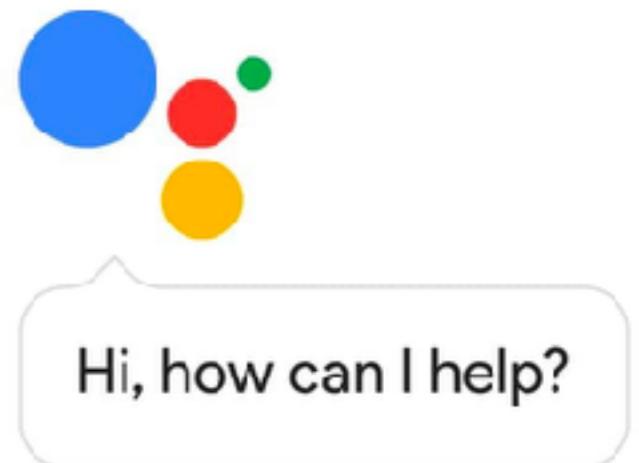
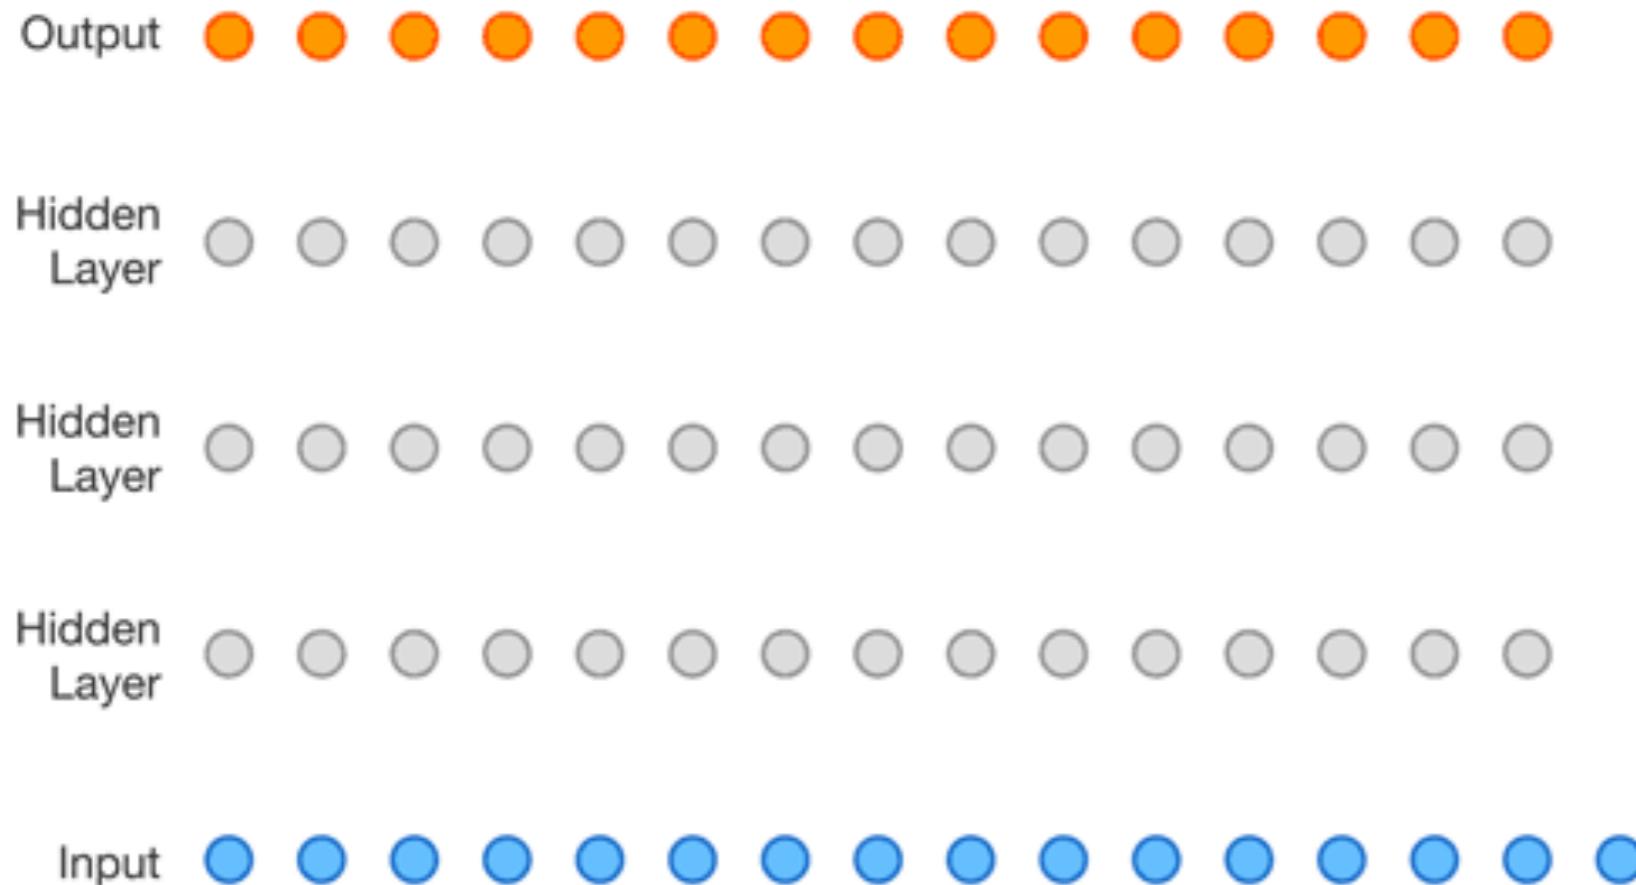
Characteristics are:

- **Probabilistic** models of data that allow for uncertainty to be captured.
- **High-dimensional** data.
- **Data distribution** is targeted.

Applications

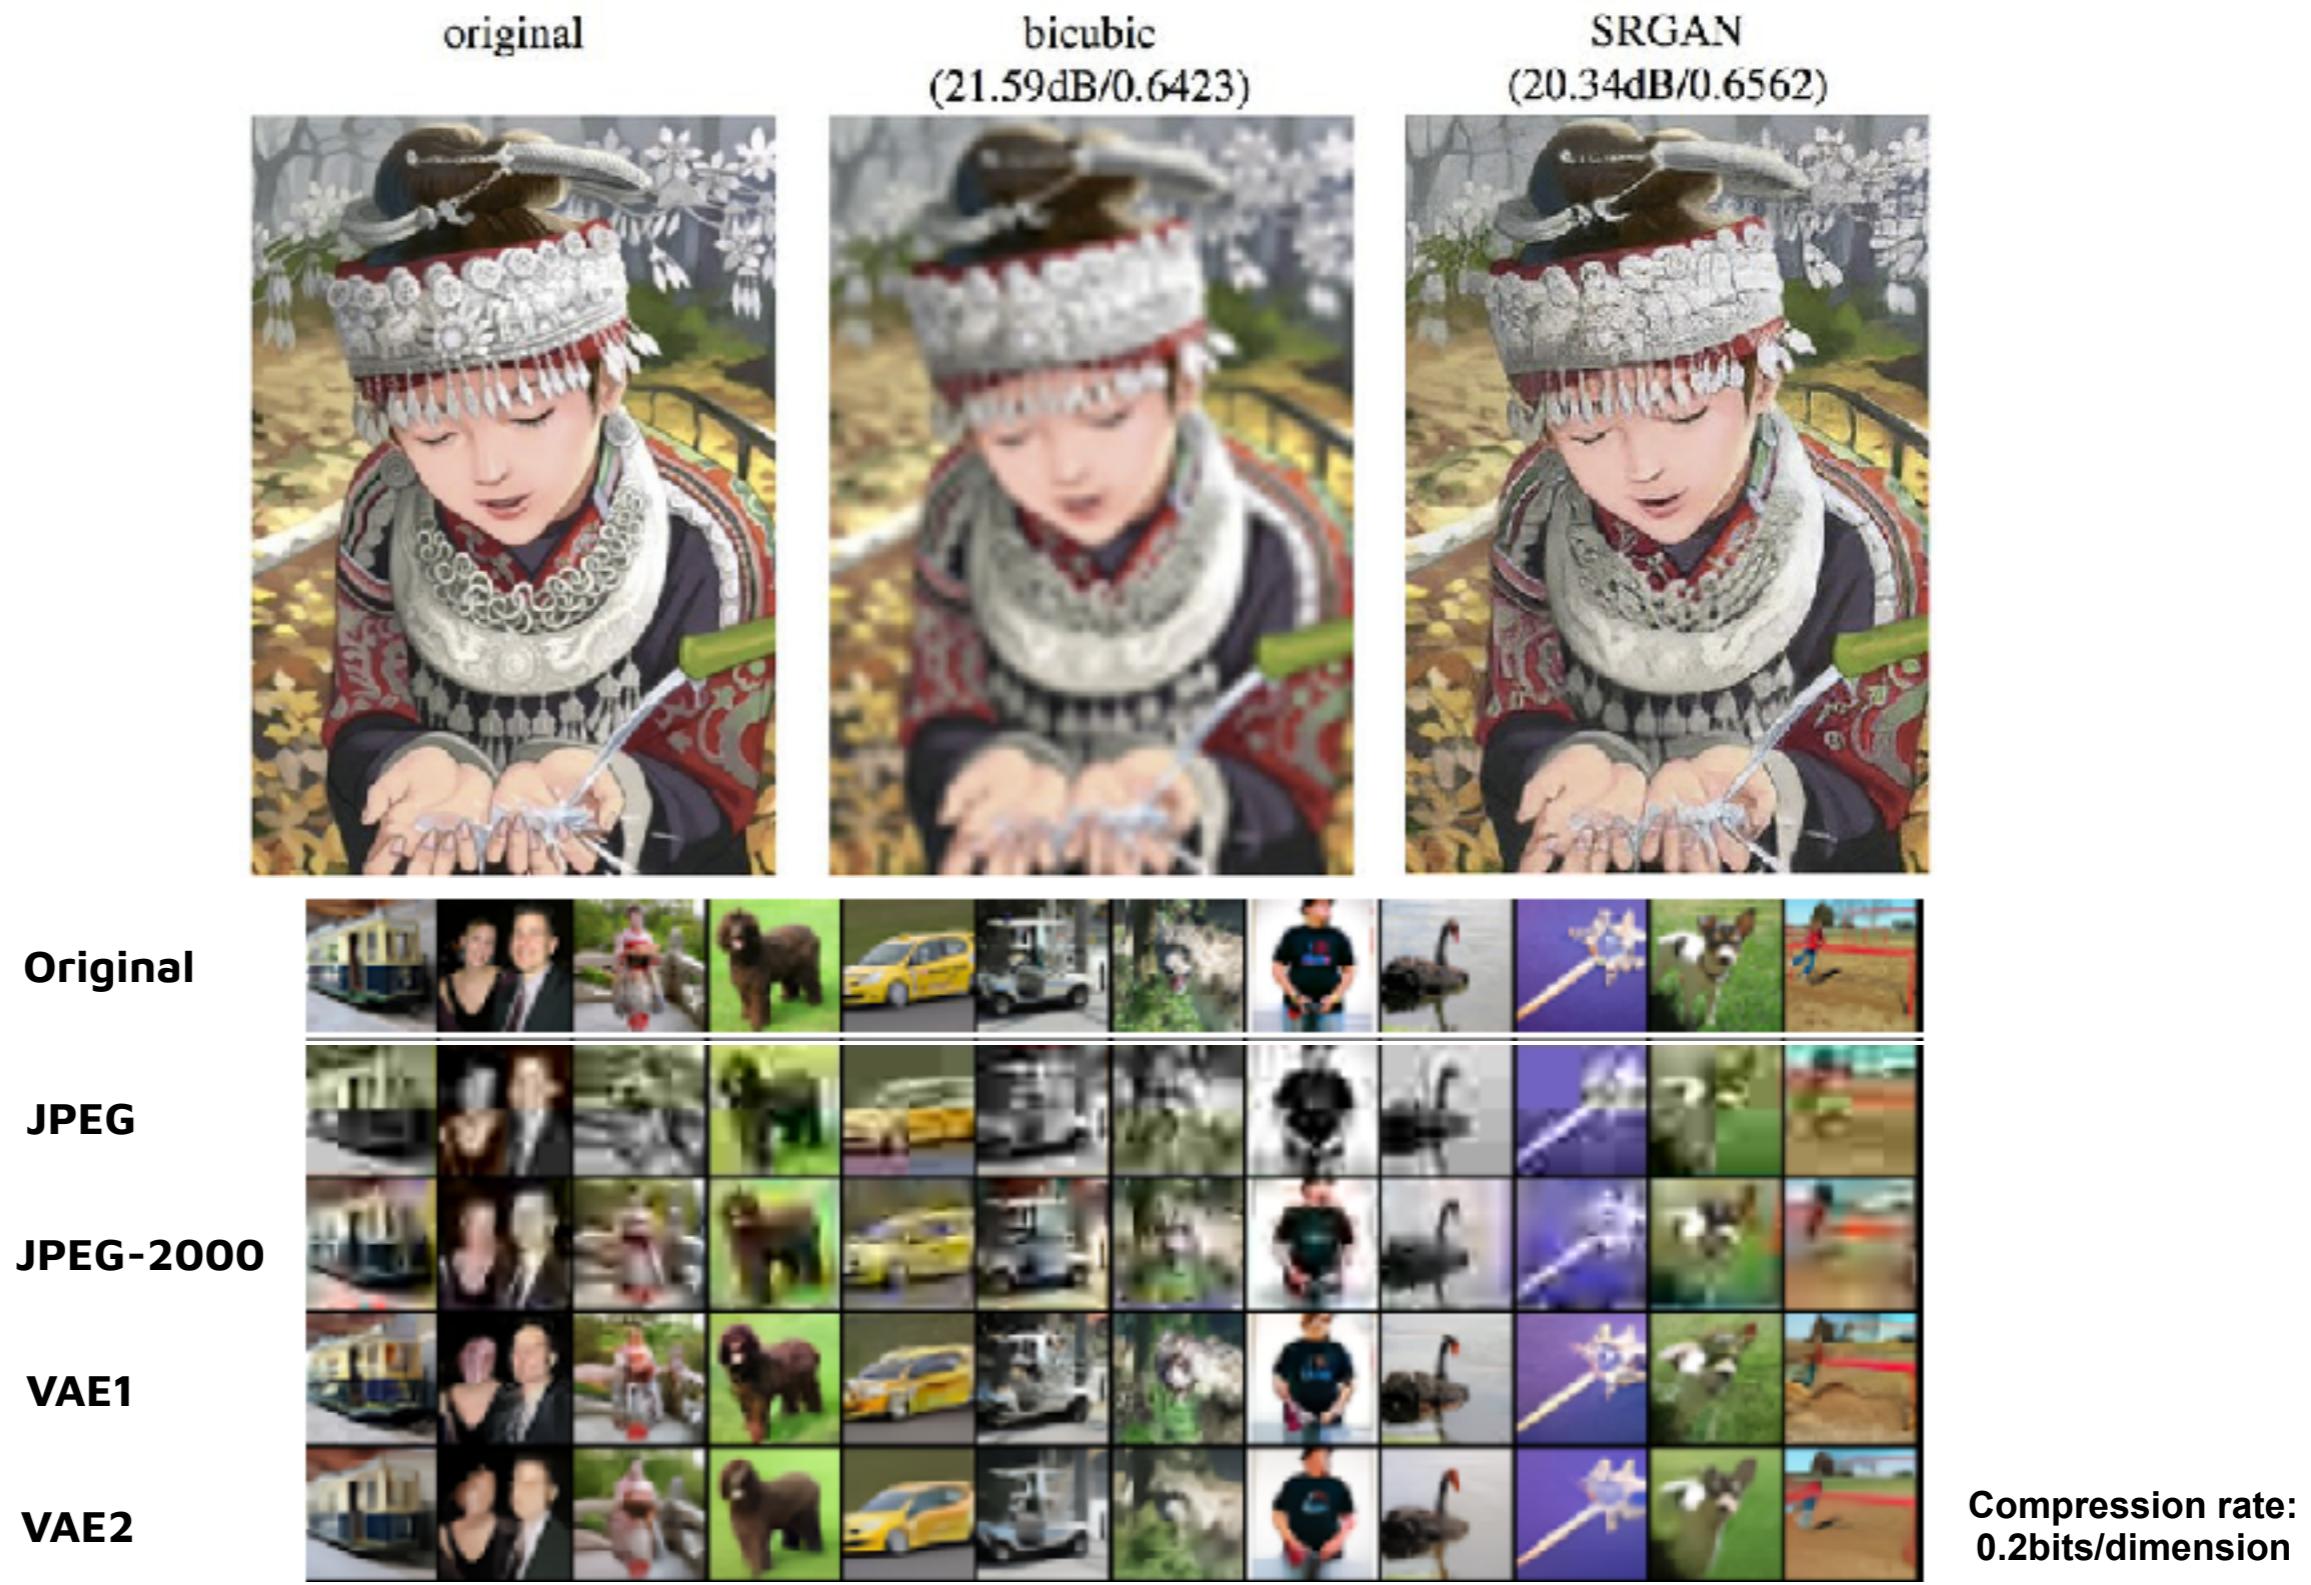


Assistive Technologies

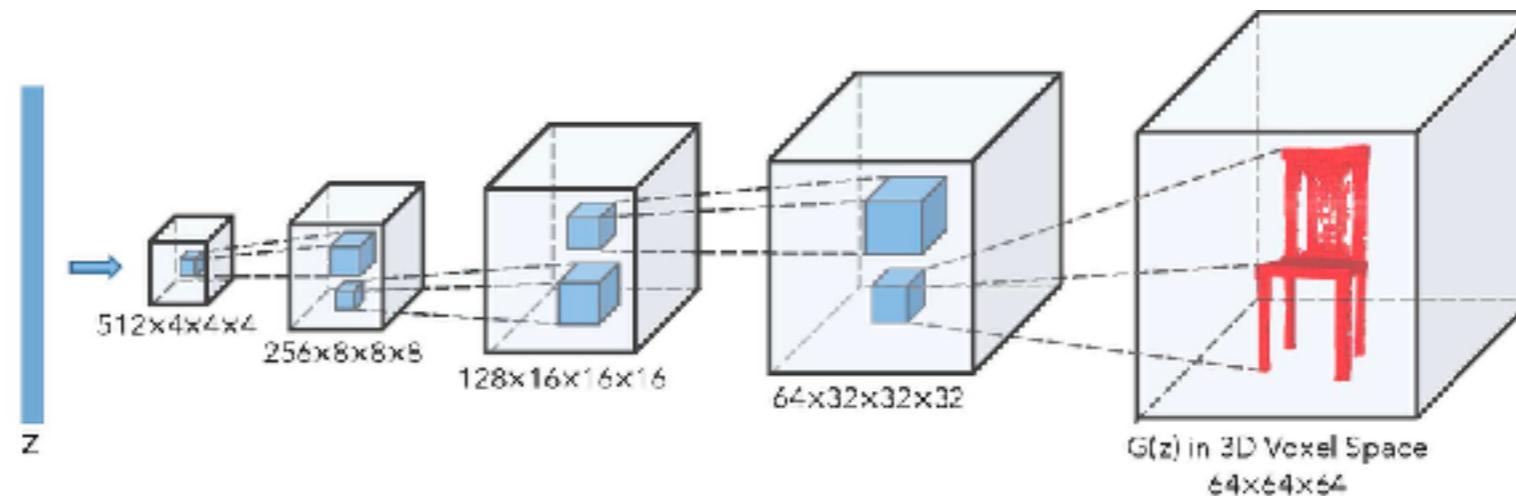


Fully-observed conditional generative model

Compression-Communication



Generative Design



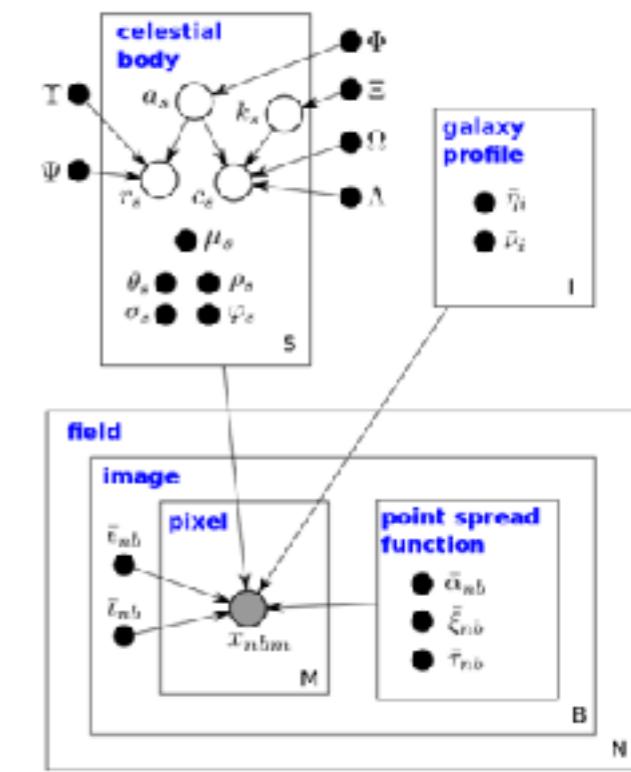
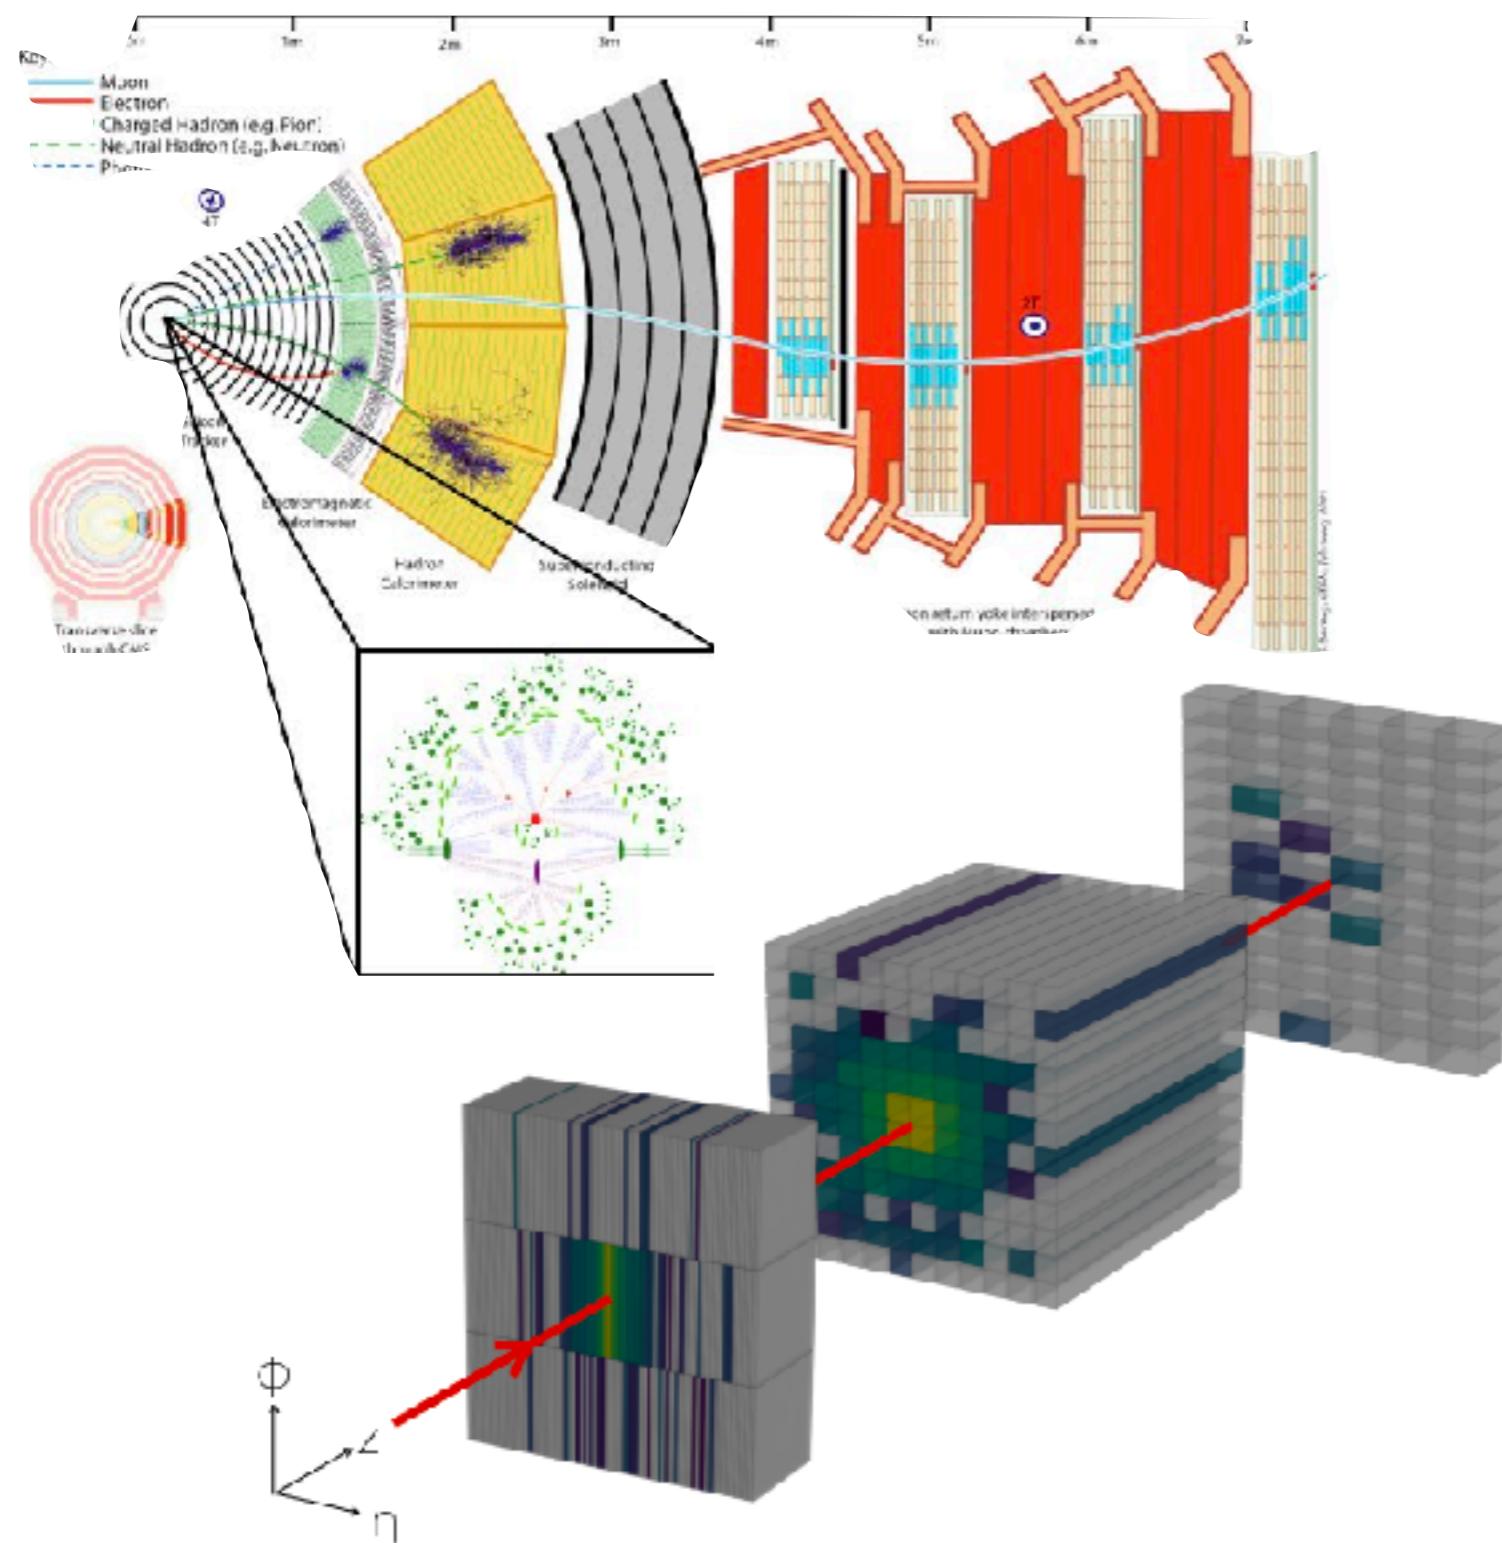


Video from work of Memo Aktem

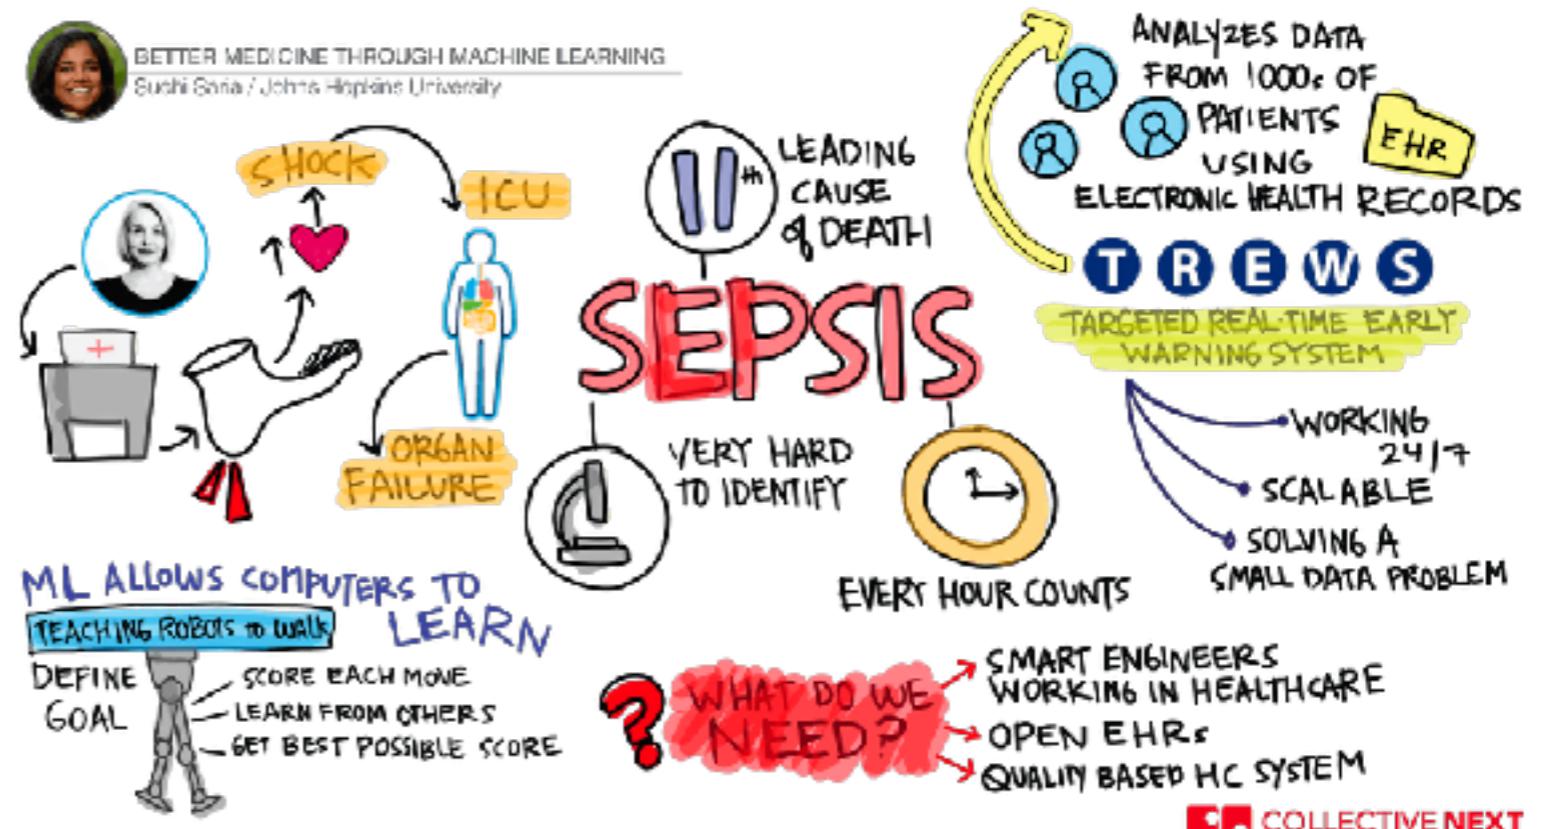
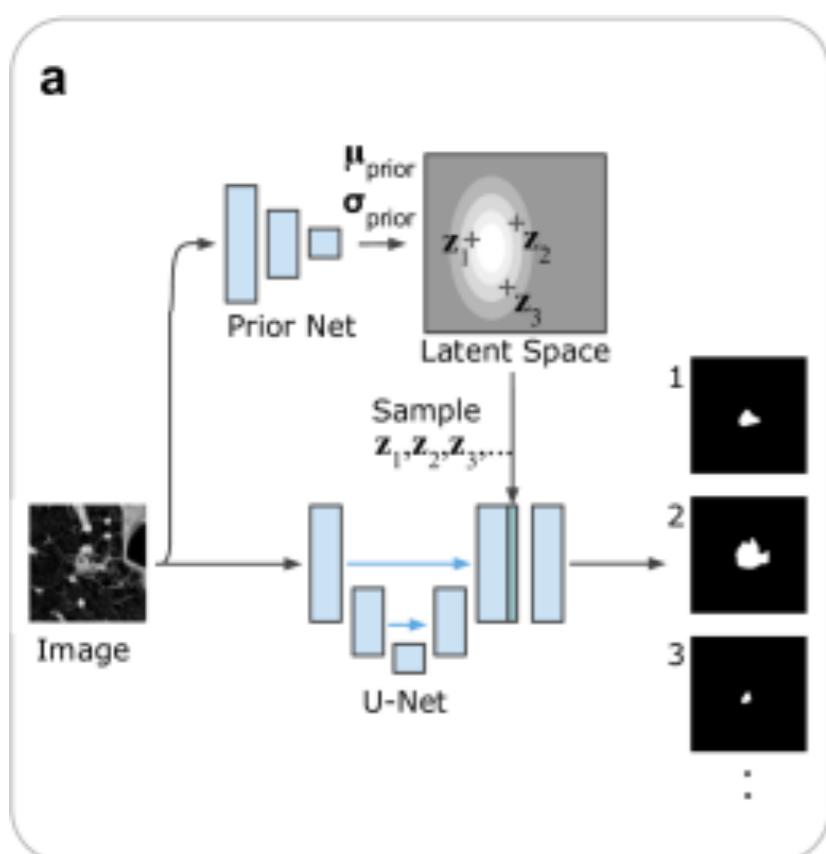
Shakir Mohamed

 DeepMind

Advancing Science

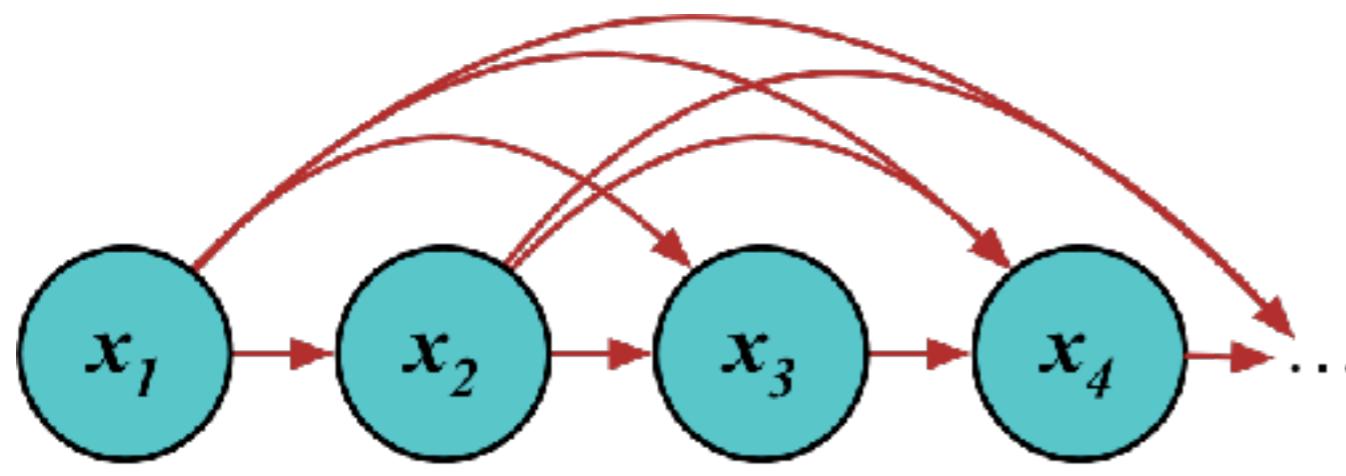


Advancing Healthcare

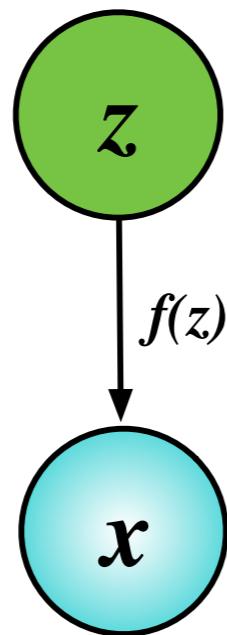


Types of Generative Models

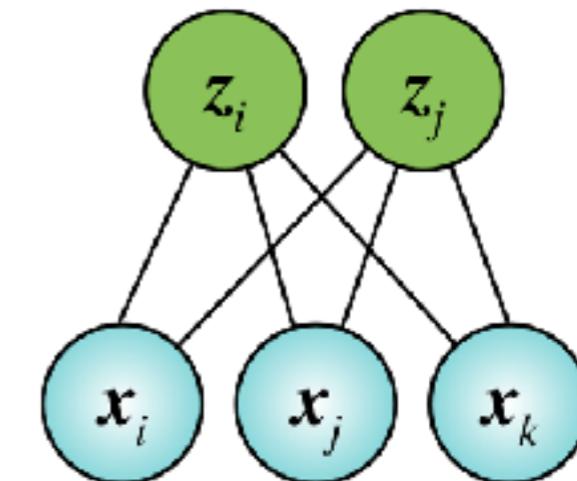
Fully-observed models



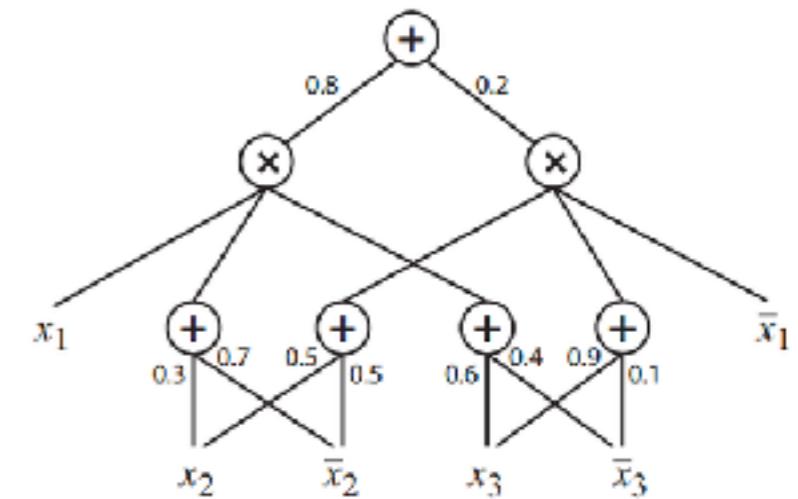
Latent variable models



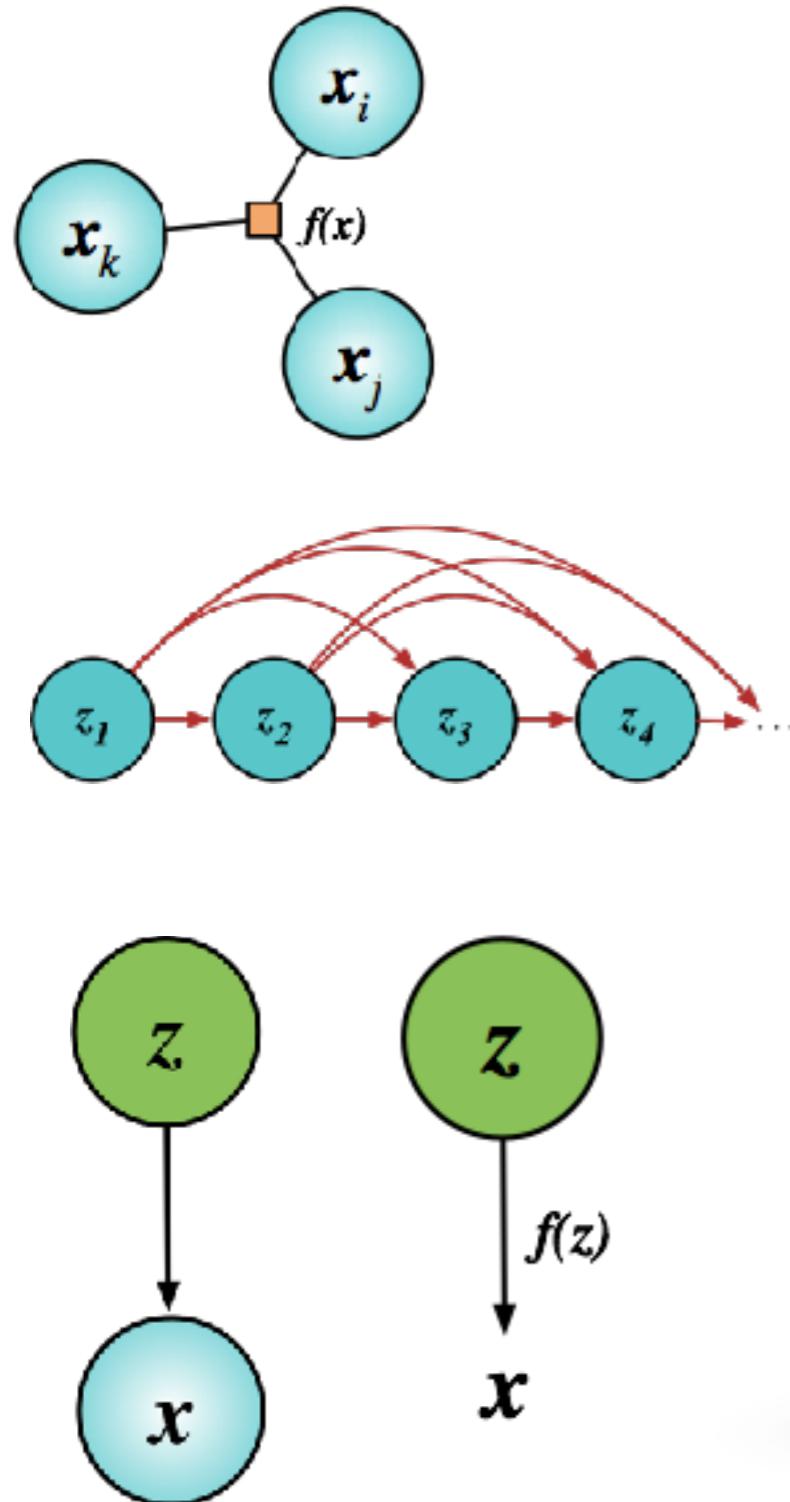
Undirected Models



Sum-Product Networks



Types of Generative Models

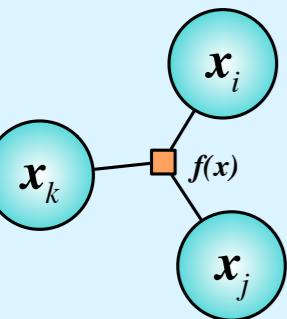


Design Dimensions

- ❖ **Data:** binary, real-valued, nominal, strings, images.
 - ❖ **Dependency:** independent, sequential, temporal, spatial.
 - ❖ **Representation:** continuous or discrete
 - ❖ **Dimension:** parametric or non-parametric
-
- ❖ Computational complexity
 - ❖ Modelling capacity
 - ❖ Bias, uncertainty, calibration
 - ❖ Interpretability

Fully-observed Models

Fully-observed models



Model observed data directly **without** introducing any new unobserved **local variables**.

Model Parameters are **global variables**.

Stochastic activations & unobserved random variables are **local variables**.

Markov Models

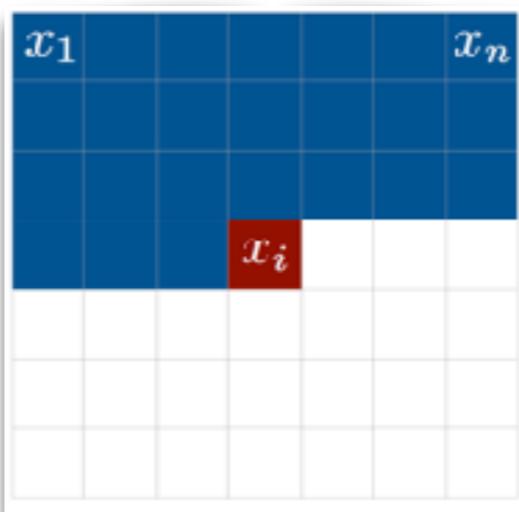
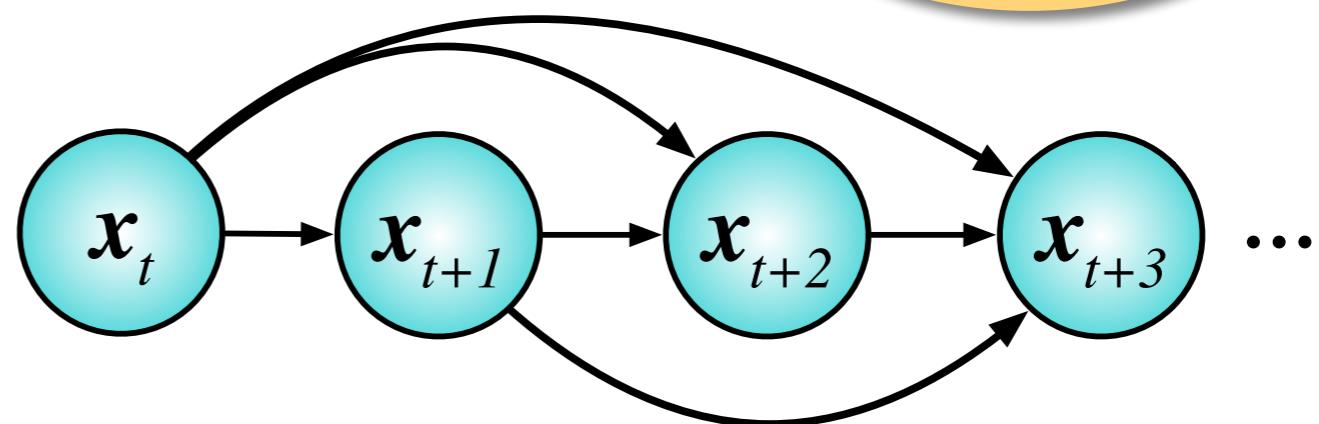
$$x_1 \sim \text{Cat}(x_1 | \pi)$$

$$x_2 \sim \text{Cat}(x_2 | \pi(\mathbf{x}_1))$$

...

$$x_i \sim \text{Cat}(x_i | \pi(\mathbf{x}_{<n}))$$

$$p(\mathbf{x}) = \prod_i p(x_i | f(\mathbf{x}_{<i}; \theta))$$



All conditional probabilities described by deep networks.

Hartebeest



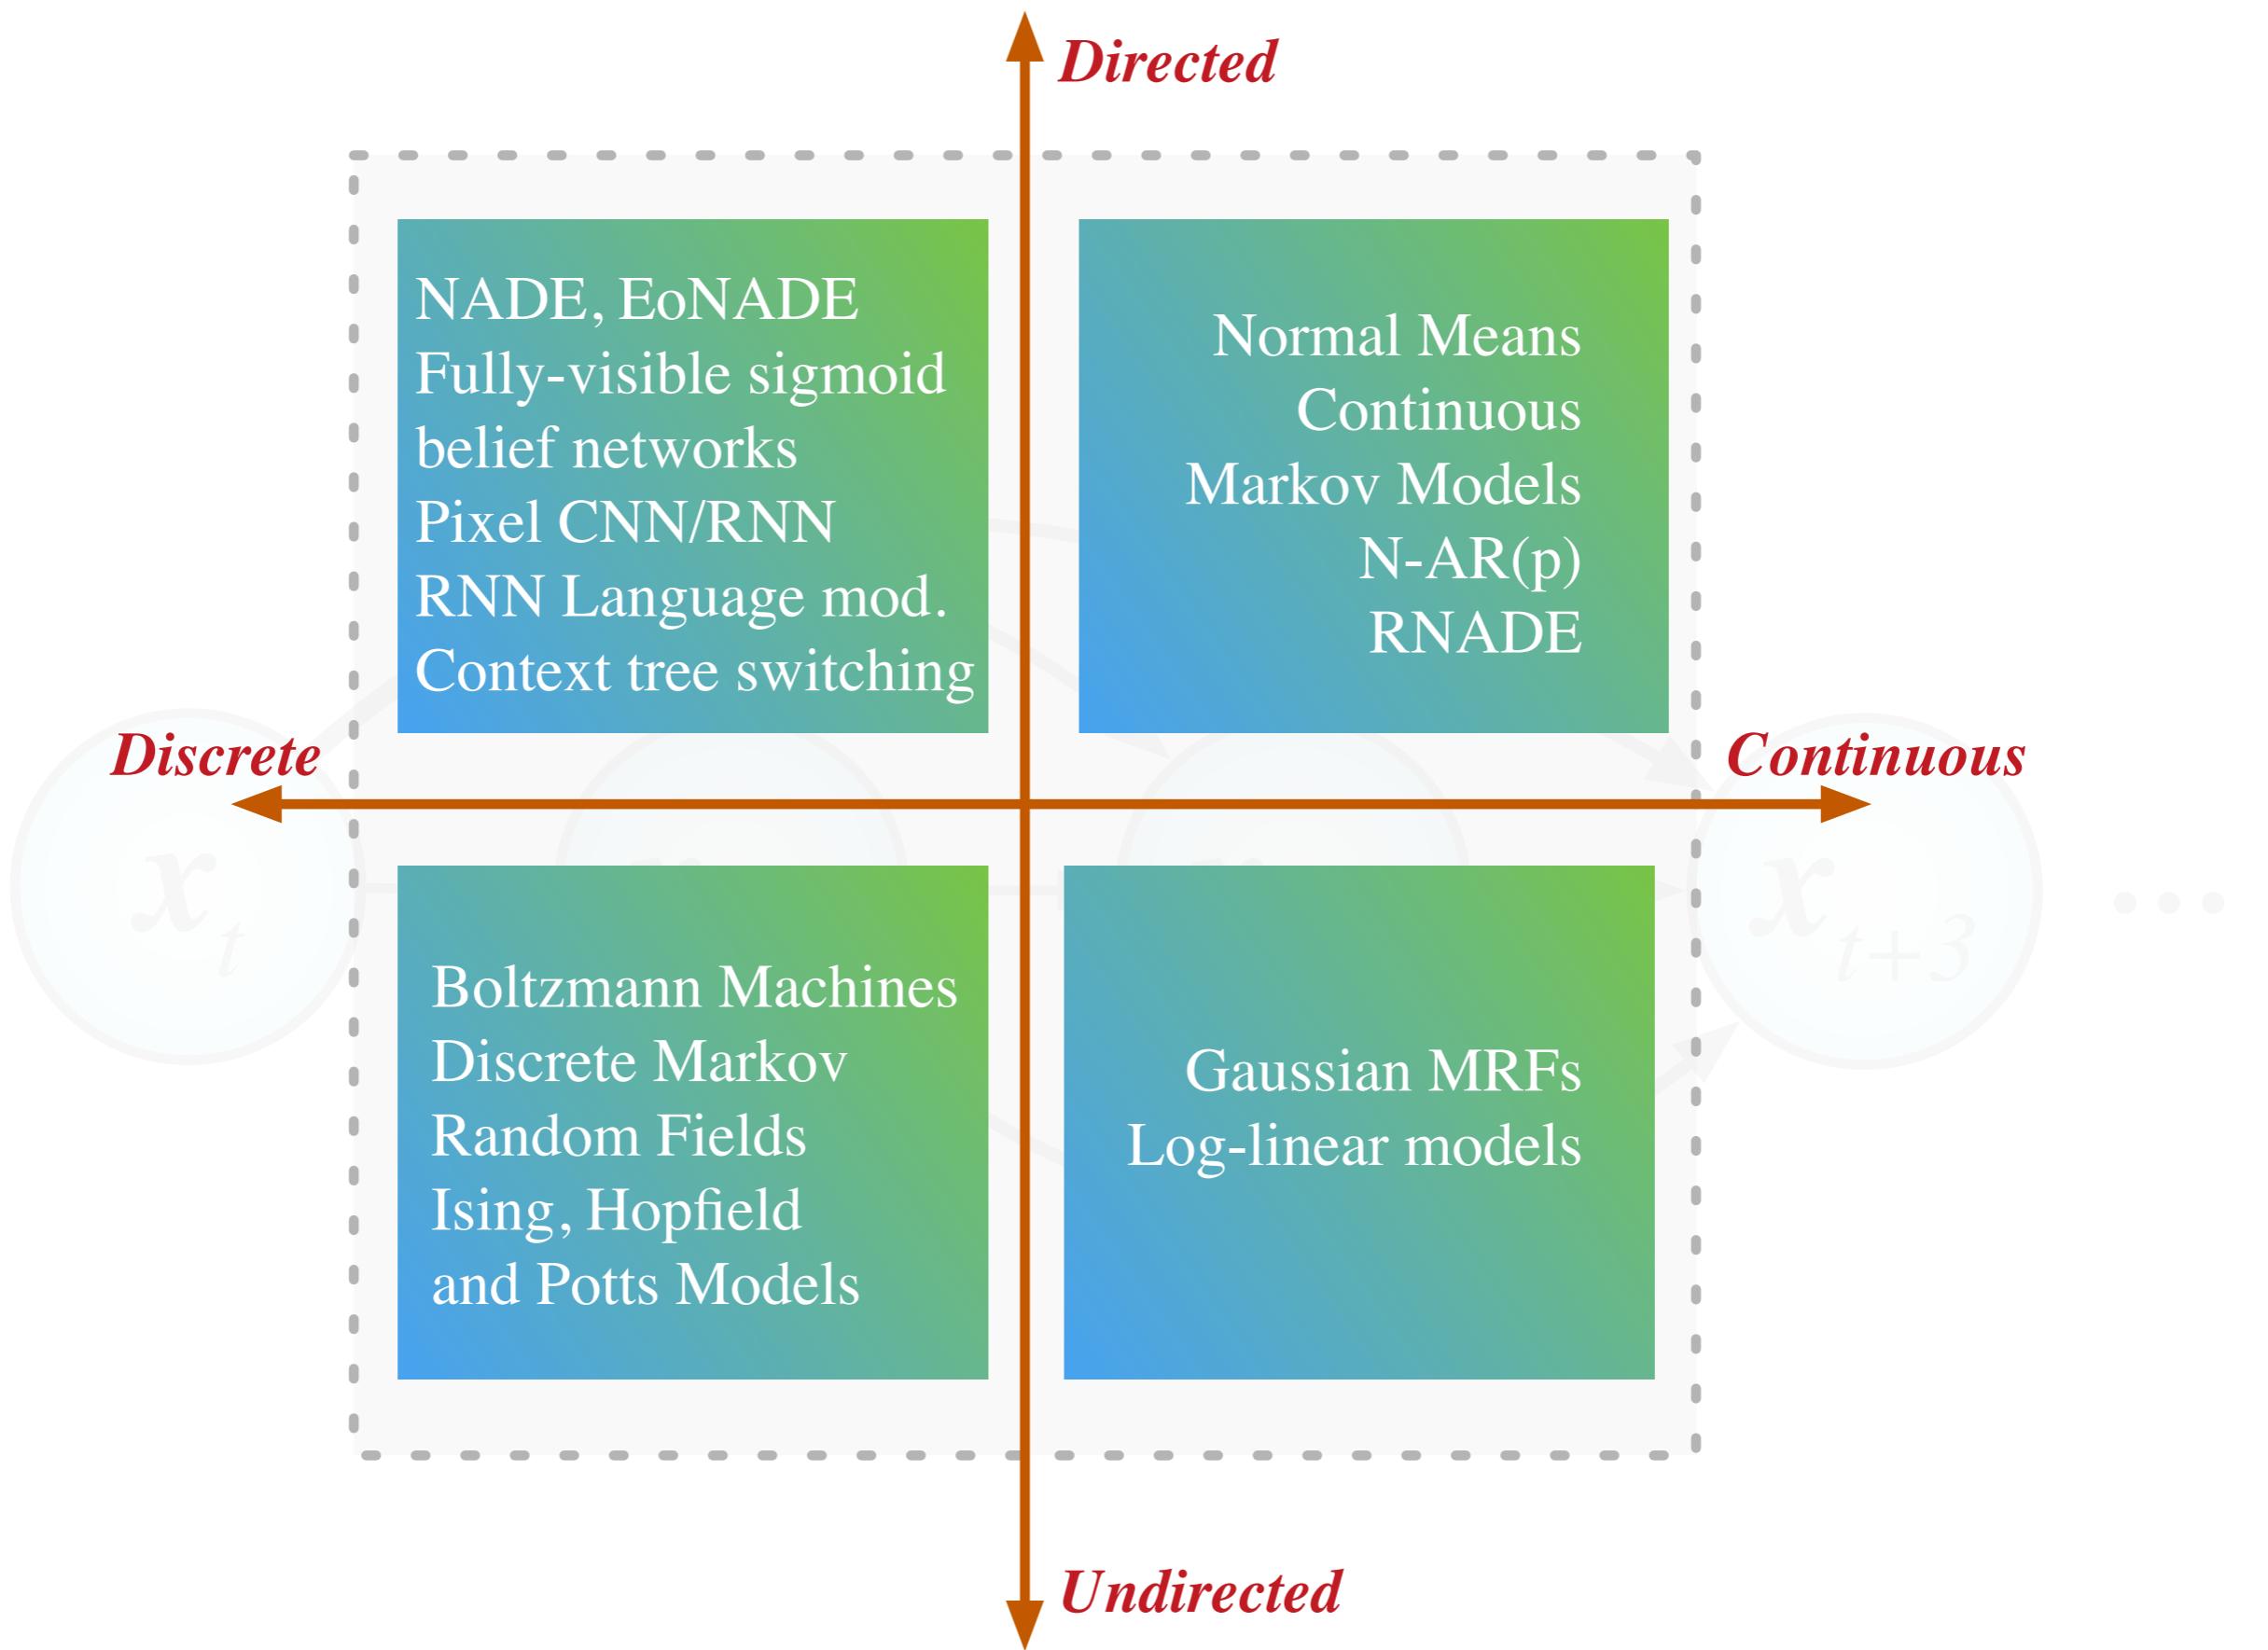
Pixel CNN

White Whale

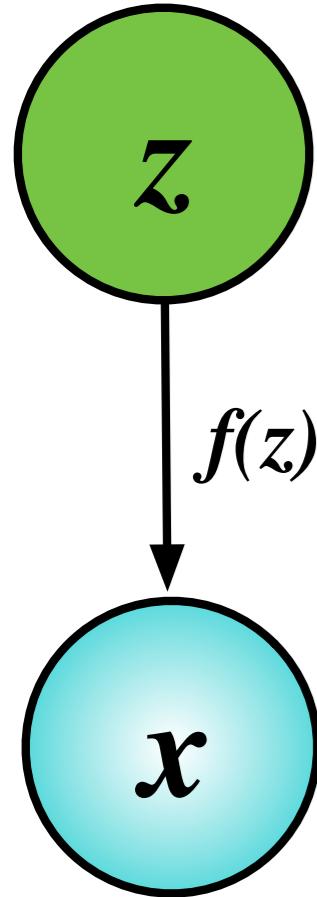


Properties

- + Can directly encode how observed points are related.
- + Any data type can be used
- + For directed graphical models:
 - + **Parameter learning simple:** Log-likelihood is directly computable, no approximation needed.
 - + Easy to scale-up to large models, many optimisation tools available.
 - Order sensitive.
- For undirected models,
 - **Parameter learning difficult:** Need to compute normalising constants.
 - **Generation can be slow:** iterate through elements sequentially, or using a Markov chain.

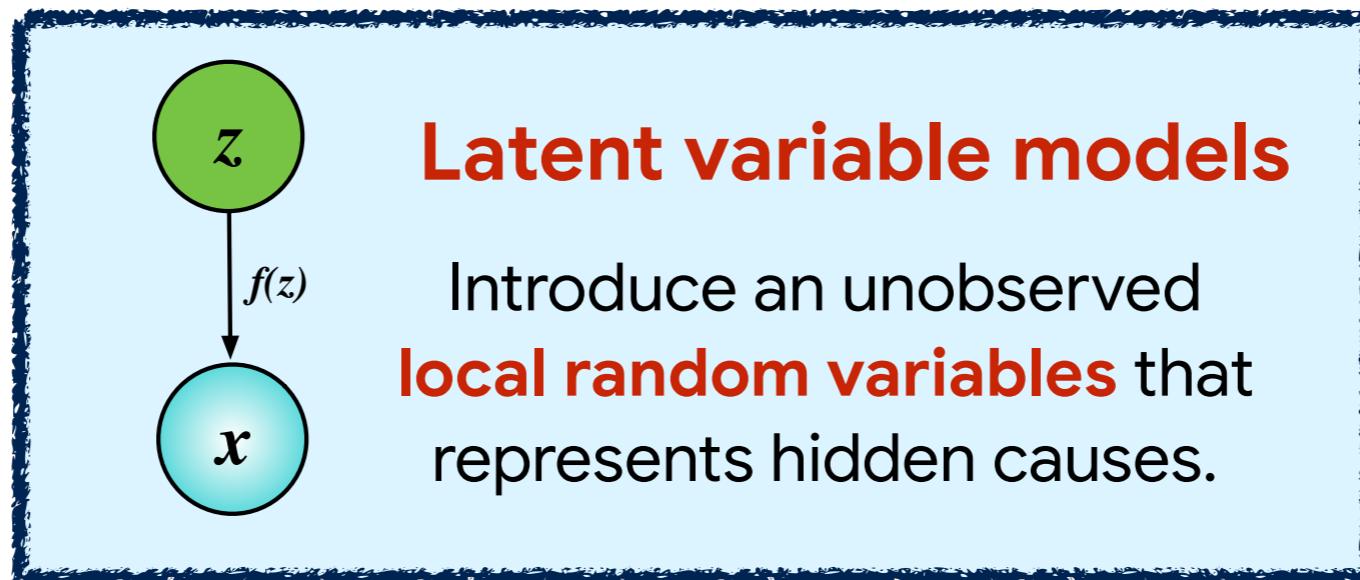


Latent Variable Models



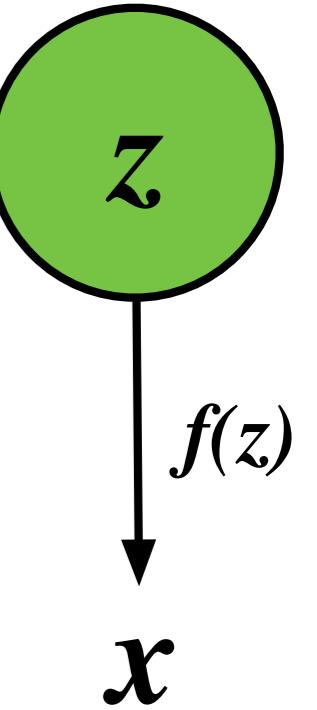
Prescribed models

Use observer likelihoods and assume observation noise.



Implicit models

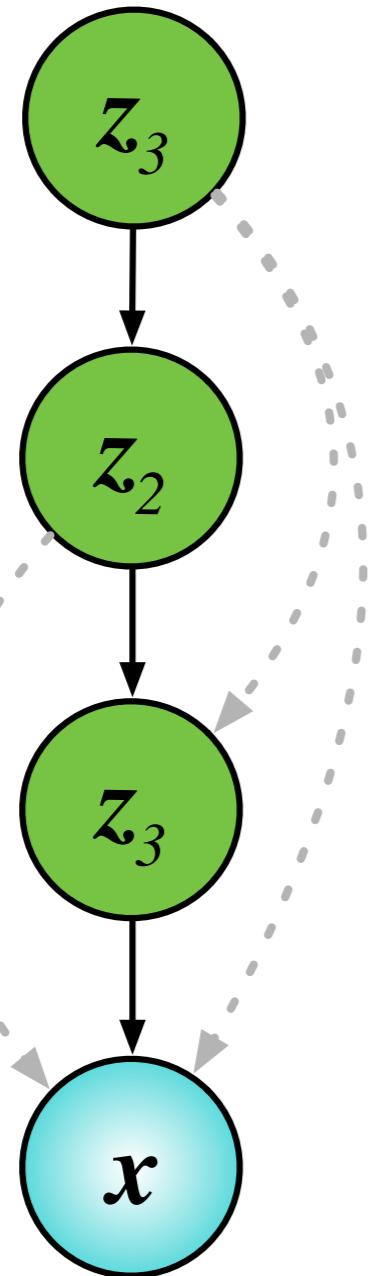
Likelihood-free or simulation-based models.



Diggle and Gratton (1984); Mohamed and Lakshminarayanan (2016)

Prescribed Models

Deep Latent Gaussian Model



$$\mathbf{z}_3 \sim \mathcal{N}(\mathbf{0}, \mathbf{I})$$

$$\mathbf{z}_2 | \mathbf{z}_3 \sim \mathcal{N}(\mu(\mathbf{z}_3), \Sigma(\mathbf{z}_3))$$

$$\mathbf{z}_1 | \mathbf{z}_2 \sim \mathcal{N}(\mu(\mathbf{z}_2), \Sigma(\mathbf{z}_2))$$

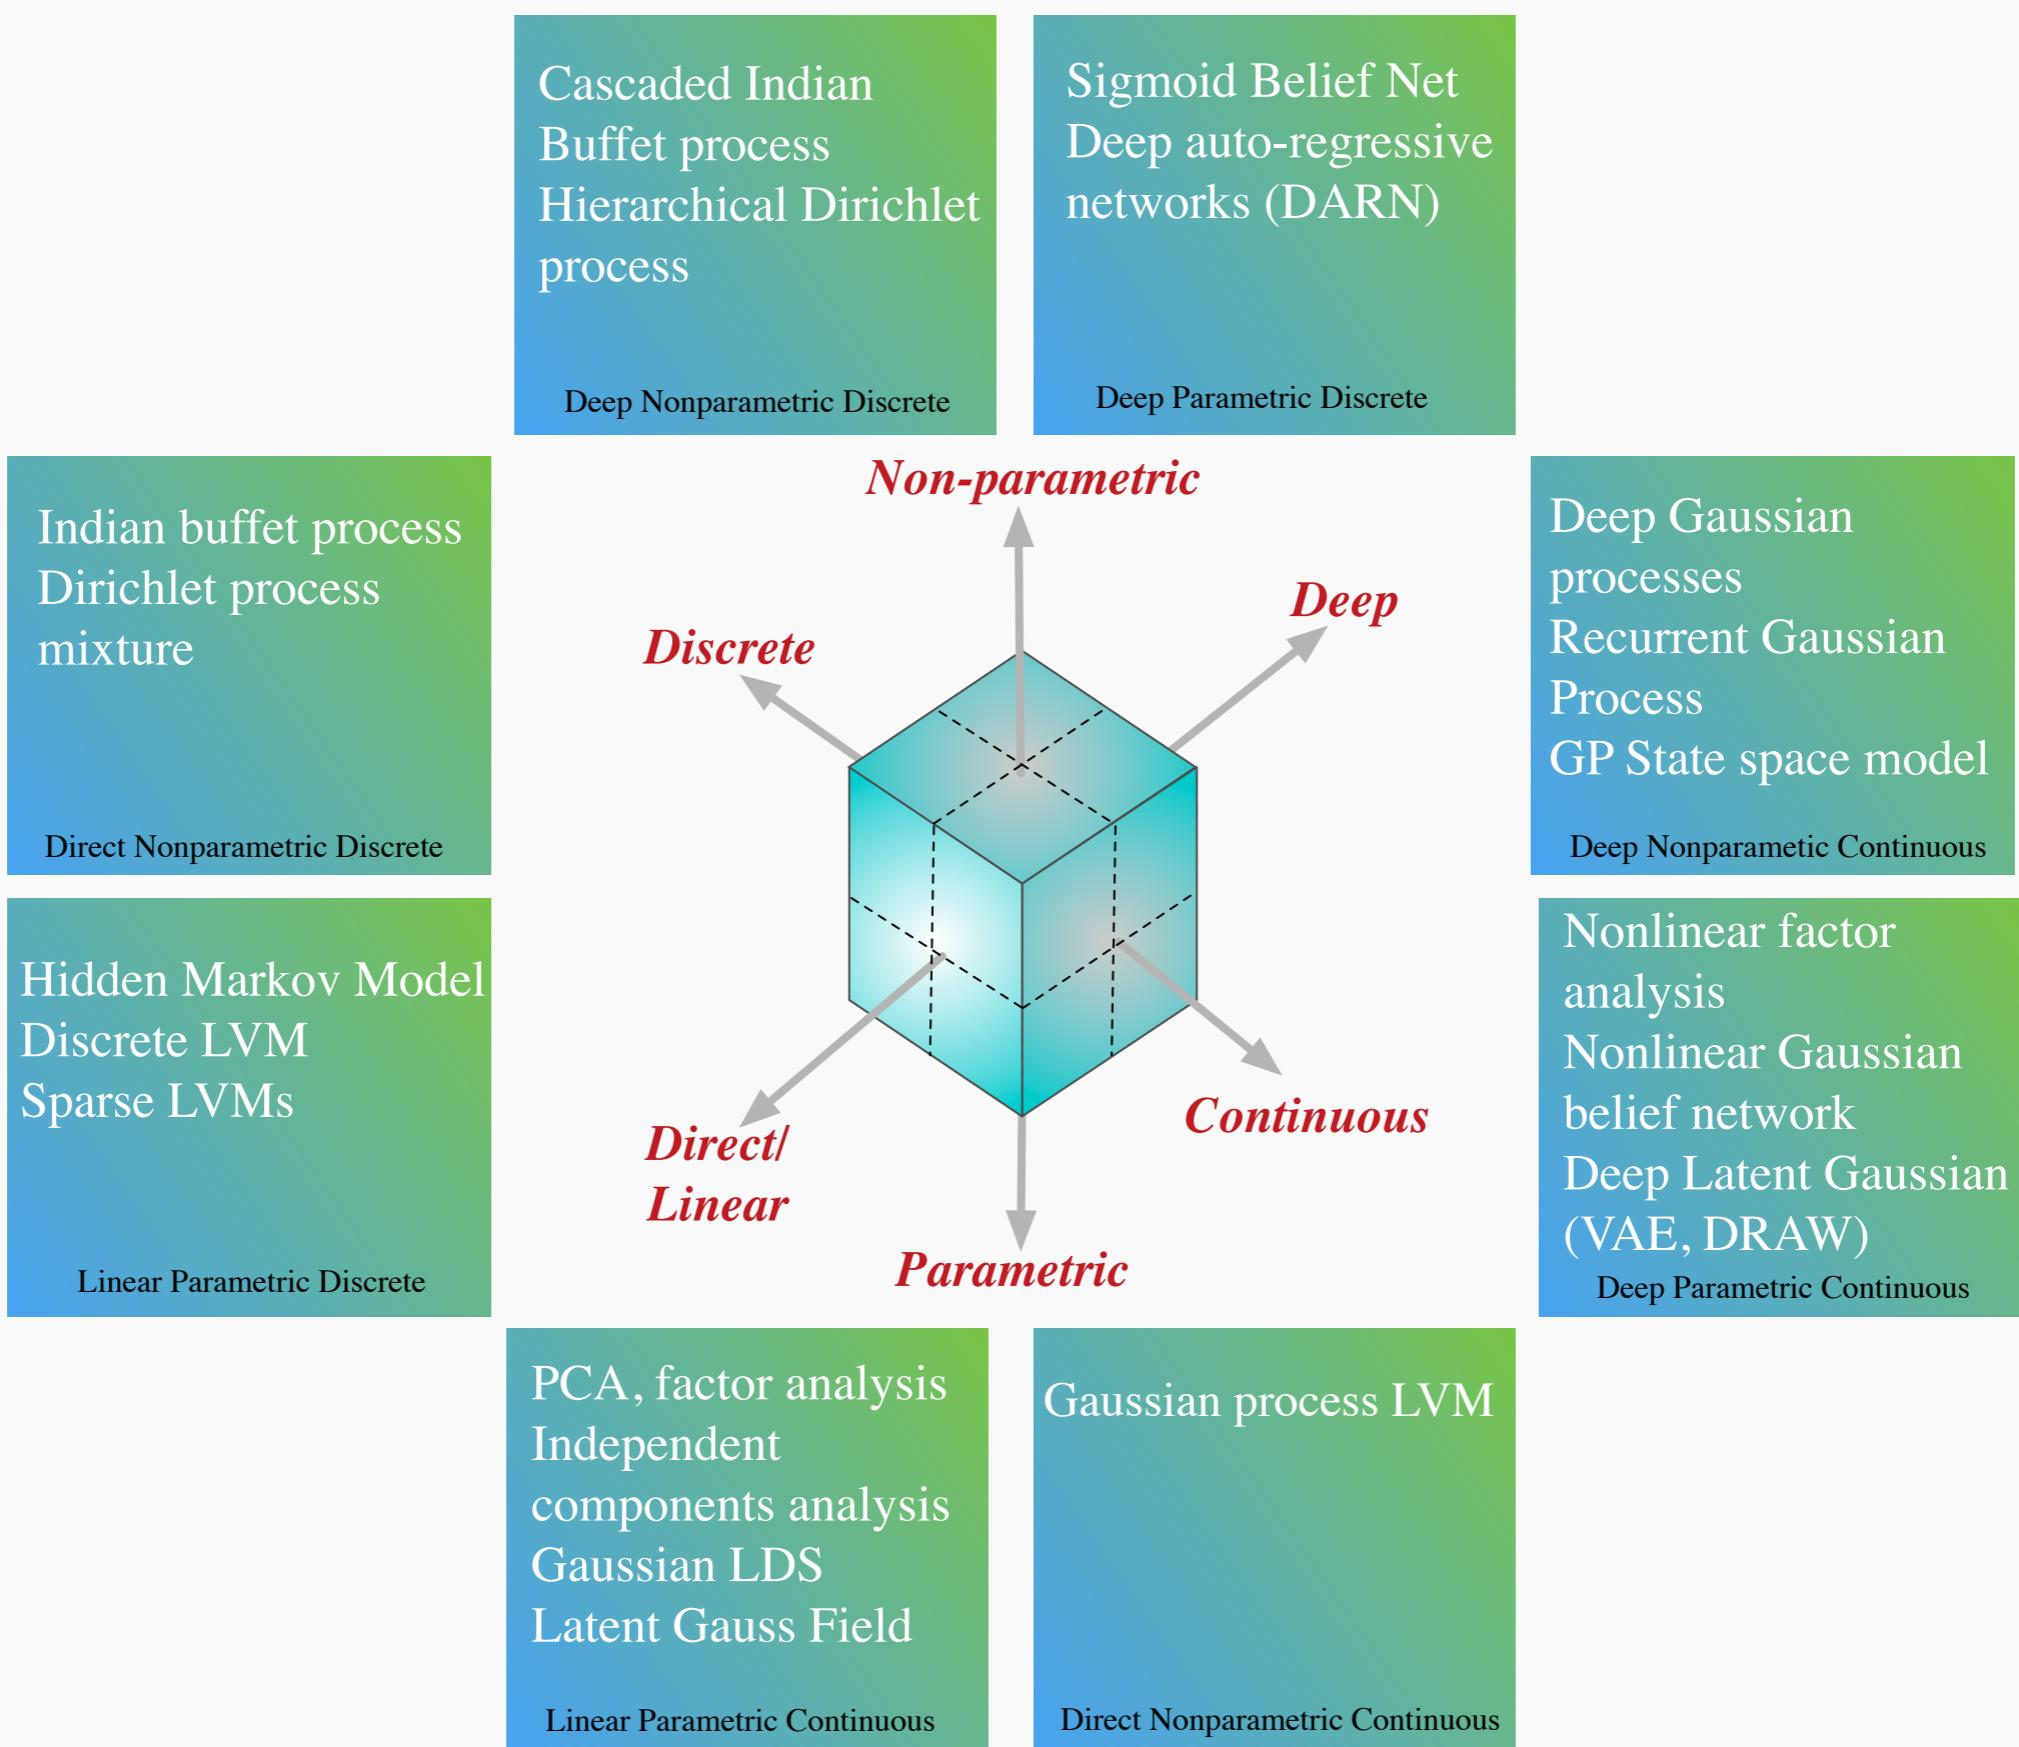
$$\mathbf{x} | \mathbf{z}_1 \sim \mathcal{N}(\mu(\mathbf{z}_1), \Sigma(\mathbf{z}_1))$$

*Convolutional
DRAW*



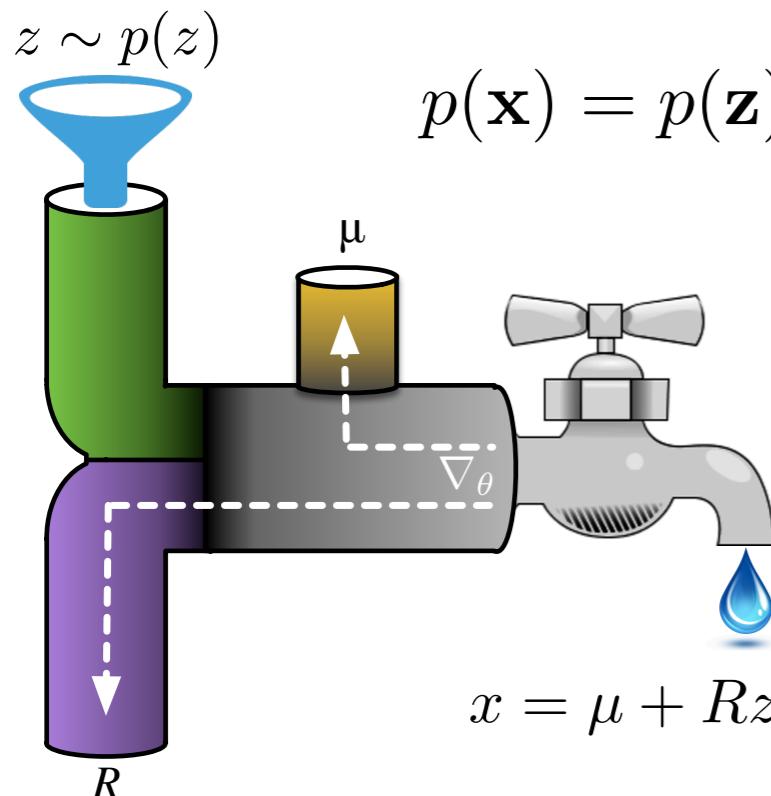
Properties

- + Easy sampling.
- + Easy way to include hierarchy and depth.
- + Easy to encode structure believed to generate the data
- + Avoids order dependency assumptions: marginalisation of latent variables induces dependencies.
- + Latents provide compression and representation the data.
- + Scoring, model comparison and selection possible using the marginalised likelihood.
- Inversion process to determine latents corresponding to a input is difficult in general
- Difficult to compute marginalised likelihood requiring approximations.
- Not easy to specify rich approximations for latent posterior distribution.



Implicit Models

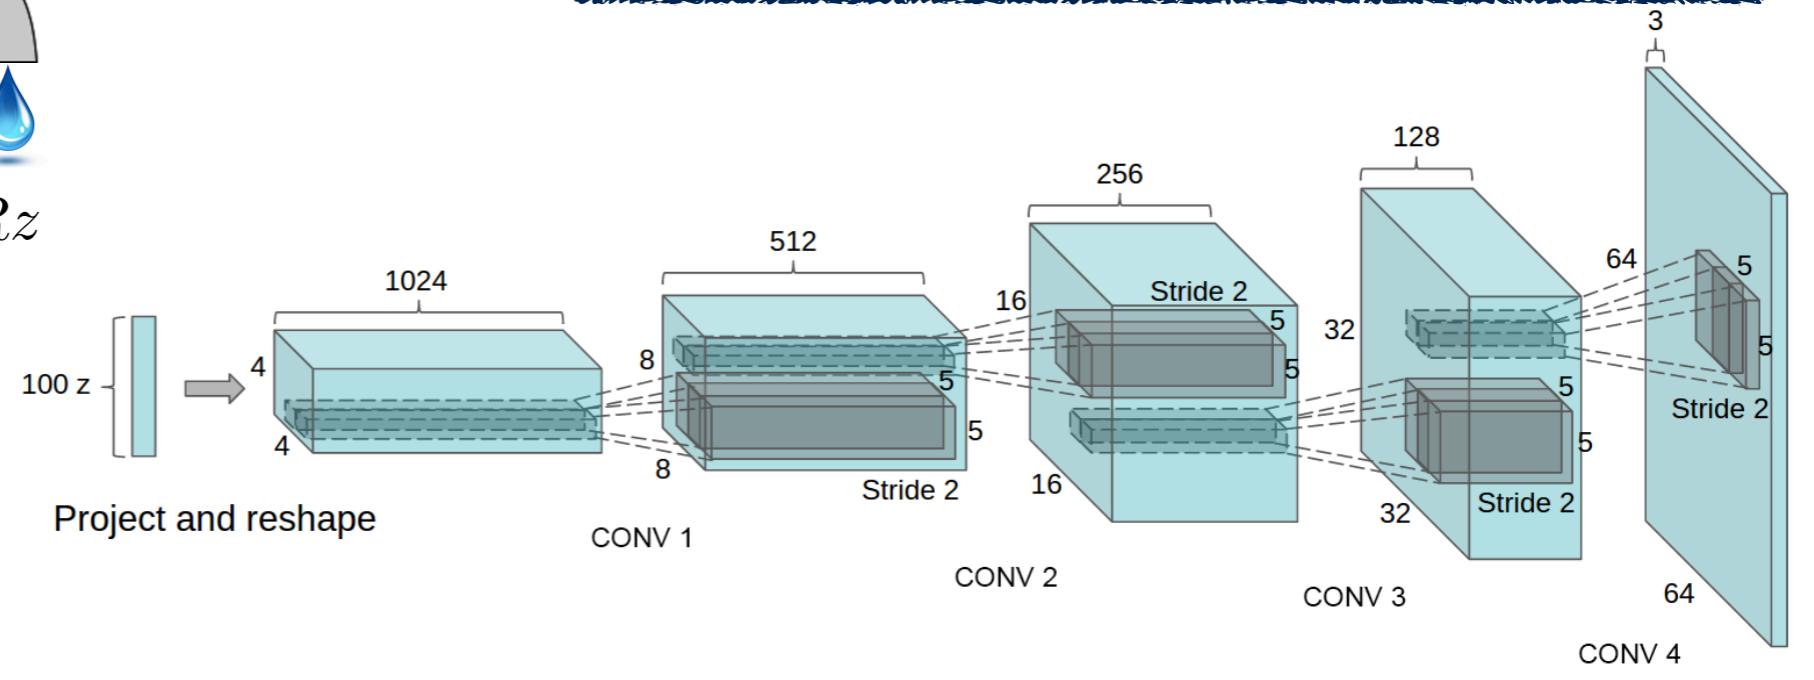
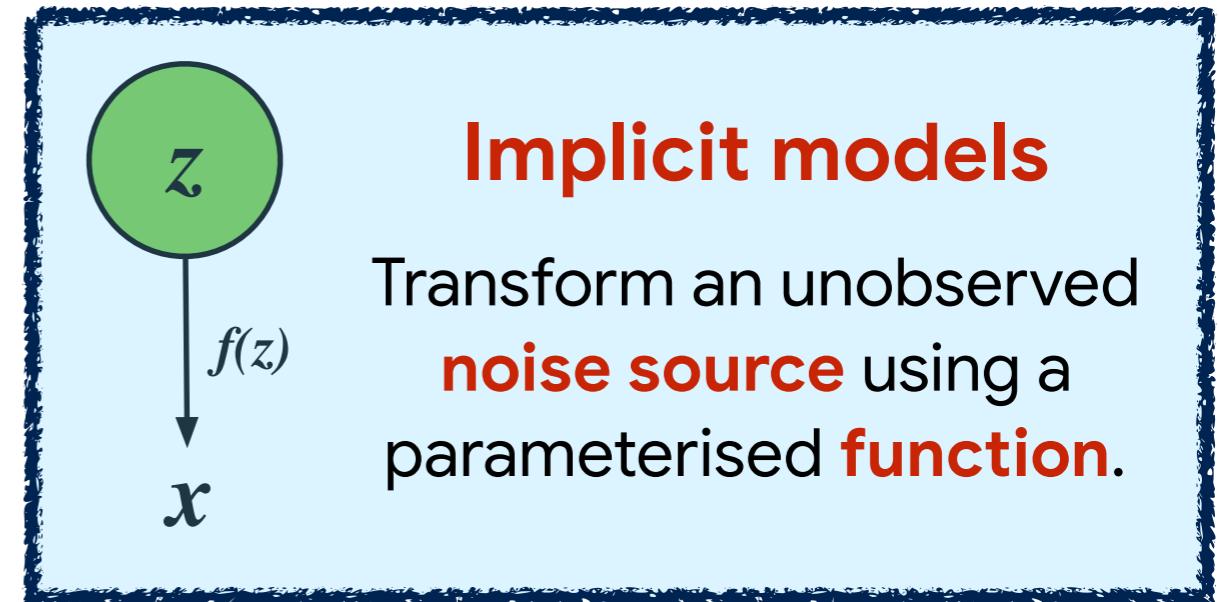
Change of variables for invertible functions



Generator Networks

$$\mathbf{z} \sim \mathcal{N}(0, \mathbf{I})$$
$$\mathbf{x} = f(\mathbf{z}; \theta)$$

$$p(\mathbf{x}) = p(\mathbf{z}) \left| \det \frac{\partial f}{\partial \mathbf{z}} \right|^{-1}$$



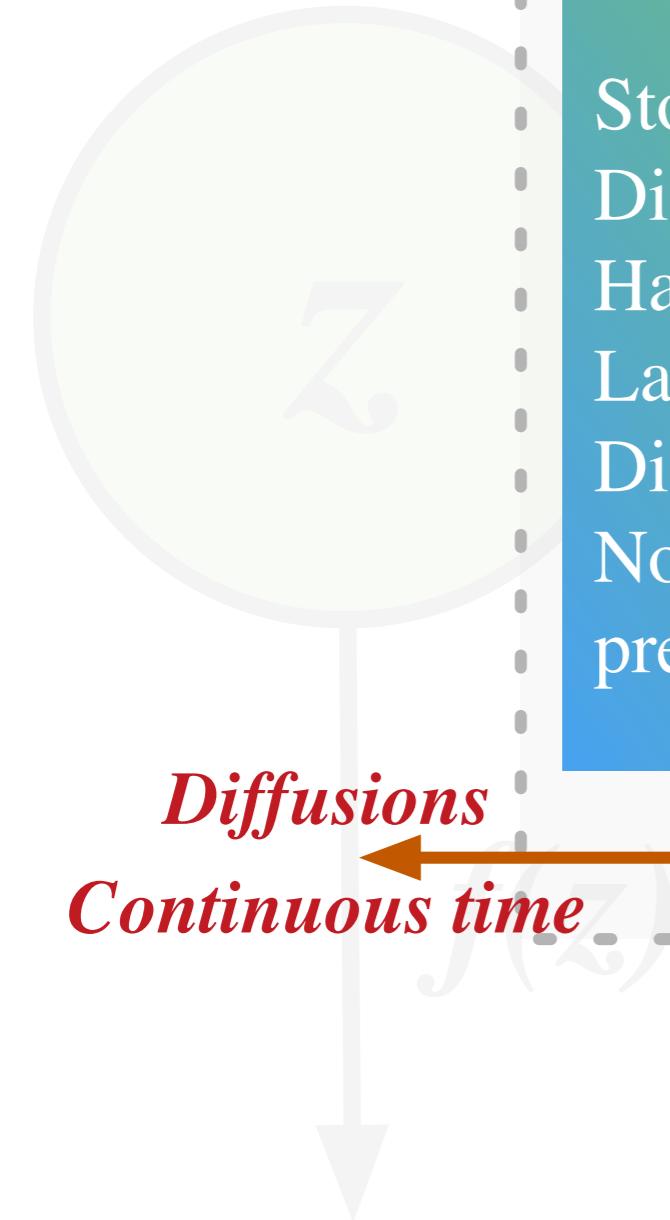
The transformation function is parameterised by a linear or deep network (fully-connected, convolutional or recurrent).

Properties

- + Easy sampling, and natural to specify.
- + Easy to compute expectations without knowing final distribution.
- + Can exploit with large-scale classifiers and convolutional networks.
- ***Difficult to satisfy constraints***: Difficult to maintain invertibility, and challenging optimisation.
- ***Lack of noise model*** (likelihood):
 - Difficult to extend to generic data types
 - Difficult to account for noise in observed data.
 - Hard to compute marginalised likelihood for model scoring, comparison and selection.

*Convolutional generative
adversarial network*

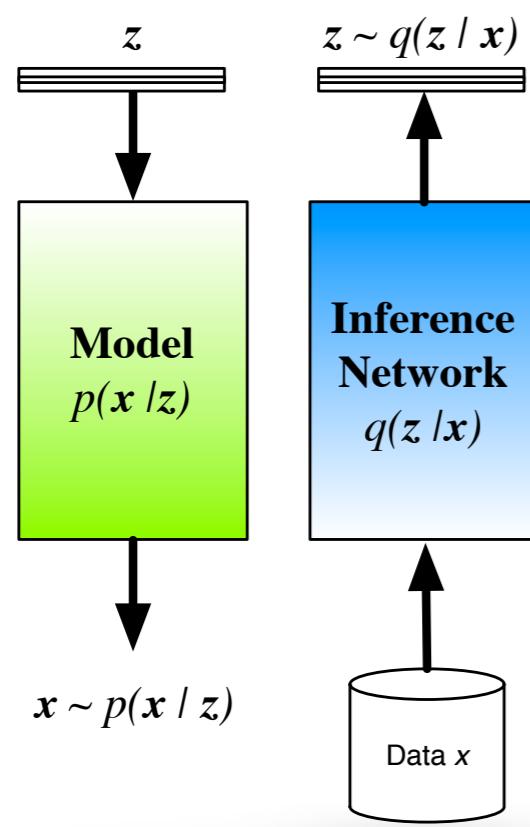




Stochastic
Differential Equations
Hamiltonian and
Langevin SDE
Diffusion Models
Non- and volume
preserving flows

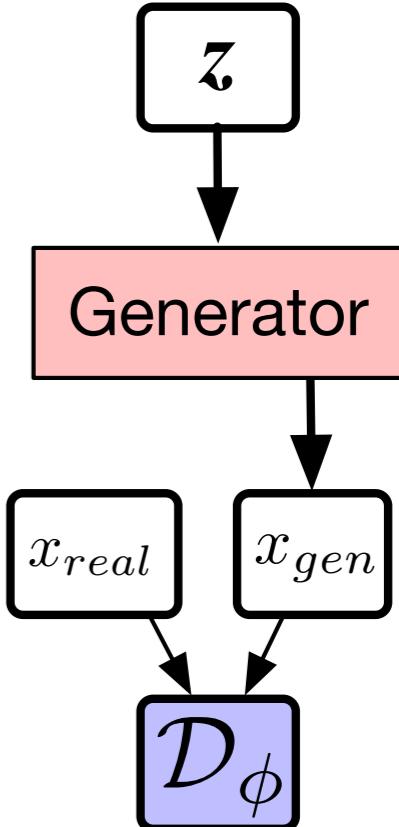
One-liners and
inverse sampling
Distrib. warping
Normalising flows
GAN generator nets
Non- and volume
preserving transforms

Model-Inference-Algorithm



Prescribed latent
variable models and
variational inference

**Variational Autoencoders
(VAEs)**



Implicit latent variable
models and estimation-
by-comparison

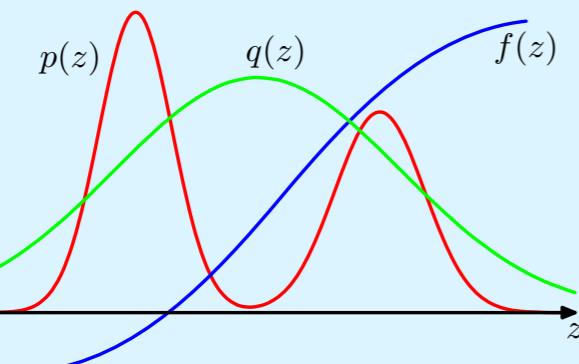
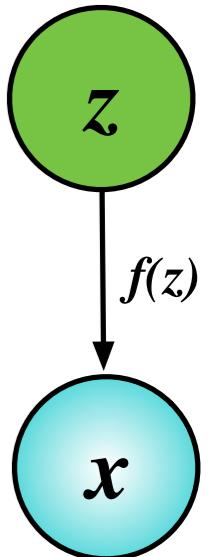
**Generative Adversarial
Networks (GANs)**

Model Evidence

Model evidence (or marginal likelihood, partition function):

Integrating out any global and local variables enables model scoring, comparison, selection, moment estimation, normalisation, posterior computation and prediction.

We take steps to improve the model evidence for given data samples.



Learning principle: Model Evidence

$$p(\mathbf{x}) = \int p(\mathbf{x}, \mathbf{z}) d\mathbf{z}$$

Integral is intractable in general and requires approximation.

Basic idea:
Transform the integral into an expectation over a simple, known distribution.

Variational Inference

$$p(\mathbf{x}) = \int p(\mathbf{x}, \mathbf{z}) d\mathbf{z}$$

$$\mathcal{F}(\mathbf{x}, q) = \mathbb{E}_{q(\mathbf{z})} [\log p(\mathbf{x}|\mathbf{z})] - KL[q(\mathbf{z}) \| p(\mathbf{z})]$$

This bound is exactly of the form we are looking for.

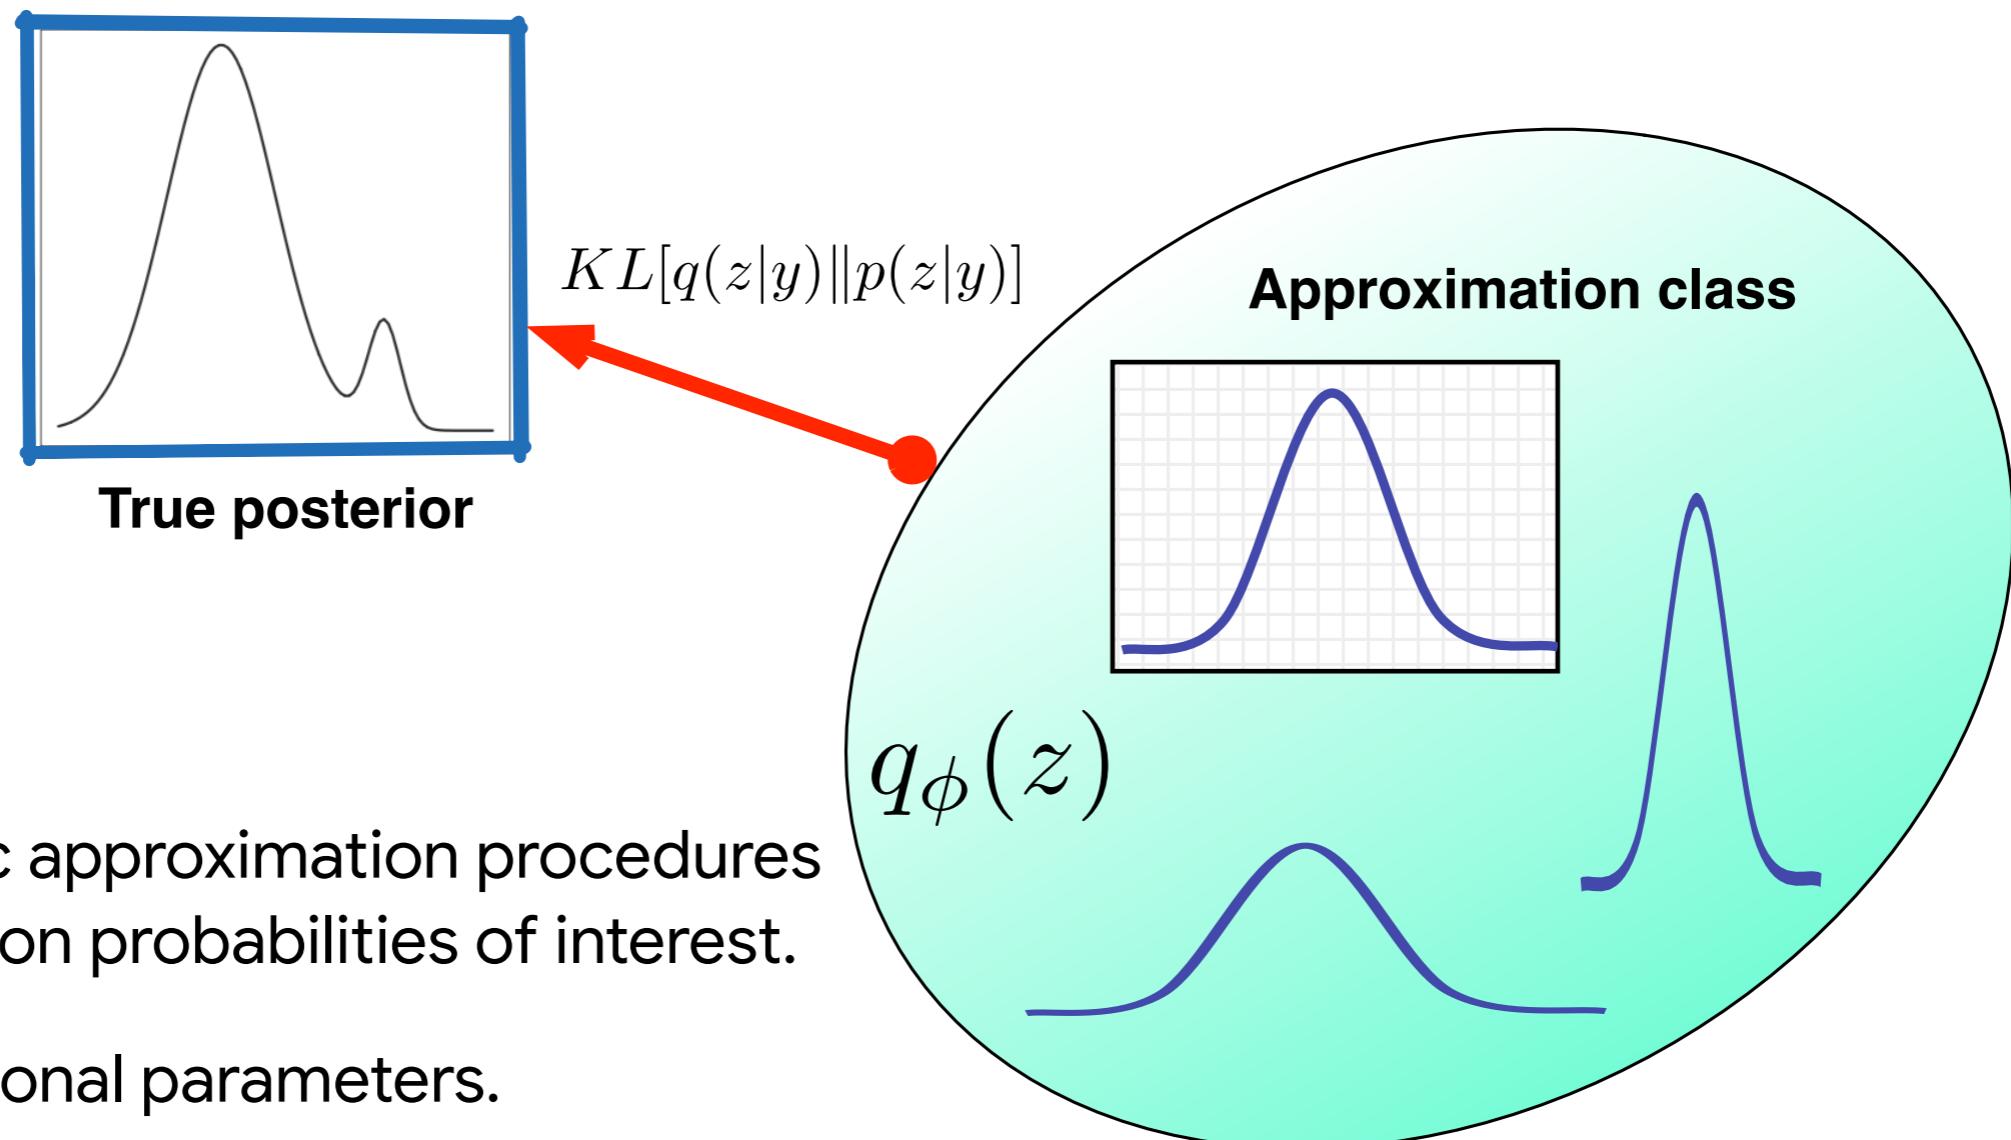
- **Variational free energy:** We obtain a functional and are free to choose the distribution $q(z)$ that best matches the true posterior.
- **Evidence lower bound (ELBO):** principled bound on the marginal likelihood, or model evidence.
- Certain choices of $q(z)$ makes this quantity easier to compute. Examples to come.



Variational Methods

Variational Principle

General family of methods for approximating complicated densities by a simpler class of densities.



Deterministic approximation procedures with bounds on probabilities of interest.

Fit the variational parameters.

Variational Bound

Interpreting the bound:

$$\mathcal{F}(\mathbf{x}, q) = \mathbb{E}_{q(\mathbf{z})} [\log p(\mathbf{x}|\mathbf{z})] - KL[q(\mathbf{z})\|p(\mathbf{z})]$$

Approx. Posterior Reconstruction Penalty

- **Approximate posterior distribution $q(\mathbf{z})$:** Best match to true posterior $p(\mathbf{z}|y)$, one of the unknown inferential quantities of interest to us.
- **Reconstruction cost:** The expected log-likelihood measure how well samples from $q(\mathbf{z})$ are able to explain the data y .
- **Penalty:** Ensures the explanation of the data $q(\mathbf{z})$ doesn't deviate too far from your beliefs $p(\mathbf{z})$. A mechanism for realising Okham's razor.

Variational Bound

$$\mathcal{F}(\mathbf{x}, q) = \mathbb{E}_{q(\mathbf{z})}[\log p(\mathbf{x}|\mathbf{z})] - KL[q(\mathbf{z})\|p(\mathbf{z})]$$

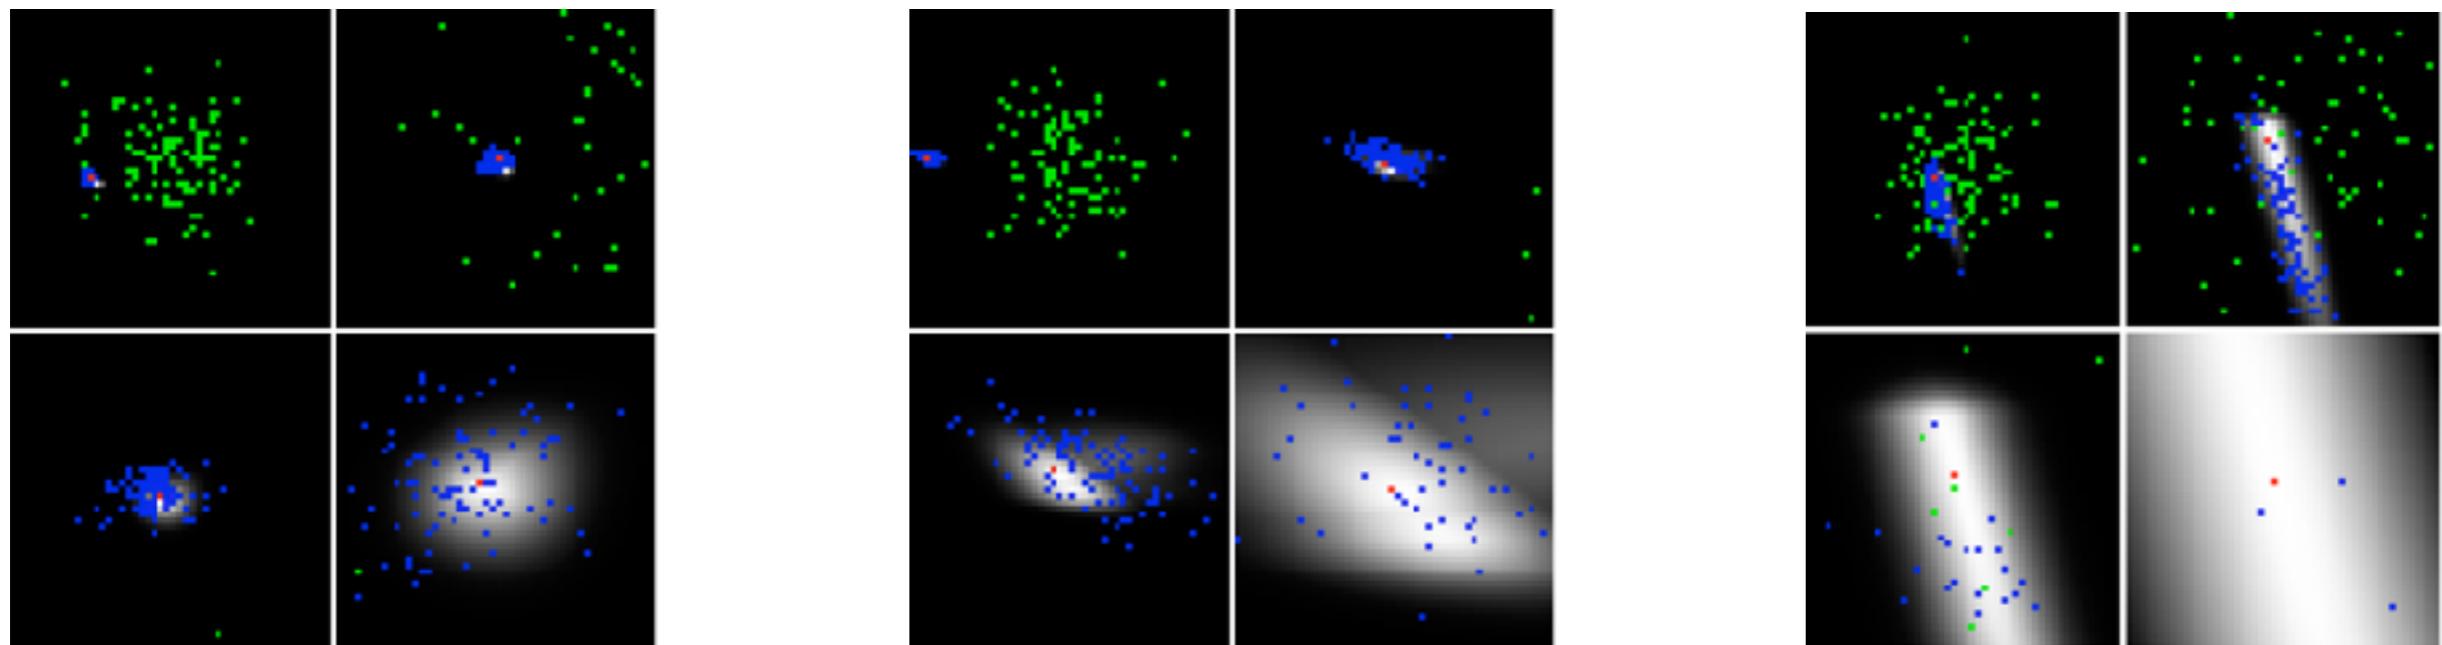
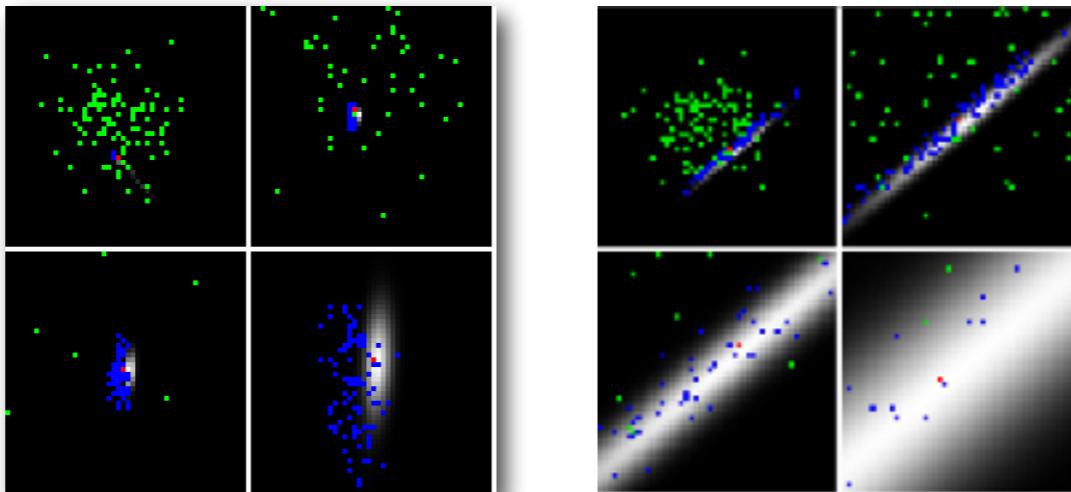
Approx. Posterior Reconstruction Penalty

Some comments on q :

- **Integration is now optimisation:** optimise for $q(z)$ directly.
 - I write $q(z)$ to simplify the notation, but it depends on the data, $q(z|x)$.
 - *Easy convergence assessment* since we wait until the free energy (loss) reaches convergence.
- **Variational parameters:** parameters of $q(z)$
 - E.g., if a Gaussian, variational parameters are mean and variance.
 - Optimisation allows us to *tighten the bound* and get as close as possible to the true marginal likelihood.

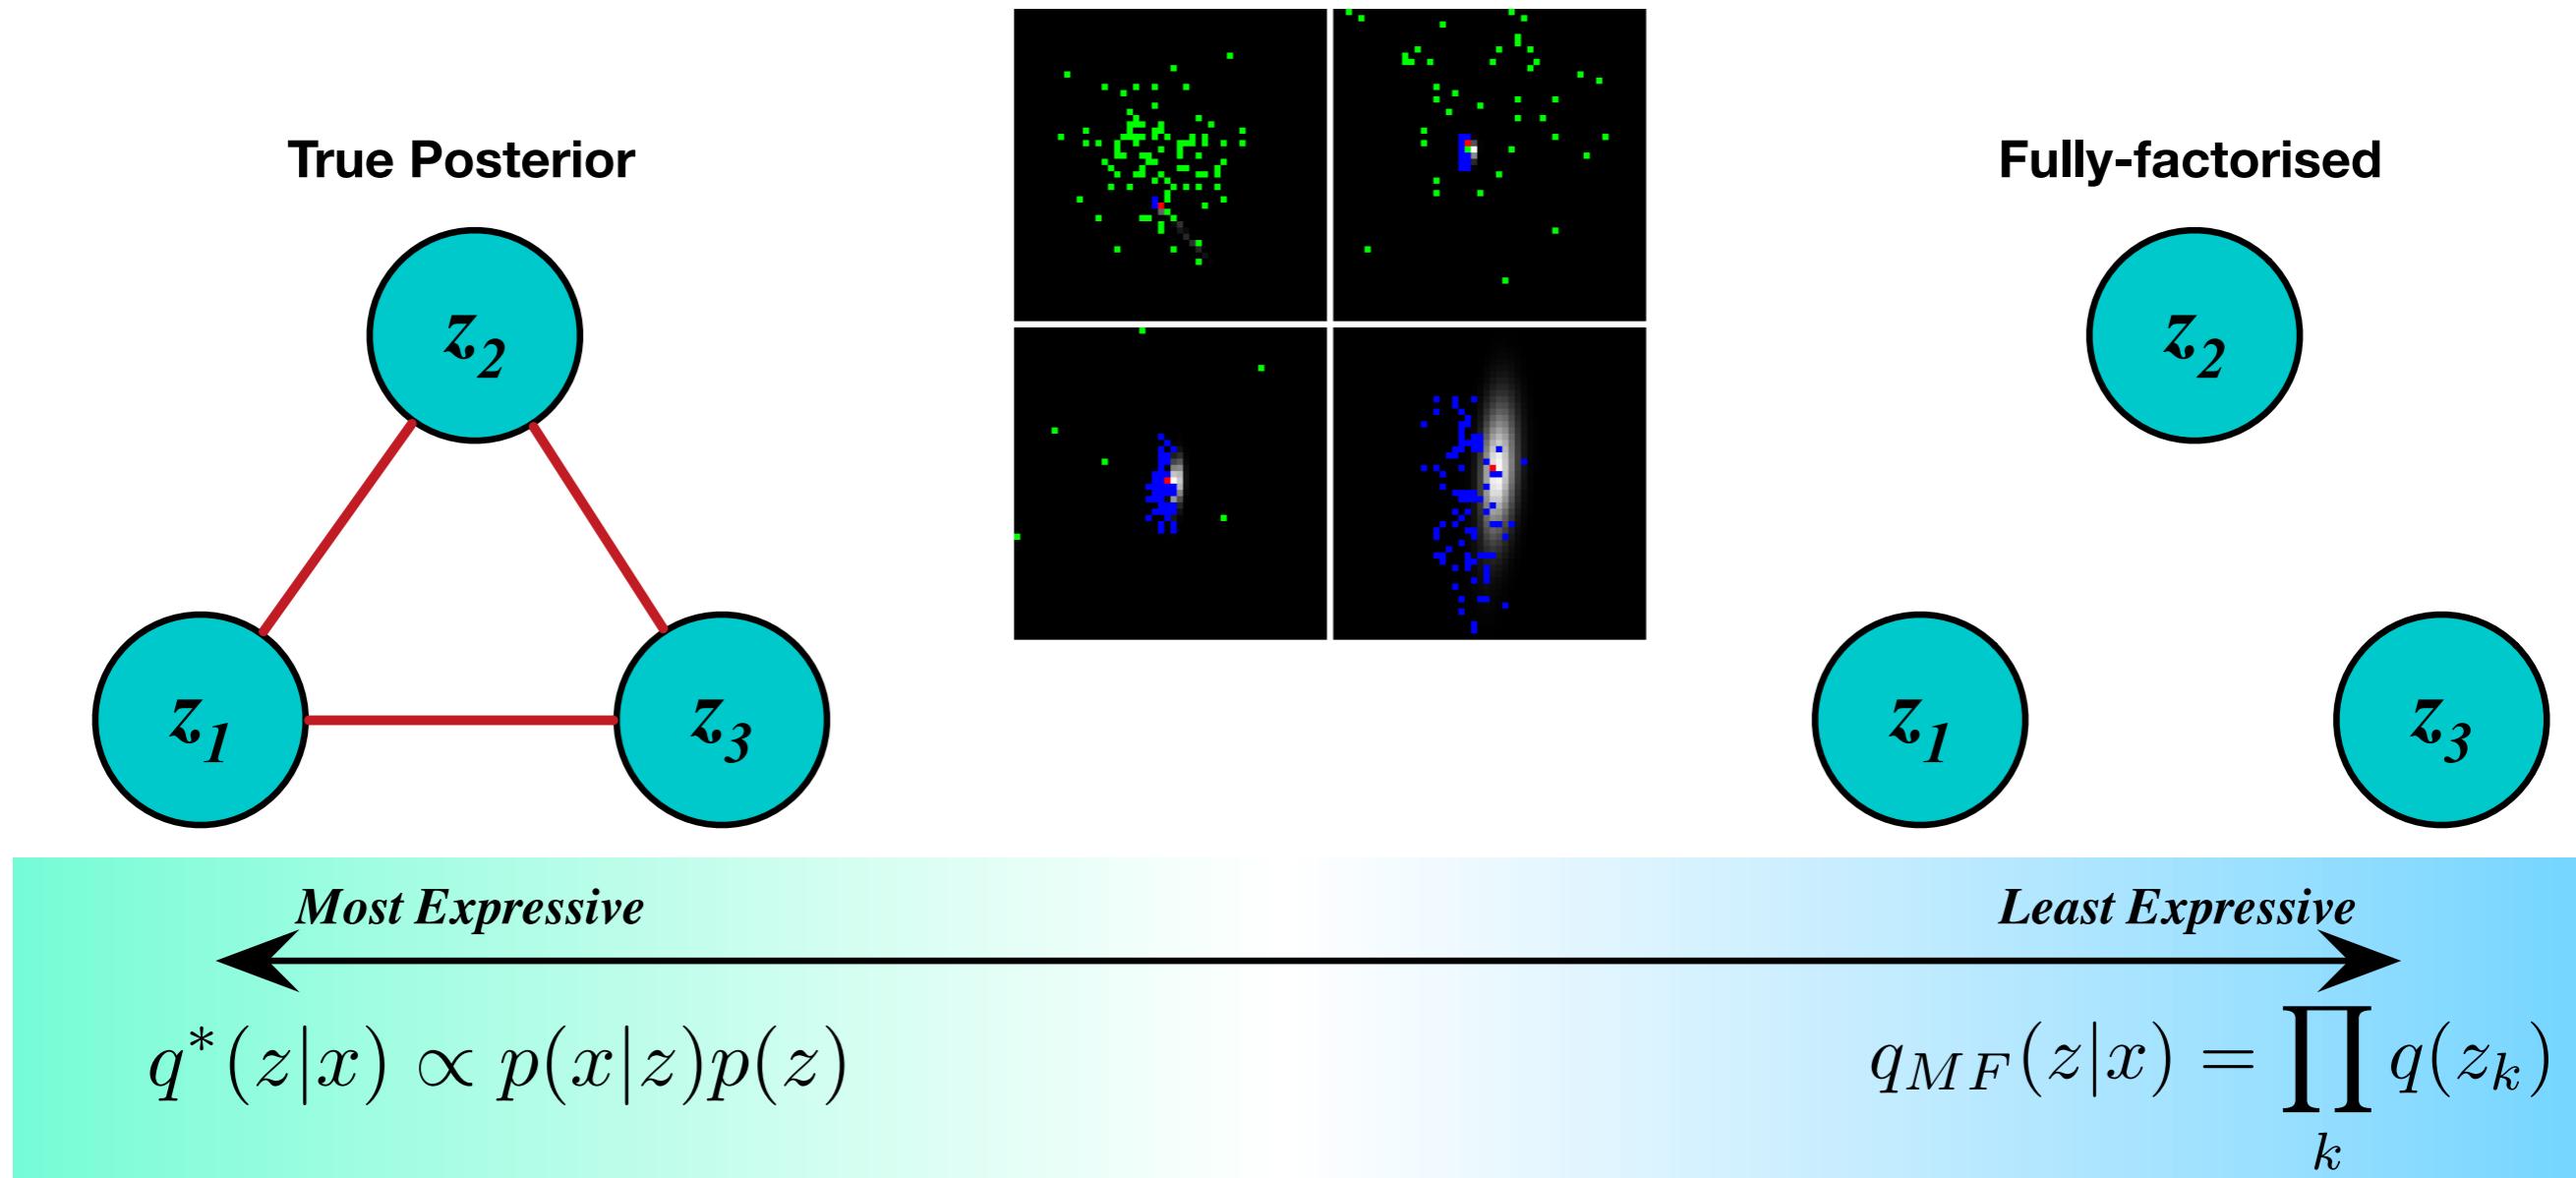
Real Posteriors

Require flexible approximations for the types of posteriors we are likely to see.



Mean-Fields

Mean-field methods assume that the distribution is factorised.

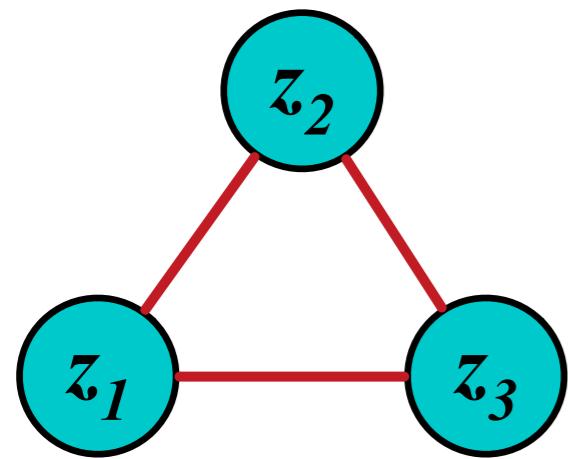


Restricted class of approximations: every dimension (or subset of dimensions) of the posterior is independent.

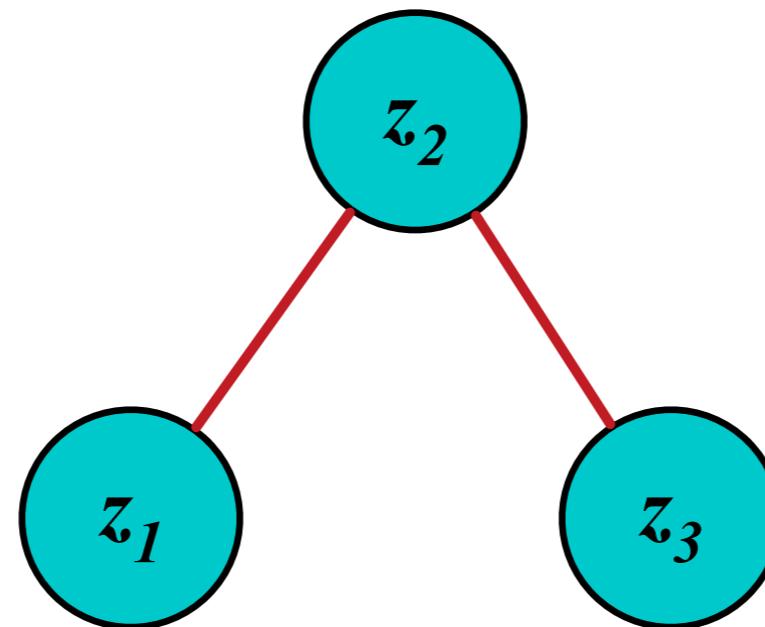
Structured Mean-field

Structured mean-field: introduce dependencies into our factorisation.

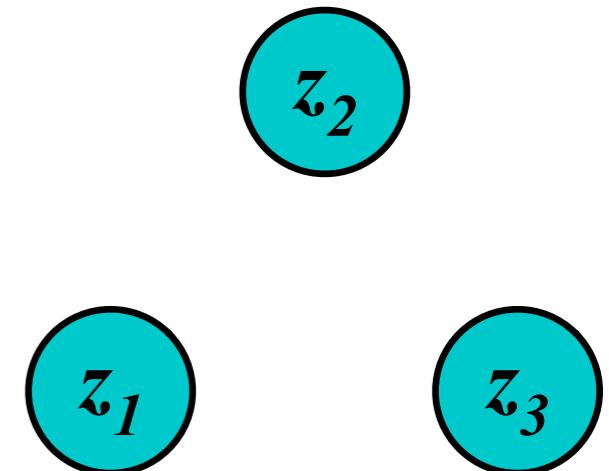
True Posterior



Structured Approx.



Fully-factorised



Most Expressive

$$q^*(z|x) \propto p(x|z)p(z)$$

$$q(z) = \prod_k q_k(z_k | \{z_j\}_{j \neq k})$$

Least Expressive

$$q_{MF}(z|x) = \prod_k q(z_k)$$

Latent Gaussian Models

Examples: GP regression or DLGM.

Probabilistic Model

$$z \sim \mathcal{N}(z|0, 1) \quad y \sim p(y|f_\theta(z))$$

Mean-field approx

$$q(z) = \prod_i \mathcal{N}(z_i|\mu_i, \sigma_i^2)$$

Variational bound

$$\mathcal{F}(y, q) = \mathbb{E}_{q(z)}[\log p(y|z)] - KL[q(z)\|p(z)]$$

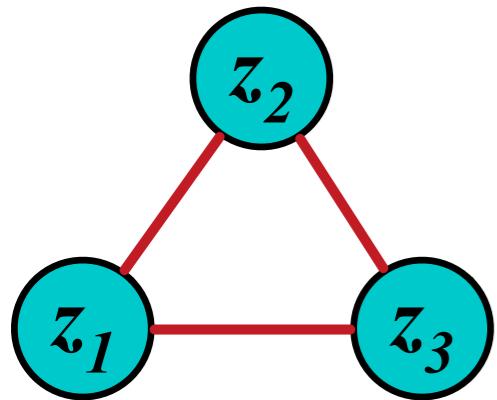
$$\mathcal{F}(y, q) = \mathbb{E}_{q(z)}[\log p(y|z)] - \sum_i KL[q(z_i)\|p(z_i)]$$

$$\mathcal{F}(y, q) = \mathbb{E}_{q(z)}[\log p(y|z)] - \sum_i KL[\mathcal{N}(z_i|\mu_i, \sigma_i^2)\|\mathcal{N}(z_i|0, 1)]$$

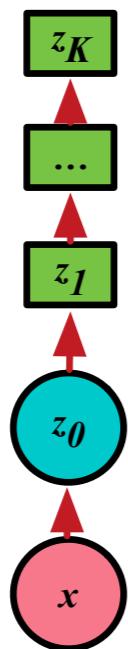
$$\mathcal{F}(y, q) = \mathbb{E}_{q(z)}[\log p(y|f_\theta(z))] - \frac{1}{2} \sum_i (\sigma_i^2 + \mu_i^2 - 1 - \ln \sigma_i^2)$$

Families of Approximations

True Posterior

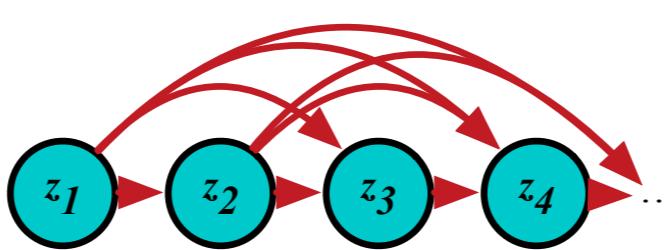


Normalising
flows

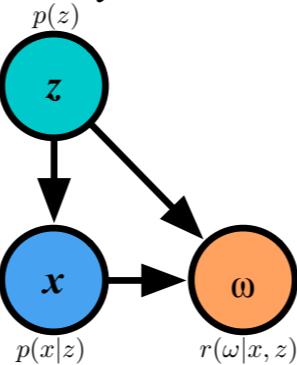


Families of Posterior Approximations

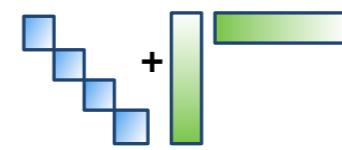
Structured mean-field



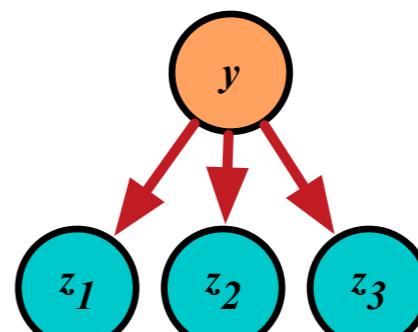
Auxiliary variables



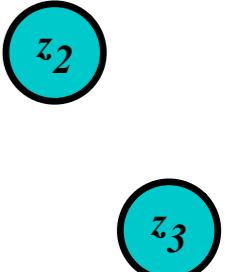
Covariance models



Mixtures



Fully-factorised



Most Expressive

Least Expressive

$$q^*(z|x) \propto p(x|z)p(z)$$

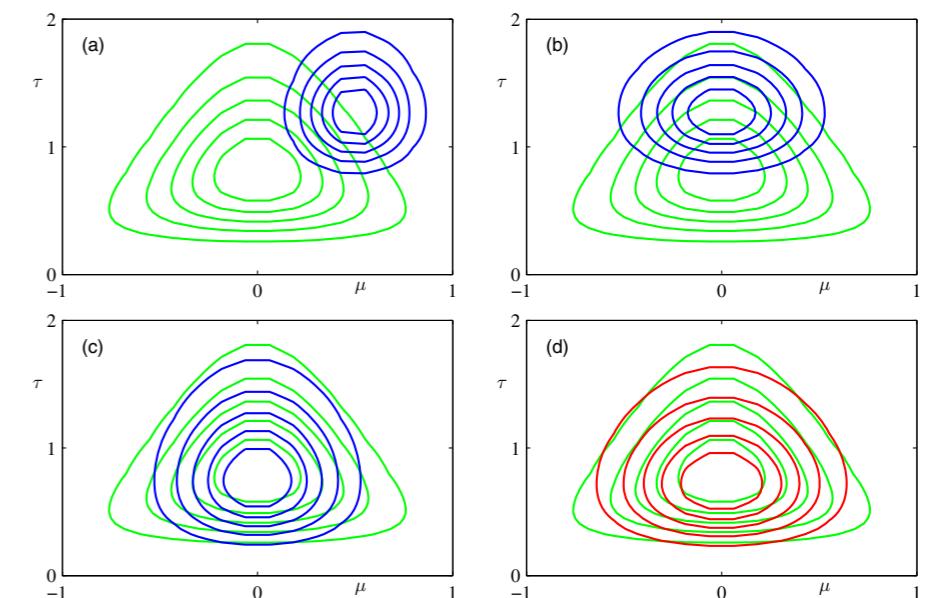
$$q_{MF}(z|x) = \prod_k q(z_k)$$

Variational Optimisation

$$\mathcal{F}(\mathbf{x}, q) = \mathbb{E}_{q(\mathbf{z})}[\log p(\mathbf{x}|\mathbf{z})] - KL[q(\mathbf{z})\|p(\mathbf{z})]$$

Approx. Posterior Reconstruction Penalty

- Variational EM
- Stochastic Variational Inference
- Doubly Stochastic Variational Inference
- Amortised Inference



Variational EM

$$\mathcal{F}(\mathbf{x}, q) = \mathbb{E}_{q(\mathbf{z})}[\log p(\mathbf{x}|\mathbf{z})] - KL[q(\mathbf{z})\|p(\mathbf{z})]$$

Alternating optimisation for the variational parameters and then model parameters (VEM).

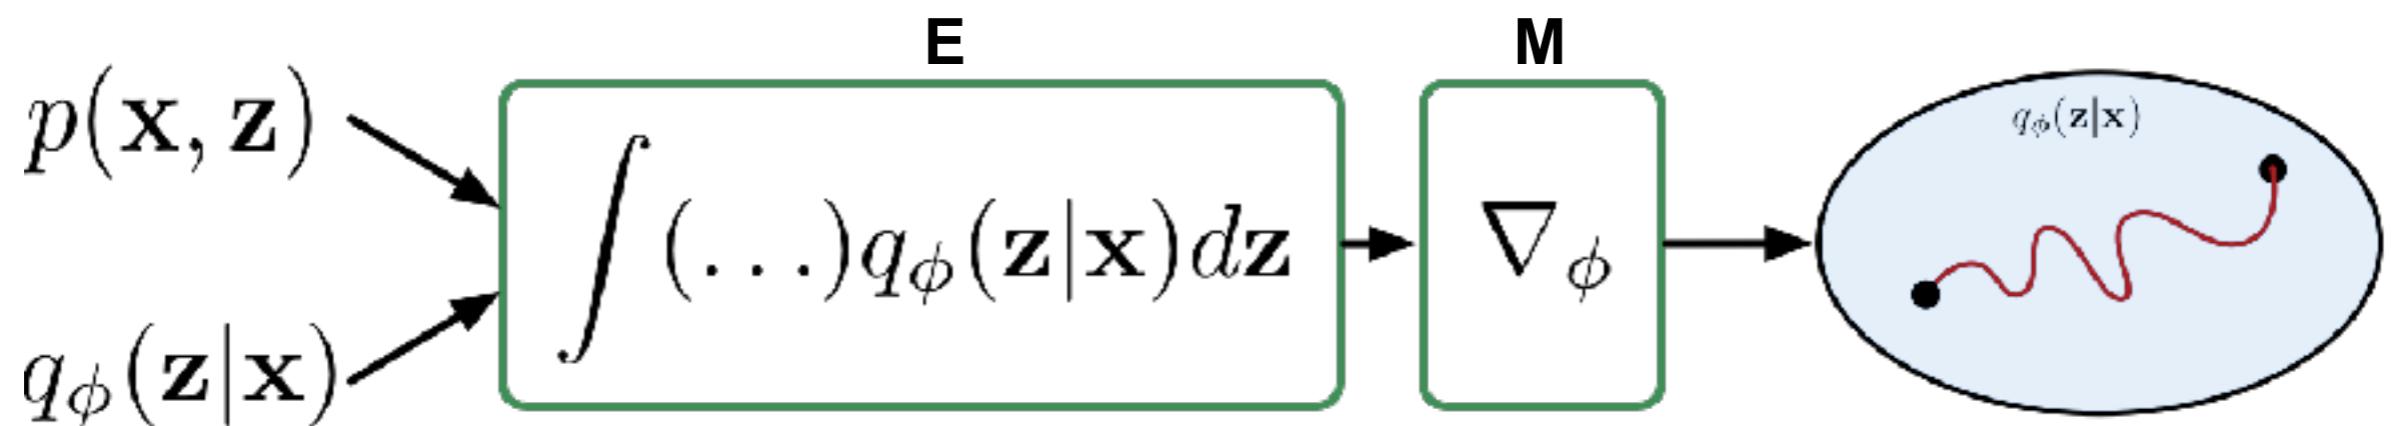
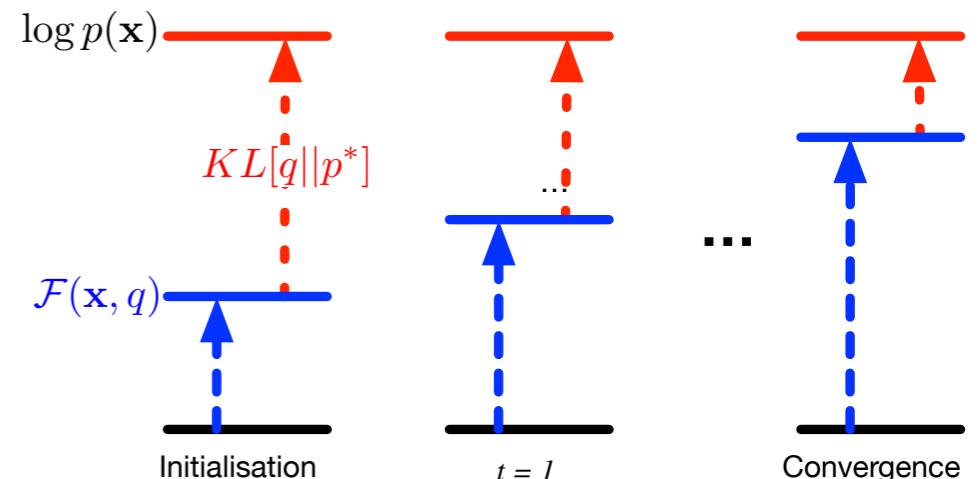
Repeat:

E-step $\phi \propto \nabla_\phi \mathcal{F}(\mathbf{x}, q)$

Var. params

M-step $\theta \propto \nabla_\theta \mathcal{F}(\mathbf{x}, q)$

Model params



Amortised Inference

Repeat:

E-step (compute q)

For $i = 1, \dots, N$

$$\phi_n \propto \nabla_\phi \mathbb{E}_{q_\phi(z)} [\log p_\theta(\mathbf{x}_n | z_n)] - \nabla_\phi KL[q(z_n) \| p(z)]$$

Instead of solving for every observation, amortise using a model.

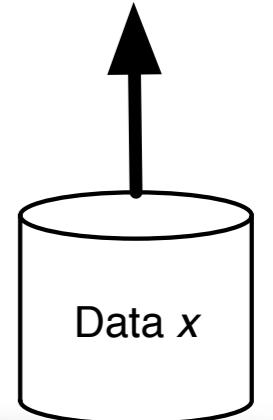
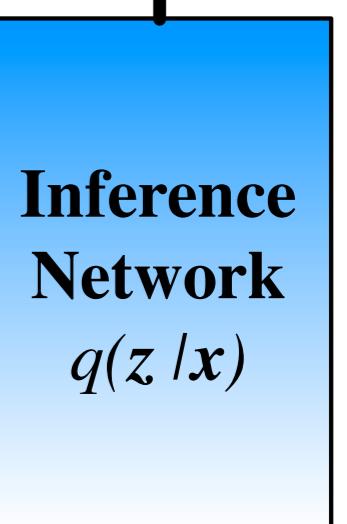
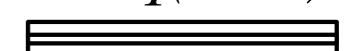
M-step

$$\theta \propto \frac{1}{N} \sum_n \mathbb{E}_{q_\phi(z)} [\nabla_\theta \log p_\theta(\mathbf{x}_n | z_n)]$$

- **Inference network:** q is an **encoder**, an **inverse model**, **recognition model**.
- Parameters of q are now a set of *global parameters* used for inference of all data points - test and train.
- **Amortise (spread) the cost of inference over all data.**
- Joint optimisation of variational and model parameters.

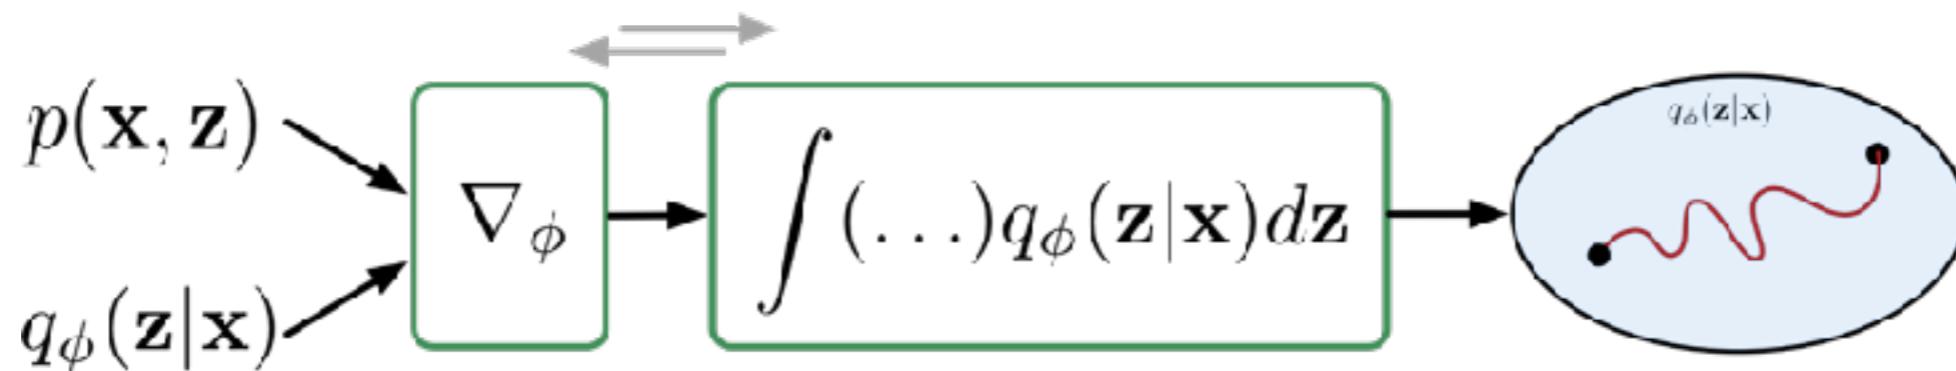
Inference networks provide an efficient mechanism for **posterior inference with memory**

$$z \sim q(z | x)$$



Stochastic Gradients

$$\nabla_{\phi} \mathbb{E}_{q_{\phi}(\mathbf{z})} [f_{\theta}(\mathbf{z})] = \nabla \int q_{\phi}(\mathbf{z}) f_{\theta}(\mathbf{z}) d\mathbf{z}$$

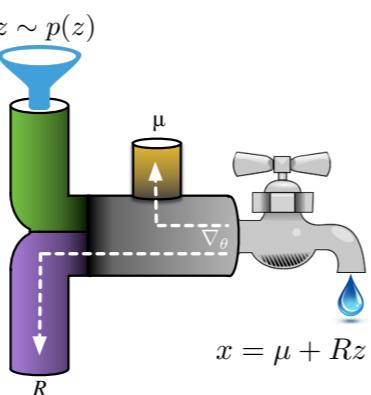


Doubly stochastic estimators

Pathwise Estimator

When easy to use transformation is available and differentiable function f .

$$= \mathbb{E}_{p(\epsilon)} [\nabla_{\phi} f_{\theta}(g(\epsilon, \phi))] \\ z \sim q_{\phi}(\mathbf{z}) \\ \mathbf{z} = g(\epsilon, \phi) \quad \epsilon \sim p(\epsilon)$$



Reparameterisation

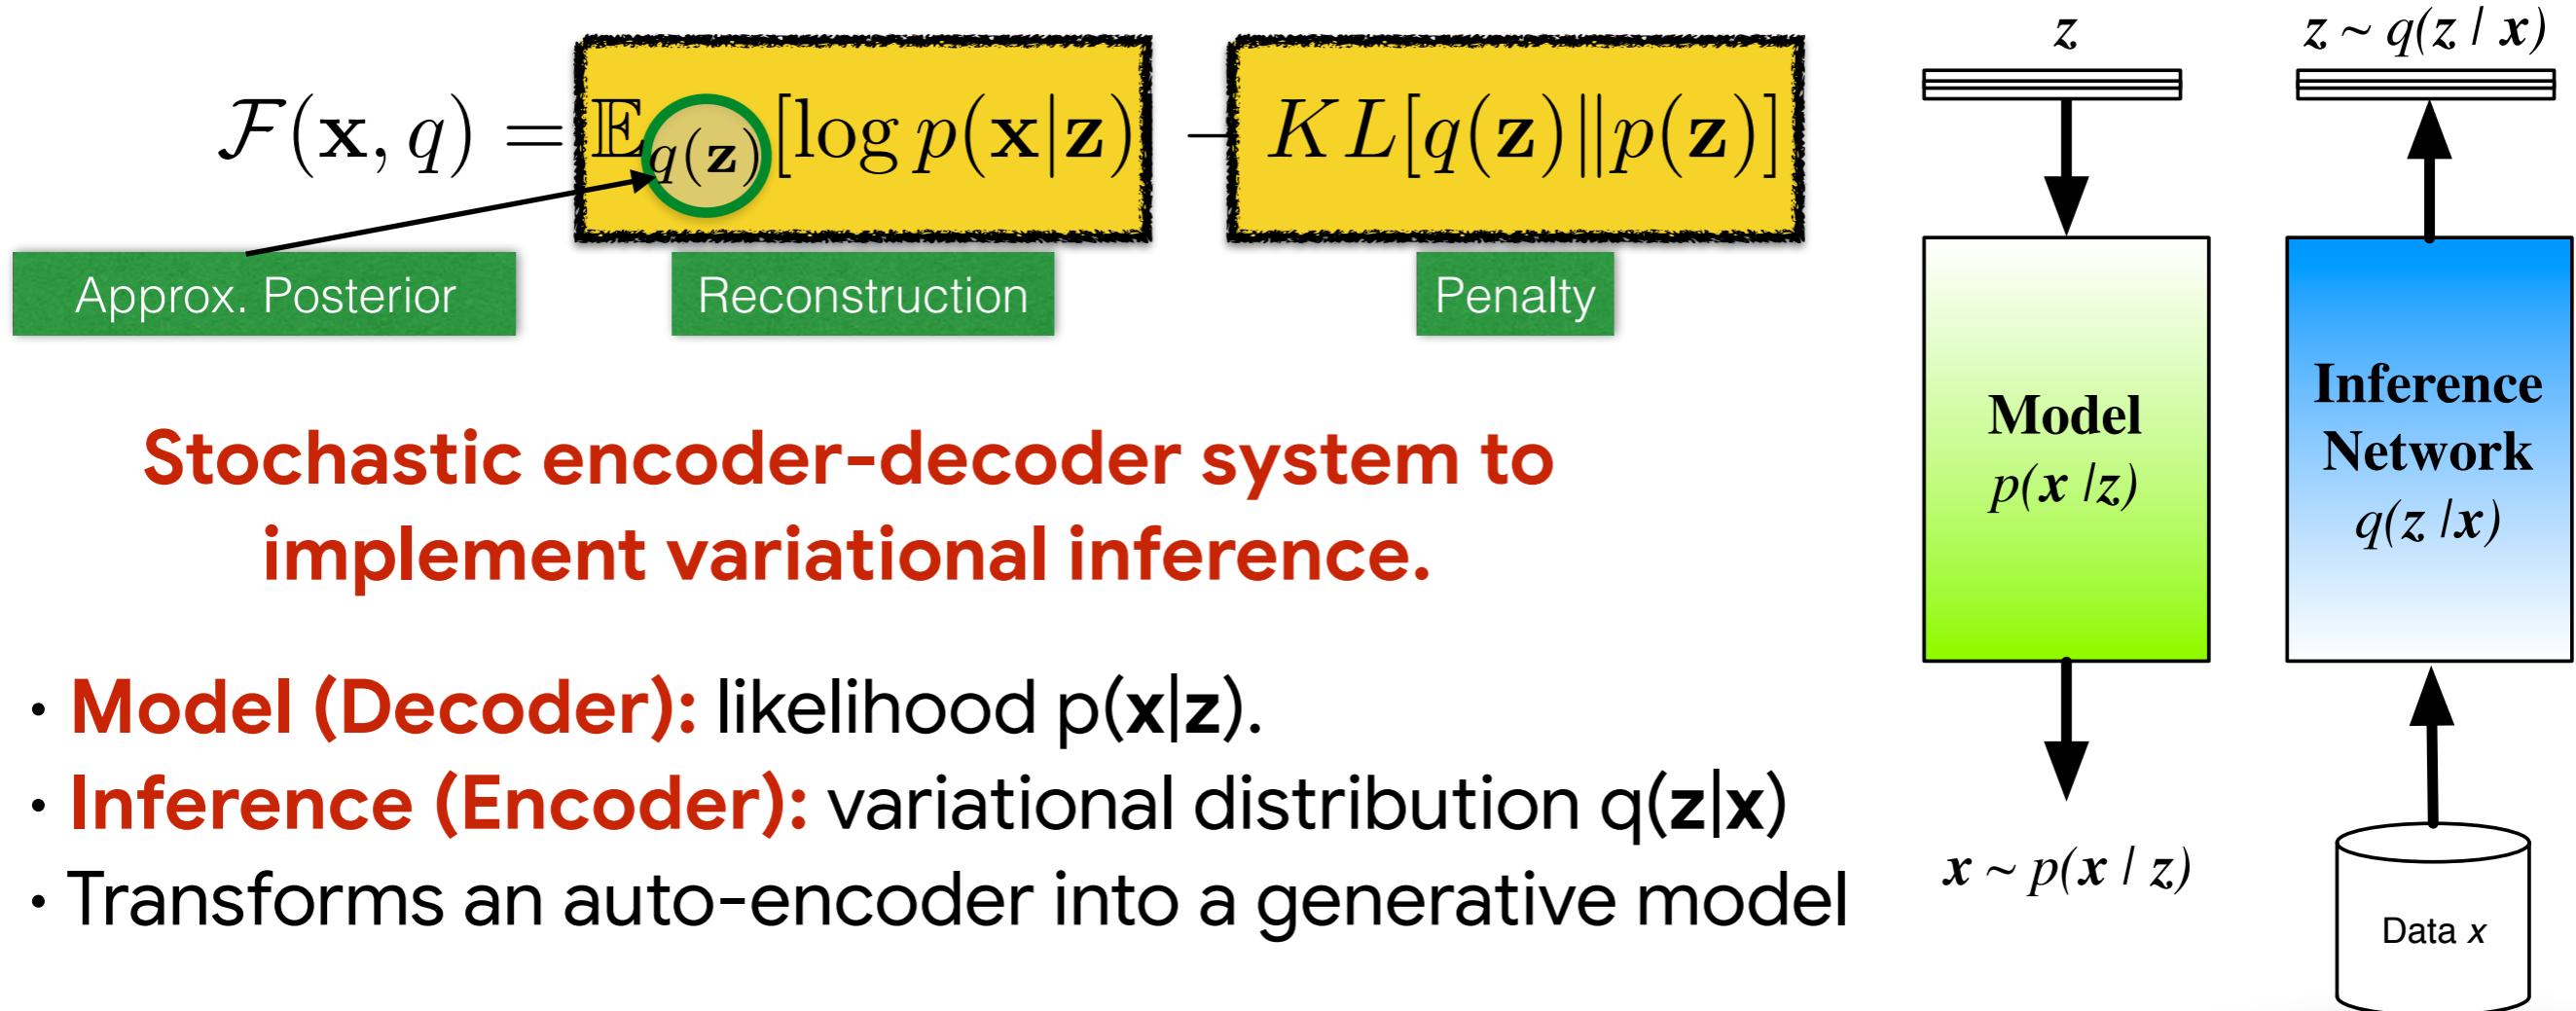
Score-function estimator
When function f non-differentiable and $q(z)$ is easy to sample from.

$$= \mathbb{E}_{q(z)} [f_{\theta}(\mathbf{z}) \nabla_{\phi} \log q_{\phi}(\mathbf{z})]$$

Identity

Log-derivative

Variational Autoencoder

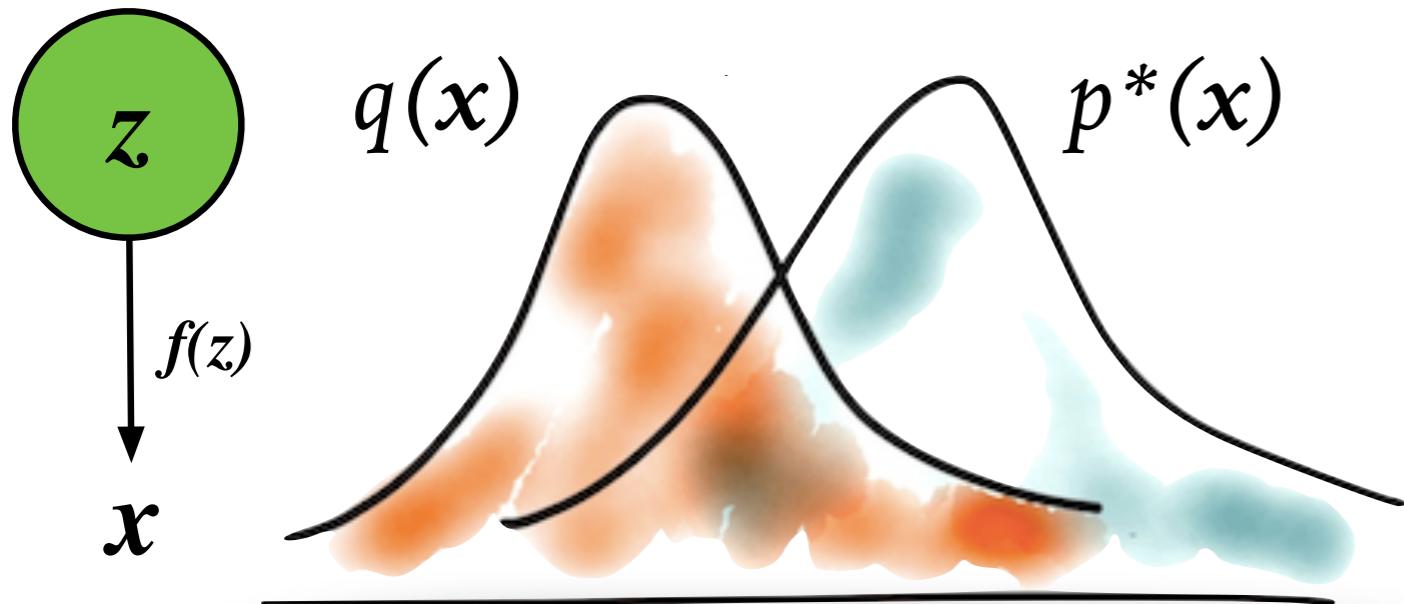


Specific combination of **variational inference** in **latent variable models** using **inference networks**
Variational Auto-encoder

But don't forget what your model is, and what inference you use.

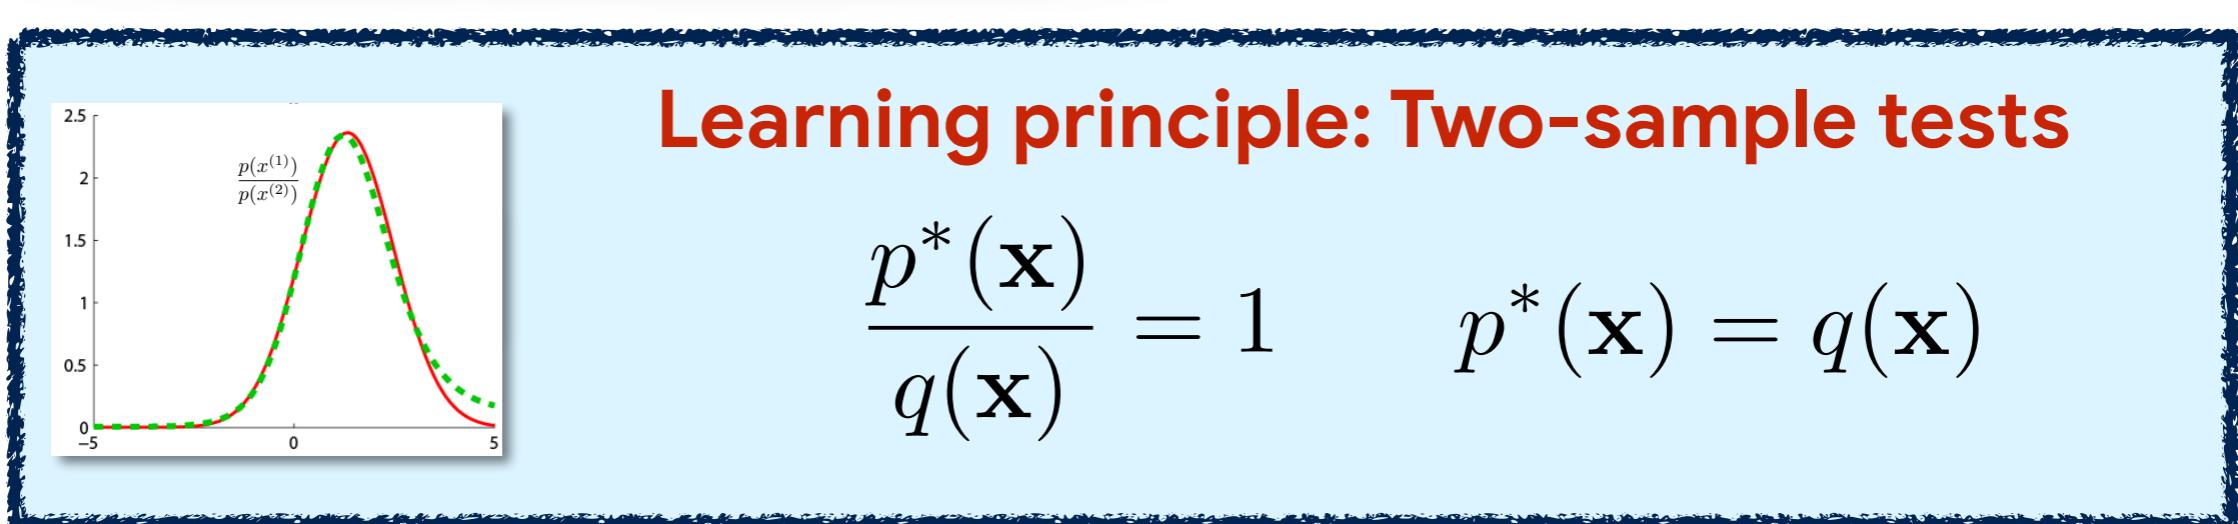
Learning by Comparison

For some models, we only have access to an unnormalised probability or partial knowledge of the distribution.



Basic idea:
Transform into
learning a model of
the density ratio.

We compare the
estimated distribution $q(x)$ to
the true distribution $p^*(x)$



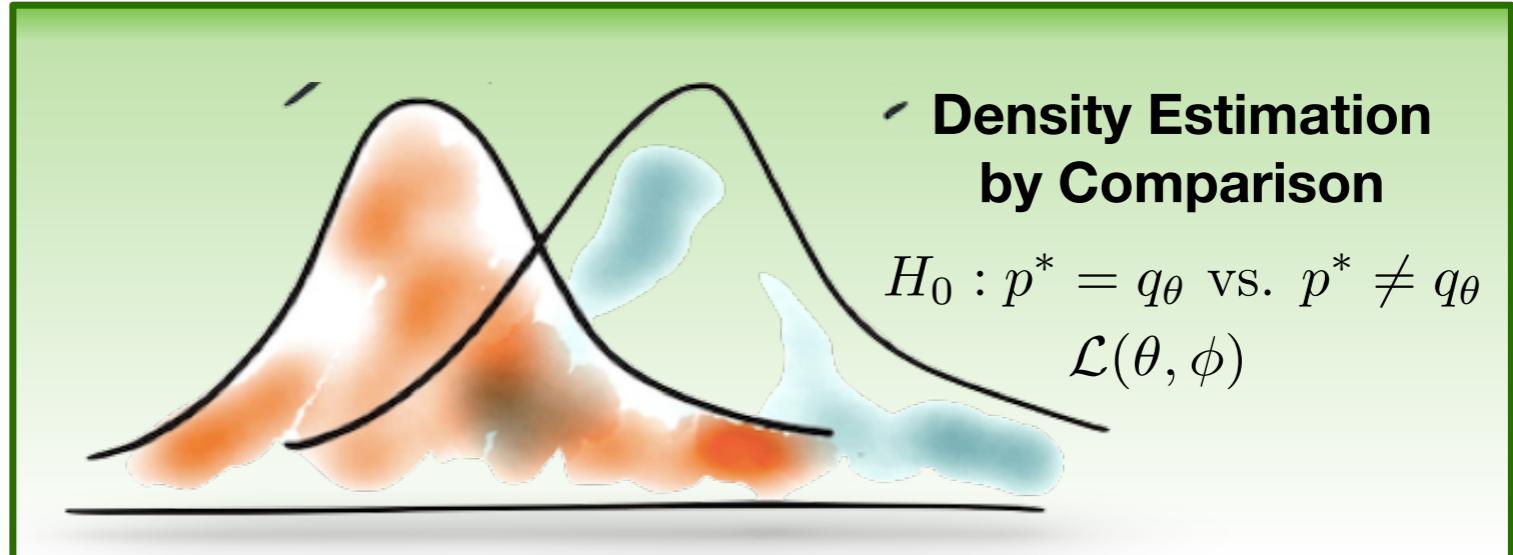
Interest is not in estimating the marginal probabilities, only in how they are related.

Estimation by Comparison

Two steps

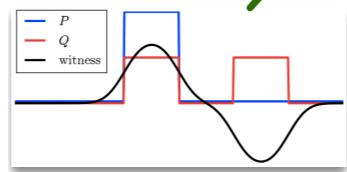
1. Use a hypothesis **test or comparison** to obtain some model to tell how data from our model differs from observed data.

2. **Adjust model** to better match the data distribution using the comparison model from step 1.

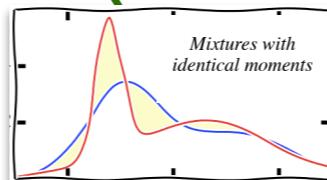


Density Difference

$$r_\phi = p^* - q_\theta$$

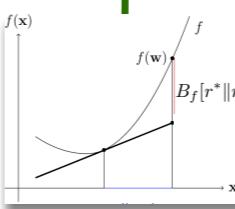


Max Mean Discrepancy

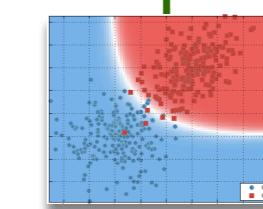


Moment Matching

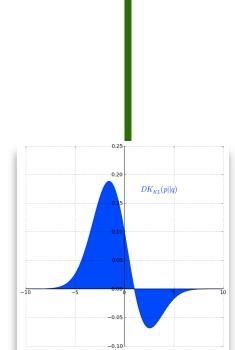
$$r_\phi = \frac{p^*}{q_\theta}$$



Bregman Divergence



Class Probability Estimation



f-Divergence

$$f(u) = u \log u - (u + 1) \log(u + 1)$$

Adversarial Learning

$$\frac{p^*(\mathbf{x})}{q(\mathbf{x})} = \frac{p(y=1|\mathbf{x})}{p(y=-1|\mathbf{x})} \quad p(y=+1|\mathbf{x}) = D_\theta(\mathbf{x}) \quad p(y=-1|\mathbf{x}) = 1 - D_\theta(\mathbf{x}) \quad \text{Scoring Function}$$

$$\mathcal{F}(\mathbf{x}, \theta, \phi) = \mathbb{E}_{p^*(x)}[\log D_\theta(\mathbf{x})] + \mathbb{E}_{q_\phi(x)}[\log(1 - D_\theta(\mathbf{x}))] \quad \text{Bernoulli Loss}$$

Generative Adversarial Networks

Alternating optimisation $\min_{\phi} \max_{\theta} \mathcal{F}(\mathbf{x}, \theta, \phi)$

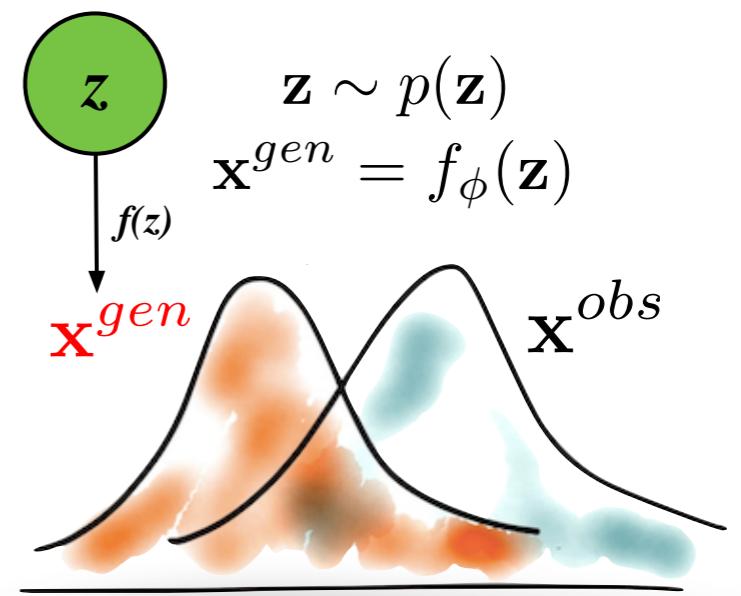
Comparison loss $\theta \propto \nabla_{\theta} \mathbb{E}_{p^*(x)}[\log D_\theta(\mathbf{x})] + \nabla_{\theta} \mathbb{E}_{q_\phi(x)}[\log(1 - D_\theta(\mathbf{x}))]$

Generative loss $\phi \propto -\nabla_{\phi} \mathbb{E}_{q(z)}[\log(1 - D_\theta(f_\phi(\mathbf{z})))]$

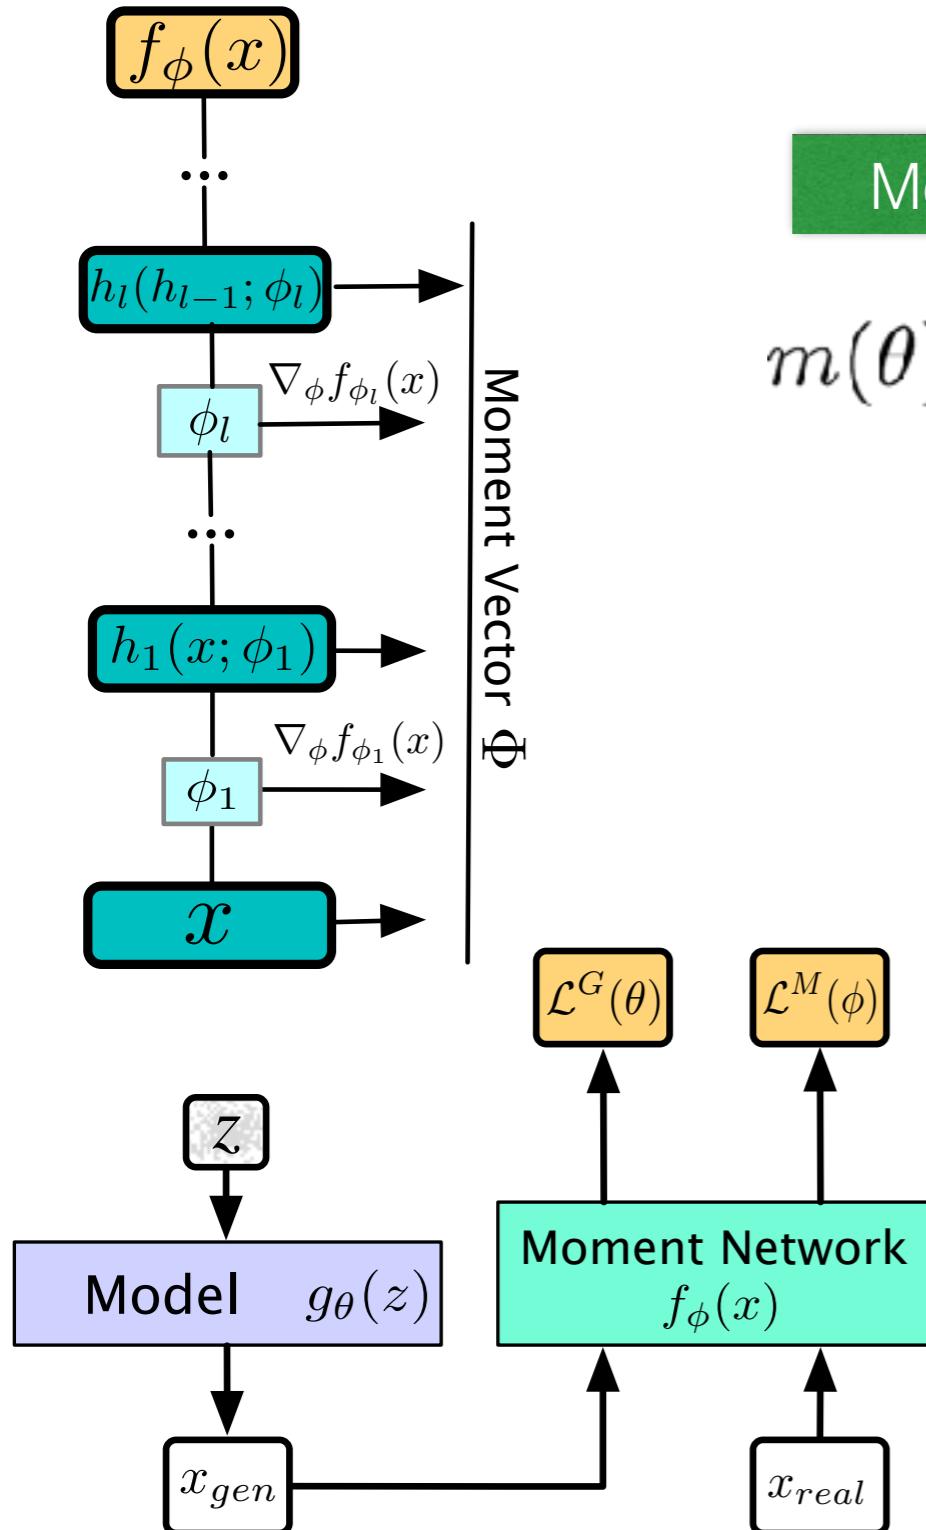
- Instances of testing and inference:**
- Unsupervised-as-supervised learning
 - Classifier ABC
 - Noise-contrastive estimation
 - Adversarial learning and GANs

Density-ratio

Reparameterisation



Method of Moments



Moment estimator

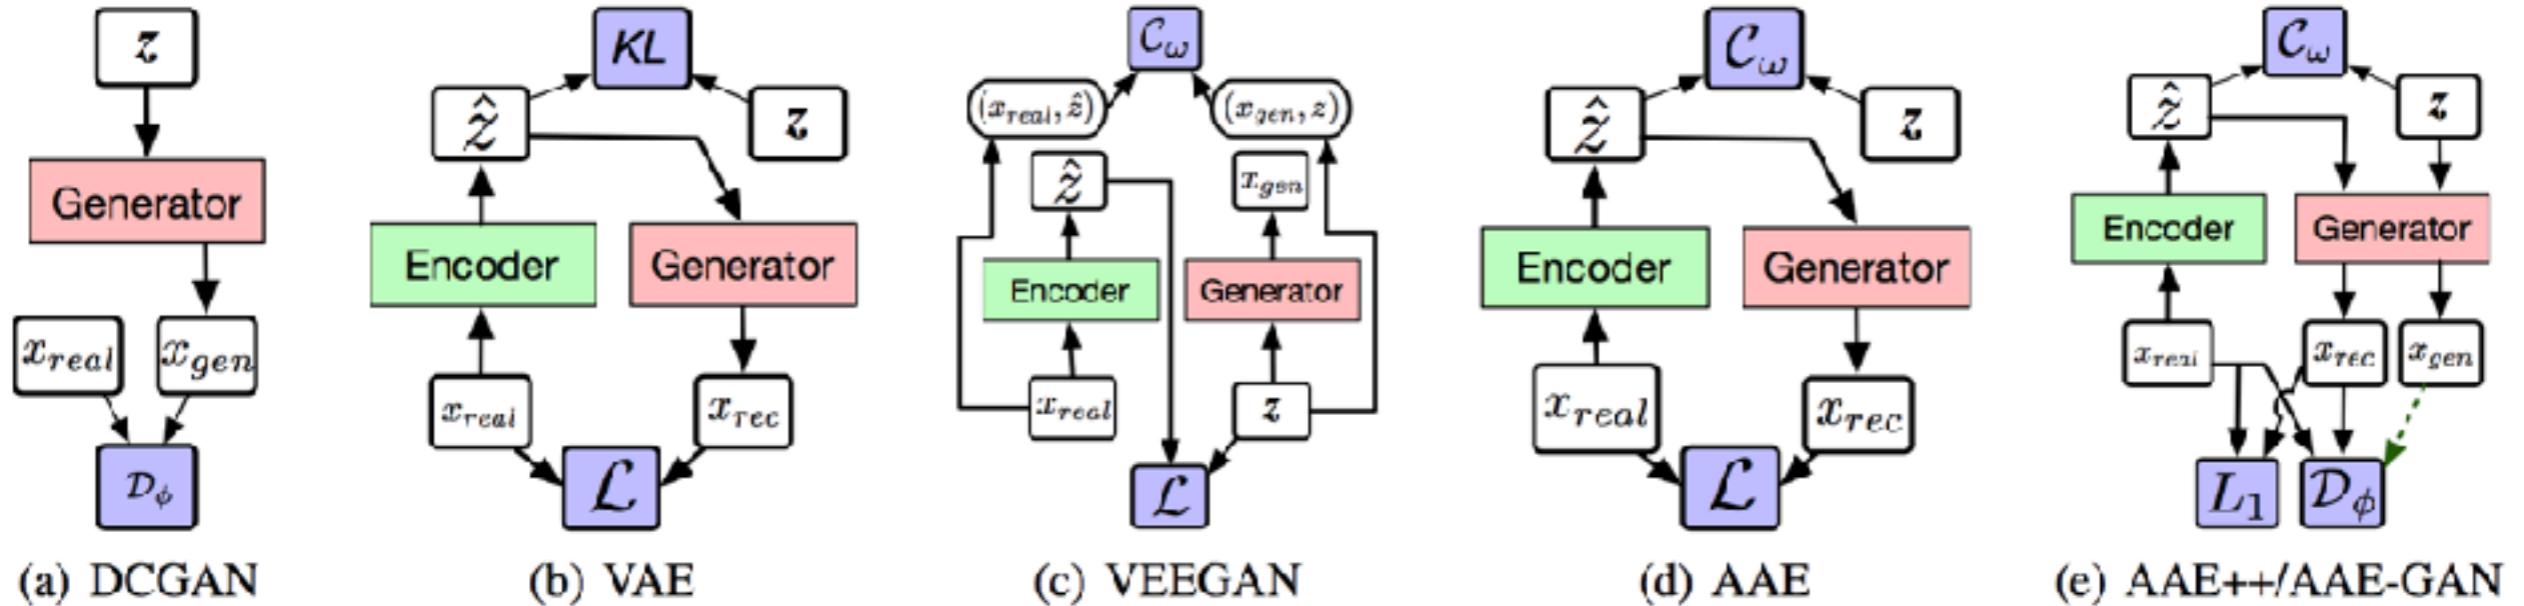
Tangent of posterior odds.

$$m(\theta) = \mathbb{E}_{p_\theta(s)}[f(s)] \quad \nabla_\phi \log \frac{p(y = +1|x)}{p(y = -1|x)} = \nabla_\phi f_\phi(x)$$

$$\mathcal{F} = \left\| \mathbb{E}_{p^*(x)}[f(x)] - \mathbb{E}_{p(z)}[f(g_\phi(z))] \right\|_2^2$$

- Consistent estimators: the number of moments is greater than the number of model parameters.
- Features should not be co-linear.
- More stable than adversarial training.
- Does not require frequent updating of the classifier.

Convergent Approaches



$$\mathcal{F}(q, \theta) = \mathbb{E}_{q_\phi(z|x)} [\log p_\theta(\mathbf{x}|\mathbf{z})] - \text{KL}[q_\phi(\mathbf{z}|\mathbf{x})||p(\mathbf{z})]$$

$$\mathbb{E}_{p^*(\mathbf{x})} \text{KL}[q(\mathbf{z}|\mathbf{x})||p(\mathbf{z})] = \text{KL}[q(\mathbf{z})||p(\mathbf{z})] + \mathbb{I}[q(\mathbf{z}|\mathbf{x}), p^*(\mathbf{x})]$$

Marginal KL

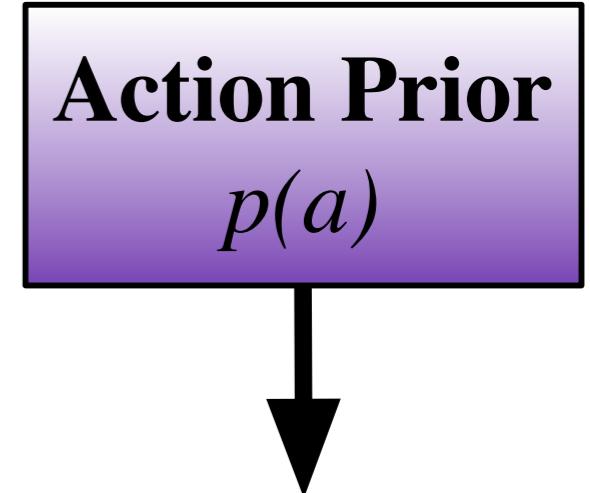
Information

- Replace density ratios by classifiers, replace posteriors with implicit models,
- Views from optimal transport and connections to integral probability metrics.

Environments and Rewards

Environment as a generative process:

- An unknown likelihood;
- Not known analytically;
- Only able to observe its outcomes.



$$a \sim p(a)$$

Prior over actions

$$u(s, a) \sim \text{Environment}(a)$$

Interaction only

$$p(R(s)|a) \propto \exp(u(s, a))$$

Long-term reward

All the key inferential questions can now be asked in this simple framework.

Planning-as-Inference

Simplest question

What is the posterior distribution over actions?
Maximising the probability of the return $\log p(R)$.

Variational inference in the hierarchical model

$$\mathcal{F}(\theta) = \mathbb{E}_{\pi(\mathbf{a}|\mathbf{s})}[R(s, a)] - \text{KL}[\pi_\theta(\mathbf{a}|s) || p(\mathbf{a})]$$

Recover policy search methods:

- Uniform prior over distributions
- Continuous policy parameters
- Can evaluate environment, but not differentiate.

Action Prior
 $p(a)$



Environment
or Model
 $p(R(s,a))$

Policy Search

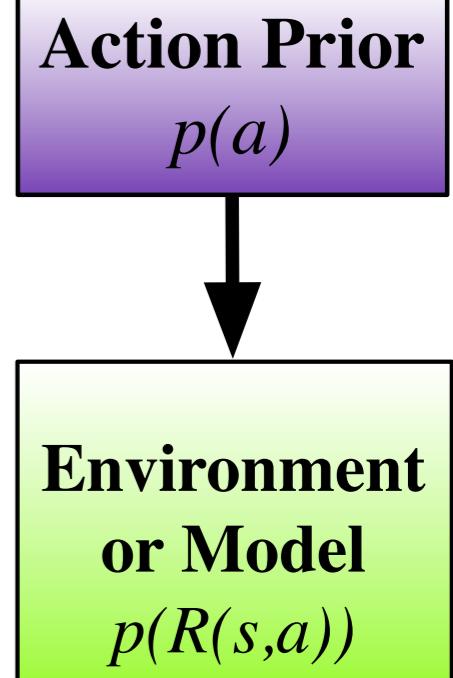
Free Energy

$$\mathcal{F}(\theta) = \mathbb{E}_{\pi(\mathbf{a}|\mathbf{s})}[R(s, a)] - \text{KL}[\pi_\theta(\mathbf{a}|s) \| p(\mathbf{a})]$$

Policy gradient using score-function gradient

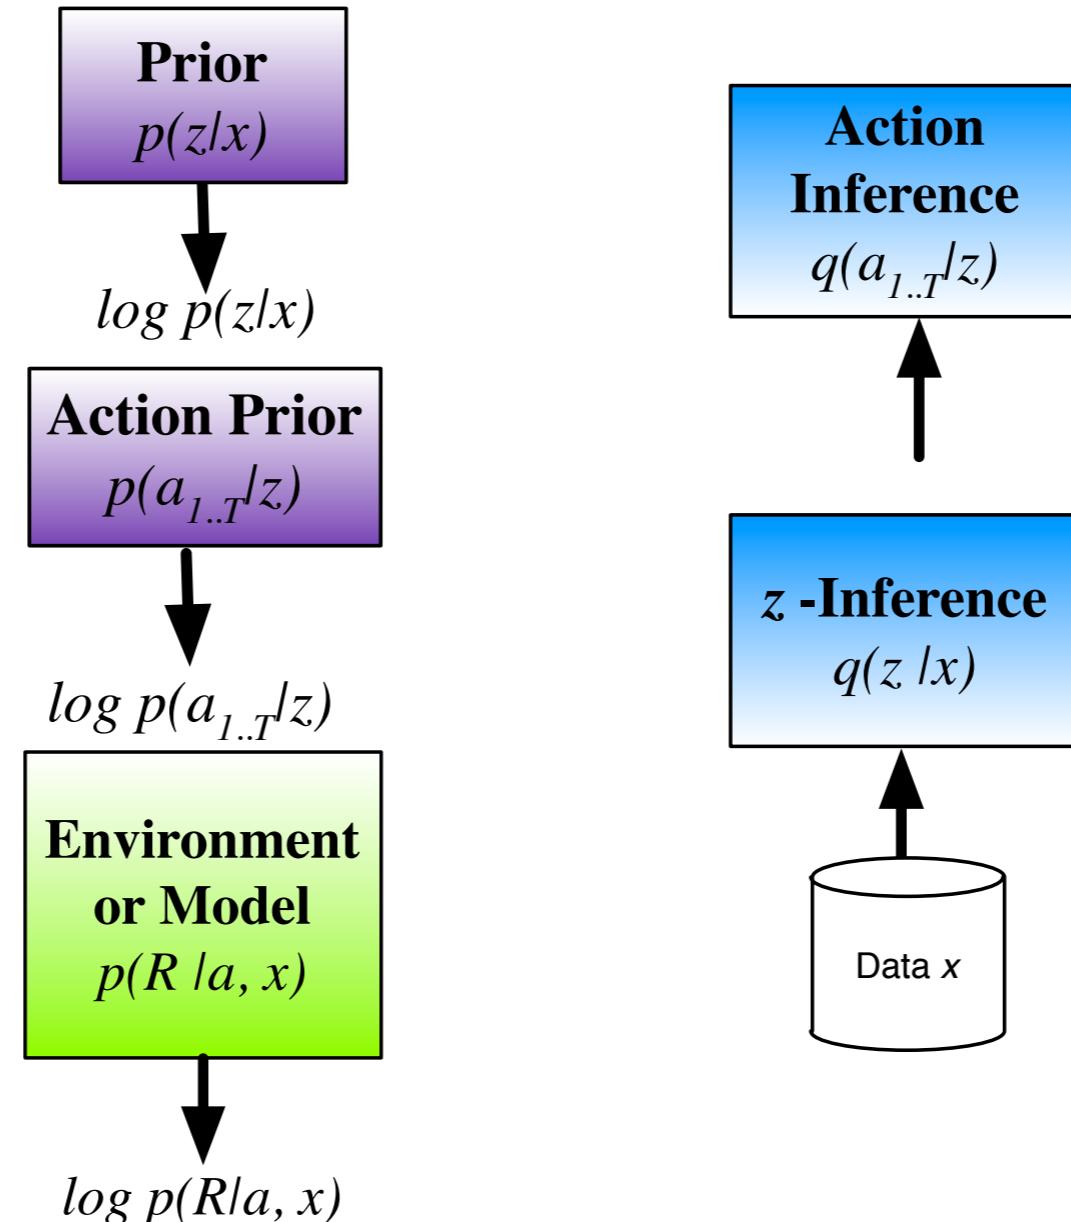
$$\nabla_\theta \mathcal{F}(\theta) = \mathbb{E}_{\pi(\mathbf{a}|\mathbf{s})}[(R(s, a) - c)\nabla_\theta \log \pi_\theta(\mathbf{a}|\mathbf{s})] + \nabla_\theta \mathbb{H}[\pi_\theta(\mathbf{a}|s)]$$

- Appearance of the entropy penalty is natural and alternative priors easy to consider.
- Can easily incorporate prior knowledge of the action space.
- Use any of the tools of probabilistic inference available.
- Easily handle stochastic and deterministic policies.



Identity | Log-derivative

Hierarchical Planning

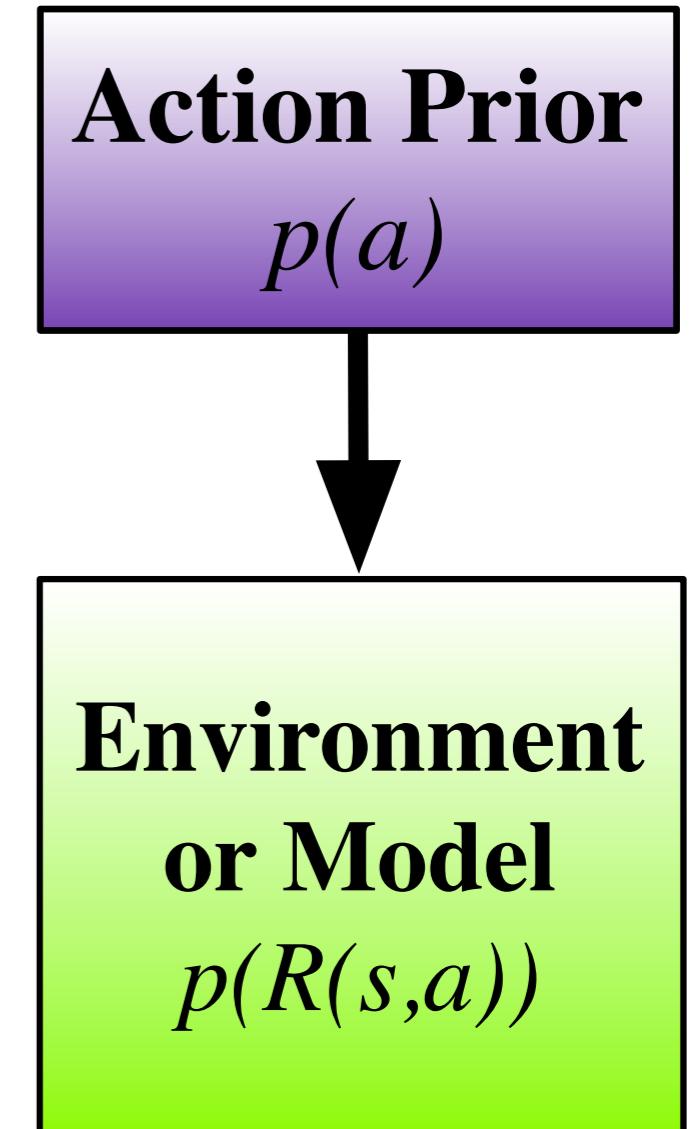


Variational MDP

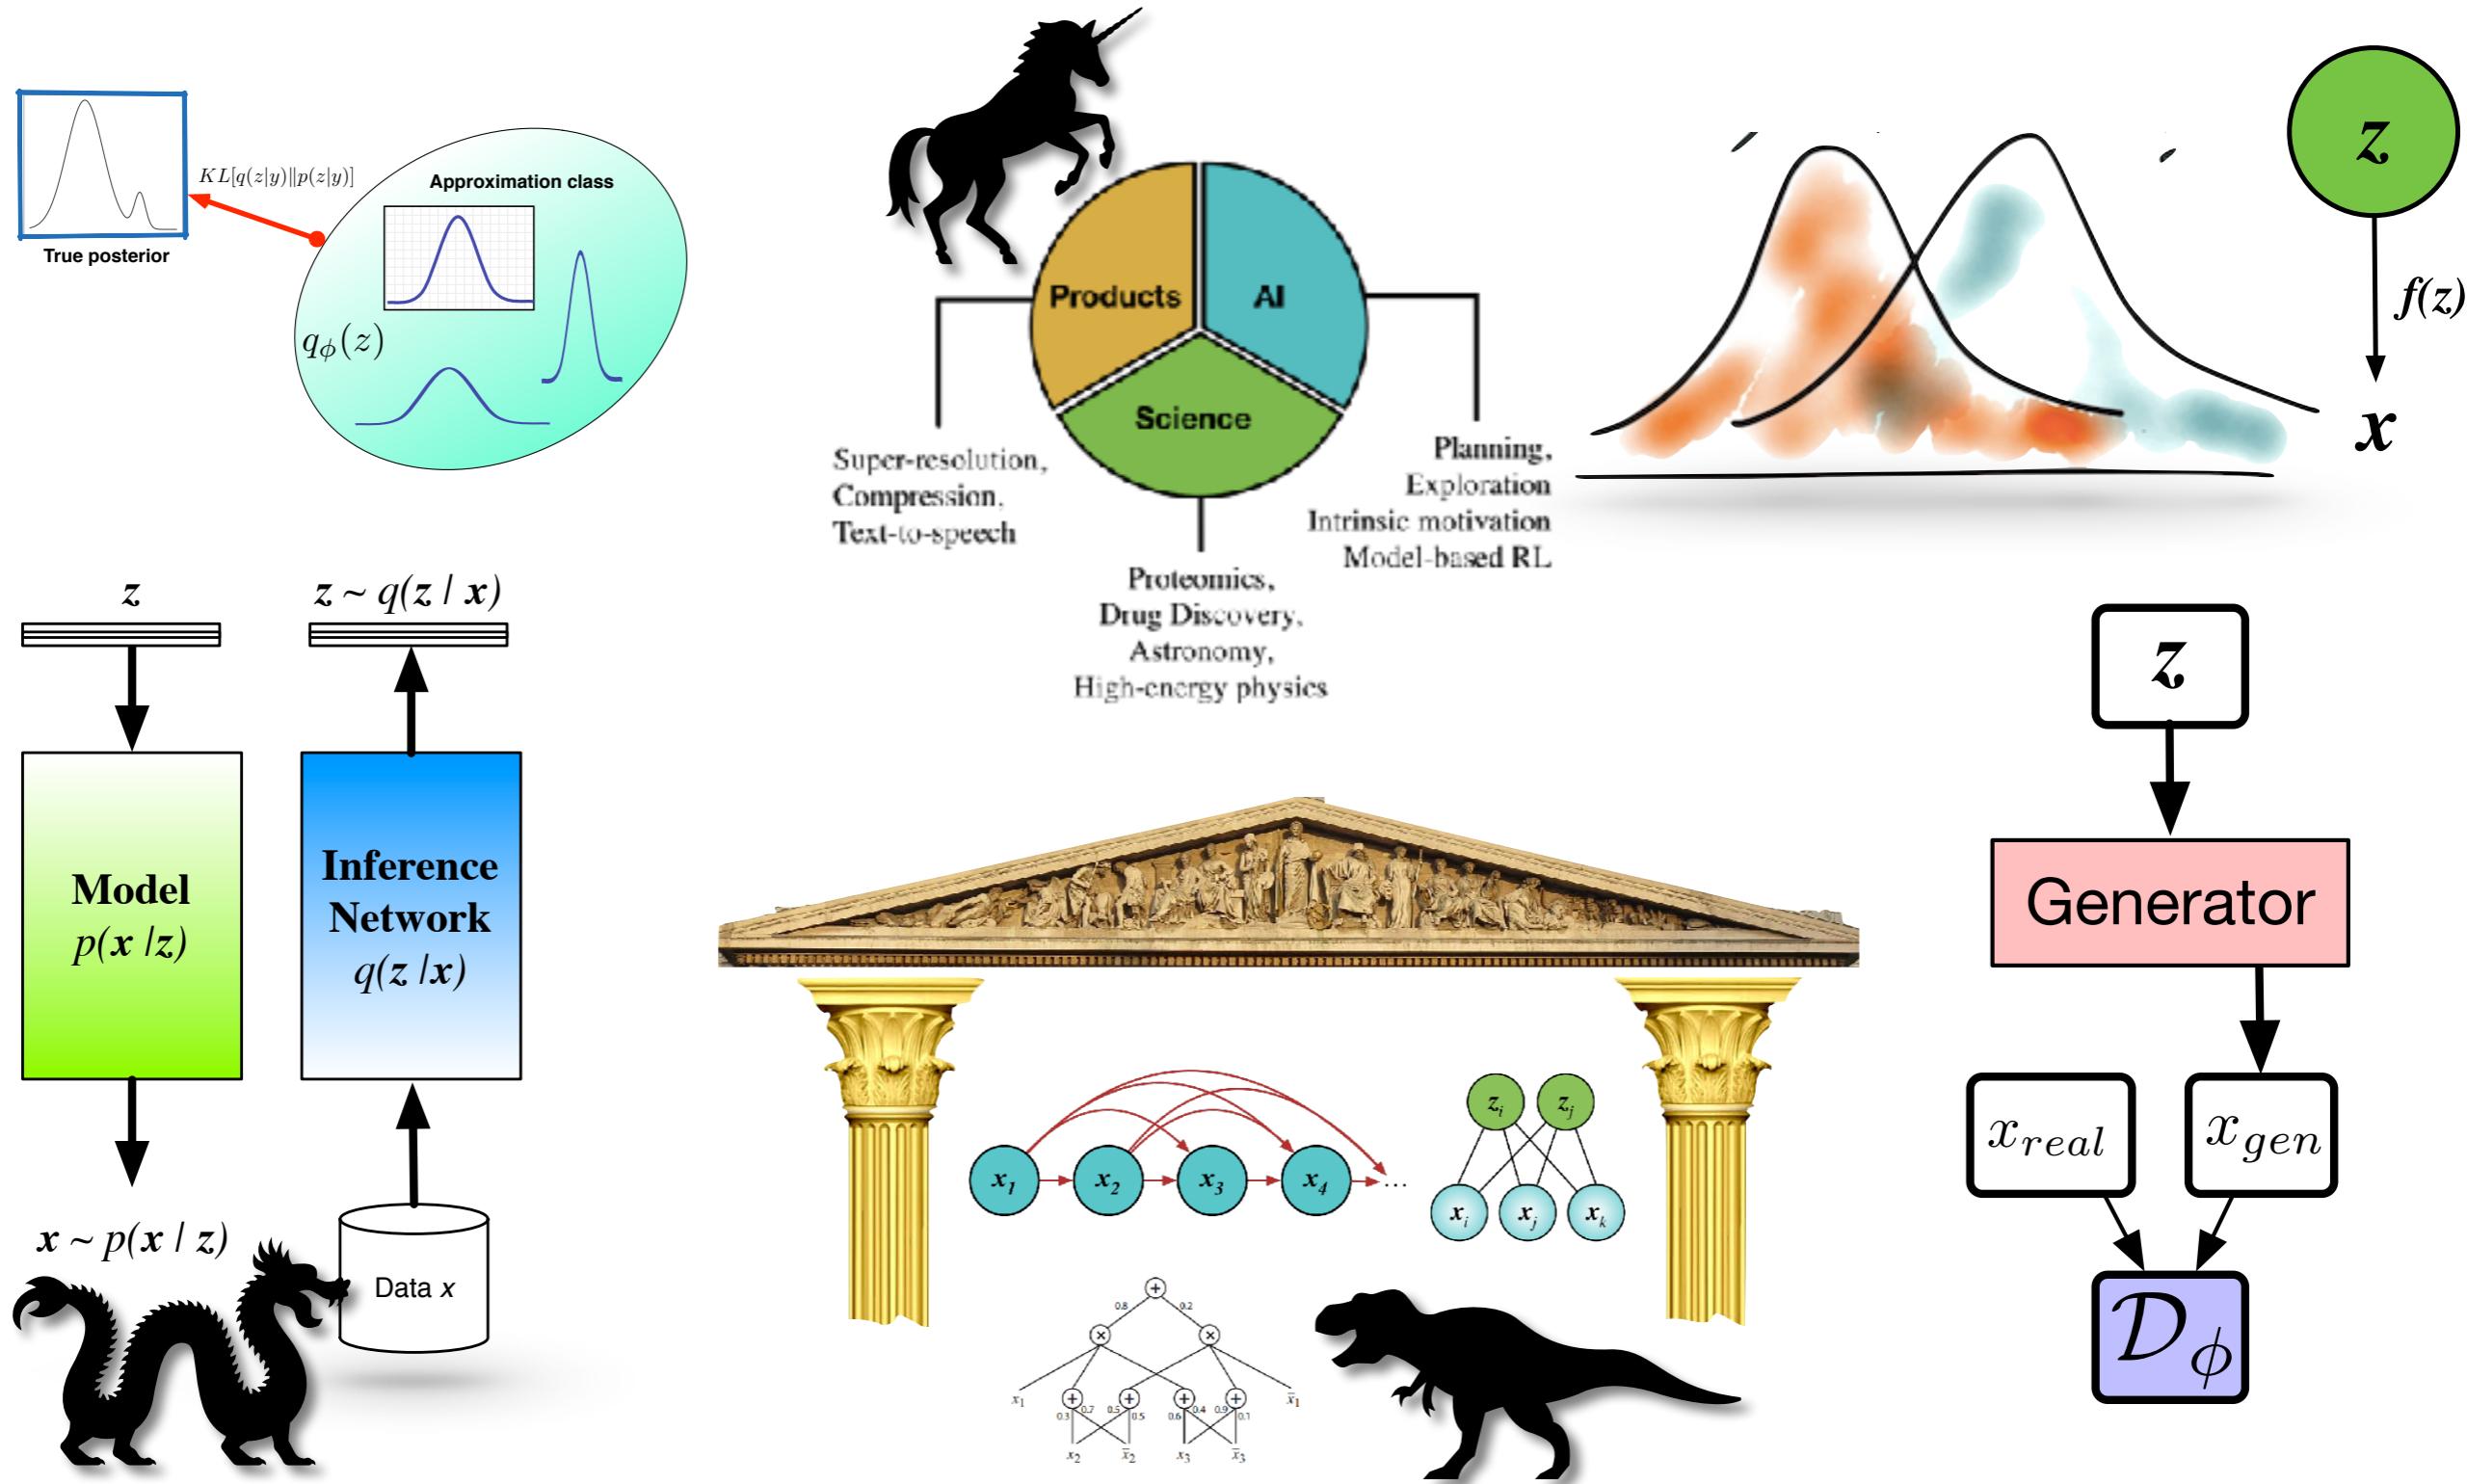
$$\mathcal{F}^\pi(\theta) = \mathbb{E}_{q(a,z|x)}[R(a|x)] - \alpha KL[q_\theta(\mathbf{z}|\mathbf{x})||p(\mathbf{z}|\mathbf{x})] + \alpha \mathbb{H}[\pi_\theta(\mathbf{a}|\mathbf{z})]$$

With a more realistic expansion as graphical model

- Derive Bellman's equation as a different writing of message passing.
- Application of the EM algorithm for policy search becomes possible.
- Easily consider other variational methods, like EP.
- Both model-free and model-based methods emerge.



Final Words



Build Communities



DEEP
LEARNING
INDABA

shakirm.com/feedback

Planting the Seeds of Probabilistic Thinking

Foundations | Tricks | Algorithms

Shakir Mohamed

Research Scientist, DeepMind

Not exhaustive list, and many references to be updated.

Probabilistic Thinking

Cheeseman, P.C., 1985, August. In Defense of Probability. In *IJCAI* (Vol. 2, pp. 1002-1009).

Murphy, K.P., 2012. Machine Learning: A Probabilistic Perspective.

Efron, B., 1982. Maximum likelihood and decision theory. *The annals of Statistics*, pp.340-356.

Stochastic Optimisation

P L'Ecuyer, Note: On the interchange of derivative and expectation for likelihood ratio derivative estimators, Management Science, 1995

Peter W Glynn, Likelihood ratio gradient estimation for stochastic systems, Communications of the ACM, 1990

Michael C Fu, Gradient estimation, Handbooks in operations research and management science, 2006

Ronald J Williams, Simple statistical gradient-following algorithms for connectionist reinforcement learning, Machine learning, 1992

Paul Glasserman, Monte Carlo methods in financial engineering, , 2003

Luc Devroye, Random variate generation in one line of code, Proceedings of the 28th conference on Winter simulation, 1996

L. Devroye, Non-uniform random variate generation, , 1986

Omiros Papaspiliopoulos, Gareth O Roberts, Martin Skold, A general framework for the parametrization of hierarchical models, Statistical Science, 2007

Michael C Fu, Gradient estimation, Handbooks in operations research and management science, 2006

Ranganath, Rajesh, Sean Gerrish, and David M. Blei. "Black Box Variational Inference." In AISTATS, pp. 814-822. 2014.

Mnih, Andriy, and Karol Gregor. "Neural variational inference and learning in belief networks." arXiv preprint arXiv:1402.0030 (2014).

Lázaro-Gredilla, Miguel. "Doubly stochastic variational Bayes for non-conjugate inference." (2014).

Wingate, David, and Theophane Weber. "Automated variational inference in probabilistic programming." arXiv preprint arXiv: 1301.1299 (2013).

Paisley, John, David Blei, and Michael Jordan. "Variational Bayesian inference with stochastic search." arXiv preprint arXiv: 1206.6430 (2012).

- Applications of Deep Generative Models
- Rezende, Danilo Jimenez, Shakir Mohamed, and Daan Wierstra. "Stochastic backpropagation and approximate inference in deep generative models." ICML 2014
- Kingma, Diederik P., and Max Welling. "Auto-encoding variational bayes." ICLR 2014
- Gregor, Karol, et al. "Towards Conceptual Compression." arXiv preprint arXiv:1604.08772 (2016).
- Eslami, S. M., Heess, N., Weber, T., Tassa, Y., Kavukcuoglu, K., & Hinton, G. E. (2016). Attend, Infer, Repeat: Fast Scene Understanding with Generative Models. arXiv preprint arXiv:1603.08575.
- Oh, Junhyuk, Xiaoxiao Guo, Honglak Lee, Richard L. Lewis, and Satinder Singh. "Action-conditional video prediction using deep networks in atari games." In Advances in Neural Information Processing Systems, pp. 2863-2871. 2015.
- Rezende, Danilo Jimenez, Shakir Mohamed, Ivo Danihelka, Karol Gregor, and Daan Wierstra. "One-Shot Generalization in Deep Generative Models." arXiv preprint arXiv:1603.05106 (2016).
- Rezende, Danilo Jimenez, S. M. Eslami, Shakir Mohamed, Peter Battaglia, Max Jaderberg, and Nicolas Heess. "Unsupervised Learning of 3D Structure from Images." arXiv preprint arXiv:1607.00662 (2016).
- Kingma, Diederik P., Shakir Mohamed, Danilo Jimenez Rezende, and Max Welling. "Semi-supervised learning with deep generative models." In Advances in Neural Information Processing Systems, pp. 3581-3589. 2014.
- Maaløe, Lars, Casper Kaae Sønderby, Søren Kaae Sønderby, and Ole Winther. "Auxiliary Deep Generative Models." arXiv preprint arXiv:1602.05473 (2016).
- Odena, Augustus. "Semi-Supervised Learning with Generative Adversarial Networks." arXiv preprint arXiv:1606.01583 (2016).
- Springenberg, Jost Tobias. "Unsupervised and Semi-supervised Learning with Categorical Generative Adversarial Networks." arXiv preprint arXiv:1511.06390 (2015).
- Blundell, Charles, Benigno Uria, Alexander Pritzel, Yazhe Li, Avraham Ruderman, Joel Z. Leibo, Jack Rae, Daan Wierstra, and Demis Hassabis. "Model-Free Episodic Control." arXiv preprint arXiv:1606.04460 (2016).
- Higgins, Irina, Loic Matthey, Xavier Glorot, Arka Pal, Benigno Uria, Charles Blundell, Shakir Mohamed, and Alexander Lerchner. "Early Visual Concept Learning with Unsupervised Deep Learning." arXiv preprint arXiv:1606.05579 (2016).
- Bellemare, Marc G., Sriram Srinivasan, Georg Ostrovski, Tom Schaul, David Saxton, and Remi Munos. "Unifying Count-Based Exploration and Intrinsic Motivation." arXiv preprint arXiv:1606.01868 (2016).
- Alexander (Sasha) Vezhnevets, Mnih, Volodymyr, John Agapiou, Simon Osindero, Alex Graves, Oriol Vinyals, and Koray Kavukcuoglu. "Strategic Attentive Writer for Learning Macro-Actions." arXiv preprint arXiv:1606.04695 (2016).
- Gregor, Karol, Ivo Danihelka, Alex Graves, Danilo Jimenez Rezende, and Daan Wierstra. "DRAW: A recurrent neural network for image generation." arXiv preprint arXiv:1502.04623 (2015).

- **Fully-observed Models**

- Oord, Aaron van den, Nal Kalchbrenner, and Koray Kavukcuoglu. "Pixel recurrent neural networks." arXiv preprint arXiv:1601.06759 (2016).
- Larochelle, Hugo, and Iain Murray. "The Neural Autoregressive Distribution Estimator." In AISTATS, vol. 1, p. 2. 2011.
- Uria, Benigno, Iain Murray, and Hugo Larochelle. "A Deep and Tractable Density Estimator." In ICML, pp. 467-475. 2014.
- Veness, Joel, Kee Siong Ng, Marcus Hutter, and Michael Bowling. "Context tree switching." In 2012 Data Compression Conference, pp. 327-336. IEEE, 2012.
- Rue, Havard, and Leonhard Held. Gaussian Markov random fields: theory and applications. CRC Press, 2005.
- Wainwright, Martin J., and Michael I. Jordan. "Graphical models, exponential families, and variational inference." Foundations and Trends® in Machine Learning 1, no. 1-2 (2008): 1-305.

- **Implicit Probabilistic Models**

- Tabak, E. G., and Cristina V. Turner. "A family of nonparametric density estimation algorithms." Communications on Pure and Applied Mathematics 66, no. 2 (2013): 145-164.
- Rezende, Danilo Jimenez, and Shakir Mohamed. "Variational inference with normalizing flows." arXiv preprint arXiv:1505.05770 (2015).
- Goodfellow, Ian, Jean Pouget-Abadie, Mehdi Mirza, Bing Xu, David Warde-Farley, Sherjil Ozair, Aaron Courville, and Yoshua Bengio. "Generative adversarial nets." In Advances in Neural Information Processing Systems, pp. 2672-2680. 2014.
- Verrelst, Herman, Johan Suykens, Joos Vandewalle, and Bart De Moor. "Bayesian learning and the Fokker-Planck machine." In Proceedings of the International Workshop on Advanced Black-box Techniques for Nonlinear Modeling, Leuven, Belgium, pp. 55-61. 1998.
- Devroye, Luc. "Random variate generation in one line of code." In Proceedings of the 28th conference on Winter simulation, pp. 265-272. IEEE Computer Society, 1996.
- Ravuri, S., Mohamed, S., Rosca, M. and Vinyals, O., 2018. Learning Implicit Generative Models with the Method of Learned Moments. *arXiv preprint arXiv:1806.11006*.

Latent variable models

- Dayan, Peter, Geoffrey E. Hinton, Radford M. Neal, and Richard S. Zemel. "The helmholtz machine." *Neural computation* 7, no. 5 (1995): 889-904.
- Hyvärinen, A., Karhunen, J., & Oja, E. (2004). *Independent component analysis* (Vol. 46). John Wiley & Sons.
- Gregor, Karol, Ivo Danihelka, Andriy Mnih, Charles Blundell, and Daan Wierstra. "Deep autoregressive networks." arXiv preprint arXiv:1310.8499 (2013).
- Ghahramani, Zoubin, and Thomas L. Griffiths. "Infinite latent feature models and the Indian buffet process." In *Advances in neural information processing systems*, pp. 475-482. 2005.
- Teh, Yee Whye, Michael I. Jordan, Matthew J. Beal, and David M. Blei. "Hierarchical dirichlet processes." *Journal of the american statistical association* (2012).
- Adams, Ryan Prescott, Hanna M. Wallach, and Zoubin Ghahramani. "Learning the Structure of Deep Sparse Graphical Models." In *AISTATS*, pp. 1-8. 2010.
- Lawrence, Neil D. "Gaussian process latent variable models for visualisation of high dimensional data." *Advances in neural information processing systems* 16.3 (2004): 329-336.
- Damianou, Andreas C., and Neil D. Lawrence. "Deep Gaussian Processes." In *AISTATS*, pp. 207-215. 2013.
- Mattos, César Lincoln C., Zhenwen Dai, Andreas Damianou, Jeremy Forth, Guilherme A. Barreto, and Neil D. Lawrence. "Recurrent Gaussian Processes." arXiv preprint arXiv:1511.06644 (2015).
- Salakhutdinov, Ruslan, Andriy Mnih, and Geoffrey Hinton. "Restricted Boltzmann machines for collaborative filtering." In *Proceedings of the 24th international conference on Machine learning*, pp. 791-798. ACM, 2007.
- Saul, Lawrence K., Tommi Jaakkola, and Michael I. Jordan. "Mean field theory for sigmoid belief networks." *Journal of artificial intelligence research* 4, no. 1 (1996): 61-76.
- Frey, Brendan J., and Geoffrey E. Hinton. "Variational learning in nonlinear Gaussian belief networks." *Neural Computation* 11, no. 1 (1999): 193-213.

Inference and Learning

- Jordan, Michael I., Zoubin Ghahramani, Tommi S. Jaakkola, and Lawrence K. Saul. "An introduction to variational methods for graphical models." *Machine learning* 37, no. 2 (1999): 183-233.
- Hoffman, Matthew D., David M. Blei, Chong Wang, and John William Paisley. "Stochastic variational inference." *Journal of Machine Learning Research* 14, no. 1 (2013): 1303-1347.
- Honkela, Antti, and Harri Valpola. "Variational learning and bits-back coding: an information-theoretic view to Bayesian learning." *IEEE Transactions on Neural Networks* 15, no. 4 (2004): 800-810.
- Burda, Yuri, Roger Grosse, and Ruslan Salakhutdinov. "Importance weighted autoencoders." arXiv preprint arXiv:1509.00519 (2015).
- Li, Yingzhen, and Richard E. Turner. "Variational Inference with R\'enyi Divergence." arXiv preprint arXiv:1602.02311 (2016).
- Borgwardt, Karsten M., and Zoubin Ghahramani. "Bayesian two-sample tests." arXiv preprint arXiv:0906.4032 (2009).
- Gutmann, Michael, and Aapo Hyv\"arinen. "Noise-contrastive estimation: A new estimation principle for unnormalized statistical models." *AISTATS*. Vol. 1. No. 2. 2010.
- Tsuboi, Yuta, Hisashi Kashima, Shohei Hido, Steffen Bickel, and Masashi Sugiyama. "Direct Density Ratio Estimation for Large-scale Covariate Shift Adaptation." *Information and Media Technologies* 4, no. 2 (2009): 529-546.
- Sugiyama, Masashi, Taiji Suzuki, and Takafumi Kanamori. *Density ratio estimation in machine learning*. Cambridge University Press, 2012.

Amortised Inference

- Gershman, Samuel J., and Noah D. Goodman. "Amortized inference in probabilistic reasoning." In *Proceedings of the 36th Annual Conference of the Cognitive Science Society*. 2014.
- Rezende, Danilo Jimenez, Shakir Mohamed, and Daan Wierstra. "Stochastic backpropagation and approximate inference in deep generative models." arXiv preprint arXiv:1401.4082 (2014).
- Heess, Nicolas, Daniel Tarlow, and John Winn. "Learning to pass expectation propagation messages." In *Advances in Neural Information Processing Systems*, pp. 3219-3227. 2013.
- Jitkrittum, Wittawat, Arthur Gretton, Nicolas Heess, S. M. Eslami, Balaji Lakshminarayanan, Dino Sejdinovic, and Zolt\'an Szab\'o. "Kernel-based just-in-time learning for passing expectation propagation messages." arXiv preprint arXiv:1503.02551 (2015).
- Korattikara, Anoop, Vivek Rathod, Kevin Murphy, and Max Welling. "Bayesian dark knowledge." arXiv preprint arXiv:1506.04416 (2015).